



BIG DATA and AI for business


Deep Learning (1)

**Decisions, Operations & Information Technologies
Robert H. Smith School of Business
Fall, 2020**



Deep learning attracts lots of attention.

- I believe you have seen lots of exciting results before.



Deep learning trends
at Google. Source:
SIGMOD/Jeff Dean

We mainly focus on the basic techniques.

Introduction to Deep Learning

Outline

Introduction of Deep Learning

“Hello World” for Deep Learning

Tips for Deep Learning

Machine Learning

≈ Looking for a Function

- Speech Recognition

$$f\left(\begin{array}{c} \text{[Speaker icon]} \\ \text{[Waveform plot]} \end{array} \right) = \text{“How are you”}$$

- Image Recognition

$$f\left(\begin{array}{c} \text{[Speaker icon]} \\ \text{[Image of a kitten]} \end{array} \right) = \text{“Cat”}$$

- Playing Go


$$\left. \begin{array}{c} f\left(\text{[Speaker icon]} \right) \\ \text{“5-5”} \end{array} \right) = \text{“5-5”} \quad (\text{next move})$$

- Dialogue System


$$\left. \begin{array}{c} f\left(\begin{array}{c} \text{“Hi”} \\ \text{(what the user said)} \end{array} \right) \\ \text{“Hello”} \end{array} \right) = \text{“Hello”} \quad (\text{system response})$$

Image Recognition:
Framework

$$f(\text{cat image}) = \text{"cat"}$$

$$f_1(\text{cat image}) = \text{"cat"}$$
$$f_2(\text{cat image}) = \text{"money"}$$
$$f_1(\text{dog image}) = \text{"dog"}$$
$$f_2(\text{dog image}) = \text{"snake"}$$

Image Recognition: Framework

$$f(\text{cat image}) = \text{"cat"}$$


Model
 $f_1, f_2 \dots$

$f_1(\text{cat image}) = \text{"cat"}$ $f_2(\text{cat image}) = \text{"money"}$

$f_1(\text{dog image}) = \text{"dog"}$ $f_2(\text{dog image}) = \text{"snake"}$

Better!

Supervised Learning

function input:




function output:
"monkey" "cat" "dog"

Framework

Image Recognition:

$$f(\text{cat}) = \text{"cat"}$$



Three Steps for Deep Learning

Step 1: define a set of function

Neural Network



Step 2: goodness of function




Step 3: pick the best function

Neural Network


Neuron

$$z = a_1 w_1 + \cdots + a_k w_k + \cdots + a_K w_K + b$$




Neural Network

Neuron




Sigmoid Function

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$




Neural Network

Different connections lead to different network structures




The neurons have different values of weights and biases.

Weights and biases are network parameters θ




Fully Connect Feedforward Network




Sigmoid Function


$$\sigma(z) = \frac{1}{1 + e^{-z}}$$



Fully Connect Feedforward Network



Fully Connect Feedforward Network



This is a function.


$$f([1@-1]) = [0.62@0.85, 0@0] = [0.51@0.85]$$

Input vector, output vector

Given parameters θ , define a function

Given network structure, define a function set

Fully Connect Feedforward Network



Deep means many hidden layers


Why Deep? Universality Theorem

Any continuous function f

$$f : R^N \rightarrow R^M$$

Can be realized by a network
with one hidden layer

(given **enough** hidden
neurons)



Reference for the reason:
[http://
neuralnetworksanddeeplearning.
com/chap4.html](http://neuralnetworksanddeeplearning.com/chap4.html)

Why “Deep” neural network not “Fat” neural
network?

Why Deep? Analogy

Logic circuits

- Logic circuits consists of **gates**
- **A two layers of logic gates** can represent **any Boolean function.**
- Using multiple layers of logic gates to build some functions are much simpler




Neural network

- Neural network consists of **neurons**
- **A hidden layer network** can represent **any continuous function.**
- Using multiple layers of neurons to represent some functions are much simpler




less gates needed



less
parameters

less
data?




More reason: https://www.youtube.com/watch?v=XsC9byQkUH8&list=PLJV_el3uVTsPy9oCRY30oBPNLCo89yu49&index=13

Deep = Many hidden layers

http://cs231n.stanford.edu/slides/winter1516_lecture8.pdf


16.4%



AlexNet (2012)

8 layers


7.3%



VGG (2014)

19 layers


6.7%



GoogleNet (2014)

22 layers

Deep = Many hidden layers



Output Layer

- Softmax layer as the output layer

Ordinary Layer

$$z_1 \rightarrow \sigma \rightarrow y_1 = \sigma(z_1)$$

$$z_2 \rightarrow \sigma \rightarrow y_2 = \sigma(z_2)$$

$$z_3 \rightarrow \sigma \rightarrow y_3 = \sigma(z_3)$$

In general, the output of network can be any value.

May not be easy to interpret


Output Layer

- Softmax layer as the output layer

Probability:

- $1 > y_i > 0$
- $\sum_i y_i = 1$


Softmax Layer




Example Application



Input




Output




Each dimension represents the confidence of a digit.

Example Application


- Handwriting Digit Recognition



Example Application



You need to decide the network structure to let a good function in your function set.



- Q: How many layers? How many neurons for each layer?

Trial and Error

+

Intuition

- Q: Can we design the network structure?

Convolutional Neural Network (CNN)
in the next lecture

- Q: Can the structure be automatically determined?
 - Yes, but not widely studied yet.

Three Steps for Deep Learning

Step 1: define a set of function



Step 2: goodness of function



Step 3: pick the best function

Training Data

- Preparing training data: images and their labels



“5”



“0”



“4”




“1”



“9”



“2”




“1”



“3”

The learning target is defined on
the training data.

Learning Target



Ink \rightarrow 1


No ink \rightarrow 0

The learning target is

Input: $\rightarrow y_1$ has the maximum value

Input: $\rightarrow y_2$ has the maximum value

A good function should make the loss of all examples as small as possible.



Loss can be **square error** or **cross entropy** between the network output and target

Total Loss

For all training data ...



⋮ ⋮ ⋮ ⋮



Total Loss:

$$L = \sum_{r=1}^R l_r$$

As small as possible

Find a function in function set that minimizes total loss L

Find the network parameters θ^{1*} that minimize total loss L

Three Steps for Deep Learning

Step 1: define a set of function



Step 2: goodness of function



Step 3: pick the best function

How to pick the best function


Find network parameters θ^{1*} that minimize total loss L

Enumerate all possible values

Network parameters $\theta = \{w^{l1}, w^{l2}, w^{l3}, \dots, b^{l1}, b^{l2}, \dots\}$

Millions of parameters


E.g. speech recognition: 8 layers and 1000 neurons each layer



Gradient Descent

Network parameters $\theta = \{w_1, w_2, \dots, b_1, b_2, \dots\}$


Find network parameters θ^* that minimize total loss L



Gradient Descent

Network parameters $\theta = \{w_1, w_2, \dots, b_1, b_2, \dots\}$


Find network parameters θ^* that minimize total loss L



Gradient Descent

Network parameters $\theta = \{w_1, w_2, \dots, b_1, b_2, \dots\}$


Find network parameters θ^* that minimize total loss L




Gradient Descent

Network parameters $\theta = \{w_1, w_2, \dots, b_1, b_2, \dots\}$

Find network parameters θ^* that minimize total loss L




Gradient Descent




Gradient Descent

Hopfully, we would reach
a minima




Local Minima



Local Minima

- Gradient descent never guarantee global minima




Gradient Descent

This is the “learning” of machines in deep learning

→ Even alpha go using this approach.

People image



Actually



I hope you are not too disappointed :p

Backpropagation

- Backpropagation: an efficient way to compute $\partial L / \partial w$ in neural network



Caffe



theano



Three Steps for Deep Learning



Deep Learning is so simple


Now If you want to find a function

If you have lots of function input/output (?) as training data


 You can use deep learning

For example, you can do


- Image Recognition



For example, you can do



For example, you can do




Outline

Introduction of Deep Learning

“Hello World” for Deep Learning

Tips for Deep Learning

Keras



Interface of
TensorFlow or
Theano



Easy to learn and use
(still have some flexibility)
You can modify it if you can write
TensorFlow or Theano

Very flexible
Need some
effort to learn

Keras

- François Chollet is the author of Keras.
- Keras means *horn* in Greek
- Documentation: <http://keras.io/>
- Example: [https://github.com/fchollet/keras/
tree/master/examples](https://github.com/fchollet/keras/tree/master/examples)


PyTorch

- Tutorial: <https://pytorch.org/tutorials>



Example Application

- Handwriting Digit Recognition



MNIST Data: <http://yann.lecun.com/exdb/mnist/>
“Hello world” for deep learning

Keras provides data sets loading function: <http://keras.io/datasets/>

Keras

Step 1:
define a set
of function

Step 2:
goodness of
function


Step 3: pick
the best
function

28x28

500

Softmax

y_1 y_2 y_{10}



```
model = Sequential()
```

```
model.add( Dense( input_dim=28*28,  
                  output_dim=500 ) )  
model.add( Activation('sigmoid') )
```

```
model.add( Dense( output_dim=500 ) )  
model.add( Activation('sigmoid') )
```


```
model.add( Dense( output_dim=10 ) )  
model.add( Activation('softmax') )
```

Keras

Step 1:
define a set
of function

Step 2:
goodness of
function

Step 3: pick
the best
function



```
model.compile(loss='mse',  
              optimizer=SGD(lr=0.1),  
              metrics=['accuracy'])
```

Keras

Step 1:
define a set
of function

Step 2:
goodness of
function

Step 3: pick
the best
function

Step 3.1: Configuration

```
model.compile(loss='mse',  
               optimizer=SGD(lr=0.1),  
               metrics=['accuracy'])
```

$$w \leftarrow w - \eta \partial L / \partial w$$

0.1

Step 3.2: Find the optimal network parameters

```
model.fit(x_train, y_train, batch_size=100, nb_epoch=20)
```

Training data
(Images)

Labels
(digits)

Keras

Step 1:
define a set
of function


Step 2:
goodness of
function

Step 3: pick
the best
function

Step 3.2: Find the optimal network parameters


```
model.fit(x_train, y_train, batch_size=100, nb_epoch=20)
```

numpy array




Number of training examples

numpy array



Number of training examples



Save and load models

<http://keras.io/getting-started/faq/#how-can-i-save-a-keras-model>

How to use the neural network (testing):

case 1:

```
score = model.evaluate(x_test, y_test)
print('Total loss on Testing Set:', score[0])
print('Accuracy of Testing Set:', score[1])
```

case 2:


```
result = model.predict(x_test)
```

PyTorch

Step 1:
define a set
of function

Step 2:
goodness of
function

Step 3: pick
the best
function



```
import torch.nn as nn
import torch.nn.functional as F

class MyNetwork(nn.Module):
    def __init__(self):
        super(MyNetwork, self).__init__()

        self.fc1 = nn.Linear(28 * 28, 500)
        self.fc2 = nn.Linear(500, 500)
        self.fc3 = nn.Linear(500, 10)


    def forward(self, x):
        x = F.sigmoid(self.fc1(x))
        x = F.sigmoid(self.fc2(x))
        x = self.fc3(x)
        return F.log_softmax(x)
```

PyTorch

Step 1:
define a set
of function

Step 2:
goodness of
function

Step 3: pick
the best
function



```
net = MyNetwork()  
optimizer = torch.optim.SGD(net.parameters(), lr=0.1, momentum=0.9)  
criterion = nn.MSELoss()
```

PyTorch

Step 1:
define a set
of function

Step 2:
goodness of
function

Step 3: pick
the best
function

Step 3.1: Training


```
for epoch in range(epochs):
    for batch_idx, (data, target) in enumerate(train_loader):
        data, target = Variable(data), Variable(target)
        # resize data from (batch_size, 1, 28, 28) to (batch_size, 28*28)
        data = data.view(-1, 28*28)

        optimizer.zero_grad()
        net_out = net(data)
        loss = criterion(net_out, target)
        loss.backward()
        optimizer.step()

        if batch_idx % log_interval == 0:
            print('Train Epoch: {} [{}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
                epoch, batch_idx * len(data), len(train_loader.dataset),
                100. * batch_idx / len(train_loader), loss.data[0]))
```



Step 3.2: Performance report



```

test_loss, correct = 0, 0
for data, target in test_loader:
    data, target = Variable(data, volatile=True), Variable(target)
    data = data.view(-1, 28 * 28)
    net_out = net(data)
    # sum up batch loss
    test_loss += criterion(net_out, target).data[0]
    pred = net_out.data.max(1)[1] # get the index of the max log-probability
    correct += pred.eq(target.data).sum()

test_loss /= len(test_loader.dataset)
print('\nTest set: Average loss: {:.4f}, Accuracy: {}/{} ({:.0f}%)\n'.format(
    test_loss, correct, len(test_loader.dataset),
    100. * correct / len(test_loader.dataset)))

```

Three Steps for Deep Learning

Step 1:
define a set
of function




Step 2:
goodness of
function



Step 3: pick
the best
function

Deep Learning is so simple



Outline




Introduction of Deep Learning



“Hello World” for Deep Learning

Tips for Deep Learning

Recipe of Deep Learning




Do not always blame Overfitting



Recipe of Deep Learning

Different approaches for different problems.

e.g. dropout for good results on testing data




Good Results on Testing Data?

Good Results on Training Data?




YES

Recipe of Deep Learning



Choosing Proper Loss



Square
Error

$$\sum_{i=1}^{10} (y_i - \hat{y}_i)^2 = 0$$

Cross
Entropy

$$-\sum_{i=1}^{10} \hat{y}_i \ln y_i = 0$$

Demo

Square Error

```
model.compile(loss='mse',  
              optimizer=SGD(lr=0.1),  
              metrics=['accuracy'])
```


Cross Entropy

```
model.compile(loss='categorical_crossentropy',  
              optimizer=SGD(lr=0.1),  
              metrics=['accuracy'])
```

Several alternatives: <https://keras.io/objectives/>

Choosing Proper Loss

When using softmax output layer,
choose cross entropy



Recipe of Deep Learning



YES

Choosing proper loss

Mini-batch

New activation function

Adaptive Learning Rate

Momentum

Good Results on
Testing Data?

YES

Good Results on
Training Data?


```
model.fit(x_train, y_train, batch_size=100, nb_epoch=20)
```

We do not really minimize total loss!


Mini-batch

➤ Randomly initialize network parameters

Mini-batch



Mini-batch



- Pick the 1st batch
 $L' = l^1 + l^{31} + \dots$
Update parameters once
- Pick the 2nd batch
 $L'' = l^2 + l^{16} + \dots$
Update parameters once
- ⋮
- Until all mini-batches have been picked


one epoch

Repeat the above process

Mini-batch

```
model.fit(x_train, y_train, batch_size=100, nb_epoch=20)
```

Mini-batch




100 examples in a mini-batch

Repeat 20 times


- Pick the 1st batch
 $L' = l^1 + l^{31} + \dots$
Update parameters once
- Pick the 2nd batch
 $L'' = l^2 + l^{16} + \dots$
Update parameters once
⋮
- Until all mini-batches have been picked

Mini-batch

Original Gradient Descent



With Mini-batch




The colors represent the total loss.

Mini-batch is Faster

Not always true with parallel computing.

Original Gradient Descent


Update after seeing all examples




Mini-batch has better performance!

With Mini-batch


If there are 20 batches, update 20 times in one epoch.




Mini-batch is Faster




Shuffle the training examples for each epoch




Recipe of Deep Learning



Hard to get the power of Deep ...



Vanishing Gradient Problem



Smaller gradients

Learn very slow


Almost random

Larger gradients

Learn very fast

Already converge

Vanishing Gradient Problem

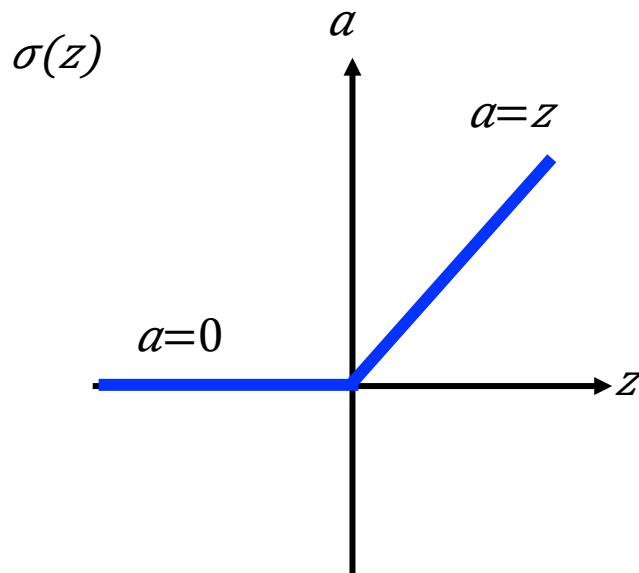


Intuitive way to compute the derivatives ...

$$\frac{\partial l}{\partial w} = ? \quad \frac{\Delta l}{\Delta w}$$

ReLU

- Rectified Linear Unit (ReLU)



Reason:

1. Fast to compute

2. Sparsity

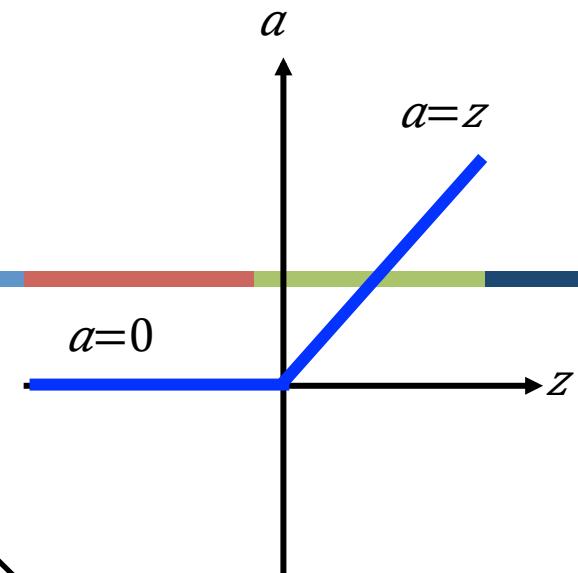
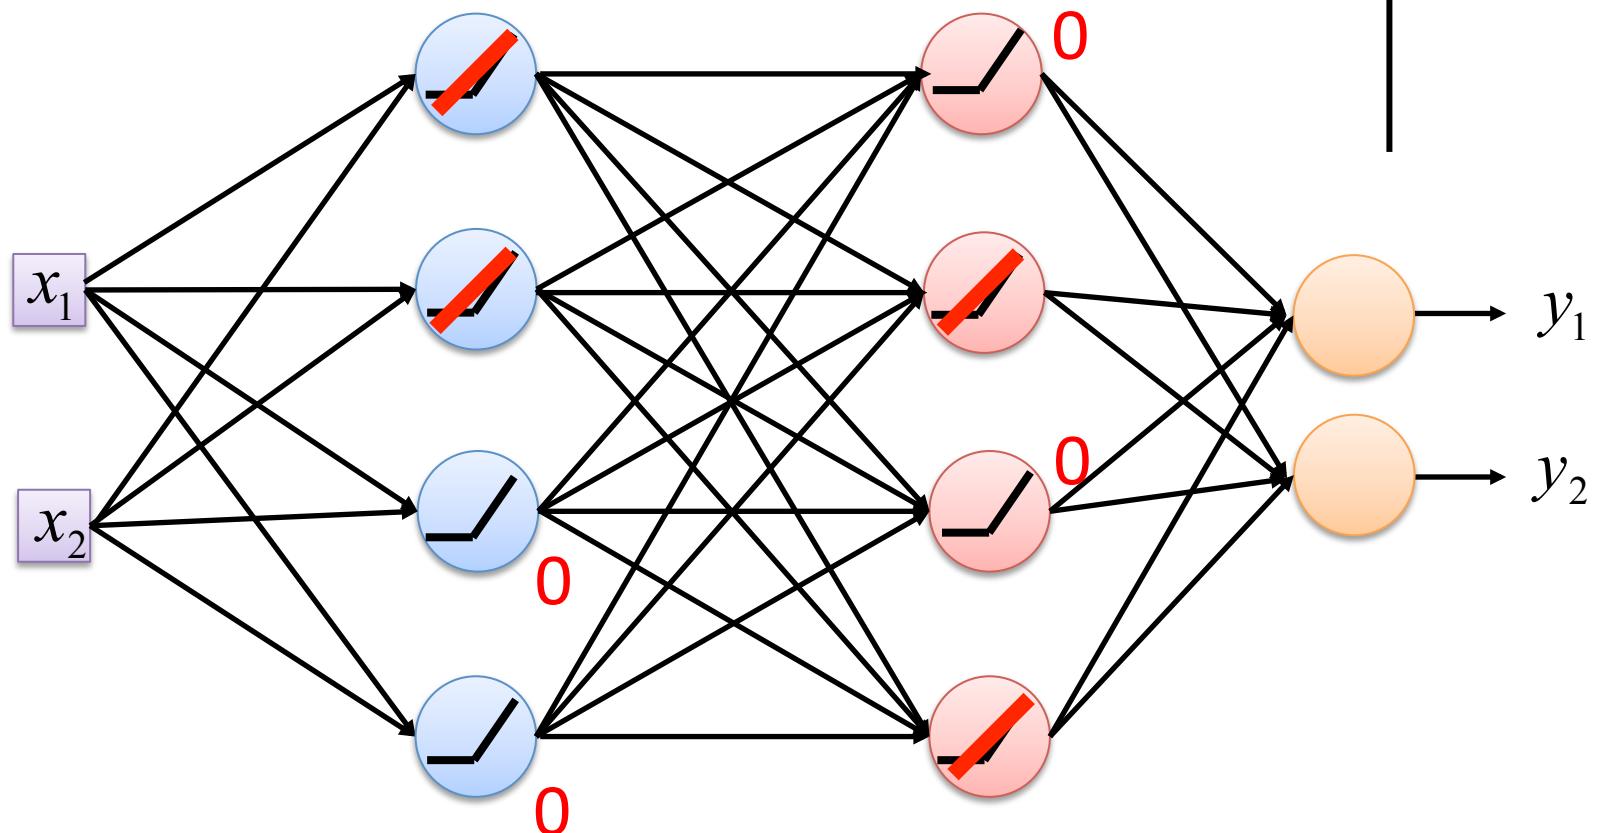
3. Vanishing gradient problem

[Xavier Glorot, AISTATS'11]


[Andrew L. Maas, ICML'13]

[Kaiming He, arXiv'15]


ReLU



ReLU

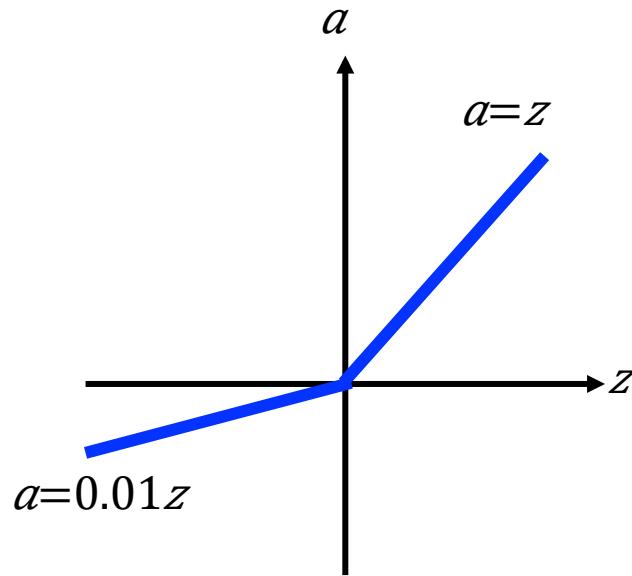


A Thinner linear network

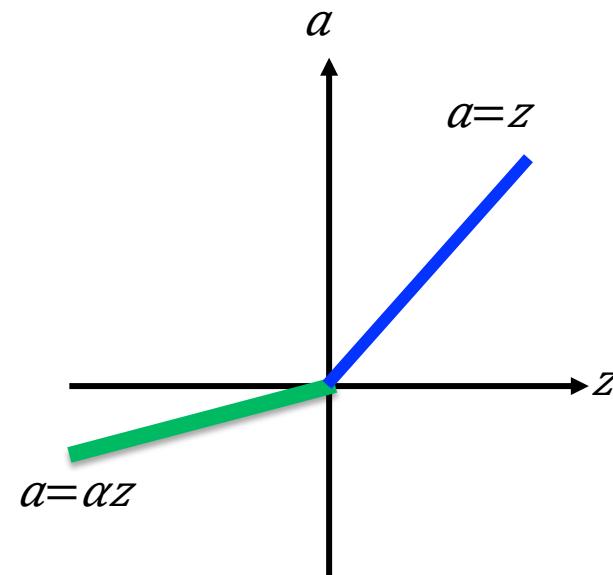


ReLU - variant

Leaky ReLU



Parametric ReLU




α also learned by
gradient descent

Maxout

ReLU is a special cases of Maxout

- Learnable activation function [Ian J. Goodfellow, ICML'13]




You can have more than 2 elements in a group.

Maxout


ReLU is a special cases of Maxout

- Learnable activation function [Ian J. Goodfellow, ICML'13]
 - Activation function in maxout network can be any piecewise linear convex function
 - How many pieces depending on how many elements in a group



2 elements in a group



3 elements in a group




Recipe of Deep Learning




Learning Rates

Set the learning rate η carefully



Learning Rates

Set the learning rate η carefully



Learning Rates

- Popular & Simple Idea: Reduce the learning rate by some factor every few epochs.
 - At the beginning, we are far from the destination, so we use larger learning rate
 - After several epochs, we are close to the destination, so we reduce the learning rate
 - E.g. 1/t decay: $\eta^t = \eta / \sqrt{t + 1}$
- Learning rate cannot be one-size-fits-all
 - Giving different parameters different learning rates

Adagrad

Original: $w \leftarrow w - \eta \partial L / \partial w$

Adagrad: $w \leftarrow w - \boxed{\eta_w} \partial L / \partial w$

Parameter dependent learning rate

$$\eta_w = \frac{\eta}{\sqrt{\sum_{i=0}^t (g^i)^2}}$$

g^i is $\partial L / \partial w$ obtained at the i-th update

Summation of the square of the previous derivatives

$$\eta_w = \frac{\eta}{\sqrt{\sum_{i=0}^t (g^i)^2}}$$

Adagrad

w_1	\mathbf{g}^0
	0.1

w_2	\mathbf{g}^0
	20.0

Learning rate:

$$\frac{\eta}{\sqrt{0.1^2}}$$

$$= \frac{\eta}{0.1}$$



$$\frac{\eta}{\sqrt{0.1^2 + 0.2^2}} = \frac{\eta}{\sqrt{0.01 + 0.04}} = \frac{\eta}{\sqrt{0.05}} = \frac{\eta}{\sqrt{0.05}} = \frac{\eta}{0.22}$$

Learning rate:

$$\frac{\eta}{\sqrt{20^2}}$$

$$= \frac{\eta}{20}$$

$$\frac{\eta}{\sqrt{20^2 + 10^2}} = \frac{\eta}{\sqrt{400 + 100}} = \frac{\eta}{\sqrt{500}} = \frac{\eta}{\sqrt{500}} = \frac{\eta}{22}$$

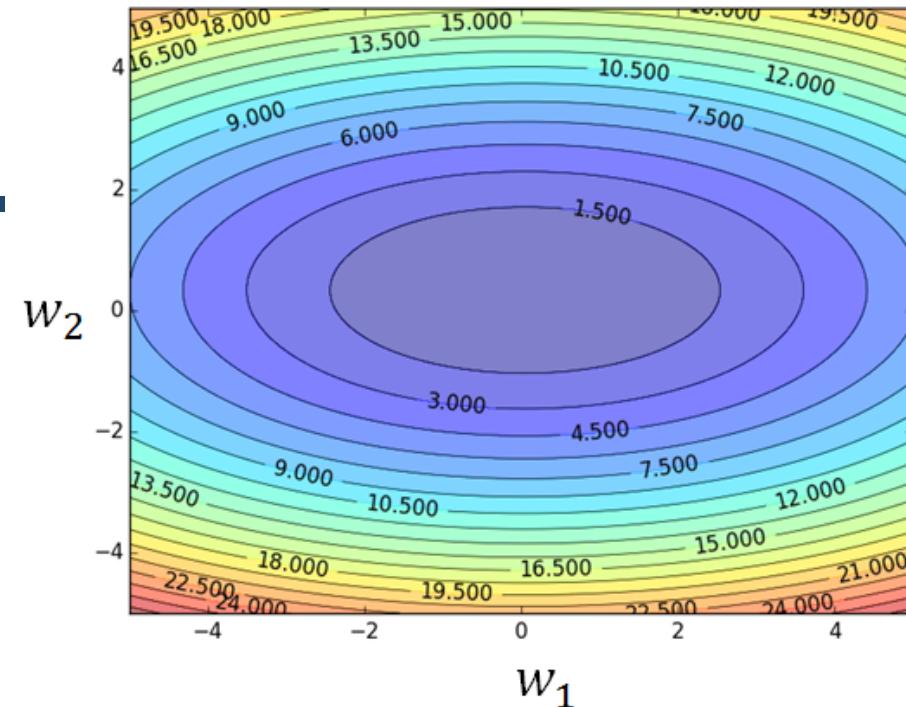


- Observation:**
1. Learning rate is smaller and smaller for all parameters
 2. Smaller derivatives, larger learning rate, and vice versa

Why?

Larger derivatives

Smaller Learning Rate



Smaller Derivatives

Larger Learning Rate



2. Smaller derivatives, larger learning rate, and vice versa

Why?


Not the whole story

- **Adagrad** [John Duchi, JMLR'11]
- **RMSprop**
 - <https://www.youtube.com/watch?v=O3sxAc4hxZU>
- **Adadelta** [Matthew D. Zeiler, arXiv'12]
- “No more pesky learning rates” [Tom Schaul, arXiv'12]
- **AdaSecant** [Caglar Gulcehre, arXiv'14]
- **Adam** [Diederik P. Kingma, ICLR'15]
- **Nadam**
 - http://cs229.stanford.edu/proj2015/054_report.pdf

Recipe of Deep Learning



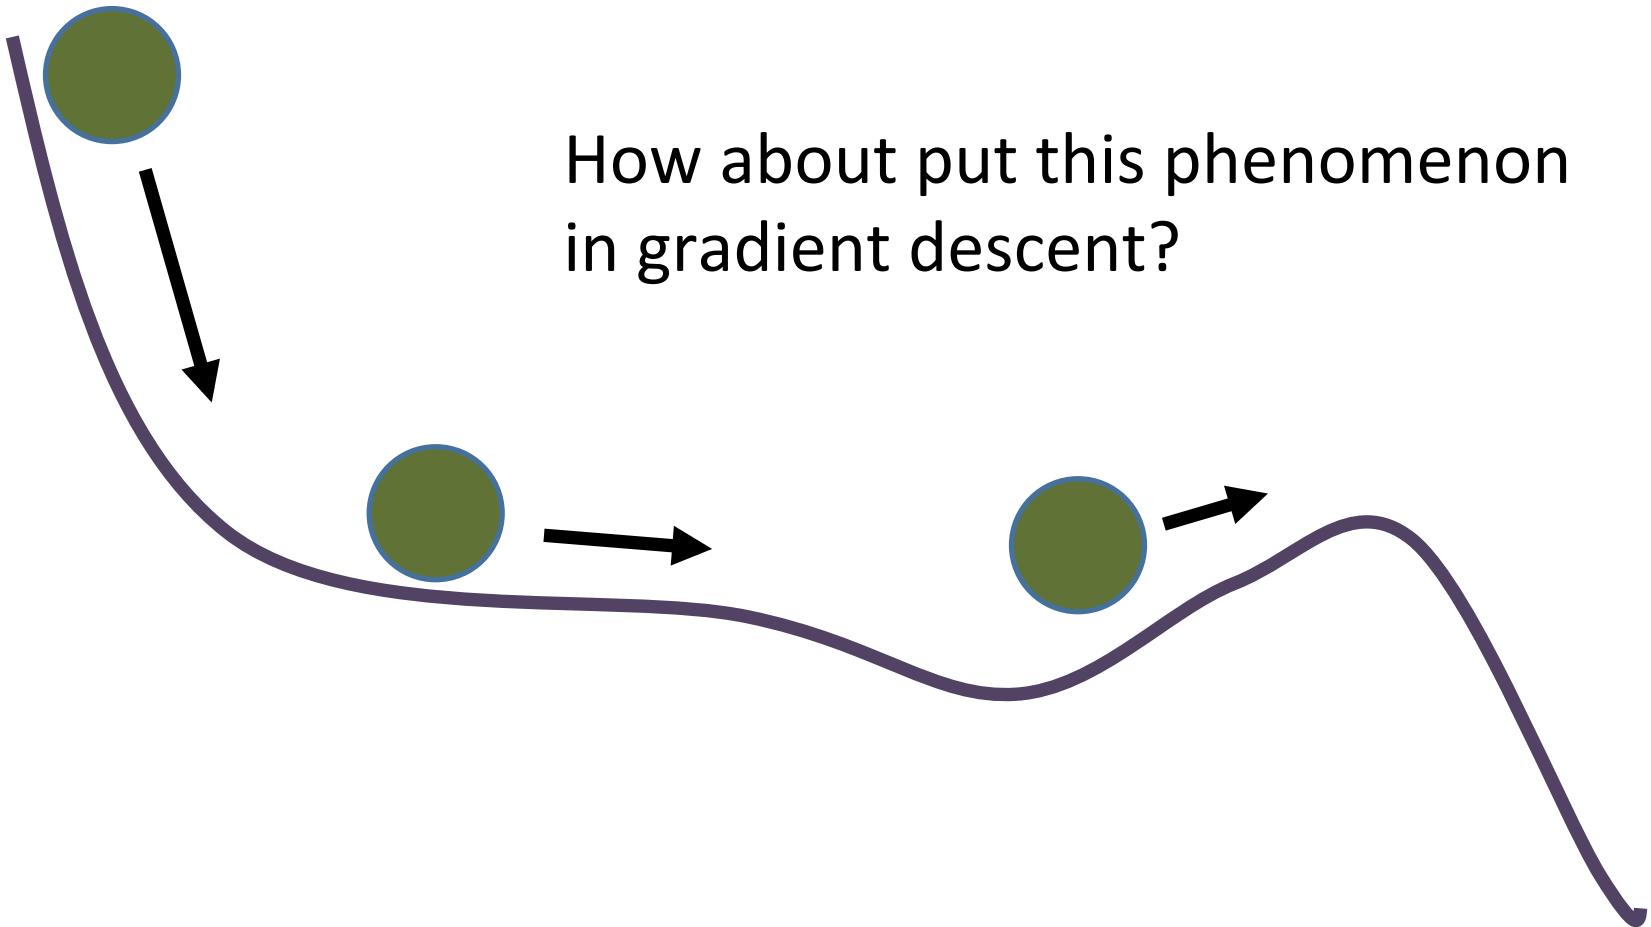
Hard to find optimal network parameters



In physical world


- Momentum

How about put this phenomenon
in gradient descent?



Momentum

Still not guarantee reaching global minima, but give some hope



RMSProp (Advanced Adagrad) + Momentum

```
model.compile(loss='categorical_crossentropy',
               optimizer=SGD(lr=0.1),
               metrics=['accuracy'])
```

```
model.compile(loss='categorical_crossentropy',
               optimizer=Adam(),
               metrics=['accuracy'])
```

Algorithm 1: *Adam*, our proposed algorithm for stochastic optimization. See section 2 for details, and for a slightly more efficient (but less clear) order of computation. g_t^2 indicates the elementwise square $g_t \odot g_t$. Good default settings for the tested machine learning problems are $\alpha = 0.001$, $\beta_1 = 0.9$, $\beta_2 = 0.999$ and $\epsilon = 10^{-8}$. All operations on vectors are element-wise. With β_1^t and β_2^t we denote β_1 and β_2 to the power t .

Require: α : Stepsize

Require: $\beta_1, \beta_2 \in [0, 1]$: Exponential decay rates for the moment estimates

Require: $f(\theta)$: Stochastic objective function with parameters θ

Require: θ_0 : Initial parameter vector

$m_0 \leftarrow 0$ (Initialize 1st moment vector)

$v_0 \leftarrow 0$ (Initialize 2nd moment vector)


$t \leftarrow 0$ (Initialize timestep)

loss = nn.CrossEntropyLoss()

optimizer = torch.optim.Adam(model.parameters(), lr=0.1)

$\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$ (Compute bias-corrected first moment estimate)
 $\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$ (Compute bias-corrected second raw moment estimate)
 $\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$ (Update parameters)
end while
return θ_t (Resulting parameters)

Recipe of Deep Learning

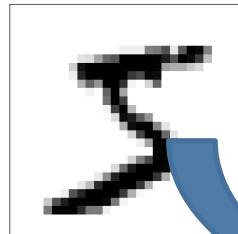


Panacea for Overfitting


- Have more training data
- *Create* more training data (?)

Handwriting recognition:

Original
Training Data:




Created
Training Data:




Shift 15 °

Recipe of Deep Learning



Dropout


Training:



- **Each time before updating the parameters**
 - Each neuron has a probability of p to dropout

Dropout

Training:

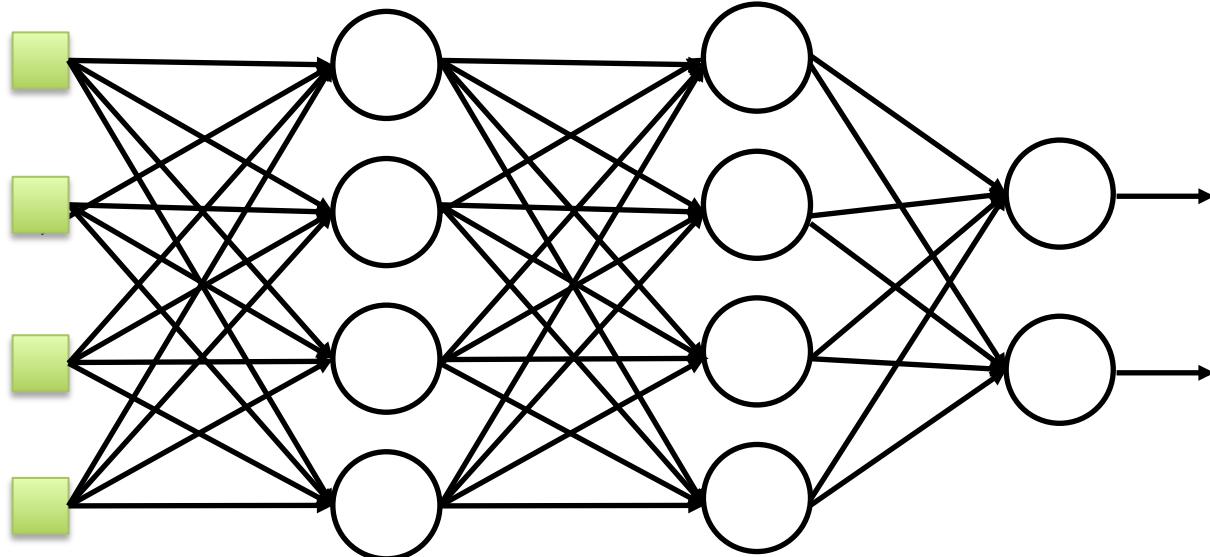


- **Each time before updating the parameters**
 - Each neuron has a probability of p to dropout
→ **The structure of the network is changed.**
 - Using the new network for training

For each mini-batch, we resample the dropout neurons

Dropout

Testing:



➤ No dropout


- If the dropout rate at training is p , all the weights times $1-p$
- Assume that the dropout rate is 50%.
If a weight $w=1$ by training, set $w=0.5$ for testing.

Dropout - Intuitive Reason

- Why the weights should multiply $(1-p)$ when testing?


Training of Dropout

Assume dropout rate is 0.5



Testing of Dropout

No dropout




Weights from training
→ $z' \approx 2z$

Weights multiply $(1-p)\%$
→ $z' \approx z$

Dropout is a kind of ensemble.


Ensemble




Train a bunch of networks with different structures

Dropout is a kind of ensemble.

Ensemble




Dropout is a kind of ensemble.



Training of
Dropout

M neurons

⋮




2^M possible
networks

- Using one mini-batch to train one network
- Some parameters in the network are shared

Dropout is a kind of ensemble.


Testing of Dropout



More about dropout

- More reference for dropout [Nitish Srivastava, JMLR'14] [Pierre Baldi, NIPS'13][Geoffrey E. Hinton, arXiv'12]
- Dropout works better with Maxout [Ian J. Goodfellow, ICML'13]
- Dropconnect [Li Wan, ICML'13]
 - Dropout delete neurons
 - Dropconnect deletes the connection between neurons
- Annealed dropout [S.J. Rennie, SLT'14]
 - Dropout rate decreases by epochs
- Standout [J. Ba, NISP'13]
 - Each neural has different dropout rate

Demo



```
model = Sequential()
```

```
model.add( Dense( input_dim=28*28,  
                  output_dim=500 ) )  
model.add( Activation('sigmoid') )
```

model.add(dropout(0.8))

```
model.add( Dense( output_dim=500 ) )  
model.add( Activation('sigmoid') )
```

model.add(dropout(0.7))


```
model.add( Dense(output_dim=10) )  
model.add( Activation('softmax') )
```

PyTorch

Step 1:
define a set
of function

Step 2:
goodness of
function


Step 3: pick
the best
function



```
import torch.nn as nn  
import torch.nn.functional as F
```

```
class MyNetwork(nn.Module):  
    def __init__(self):  
        super(MyNetwork, self).__init__()  
  
        self.fc1 = nn.Linear(28 * 28, 500)  
        self.fc2 = nn.Linear(500, 500)  
        self.fc3 = nn.Linear(500, 10)  
        self.do1 = nn.Dropout(0.8)  
        self.do2 = nn.Dropout(0.7)  
  
    def forward(self, x):  
        x = F.sigmoid(self.fc1(x))  
        x = self.do1(x)  
        x = F.sigmoid(self.fc2(x))  
        x = self.do2(x)  
        return F.log_softmax(self.fc3(x))
```

Recipe of Deep Learning



Concluding Remarks

Recipe of Deep Learning

