



BIG DATA and AI for business

Lecture 4 (02/14, 02/18): Deep Learning (2)

**Decisions, Operations & Information Technologies
Robert H. Smith School of Business
Spring, 2019**



Variants of Neural Networks

Convolutional Neural
Network (CNN)

Widely used in
image processing

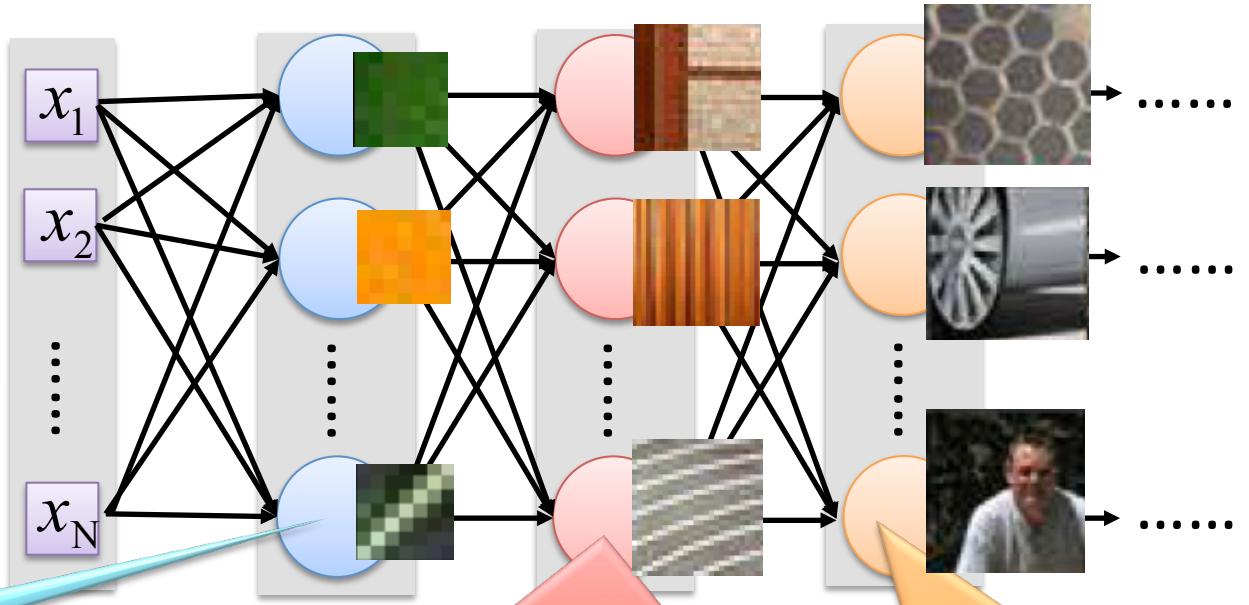
Recurrent Neural Network
(RNN)

Why CNN for Image?

[Zeiler, M. D., ECCV 2014]



Represented
as pixels



The most basic
classifiers

Use 1st layer as module
to build classifiers

Use 2nd layer as
module

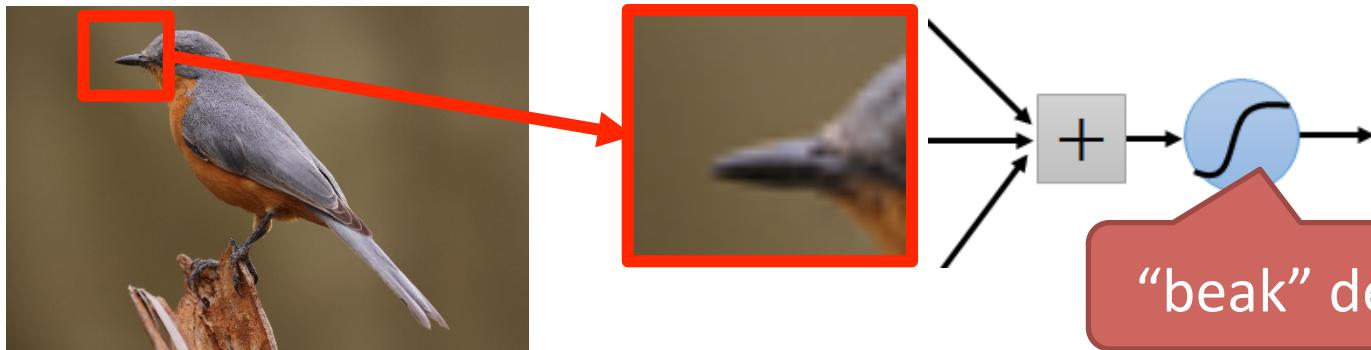
Can the network be simplified by
considering the properties of images?

Why CNN for Image

- Some patterns are much smaller than the whole image

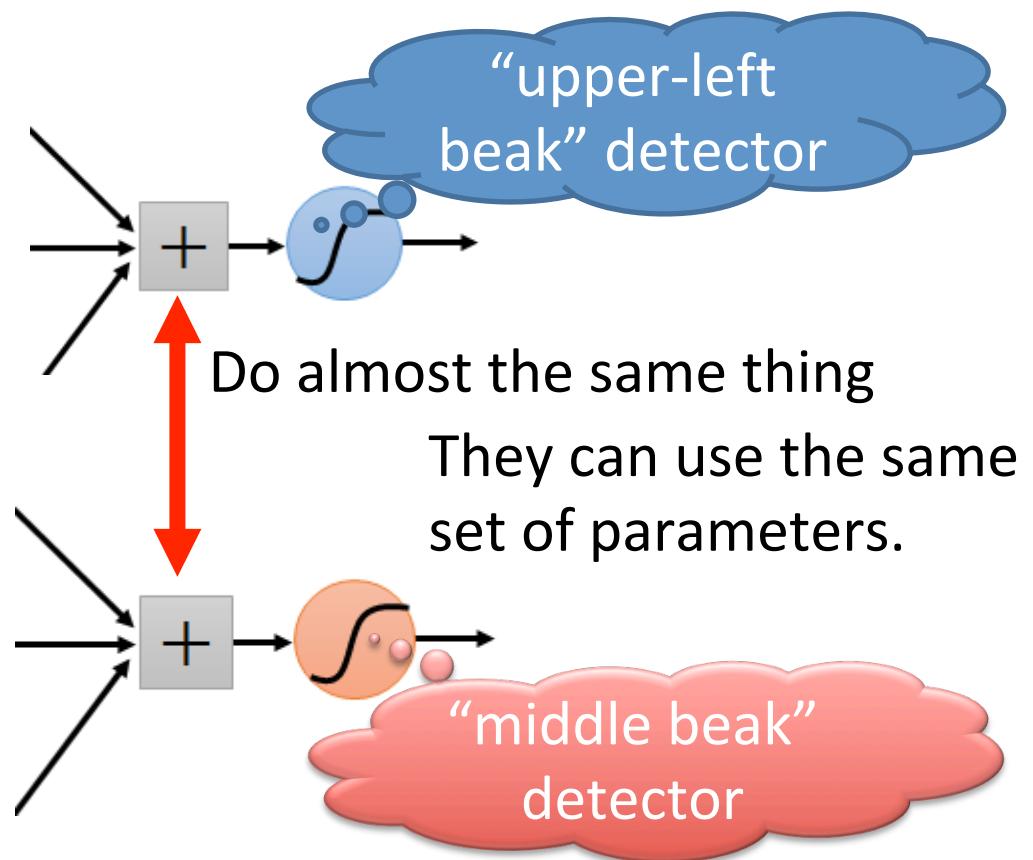
A neuron does not have to see the whole image to discover the pattern.

Connecting to small region with less parameters



Why CNN for Image

- The same patterns appear in different regions.



Why CNN for Image

- Subsampling the pixels will not change the object bird



subsampling



bird

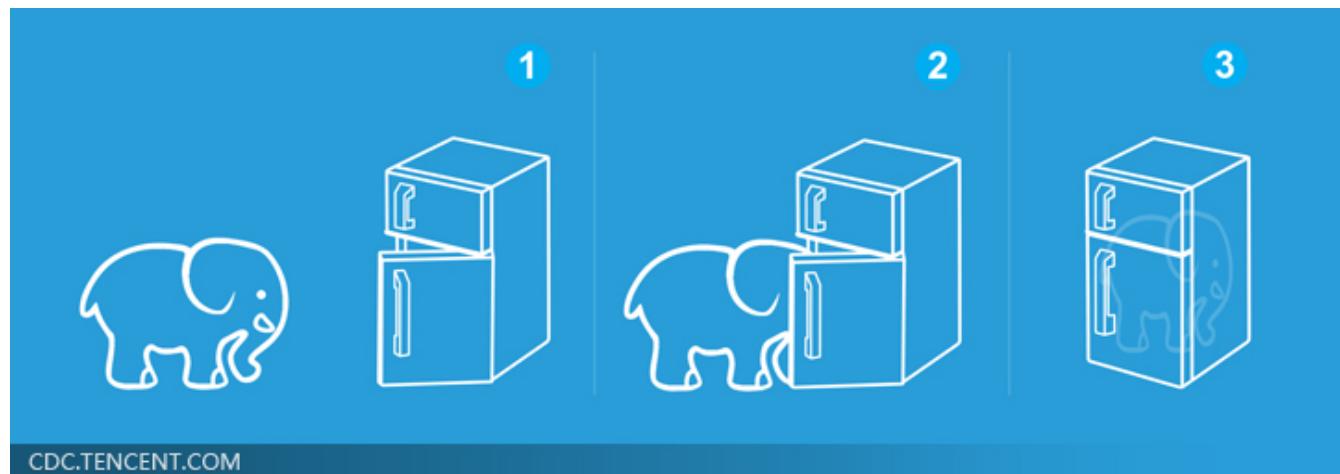
We can subsample the pixels to make image smaller

→ Less parameters for the network to process the image

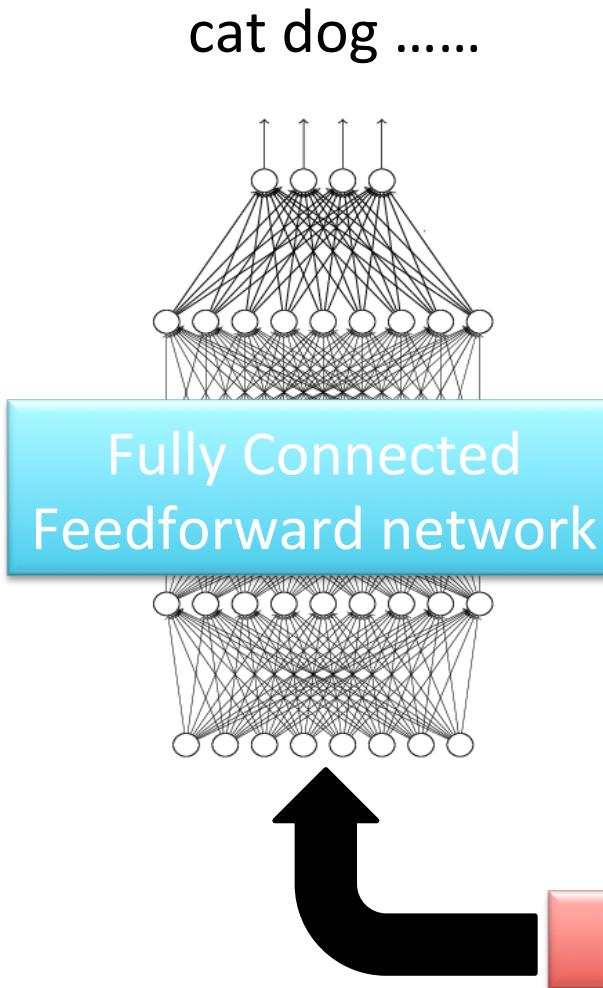
Three Steps for Deep Learning



Deep Learning is so simple



The whole CNN



Convolution

Max Pooling

Convolution

Max Pooling

Flatten

Can repeat
many times

The whole CNN

Property 1

- Some patterns are much smaller than the whole image

Property 2

- The same patterns appear in different regions.

Property 3

- Subsampling the pixels will not change the object



Convolution

Max Pooling

Convolution

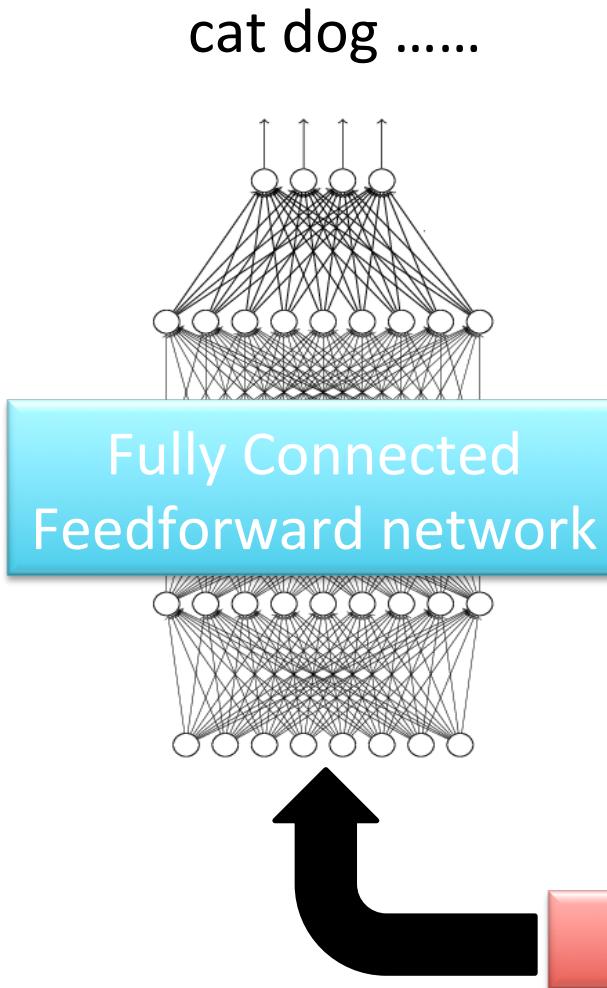
Max Pooling

Flatten



Can repeat
many times

The whole CNN



Flatten

Can repeat
many times

CNN – Convolution

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

Those are the network parameters to be learned.

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1
Matrix

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2
Matrix

⋮

Property 1

Each filter detects a small pattern (3 x 3).

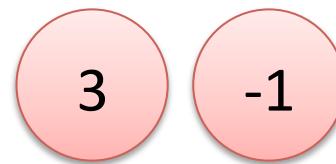
CNN – Convolution

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0



6 x 6 image

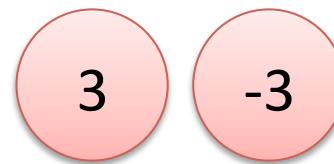
CNN – Convolution

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

If stride=2

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

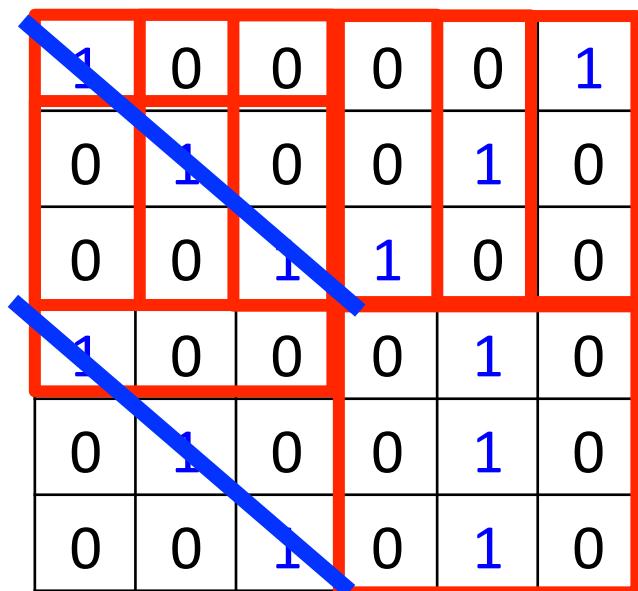


We set stride=1 below

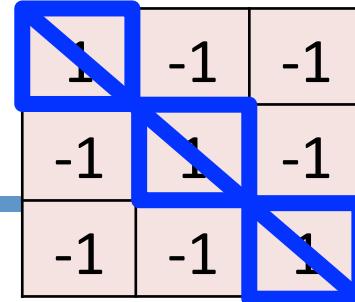
6 x 6 image

CNN – Convolution

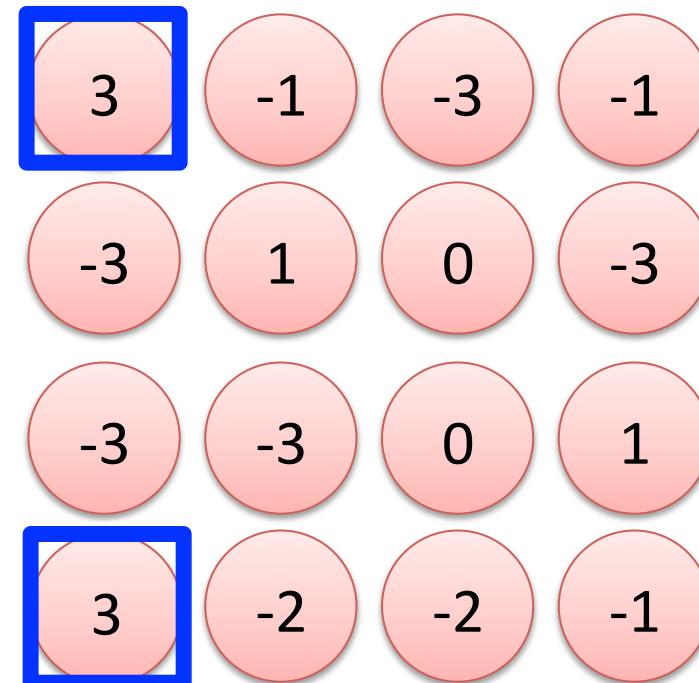
stride=1



6 x 6 image



Filter 1



Property 2

CNN – Convolution

stride=1

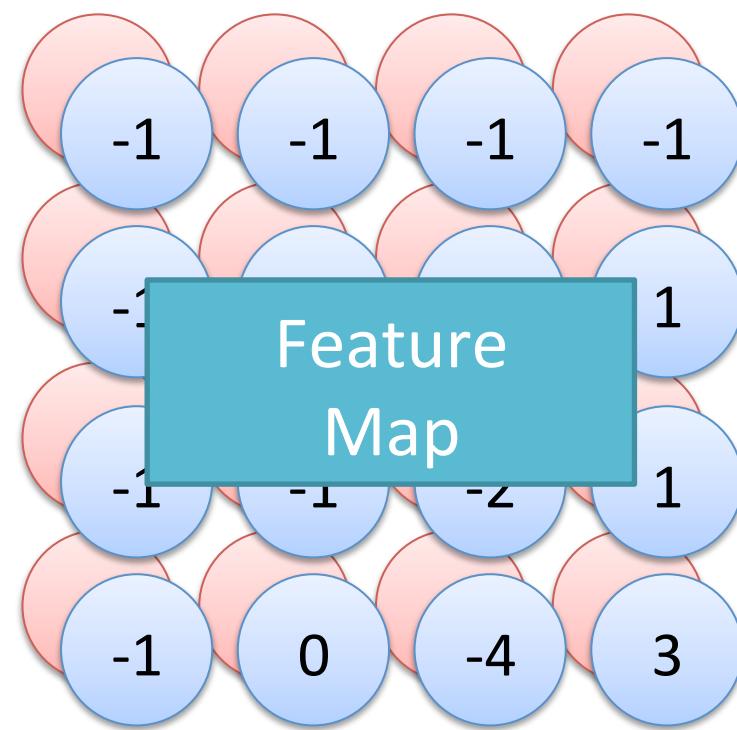
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

Do the same process for every filter

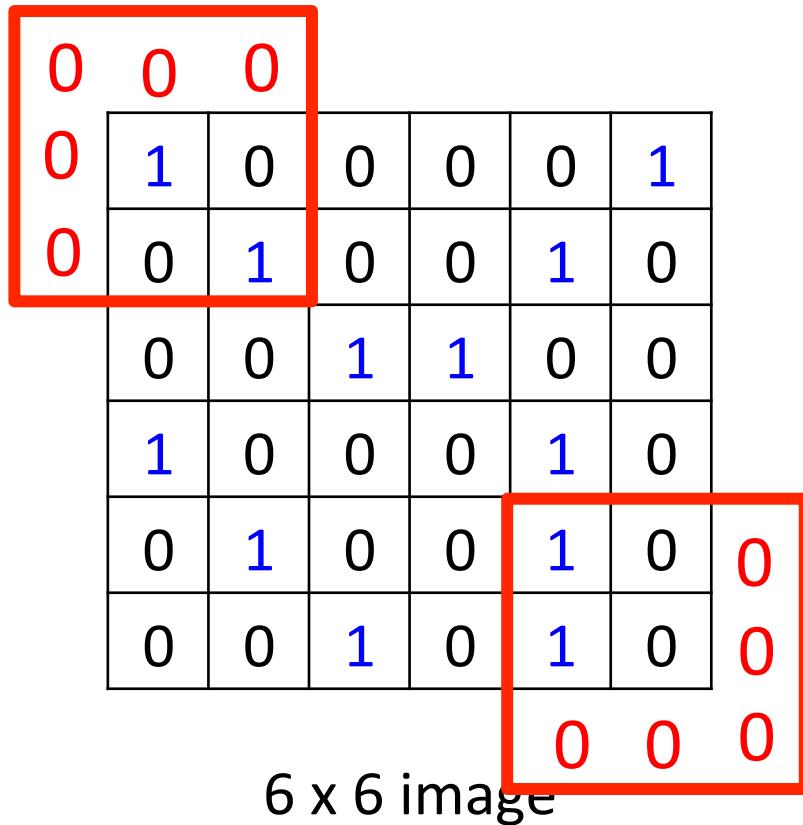


4 x 4 image

CNN – Zero Padding

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1



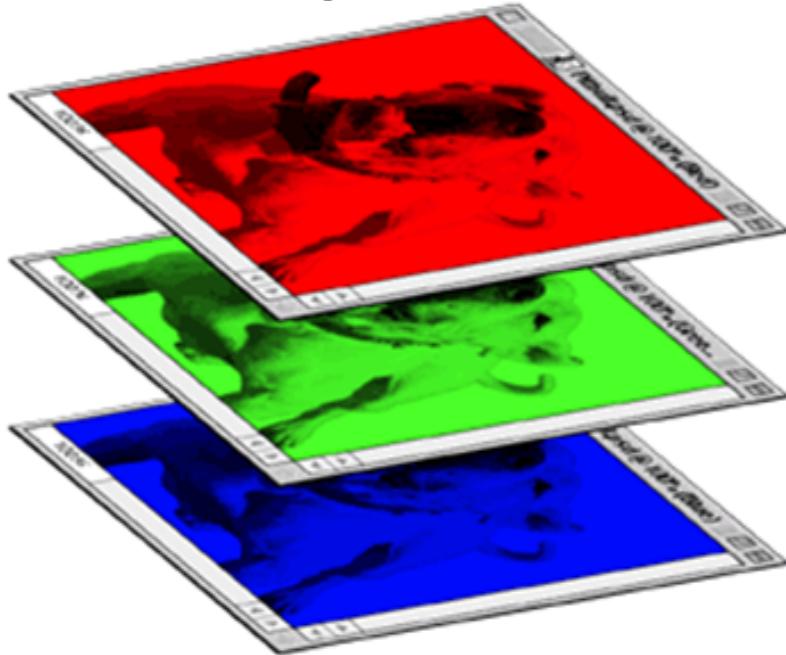
You will get another 6×6 images in this way



Zero padding

CNN – Colorful image

Colorful image



$$\begin{matrix} & & & \\ & & & \\ \begin{matrix} 1 & -1 & -1 \\ -1 & 1 & -1 \\ -1 & -1 & 1 \end{matrix} & & & \\ & & & \end{matrix}$$

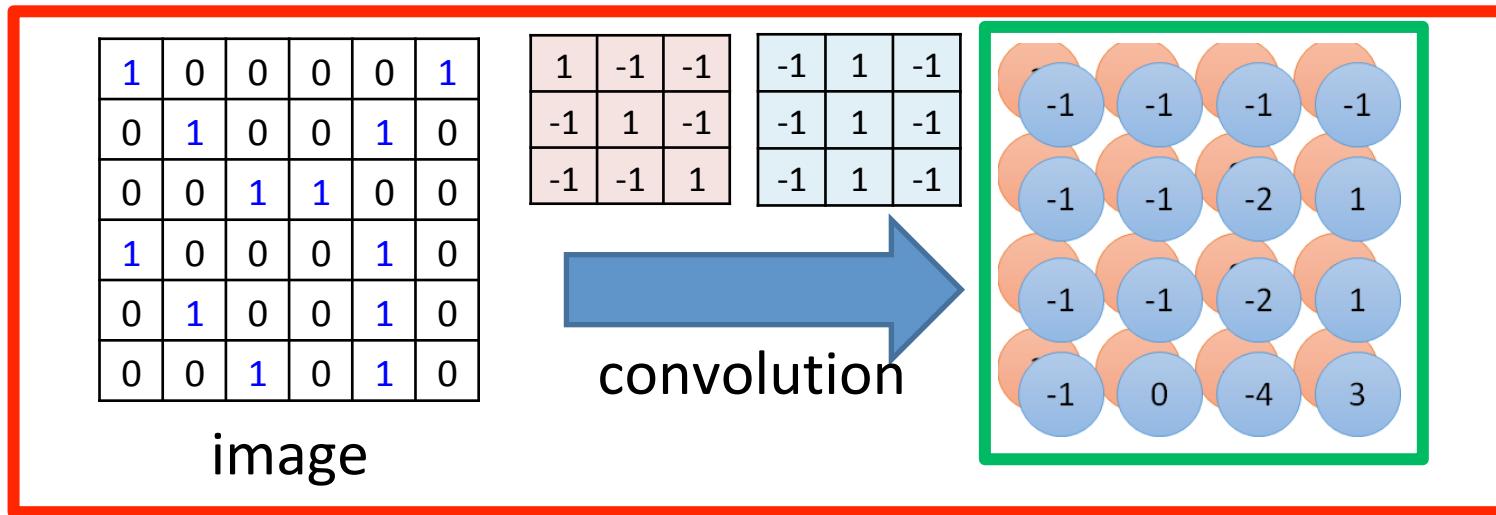
Filter 1

$$\begin{matrix} & & & \\ & & & \\ \begin{matrix} -1 & 1 & -1 \\ -1 & 1 & -1 \\ -1 & 1 & -1 \end{matrix} & & & \\ & & & \end{matrix}$$

Filter 2

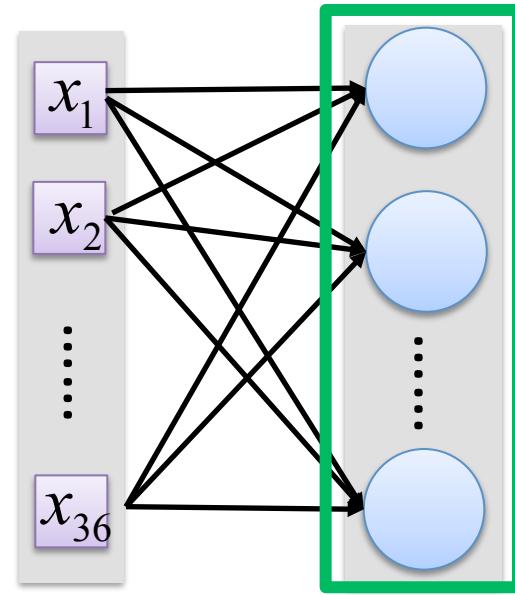
$$\begin{matrix} & & & & & \\ & & & & & \\ \begin{matrix} 1 & 0 & 0 & 0 & 0 & 1 \\ 0 & 1 & 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 & 1 & 0 \\ 0 & 1 & 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 & 1 & 0 \end{matrix} & & & & & \\ & & & & & \end{matrix}$$

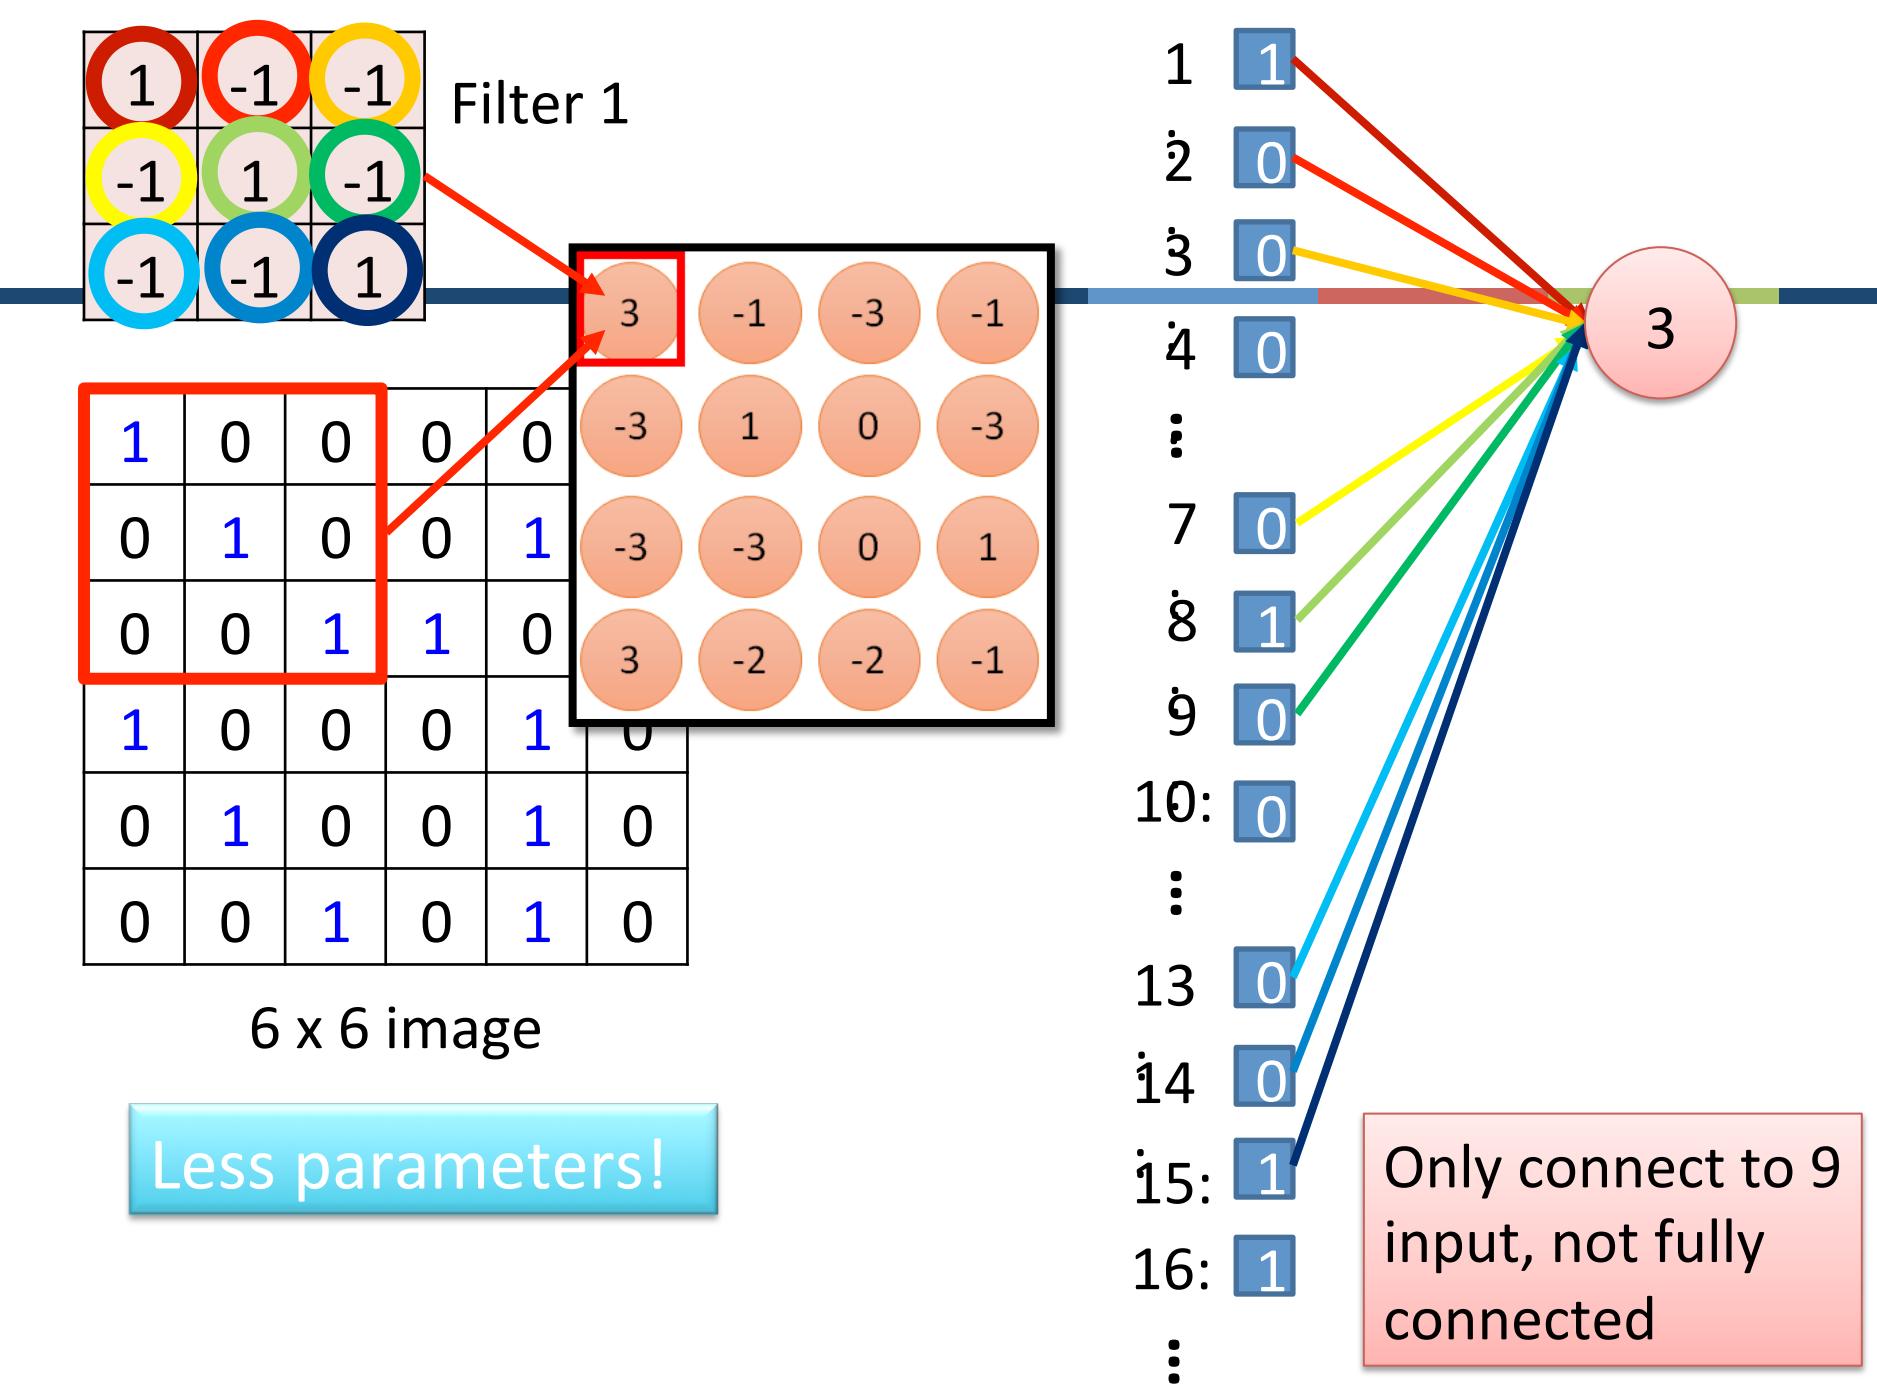
Convolution v.s. Fully Connected

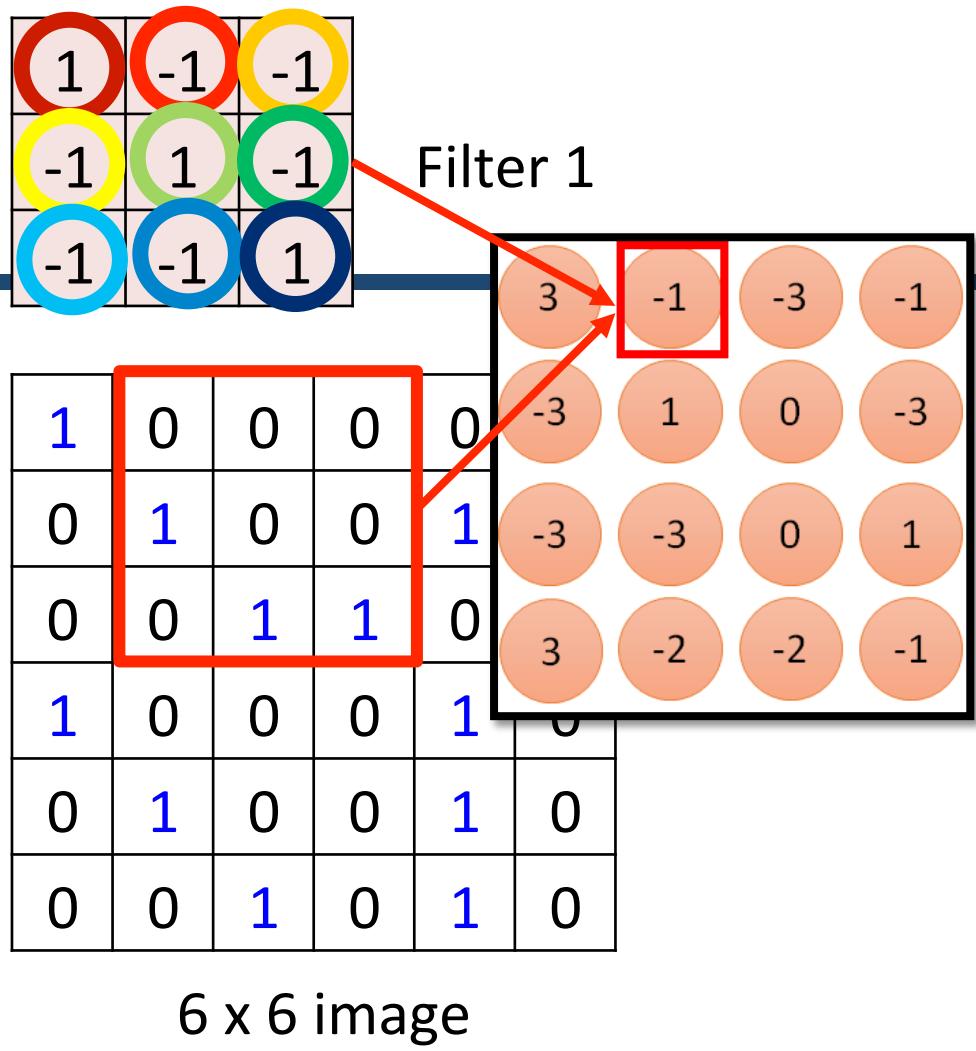


Fully-
connected

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

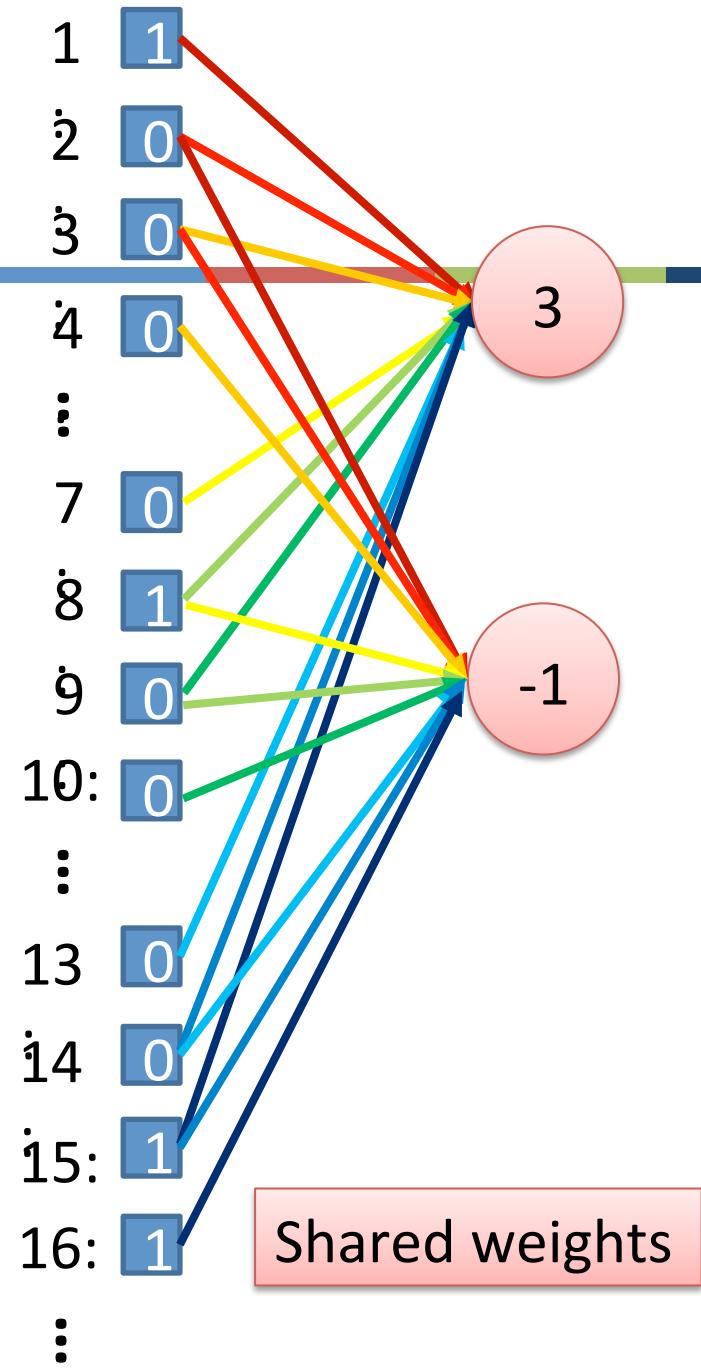




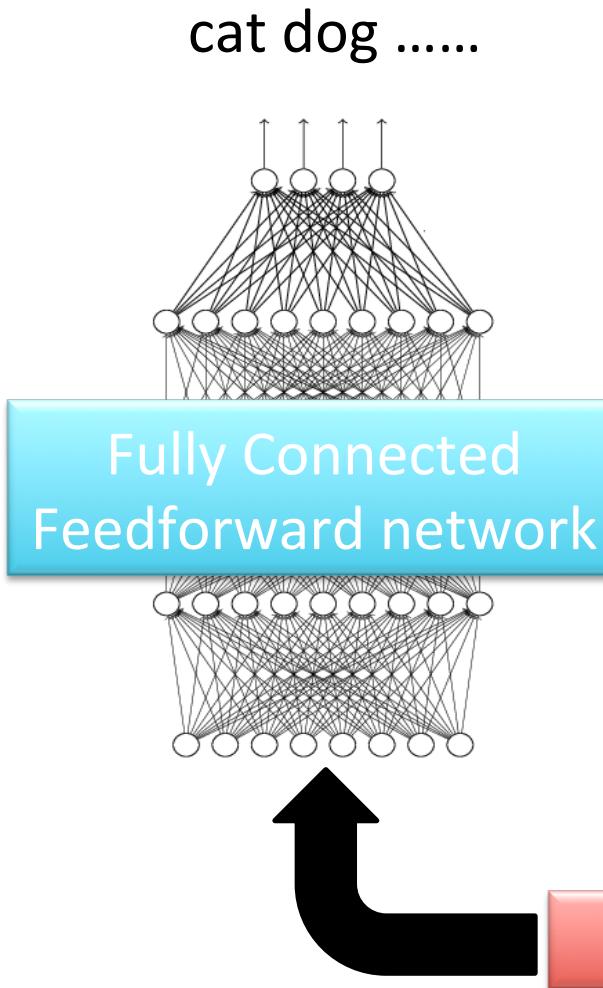


Less parameters!

Even less parameters!



The whole CNN



Convolution

Max Pooling

Convolution

Max Pooling

Flatten

Can repeat
many times

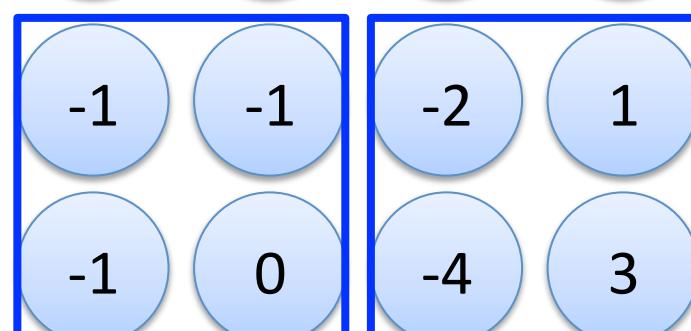
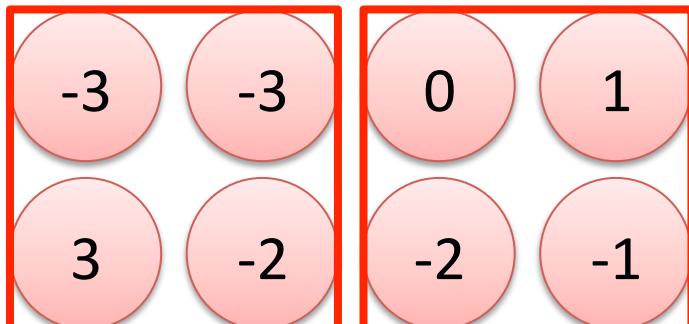
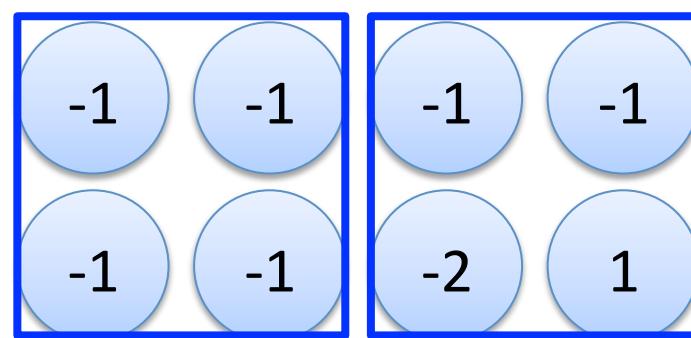
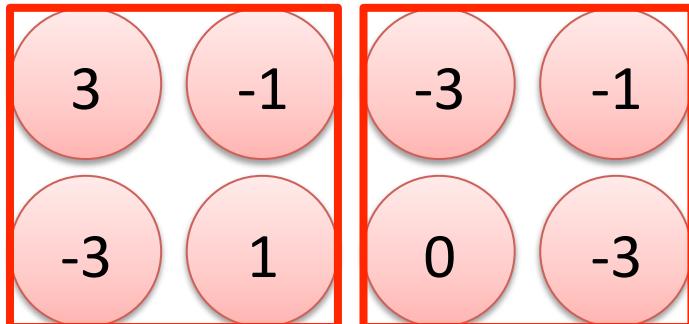
CNN – Max Pooling

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

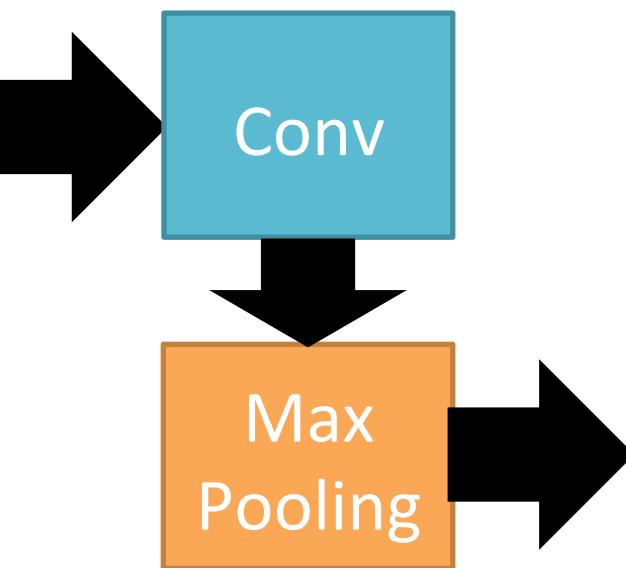
-1	1	-1
-1	1	-1
-1	1	-1

Filter 2



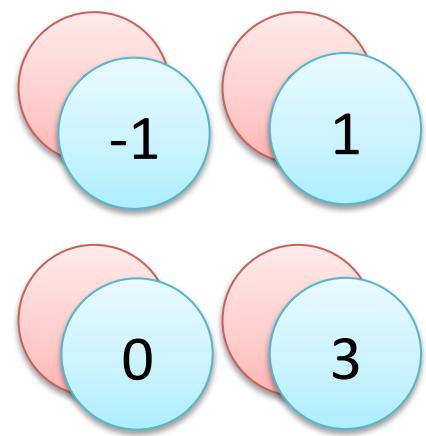
CNN – Max Pooling

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0



6 x 6 image

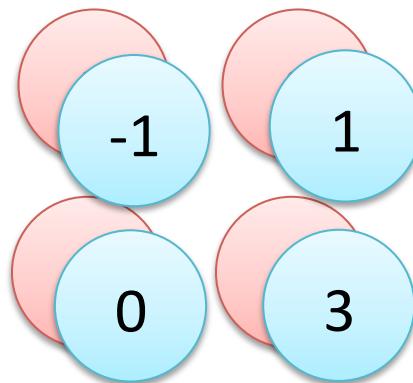
New image
but smaller



2 x 2 image

Each filter
is a channel

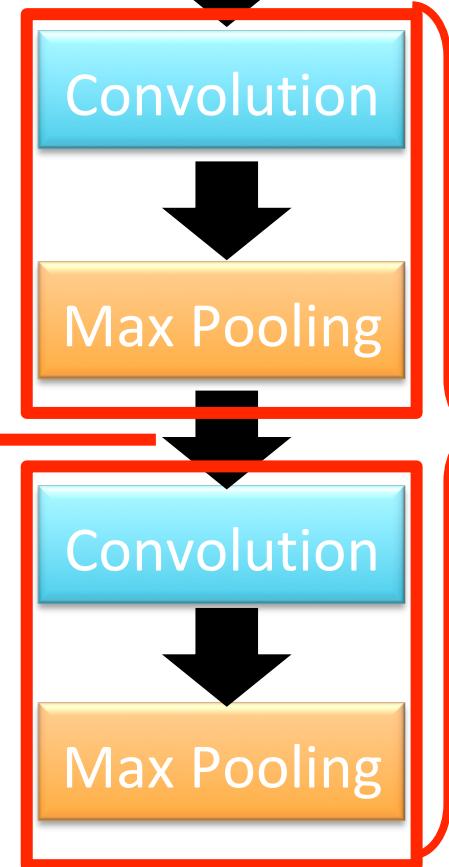
The whole CNN



A new image

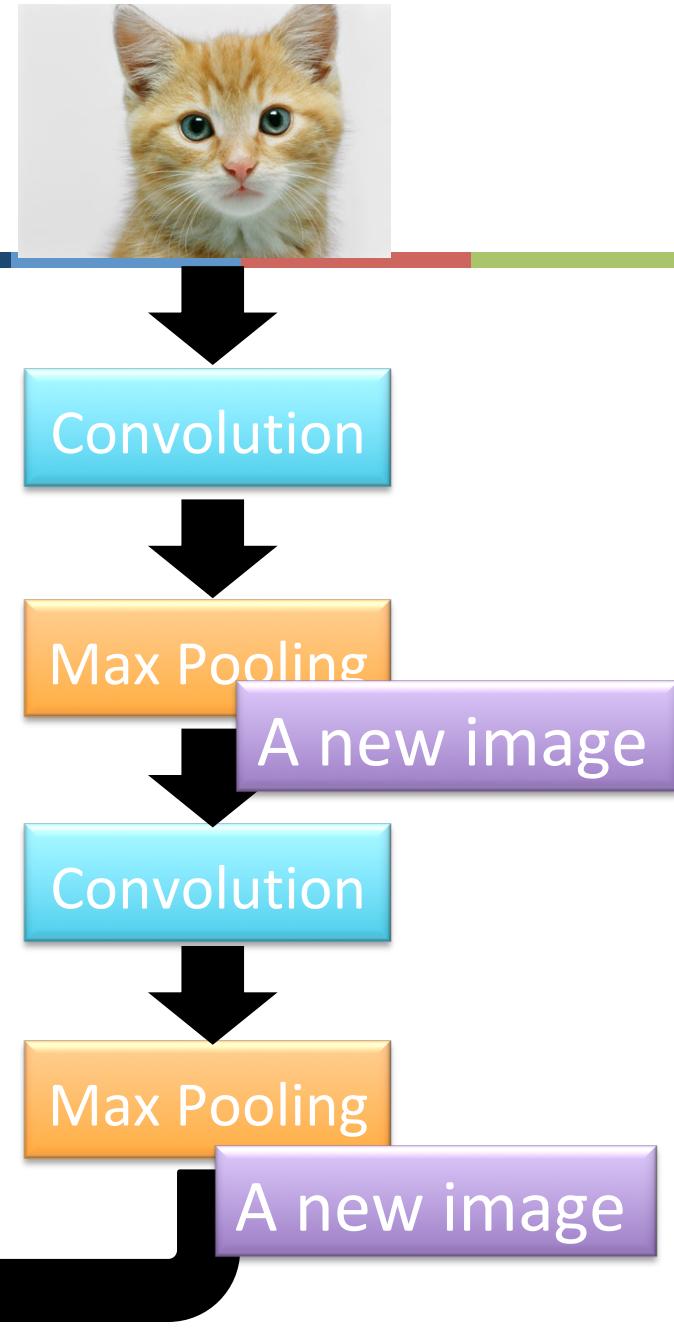
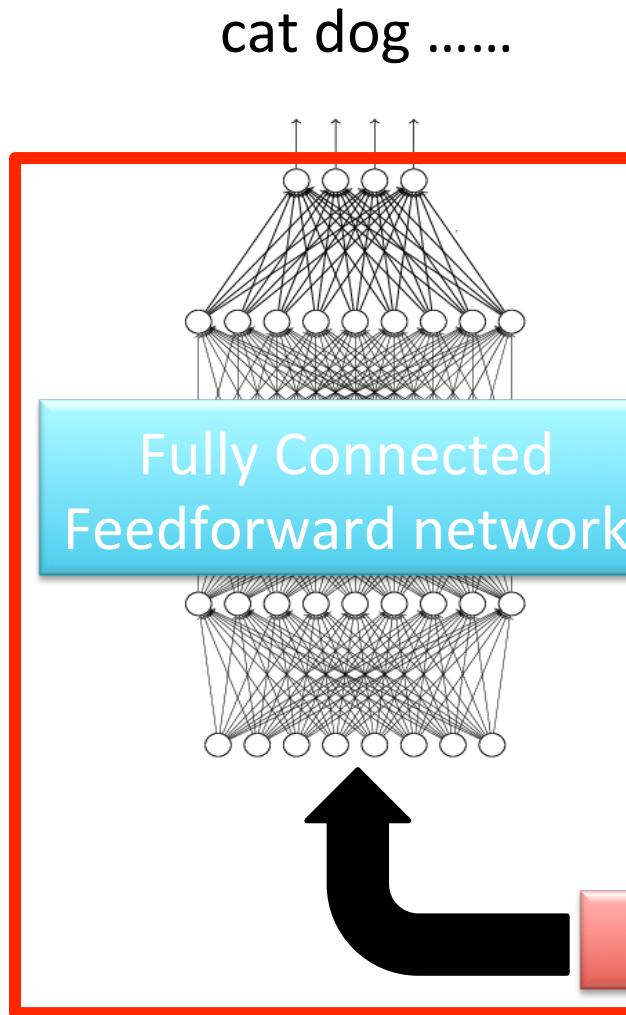
Smaller than the original image

The number of the channel is the number of filters

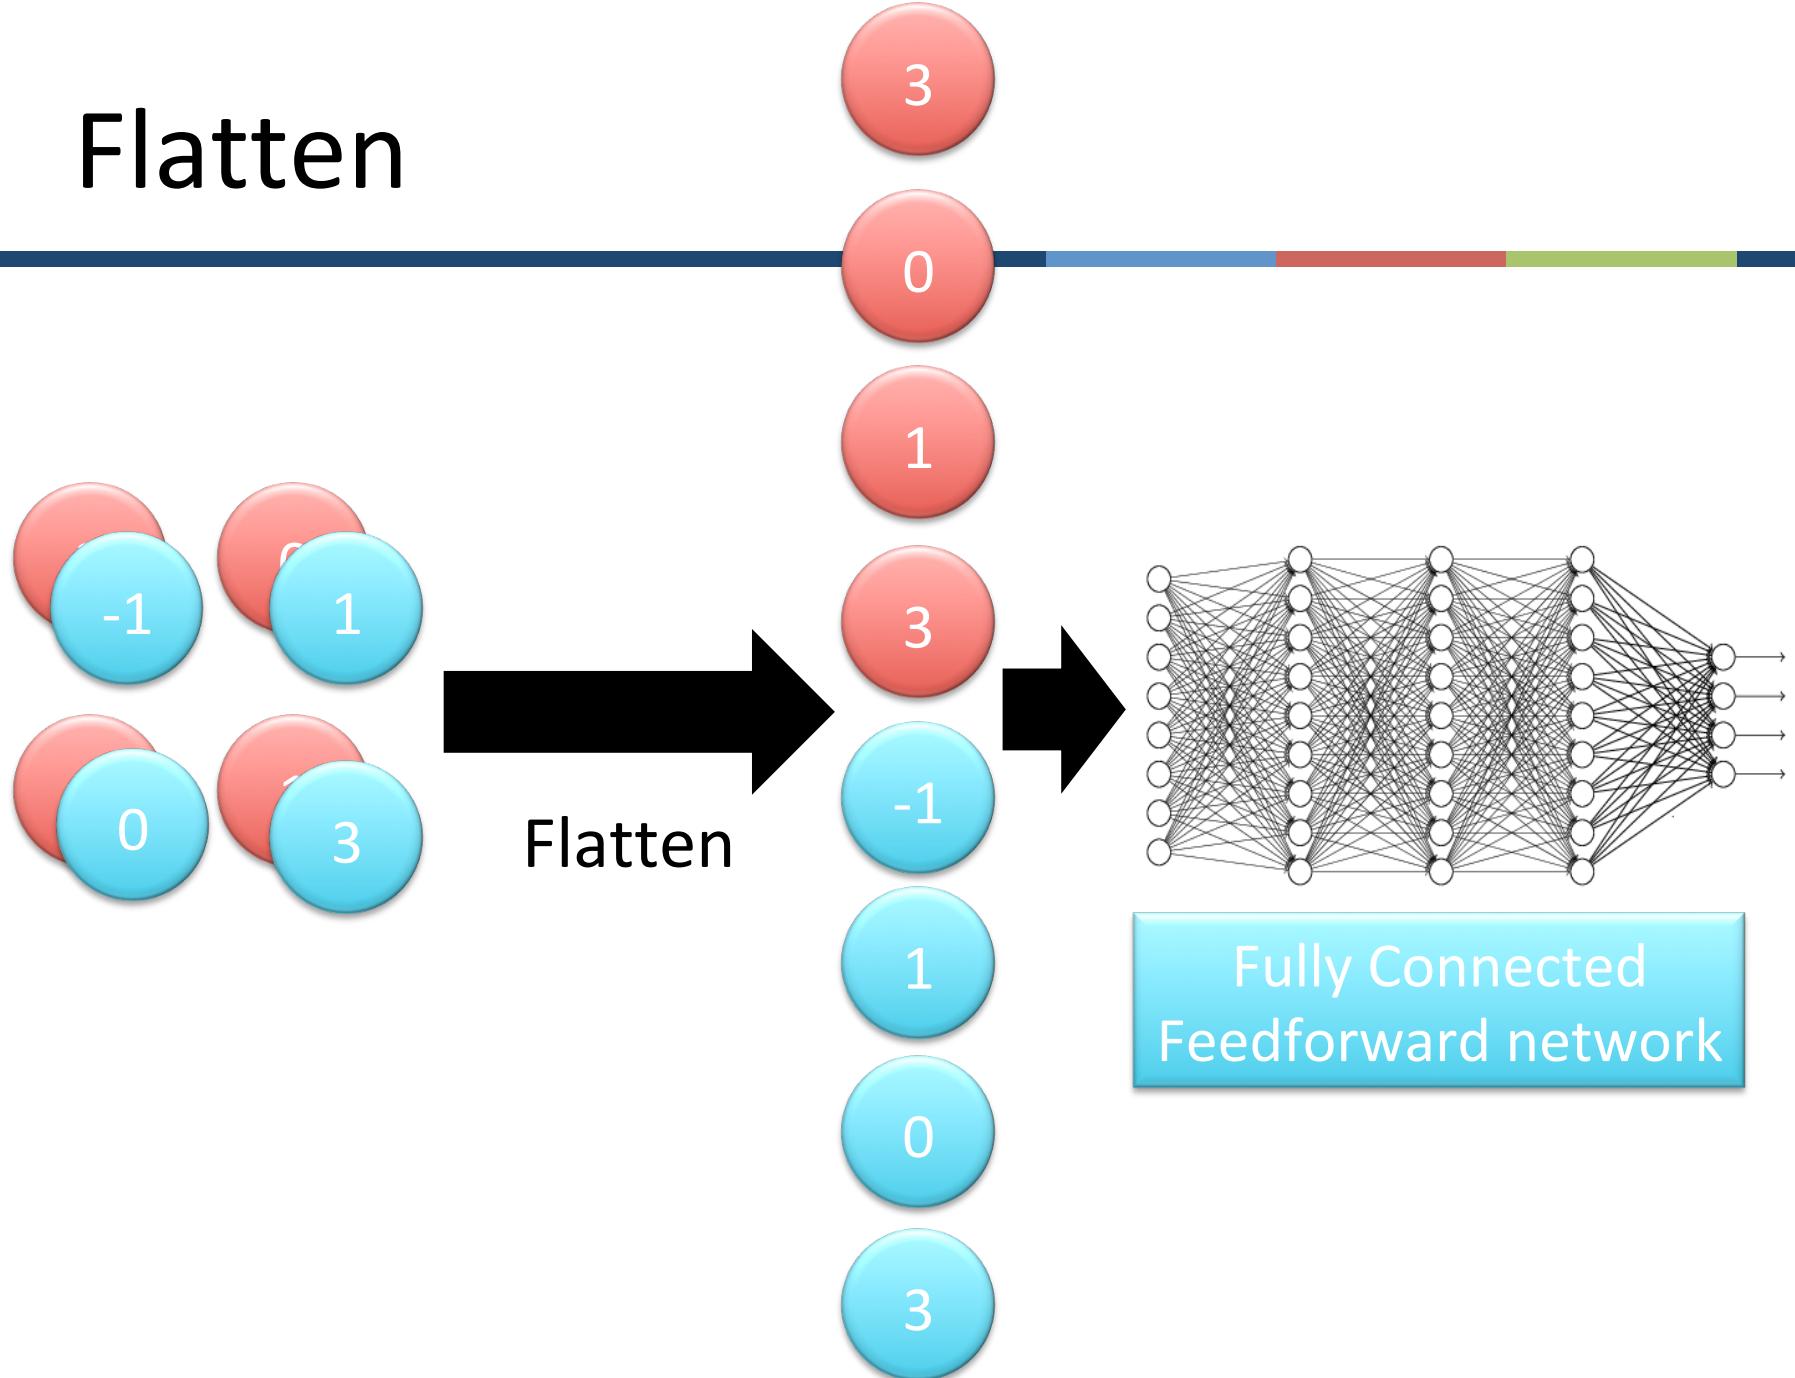


Can repeat
many times

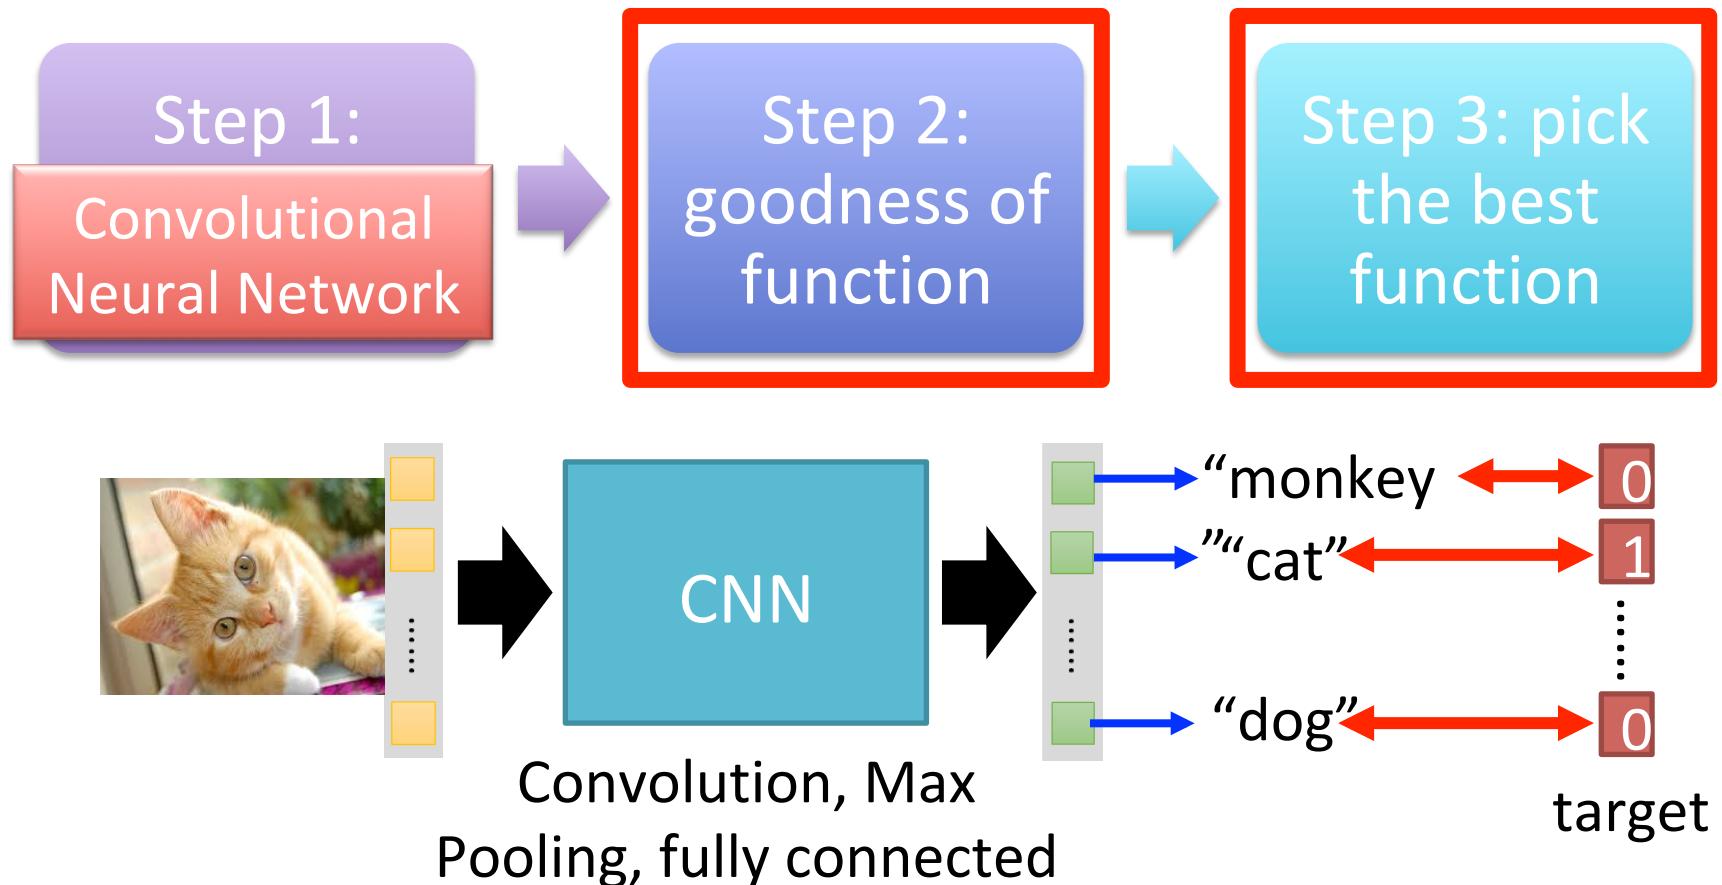
The whole CNN



Flatten



Convolutional Neural Network

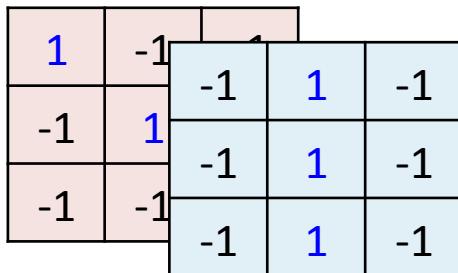


Learning: Nothing special, just gradient descent

CNN in Keras

Only modified the *network structure* and *input format (vector -> 3-D tensor)*

```
model2.add( Convolution2D( 25, 3, 3,  
                           input_shape=(1, 28, 28) ) )
```

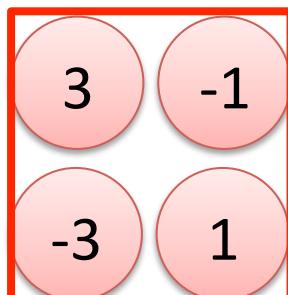


.....
There are 25
3x3 filters.

Input_shape = (1, 28, 28)

1: black/weight, 3: RGB 28 x 28 pixels

```
model2 .add (MaxPooling2D ( (2,2) ))
```



input
↓

Convolution
↓

Max Pooling
↓

Convolution
↓

Max Pooling

CNN in Keras

Only modified the *network structure* and *input format (vector -> 3-D tensor)*

How many parameters
for each filter?

```
model2.add( Convolution2D( 25, 3, 3,  
    input_shape=(1,28,28) ) )
```

9

25 x 26 x 26

```
model2.add(MaxPooling2D( (2,2) ))
```

25 x 13 x 13

```
model2.add(Convolution2D(50, 3, 3))
```

225

50 x 11 x 11

```
model2.add(MaxPooling2D( (2,2) ))
```

50 x 5 x 5

input
↓

Convolution
↓

Max Pooling
↓

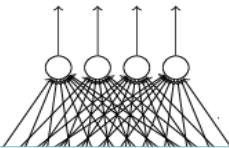
Convolution
↓

Max Pooling

CNN in Keras

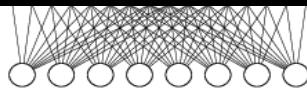
Only modified the ***network structure*** and ***input format (vector -> 3-D tensor)***

output



Fully Connected
Feedforward network

```
model2.add(Dense(output_dim=100))  
model2.add(Activation('relu'))  
model2.add(Dense(output_dim=10))  
model2.add(Activation('softmax'))
```



1250

Flatten

```
model2.add(Flatten())
```

input

$1 \times 28 \times 28$

Convolution

$25 \times 26 \times 26$

Max Pooling

$25 \times 13 \times 13$

Convolution

$50 \times 11 \times 11$

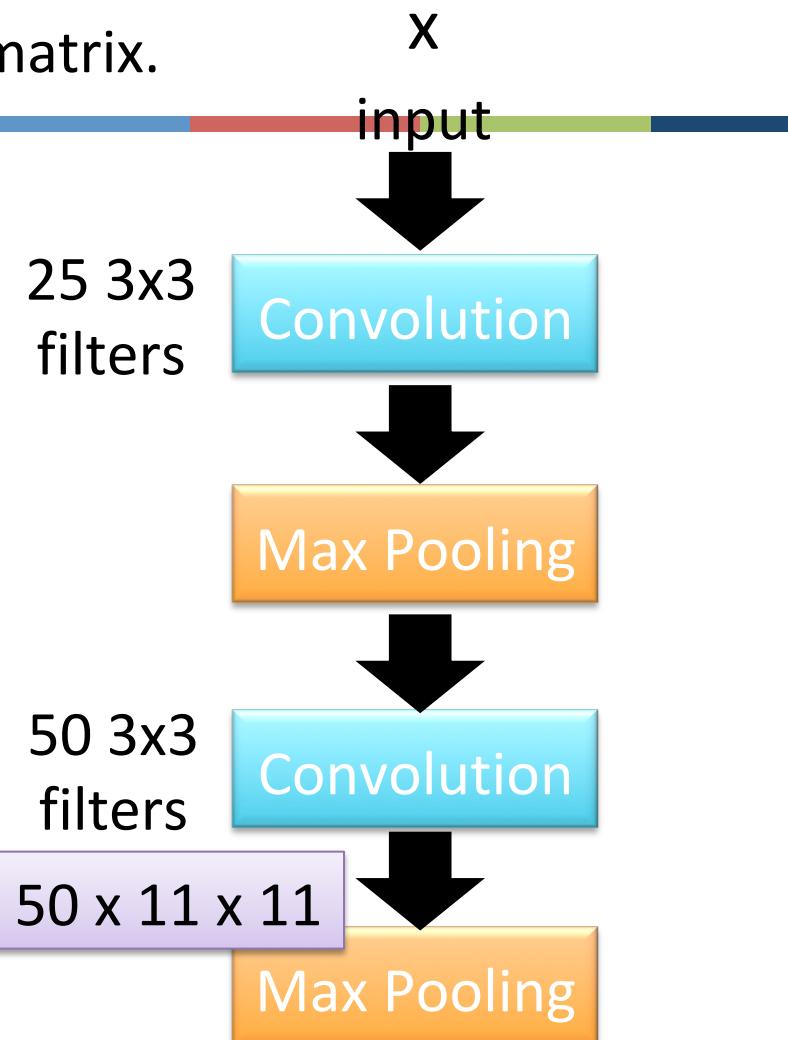
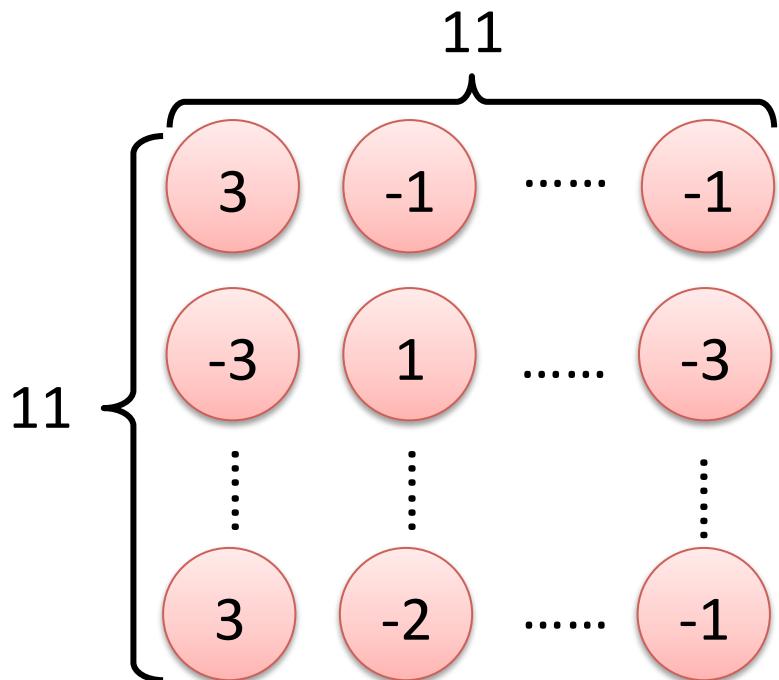
Max Pooling

$50 \times 5 \times 5$



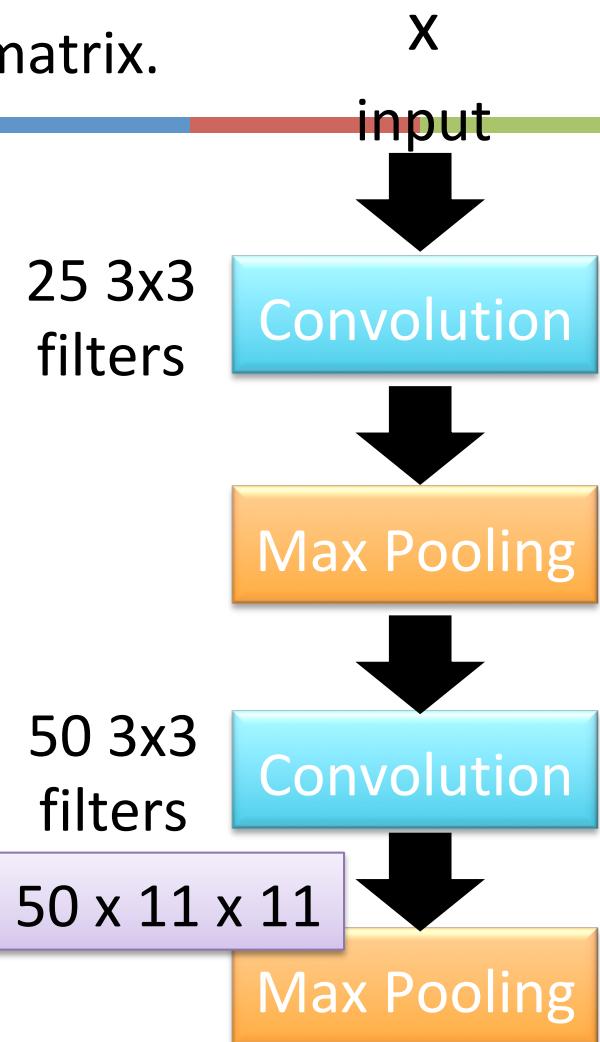
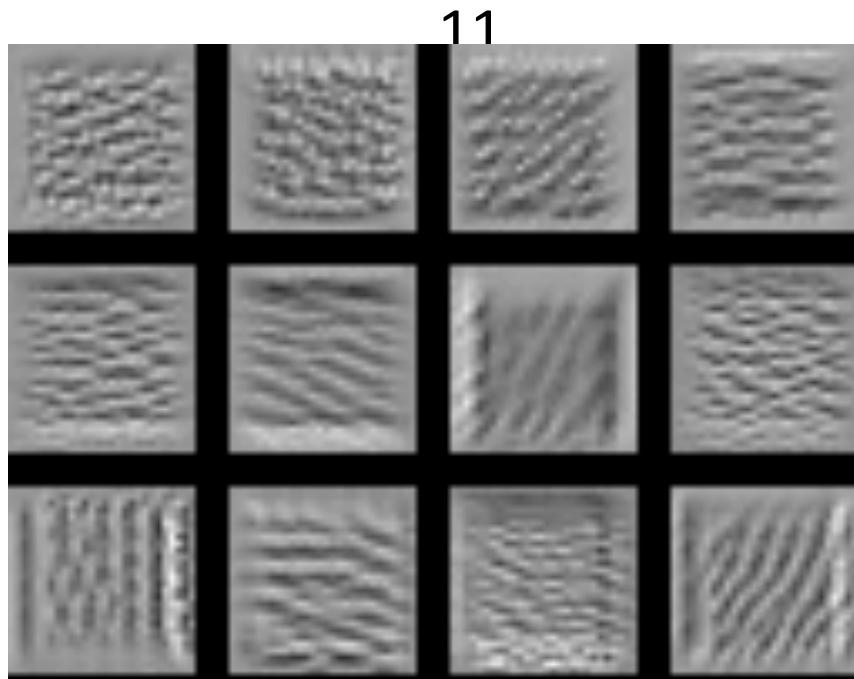
What does CNN learn?

The output of the k-th filter is a 11×11 matrix.

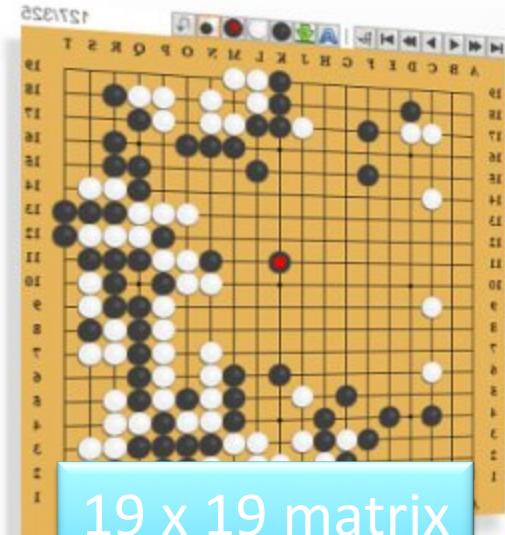


What does CNN learn?

The output of the k-th filter is a 11×11 matrix.

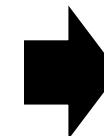
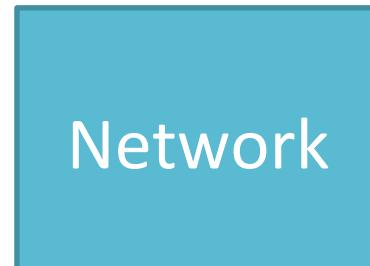


More Application: Playing Go



19 x 19 matrix
(image)

Black: 1
white: -1
none: 0



Next move
(19 x 19
positions)

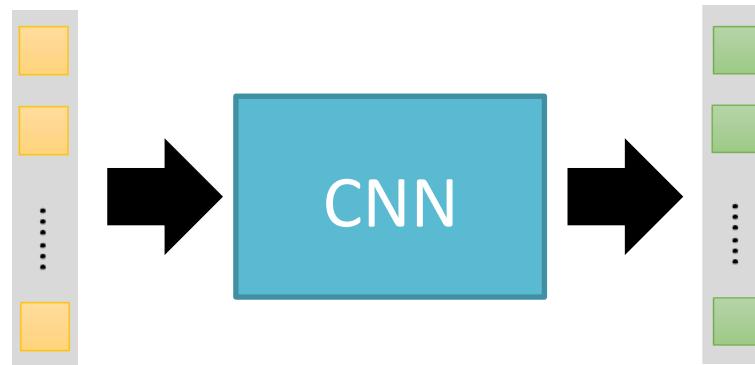
19 x 19 vector

Fully-connected feedforward
network can be used

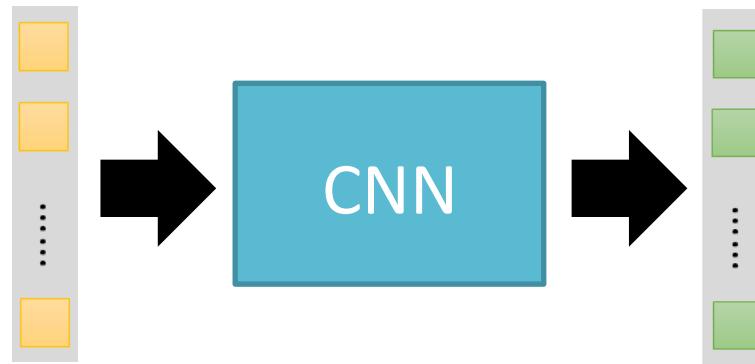
But CNN performs much better.

More Application: Playing Go

Training: record of previous plays 黑: 5之五 → 白: 天元 → 黑: 五之5 ...



Target:
“天元” = 1
else = 0

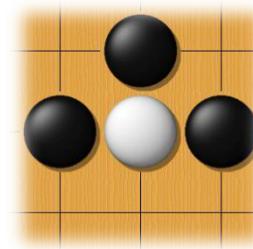


Target:
“五之5” = 1
else = 0

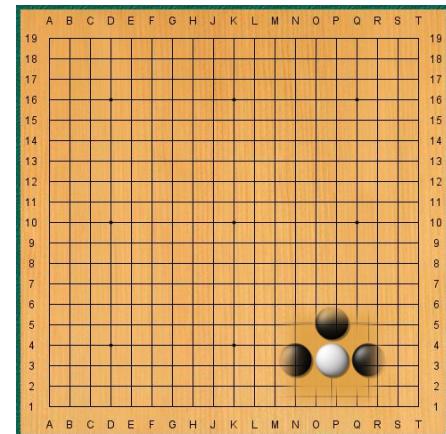
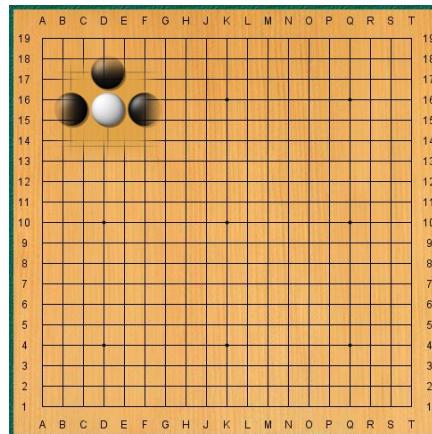
Why CNN for playing Go?

- Some patterns are much smaller than the whole image

Alpha Go uses 5×5 for first layer



- The same patterns appear in different regions.



Why CNN for playing Go?

- Subsampling the pixels will not change the object

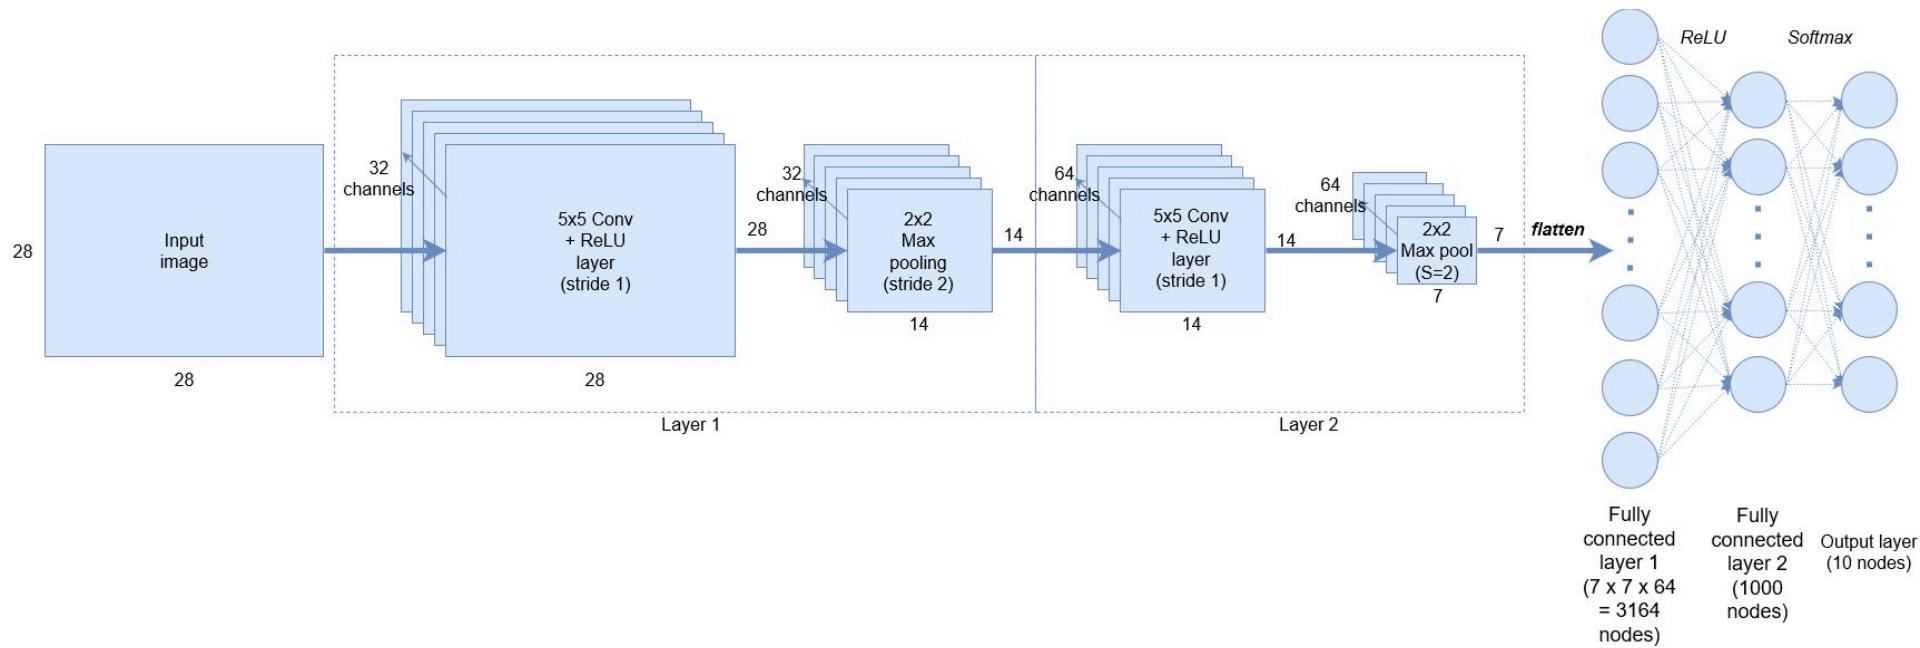


Max Pooling

How to explain this???

Neural network architecture. The input to the policy network is a $19 \times 19 \times 48$ image stack consisting of 48 feature planes. The first hidden layer zero pads the input into a 23×23 image, then convolves k filters of kernel size 5×5 with stride 1 with the input image and applies a rectifier nonlinearity. Each of the subsequent hidden layers 2 to 12 zero pads the respective previous hidden layer into a 21×21 image, then convolves k filters of kernel size 3×3 with stride 1, again followed by a rectifier nonlinearity. The final layer convolves 1 filter of kernel size 1×1 with stride 1 with a different bias for each position and applying a softmax function. The Alpha Go does not use Max Pooling Extended Data Table 3 additionally show the results of training with $k = 128, 256$ and 384 filters.

Example



```
model = Sequential()

model.add(Conv2D(32, kernel_size=(5, 5), strides=(1, 1), activation='relu',
    input_shape = (28,28,1) ))

model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))

model.add(Conv2D(64, (5, 5), activation='relu'))

model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())

model.add(Dense(1000, activation='relu'))

model.add(Dense(10, activation='softmax'))
```

```
model.compile(loss=keras.losses.categorical_crossentropy,  
              optimizer=keras.optimizers.SGD(lr=0.01),  
              metrics=['accuracy'])
```

```
model.fit(x_train, y_train,  
          batch_size=64,  
          epochs=10,  
          verbose=1,  
          validation_data=(x_test, y_test)  
        )
```

```
score = model.evaluate(x_test, y_test, verbose=0)  
print('Test loss:', score[0])  
print('Test accuracy:', score[1])
```

Variants of Neural Networks

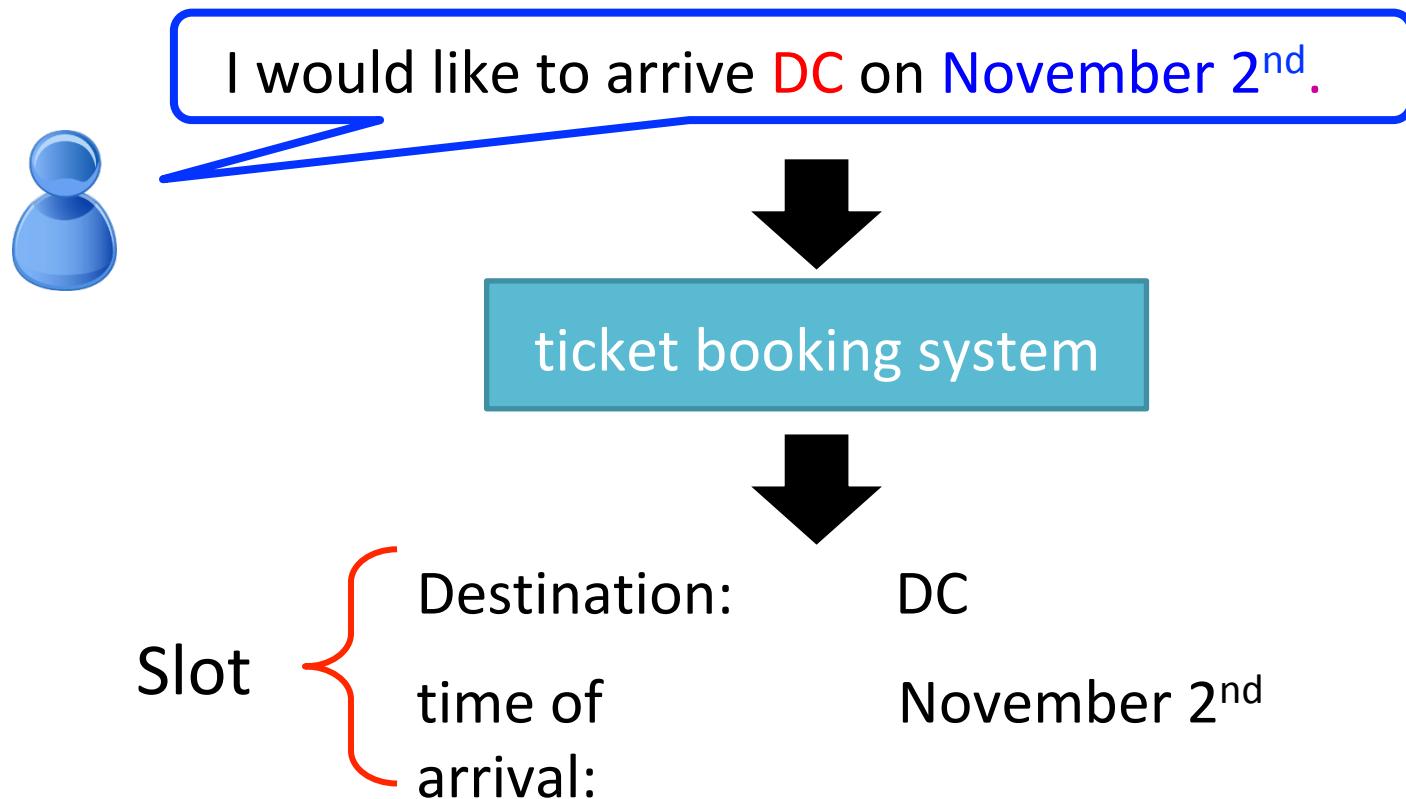
Convolutional Neural
Network (CNN)

Recurrent Neural Network
(RNN)

Neural Network with Memory

Example Application

- Slot Filling

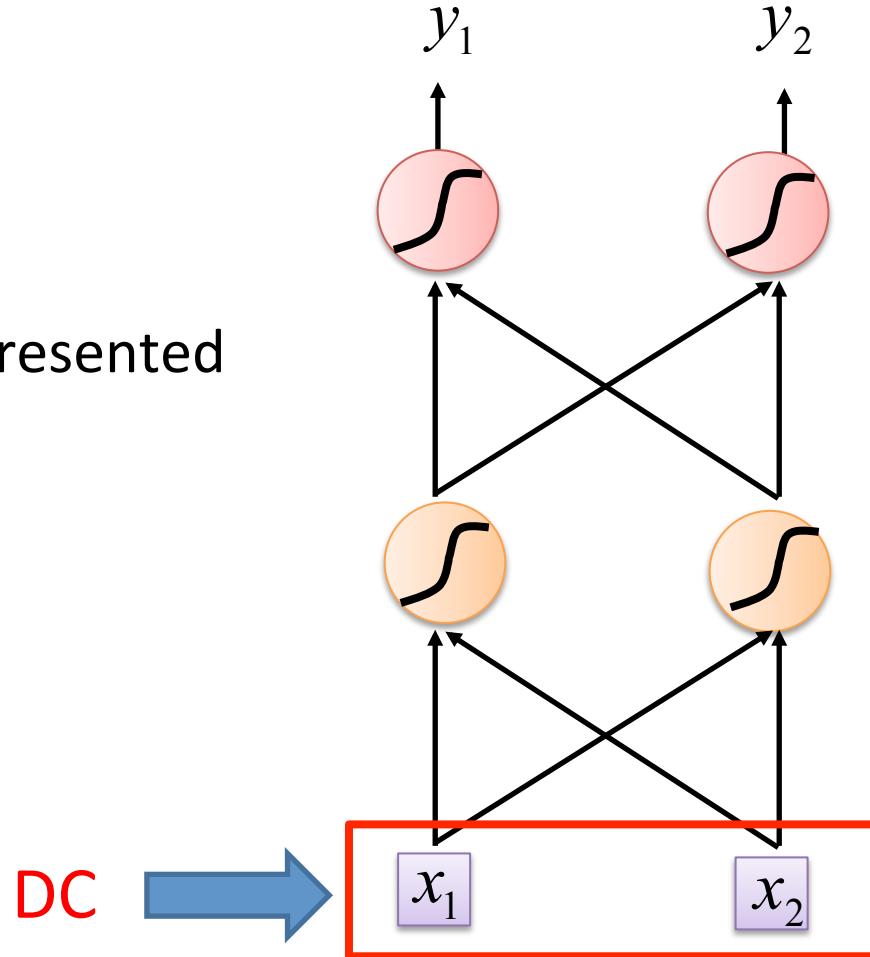


Example Application

Solving slot filling by
Feedforward network?

Input: a word

(Each word is represented
as a vector)



1-of-N encoding

How to represent each word as a vector?

1-of-N Encoding lexicon = {apple, bag, cat, dog, elephant}

The vector is lexicon size.

$$\text{apple} = [1 \ 0 \ 0 \ 0 \ 0]$$

Each dimension corresponds
to a word in the lexicon

$$\text{bag} = [0 \ 1 \ 0 \ 0 \ 0]$$

The dimension for the word
is 1, and others are 0

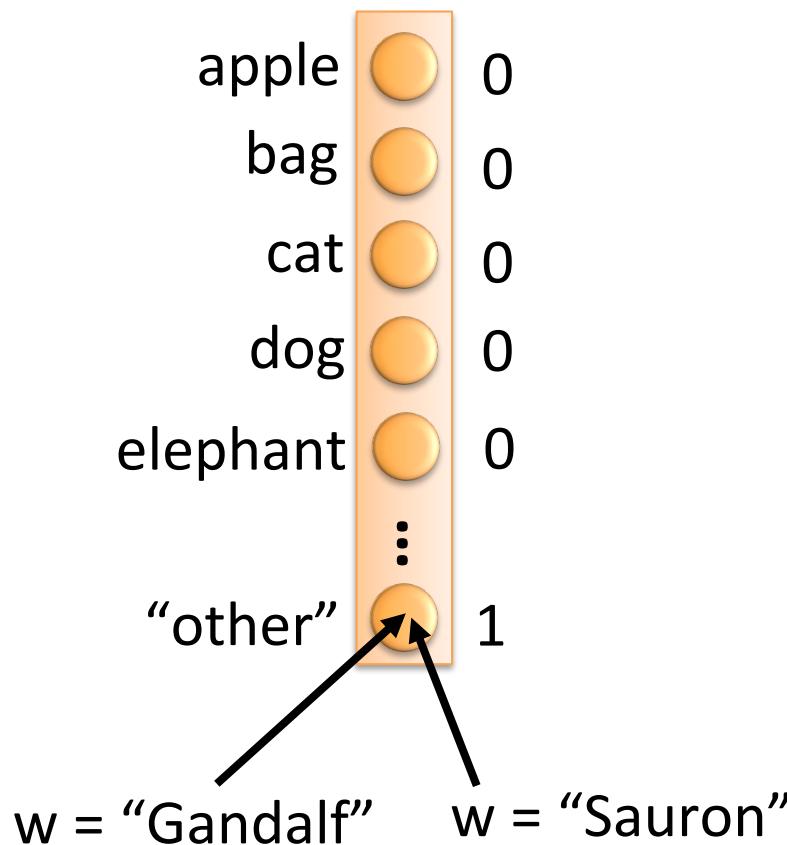
$$\text{cat} = [0 \ 0 \ 1 \ 0 \ 0]$$

$$\text{dog} = [0 \ 0 \ 0 \ 1 \ 0]$$

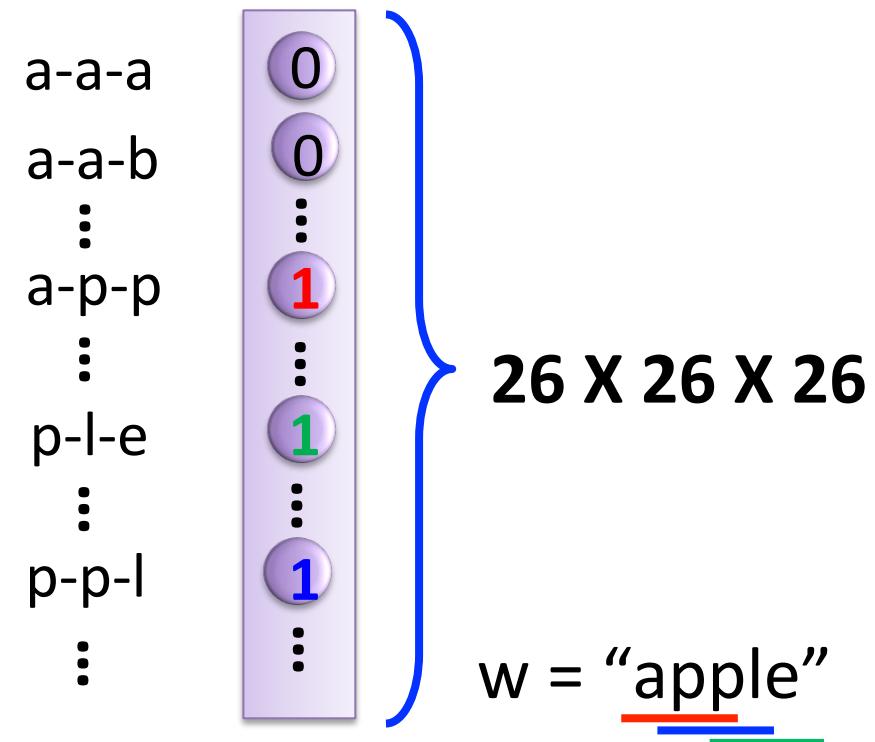
$$\text{elephant} = [0 \ 0 \ 0 \ 0 \ 1]$$

Beyond 1-of-N encoding

Dimension for “Other”



Word hashing



Example Application

dest

time of
departure

Solving slot filling by
Feedforward network?

Input: a word

(Each word is represented
as a vector)

Output:

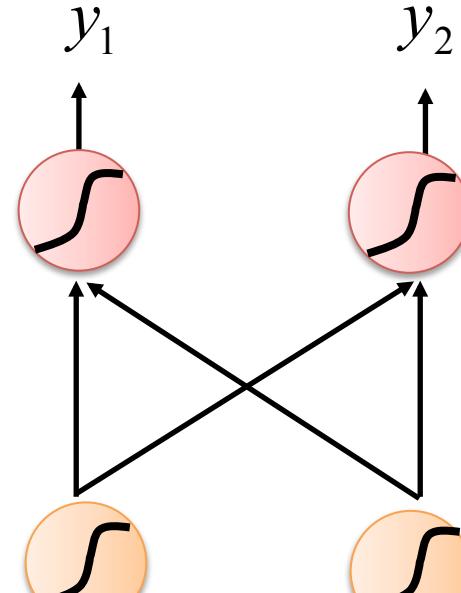
Probability distribution that
the input word belonging to
the slots

DC

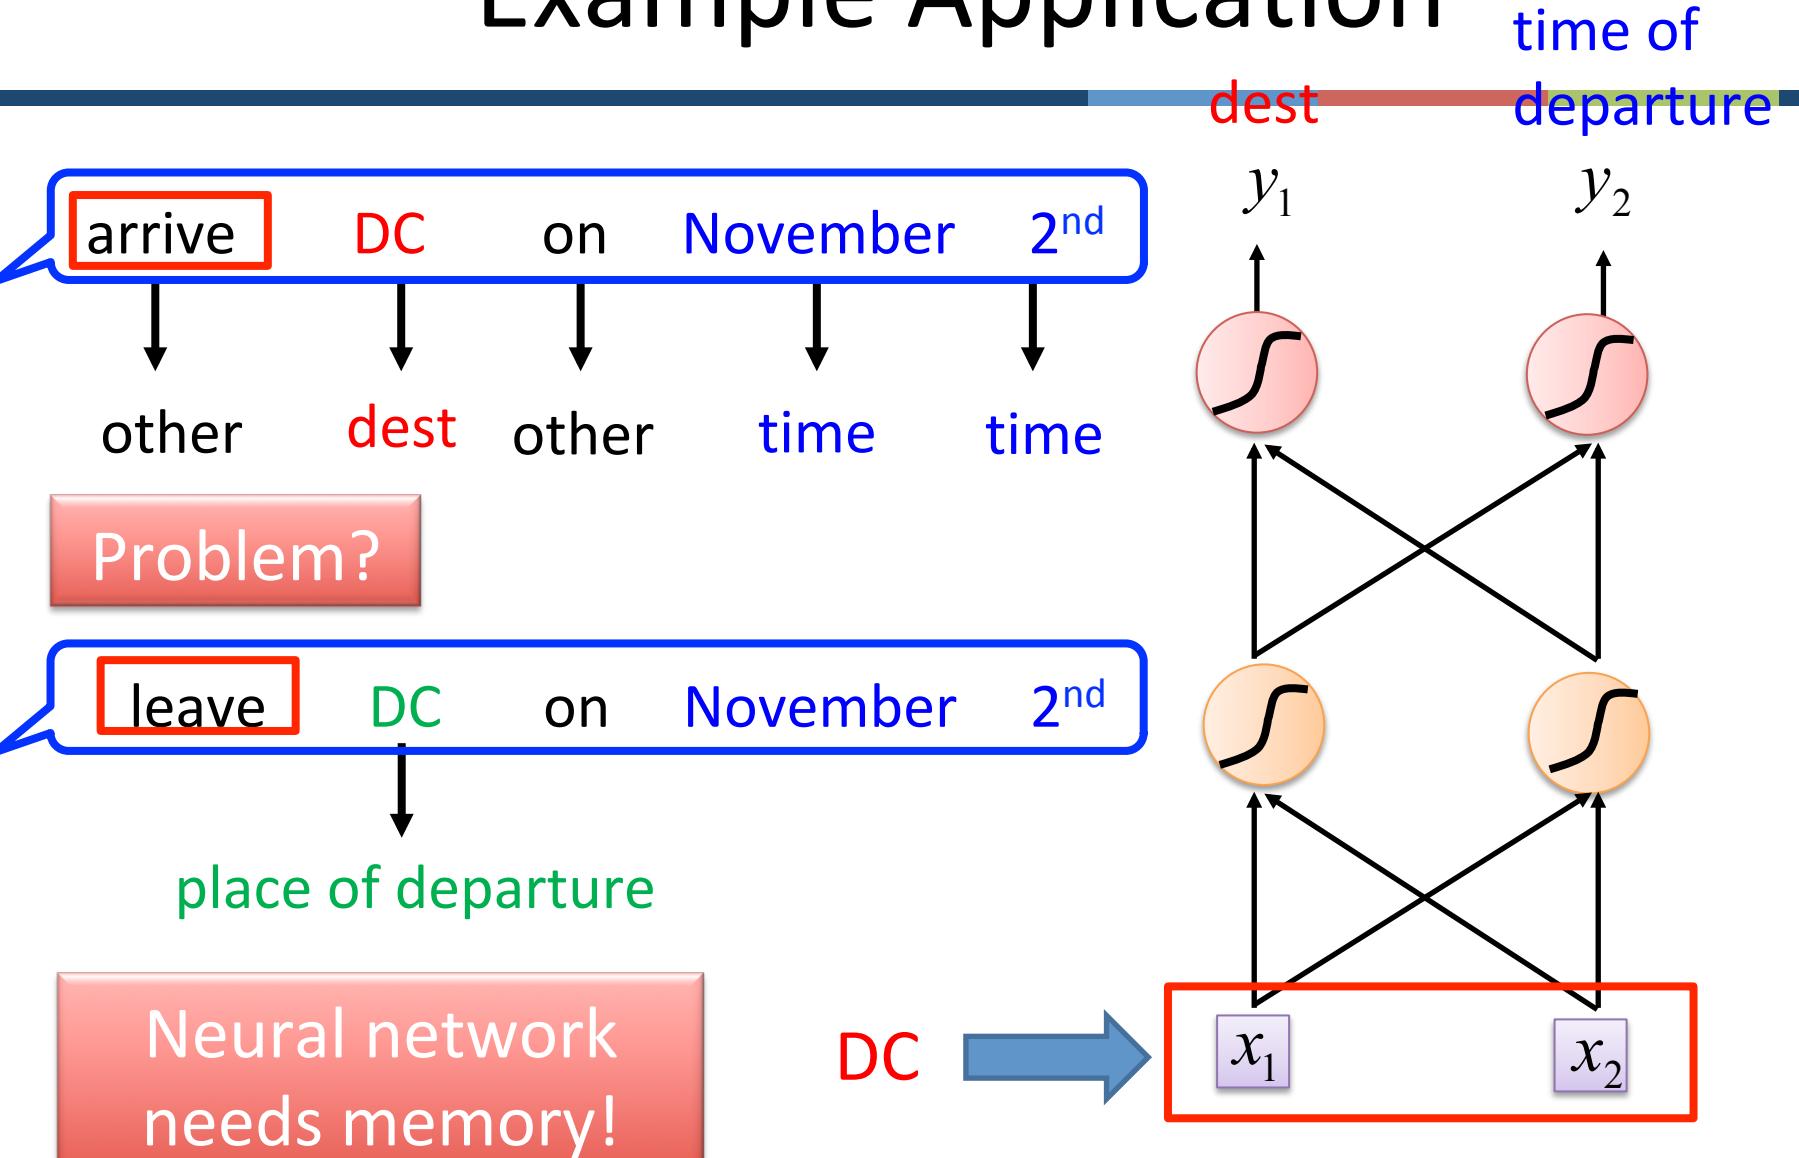


x_1

x_2



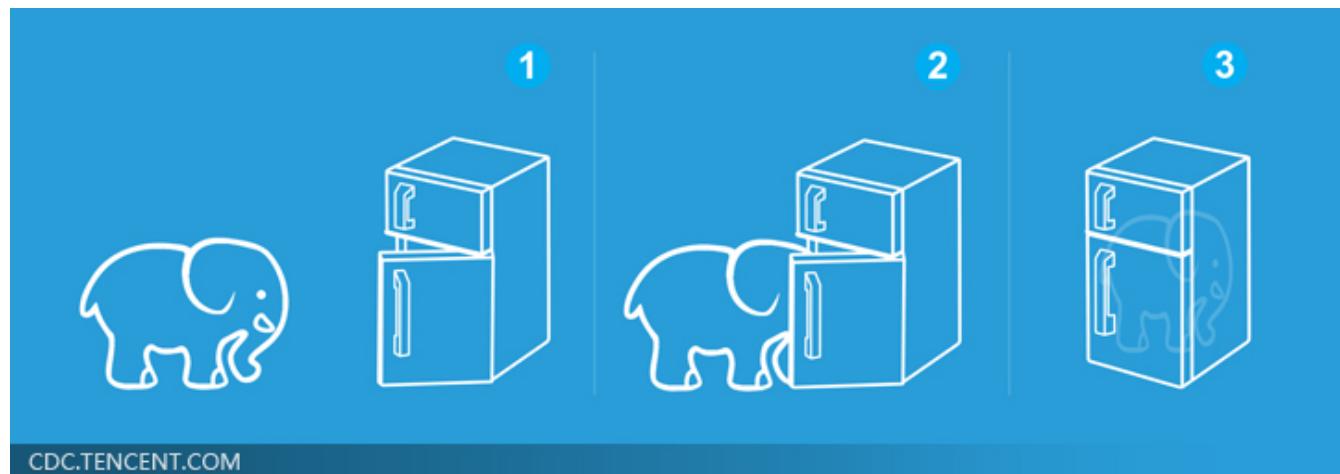
Example Application



Three Steps for Deep Learning

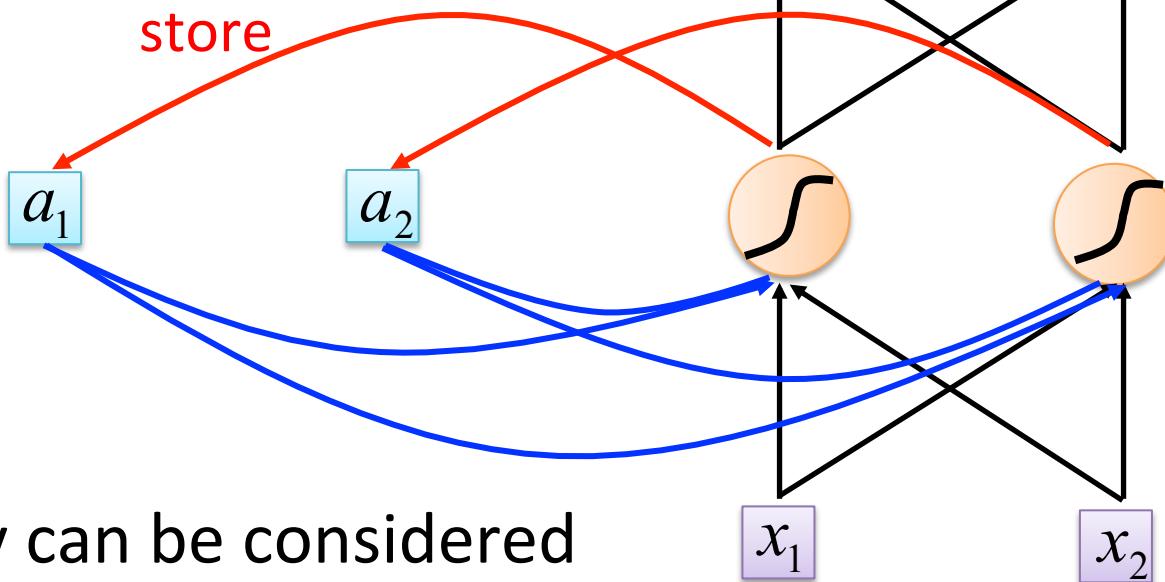


Deep Learning is so simple



Recurrent Neural Network (RNN)

The output of hidden layer
are stored in the memory.



Memory can be considered
as another input.

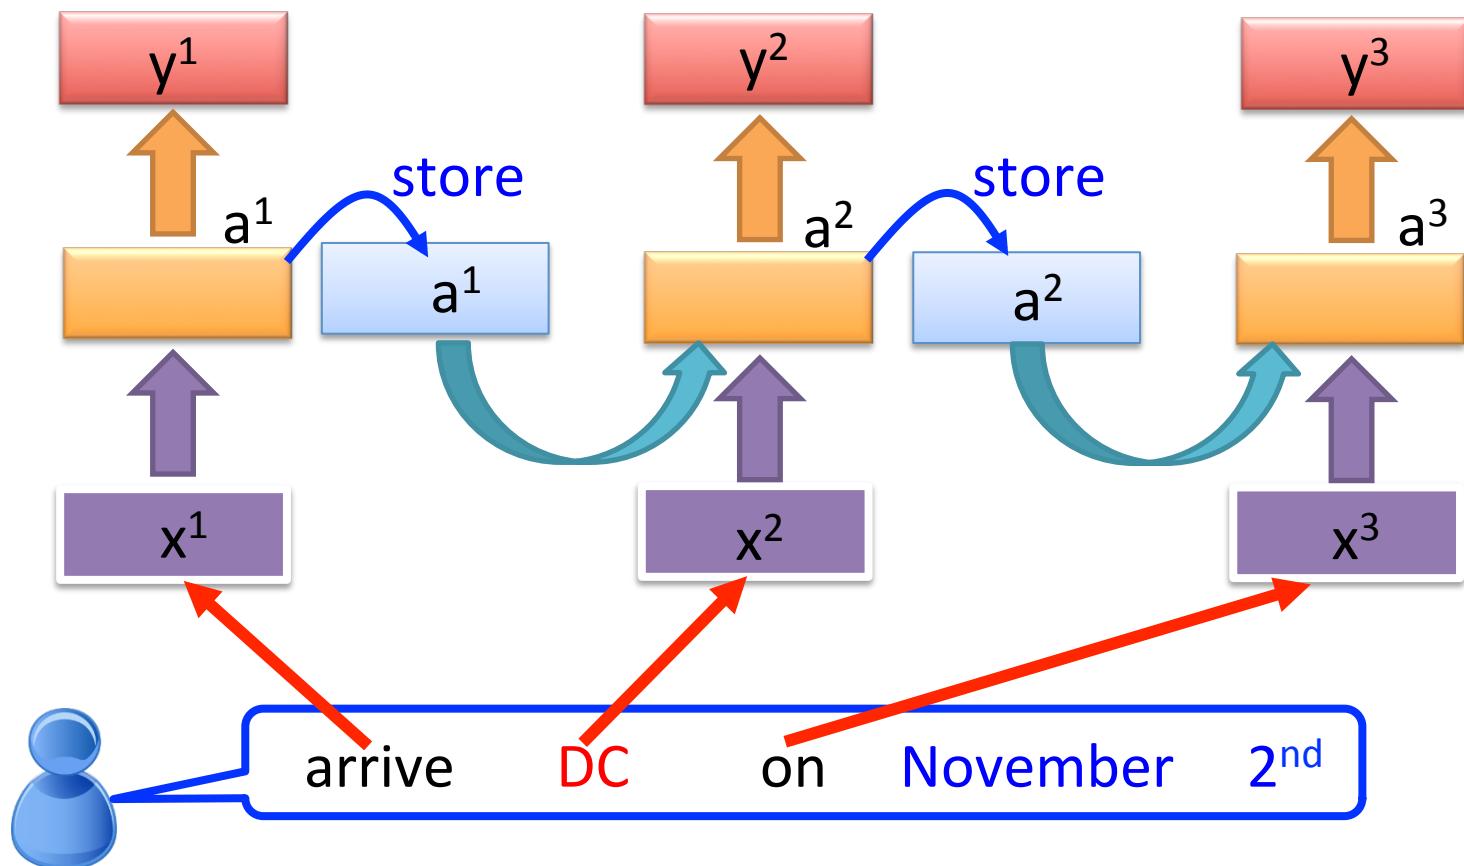
RNN

The same network is used again and again.

Probability of
“arrive” in each slot

Probability of “DC”
in each slot

Probability of
“on” in each slot



RNN

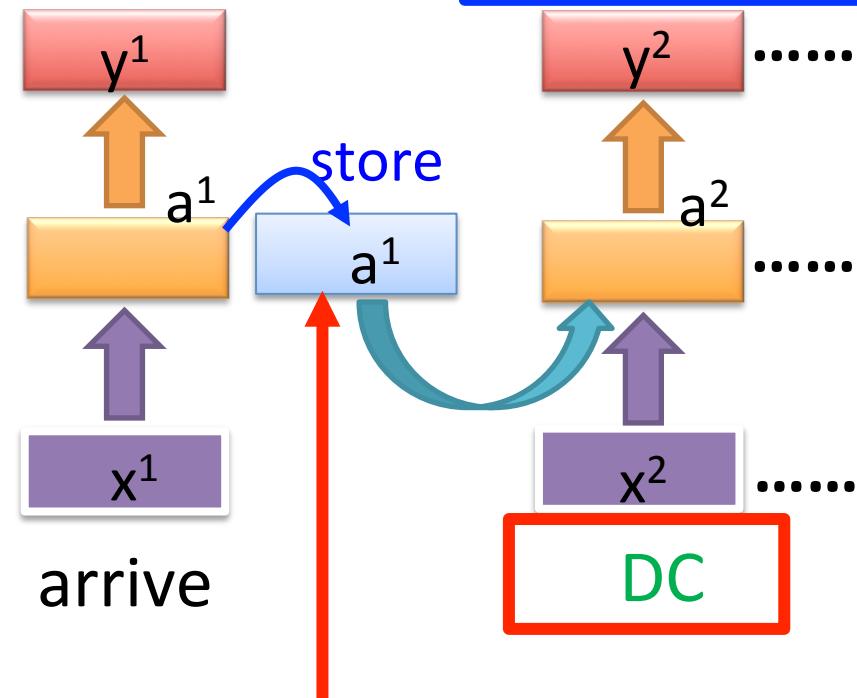
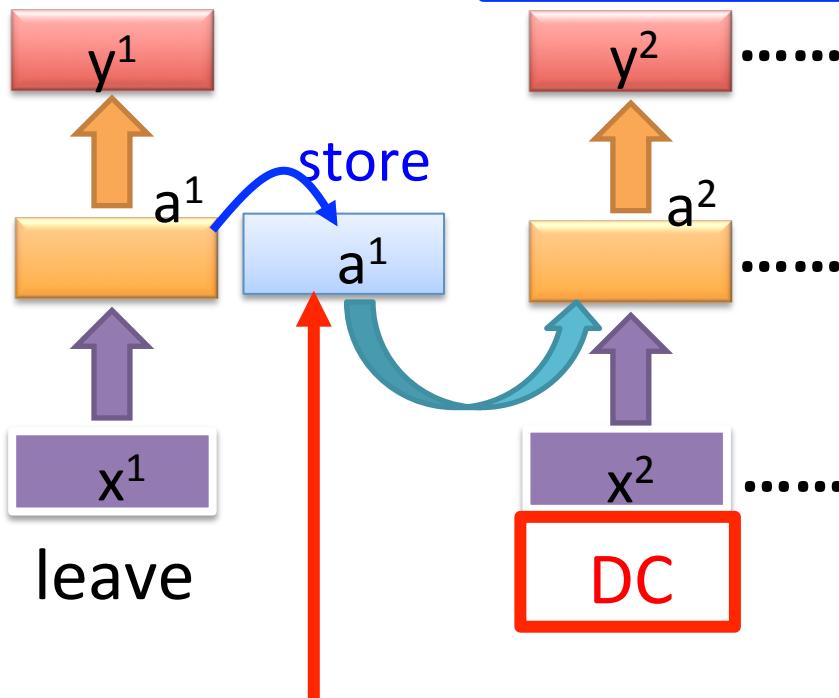
Different

Prob of “leave”
in each slot

Prob of “DC” in
each slot

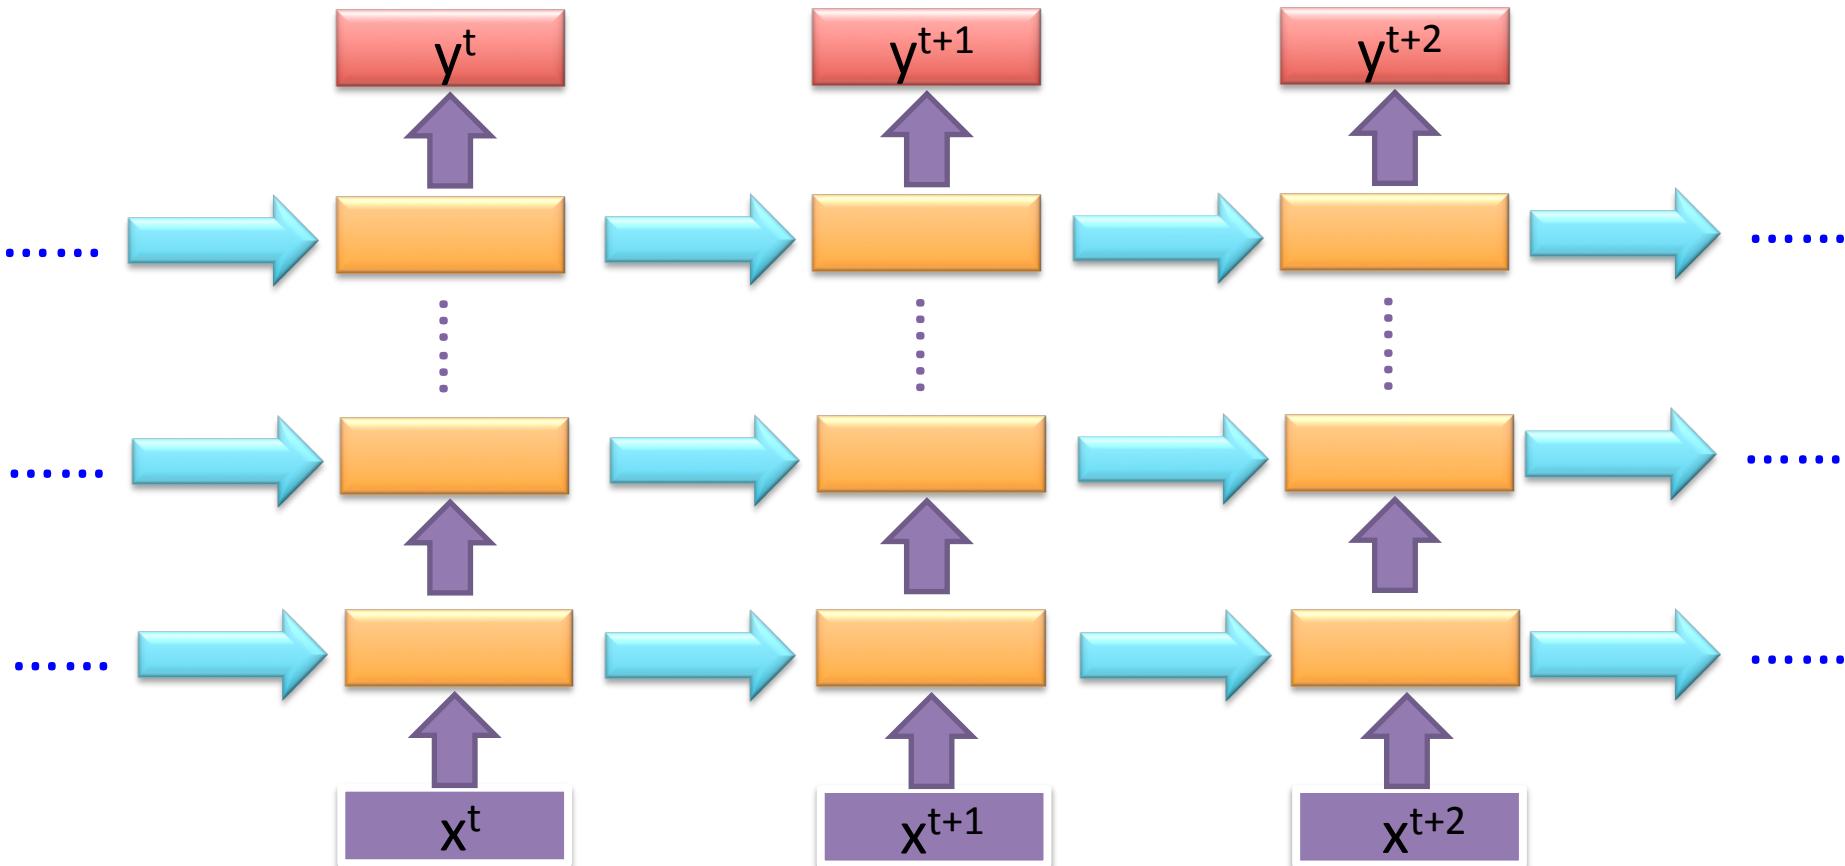
Prob of “arrive”
in each slot

Prob of “DC” in
each slot

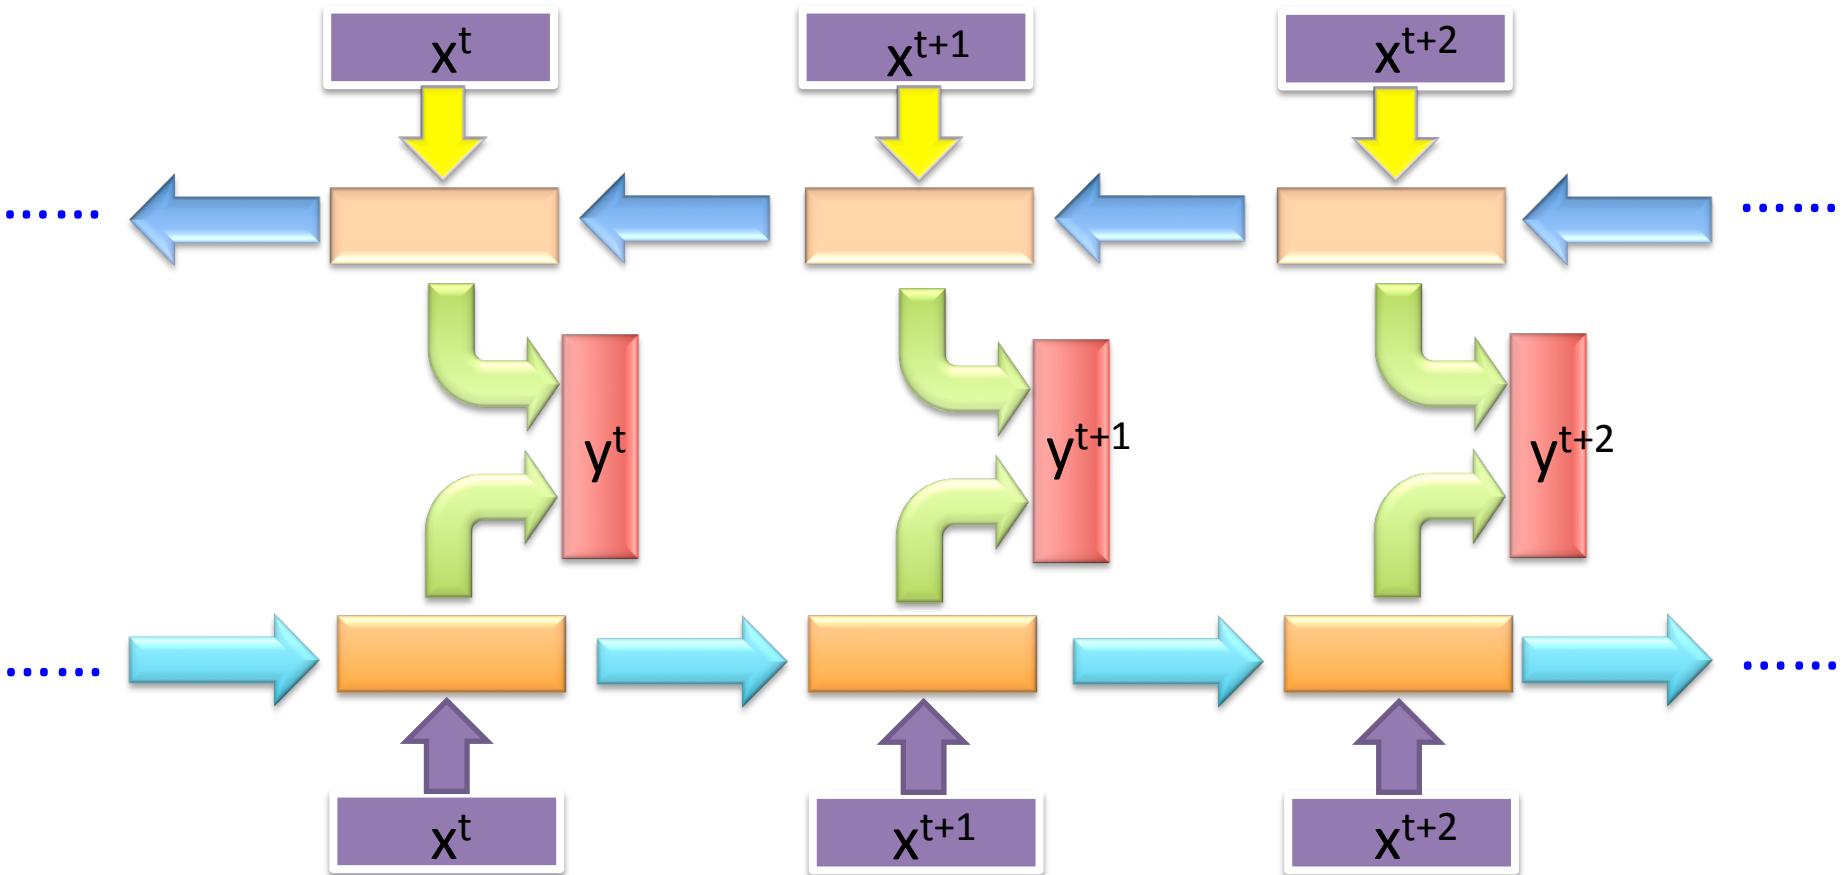


The values stored in the memory is different.

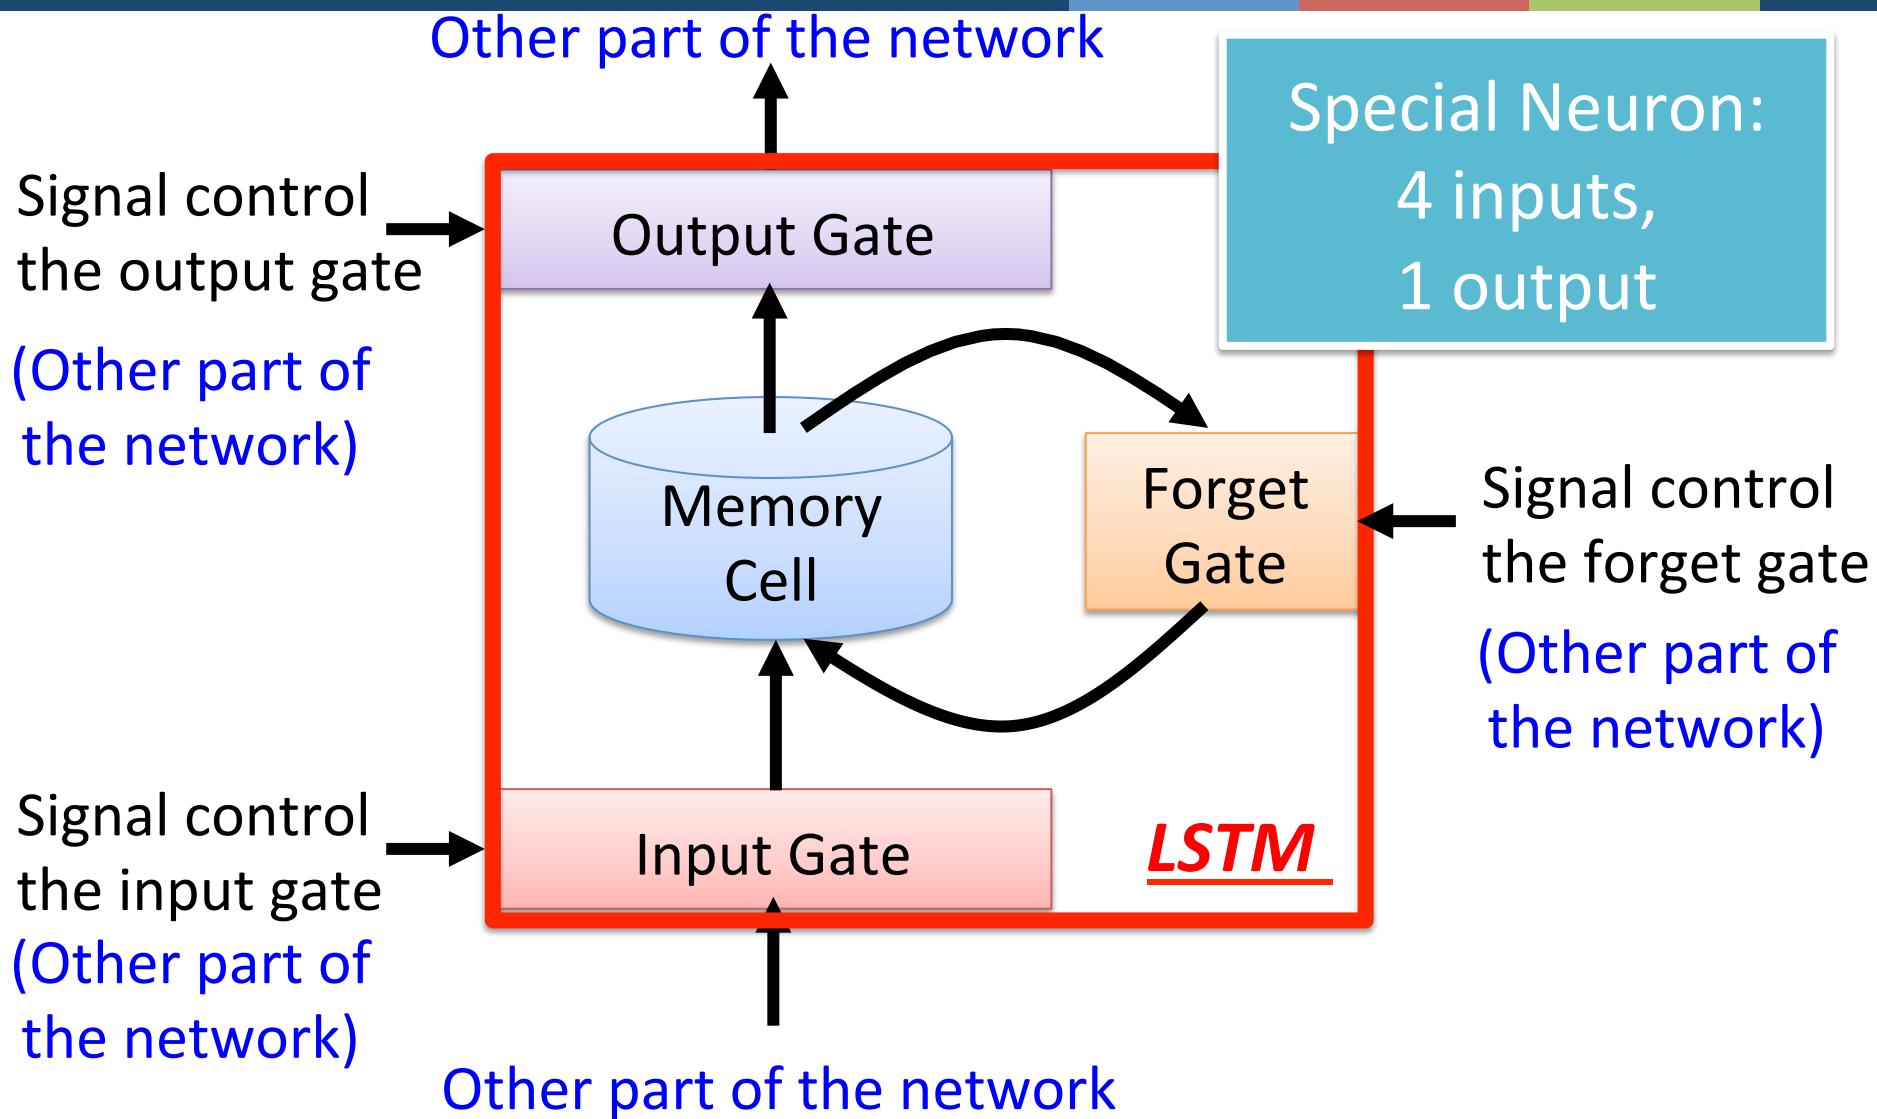
Of course it can be deep ...

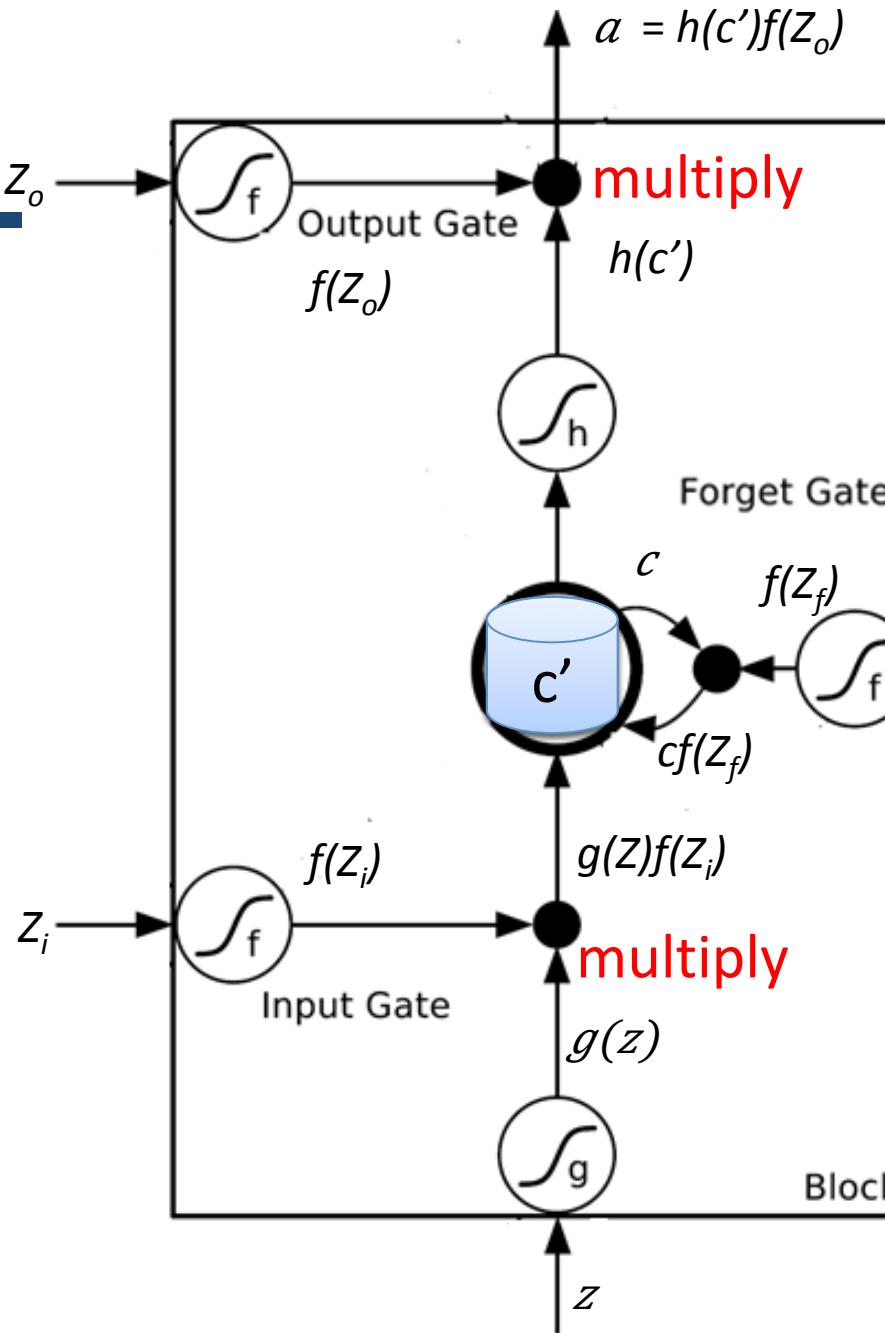


Bidirectional RNN



Long Short-term Memory (LSTM)



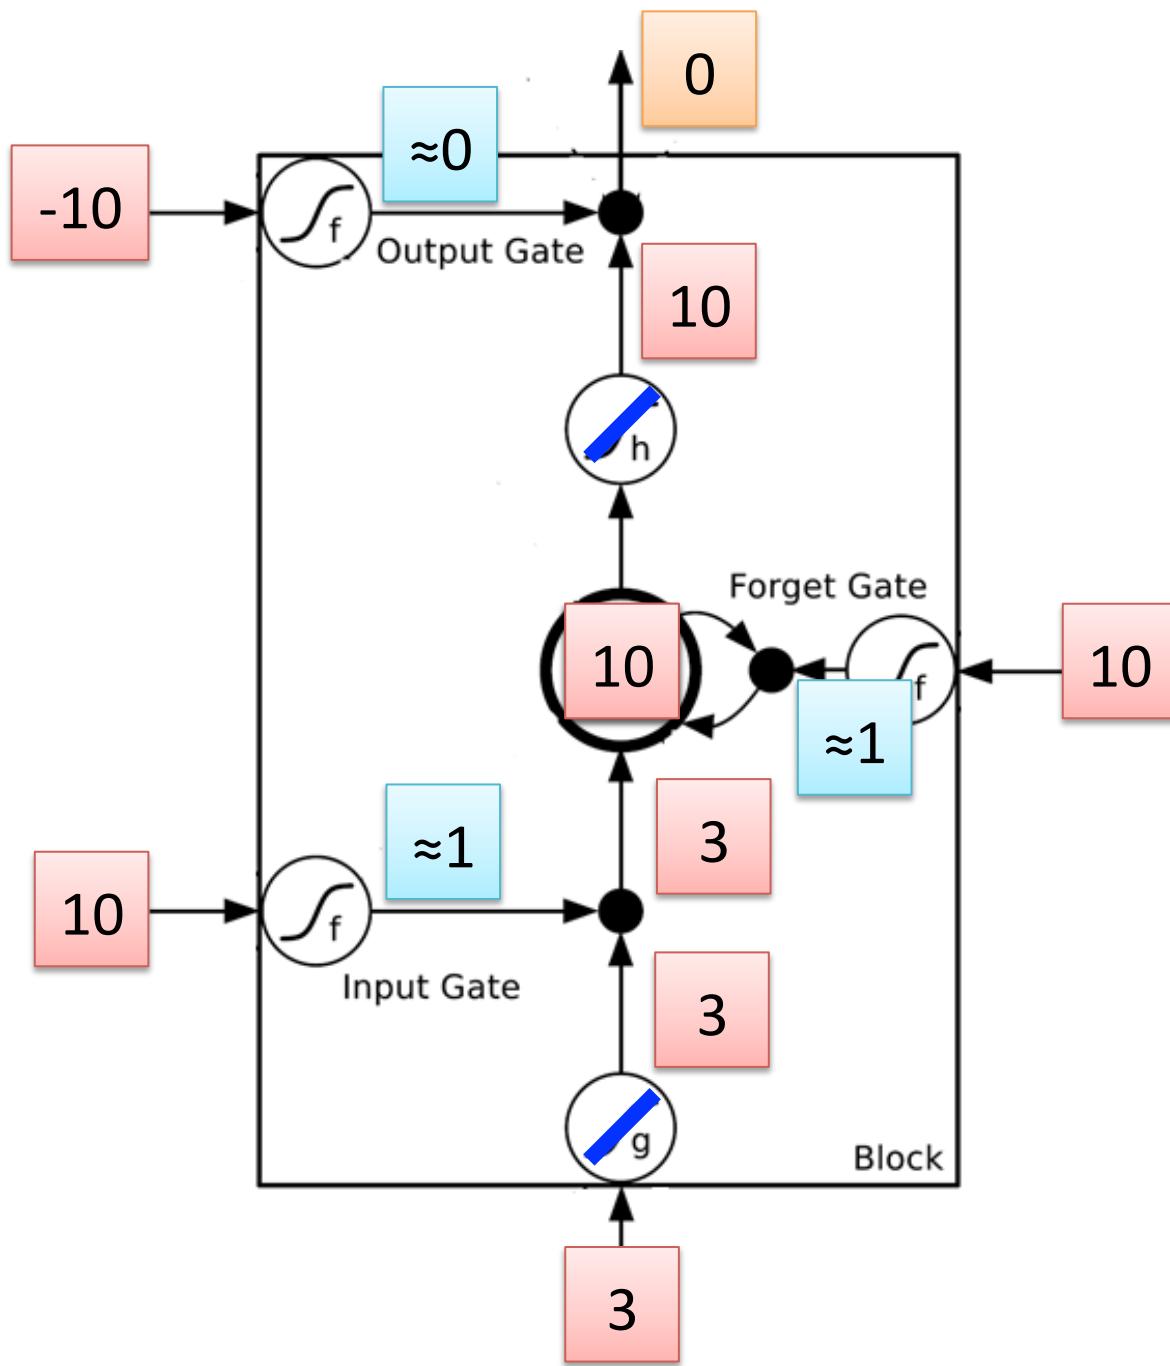


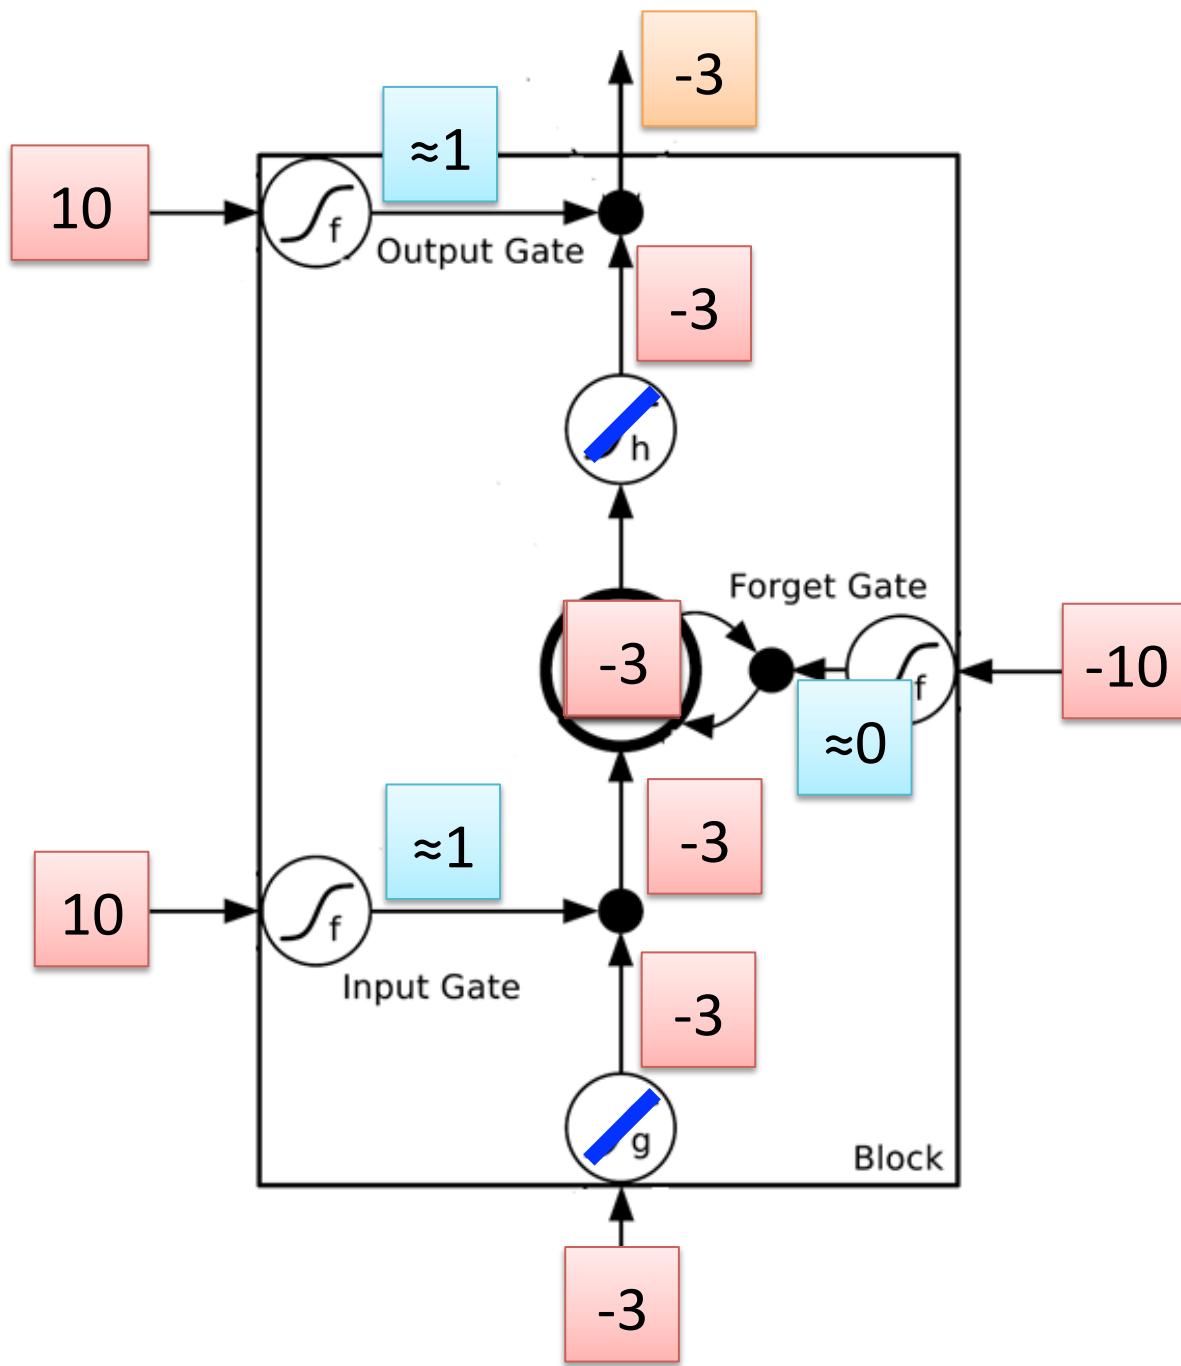
Activation function f is usually a sigmoid function

Between 0 and 1

Mimic open and close gate

$$c' = g(z)f(Z_i) + cf(Z_f)$$

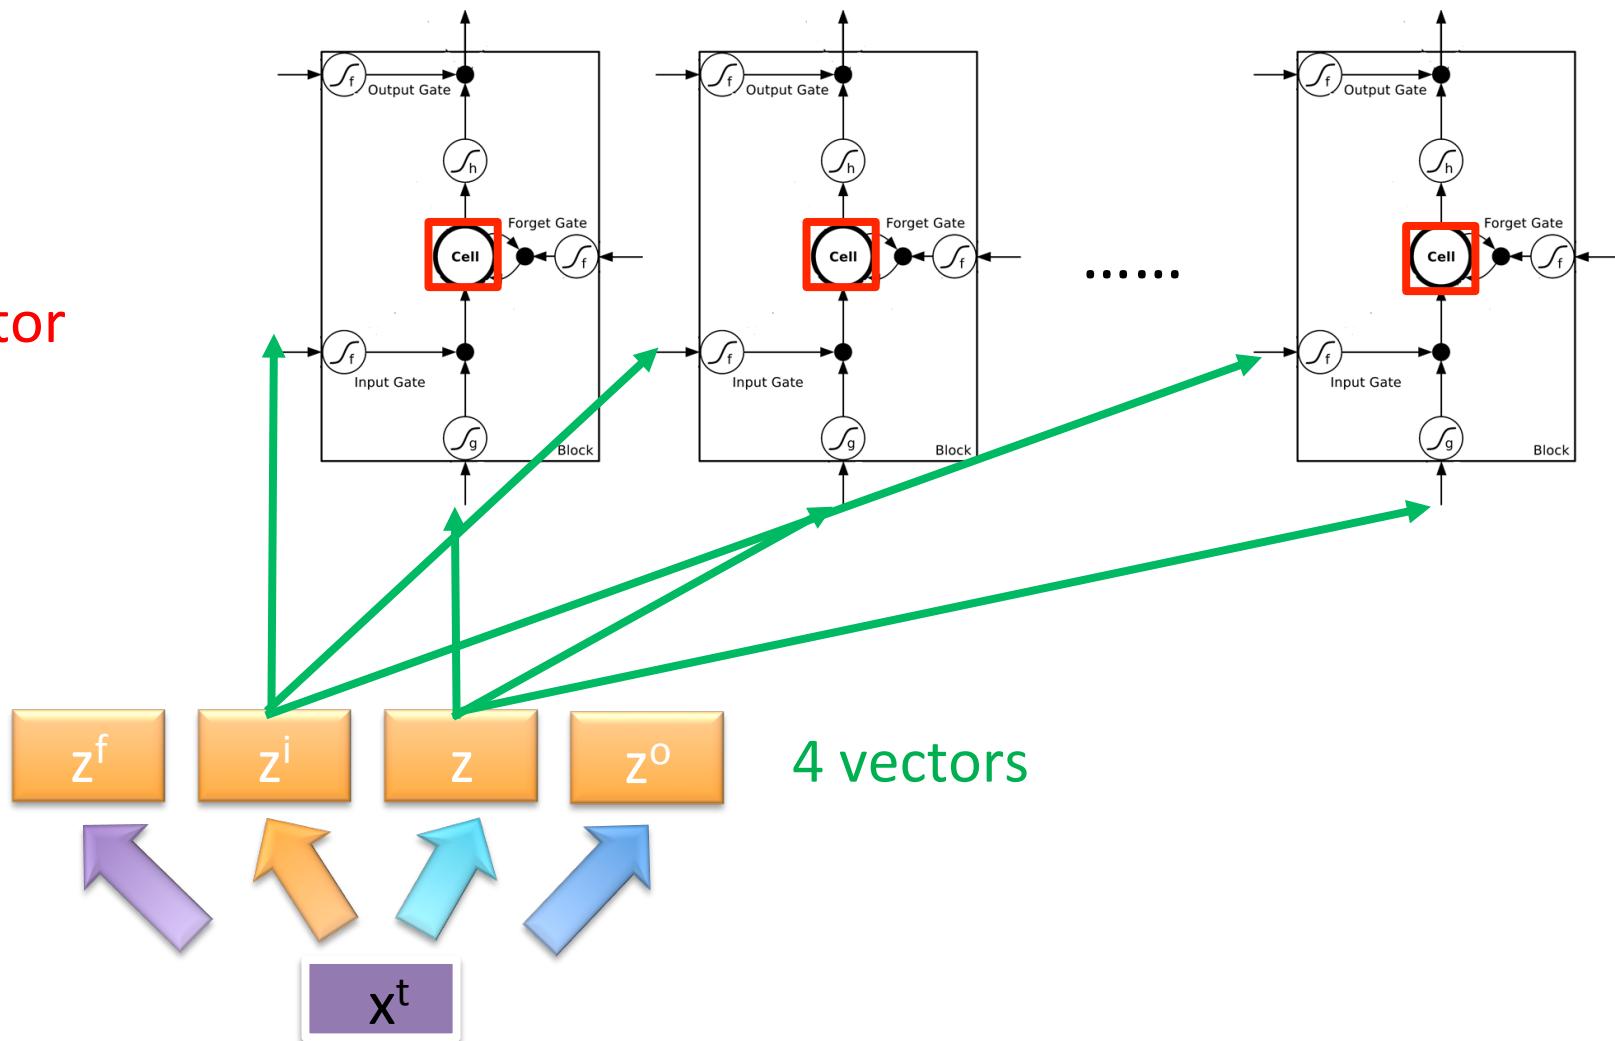




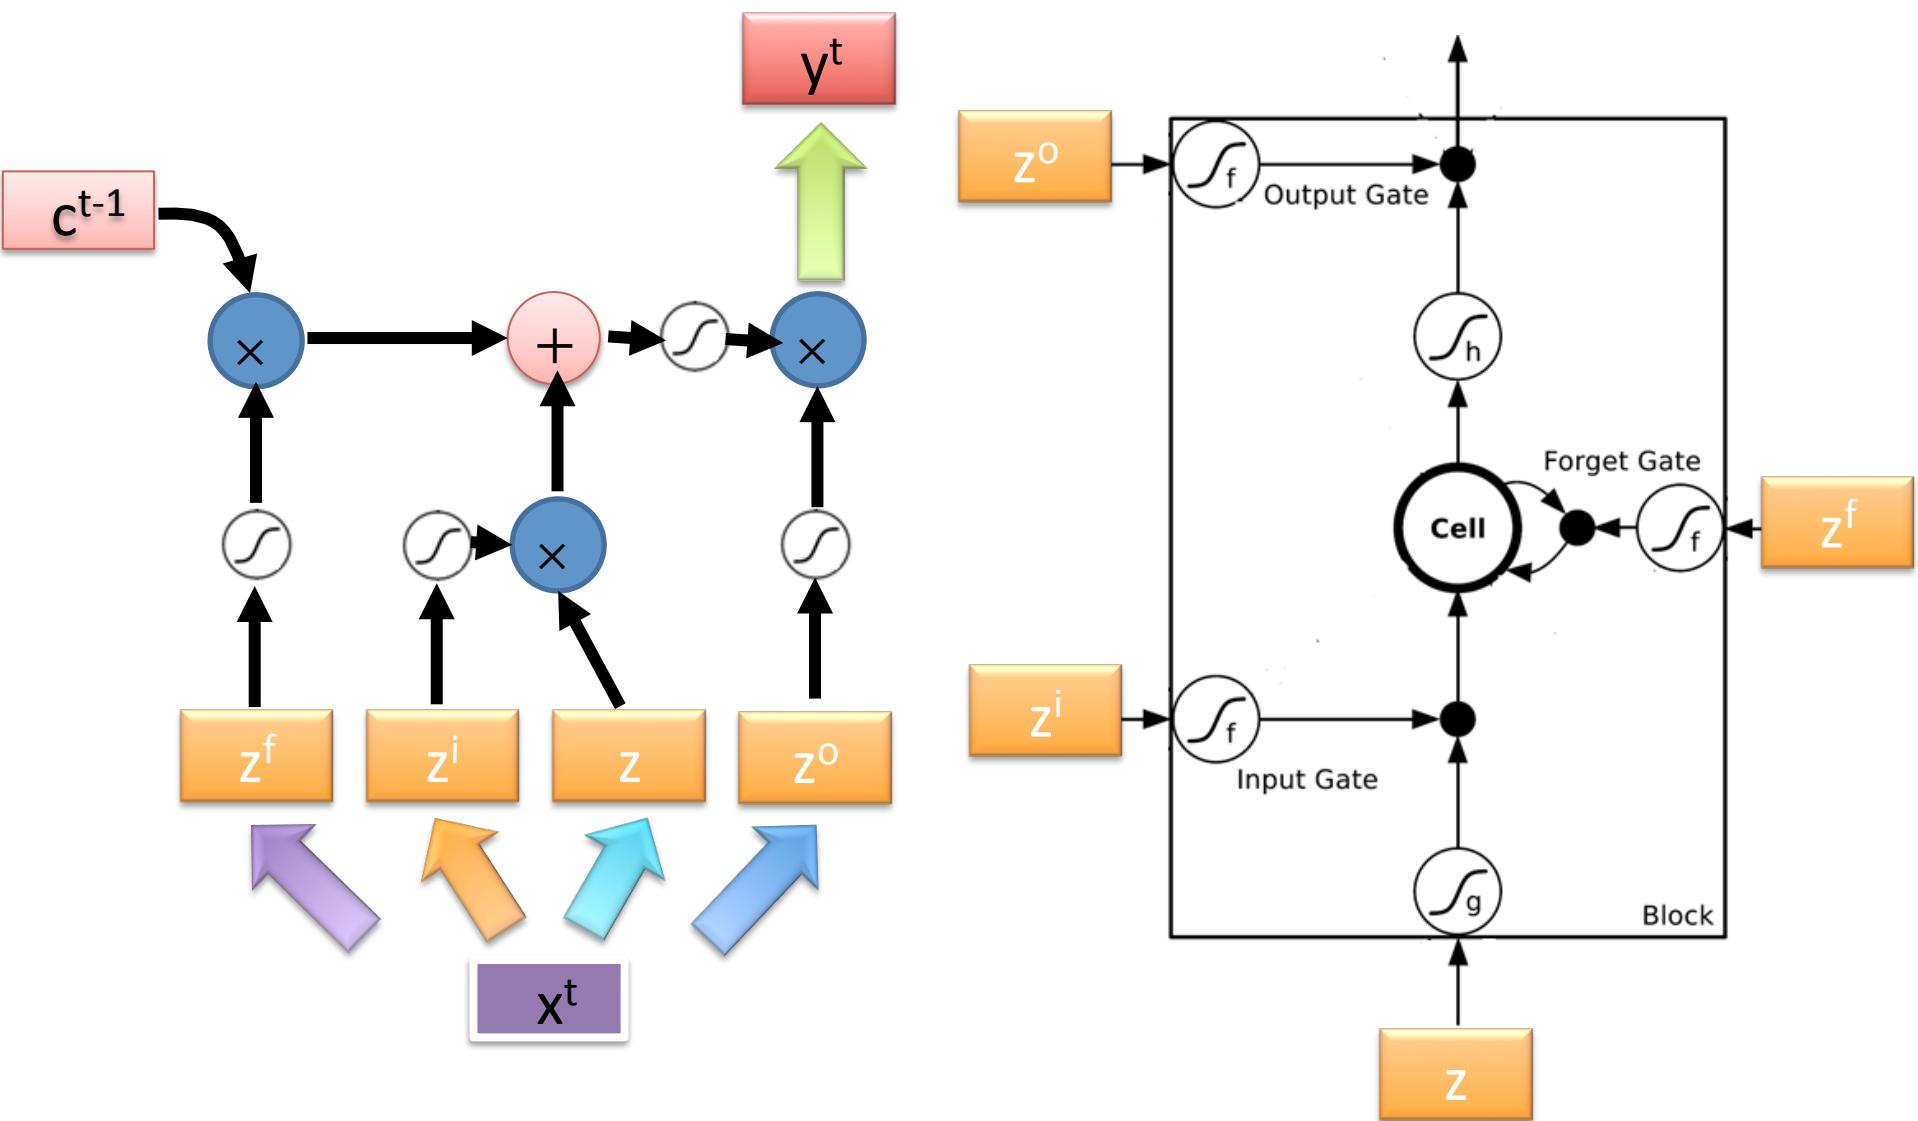
LSTM

C^{t-1}

vector

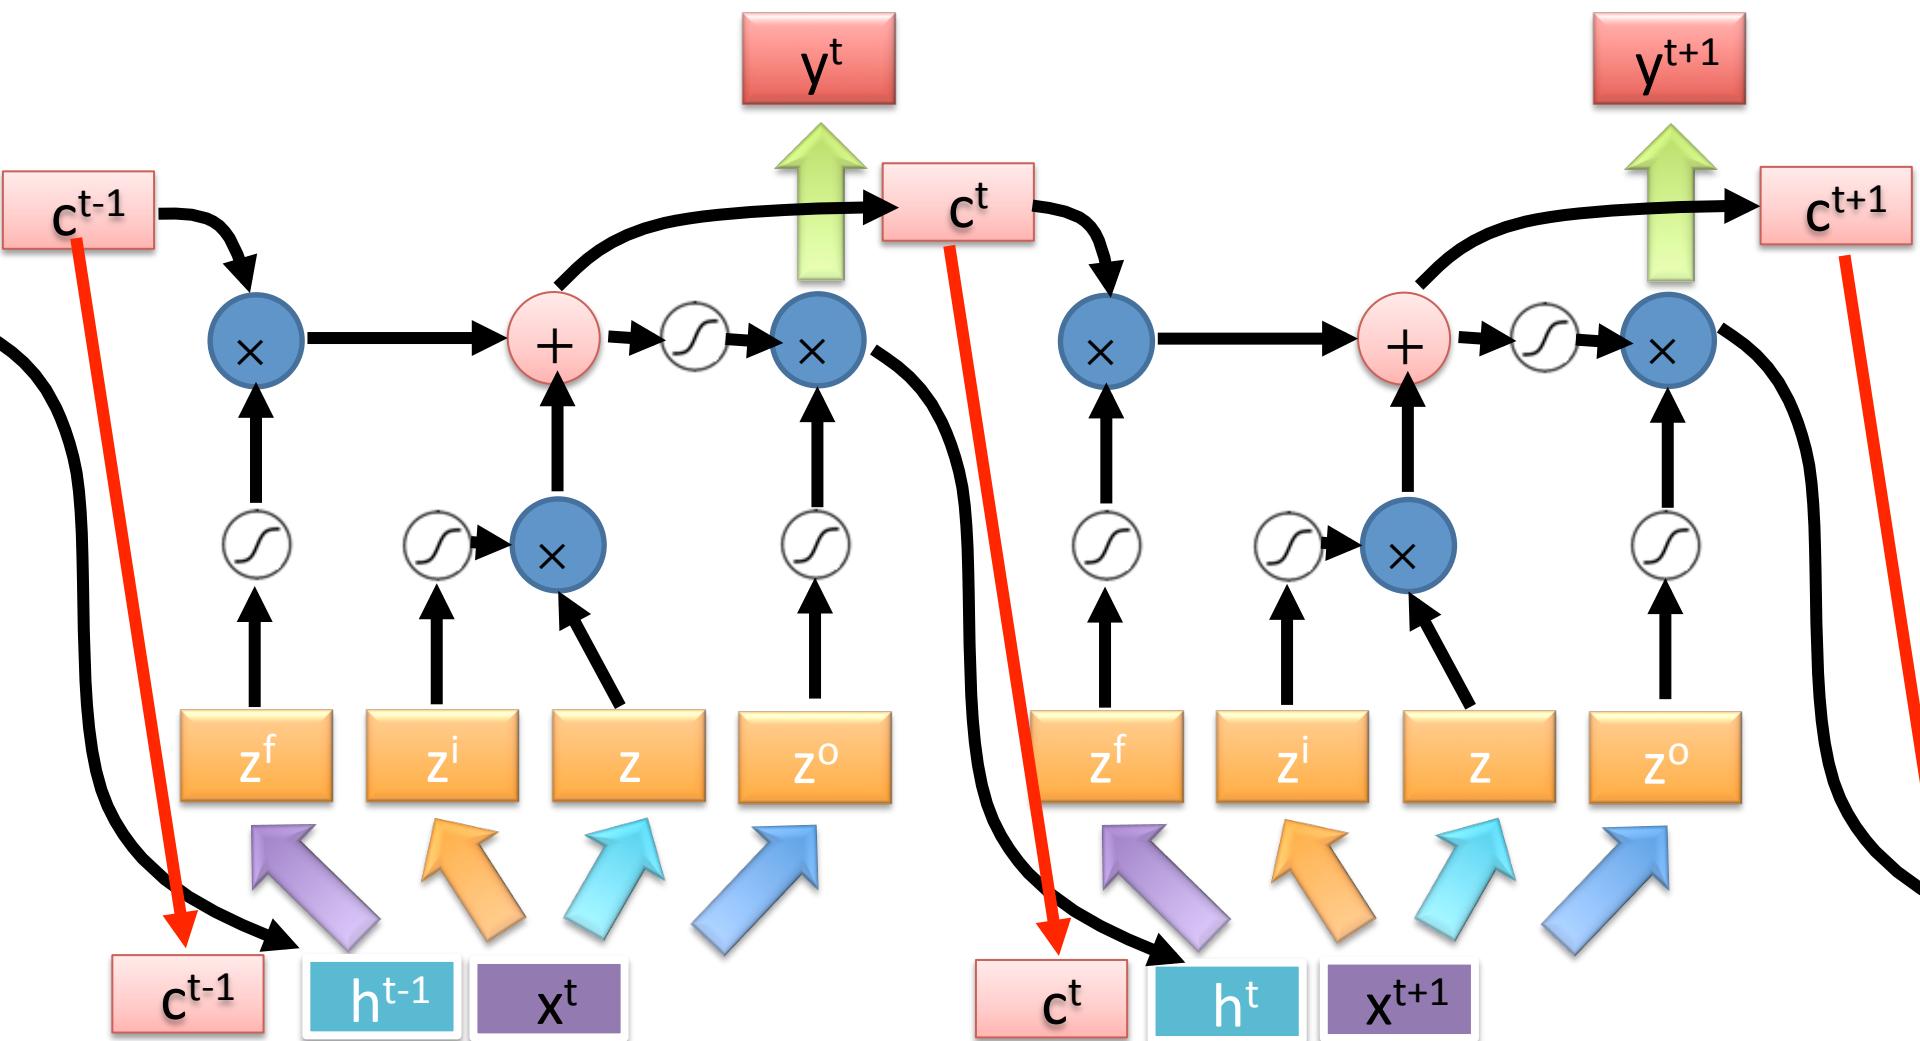


LSTM

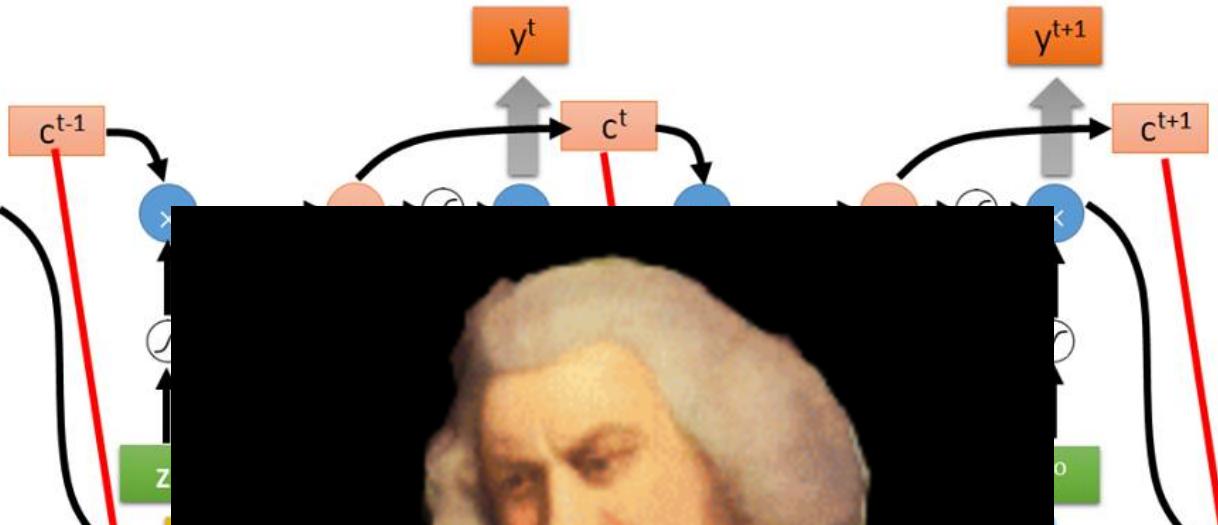


LSTM

Extension: “peephole”



Multiple-layer LSTM



Don't worry if you cannot understand this.
Keras can handle it.

Keras supports
“LSTM”, “GRU”, “SimpleRNN” layers

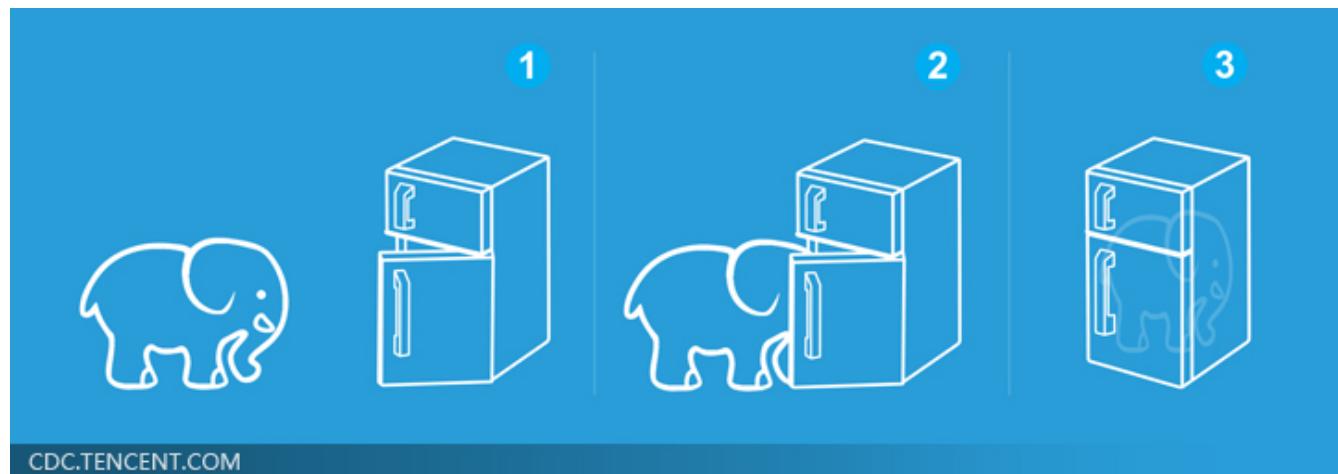
This is quite
standard now.



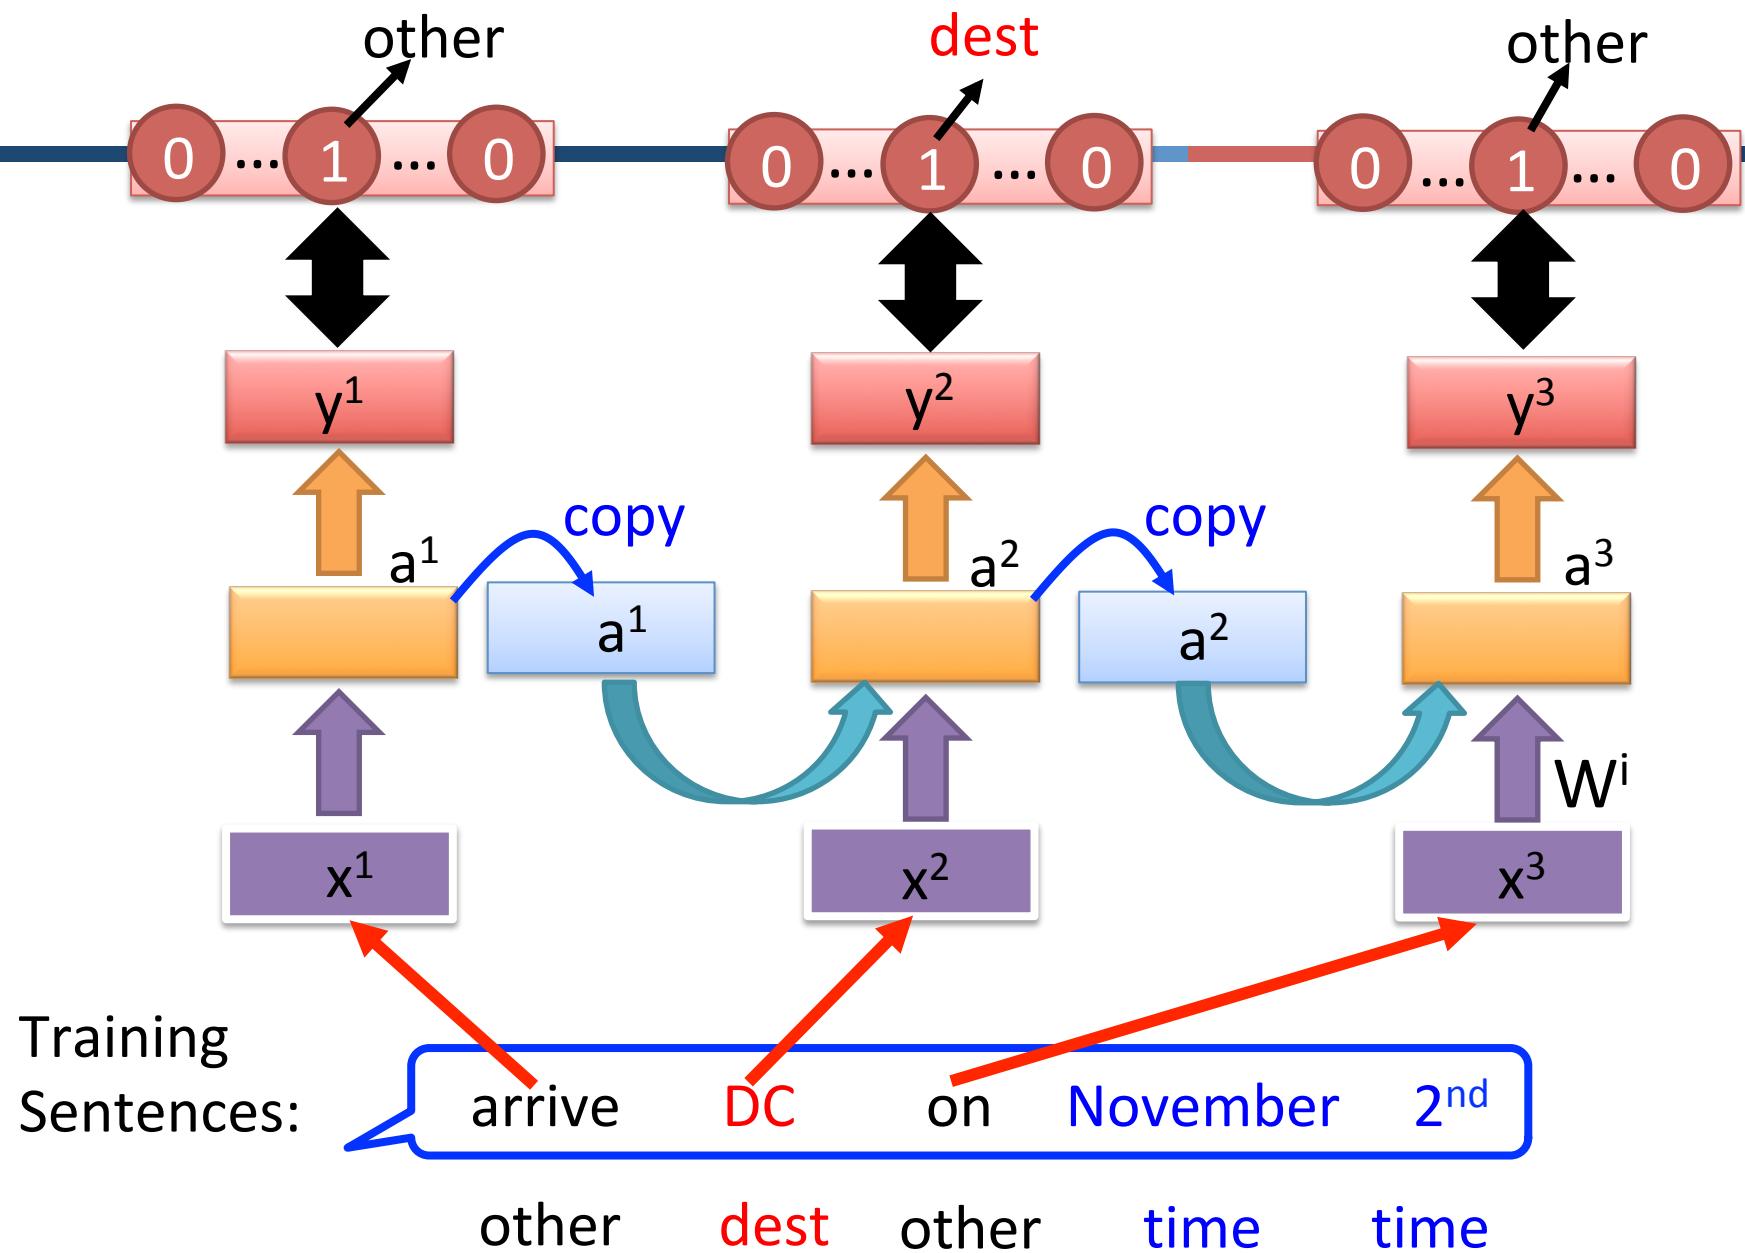
Three Steps for Deep Learning



Deep Learning is so simple



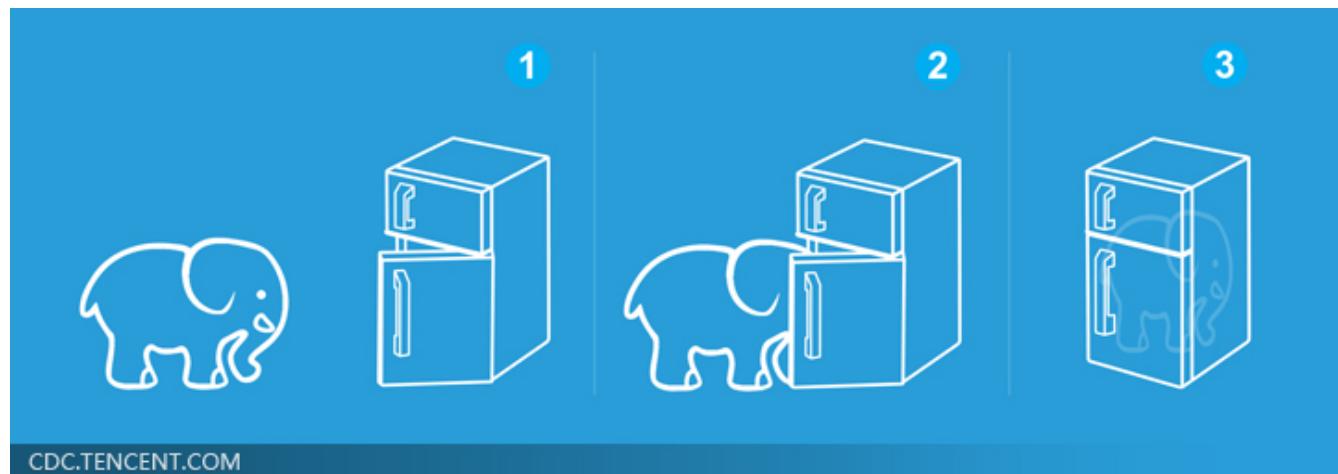
Learning Target



Three Steps for Deep Learning

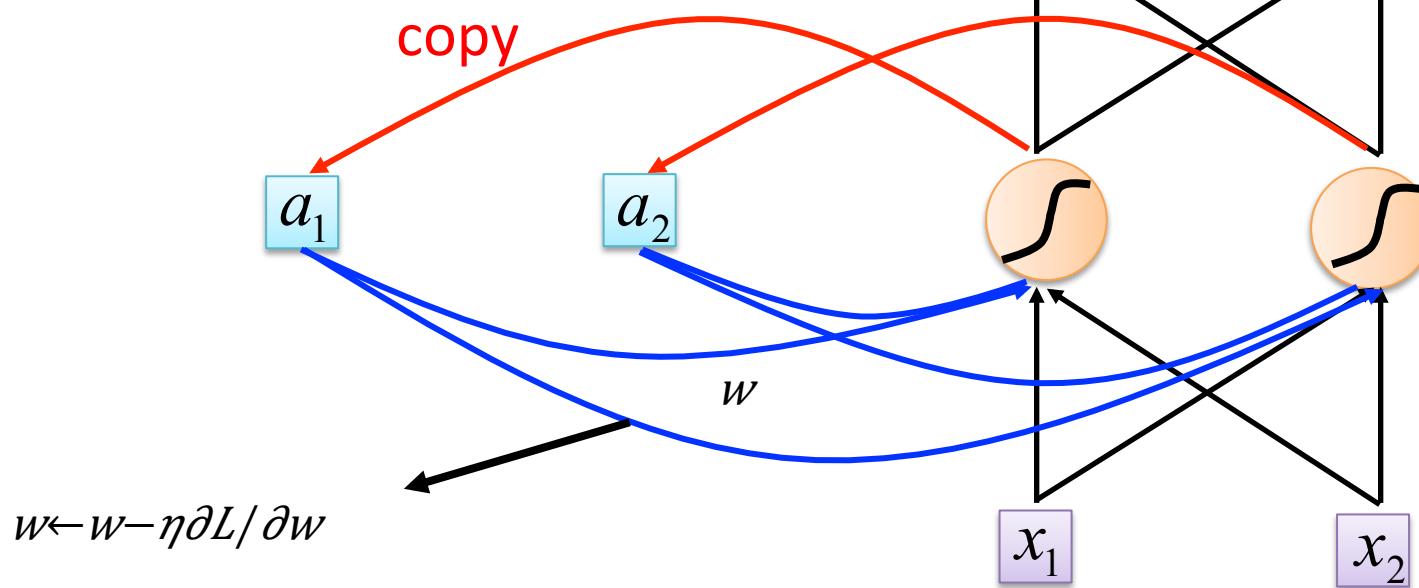


Deep Learning is so simple



Learning

Backpropagation
through time (BPTT)

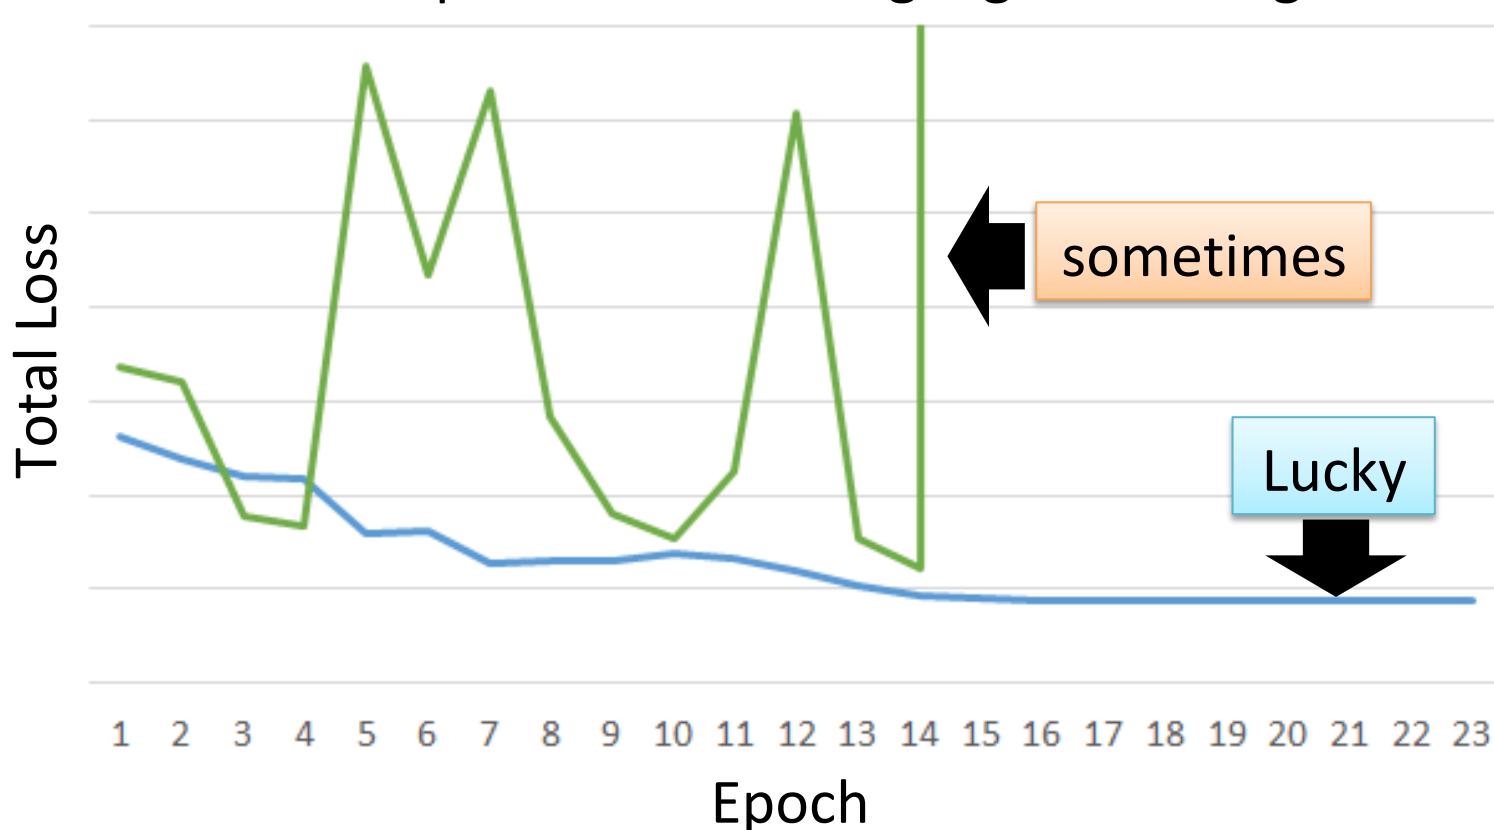


RNN Learning is very difficult in practice.

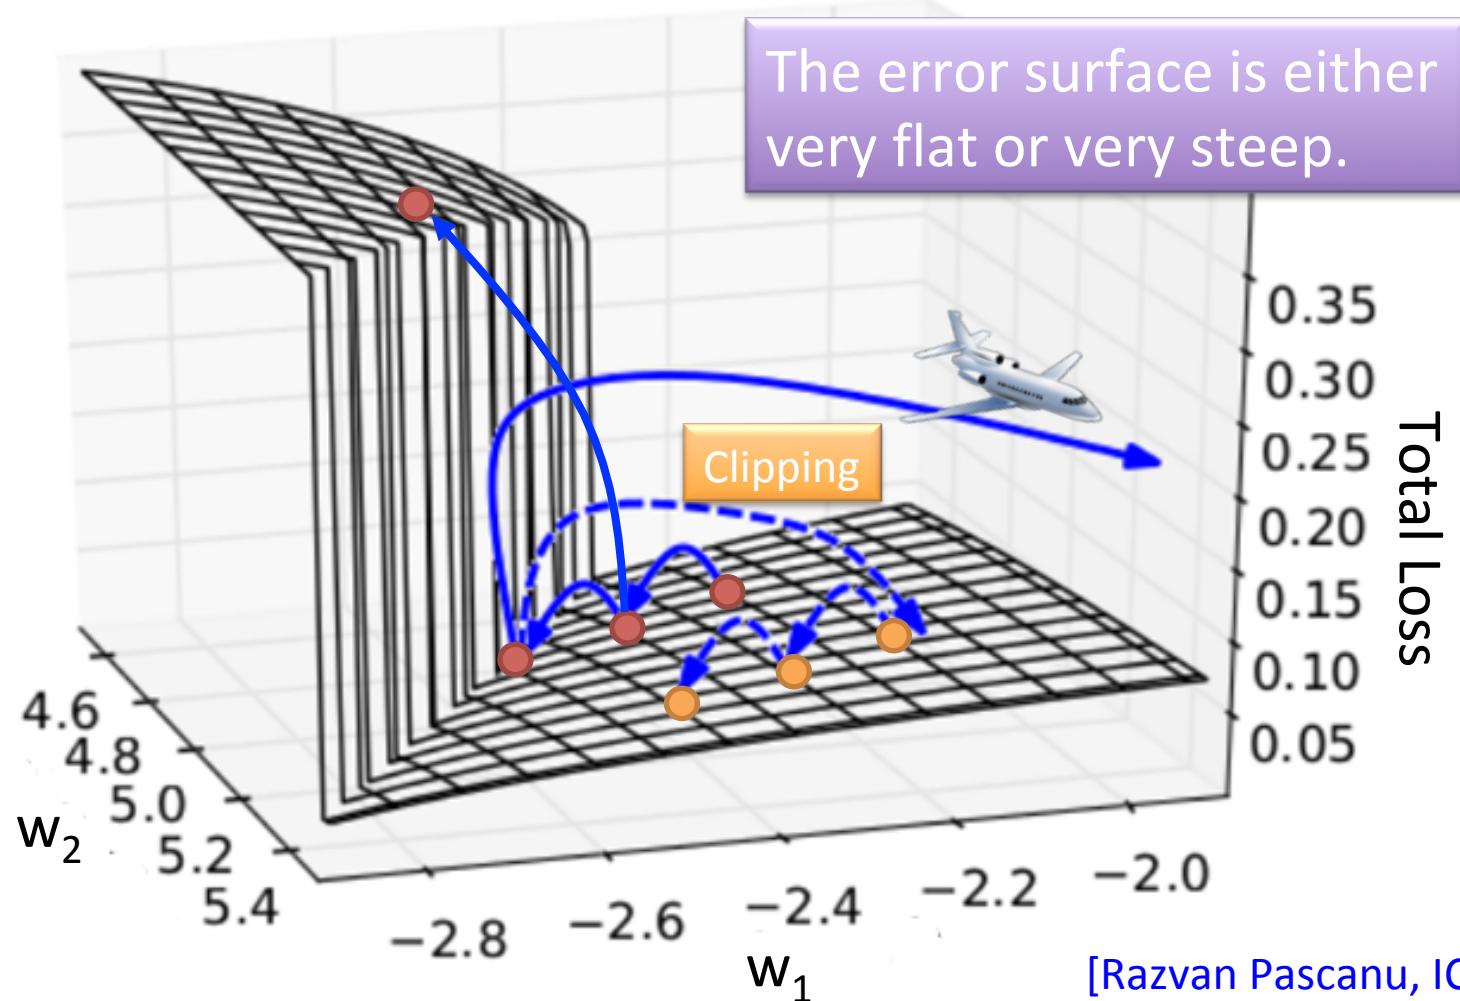
Unfortunately

- RNN-based network is not always easy to learn

Real experiments on Language modeling



The error surface is rough.



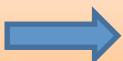
Why?

$$w=1$$



$$y^{1000} = 1$$

$$w=1.01$$



$$y^{1000} \approx 20000$$

$$w=0.99$$



$$y^{1000} \approx 0$$

$$w=0.01$$



$$y^{1000} \approx 0$$

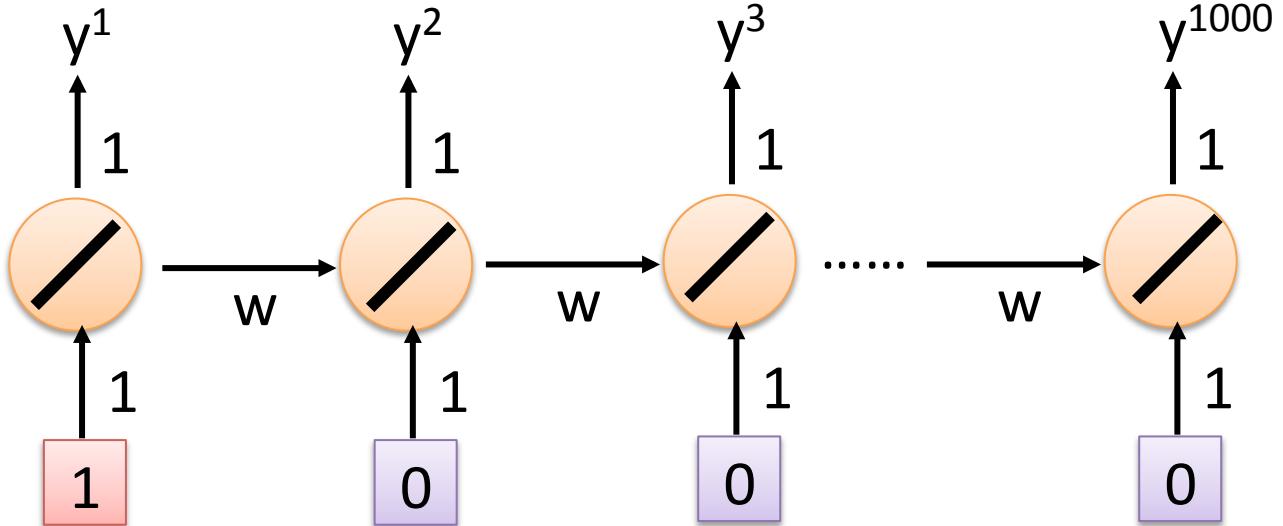
Large
 $\partial L/\partial w$

Small Learning
rate?

small
 $\partial L/\partial w$

Large Learning
rate?

Toy Example



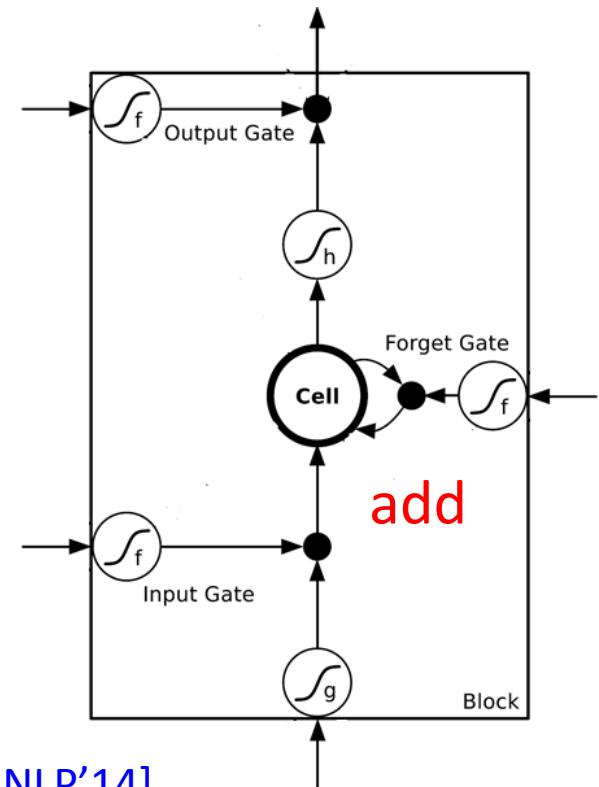
Helpful Techniques

• Long Short-term Memory (LSTM)

- Can deal with gradient vanishing (not gradient explode)
- Memory and input are added
- The influence never disappears unless forget gate is closed
- No Gradient vanishing
(If forget gate is open)

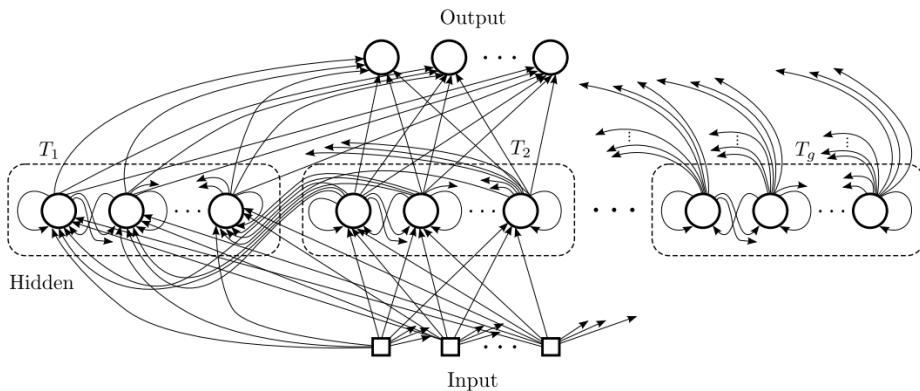
Gated Recurrent Unit (GRU):
simpler than LSTM

[Cho, EMNLP'14]



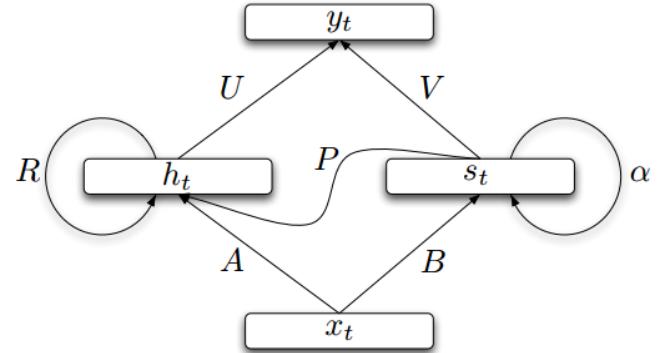
Helpful Techniques

Clockwise RNN



[Jan Koutnik, JMLR'14]

Structurally Constrained
Recurrent Network (SCRN)



[Tomas Mikolov, ICLR'15]

Vanilla RNN Initialized with Identity matrix + ReLU activation function [Quoc V. Le, arXiv'15]

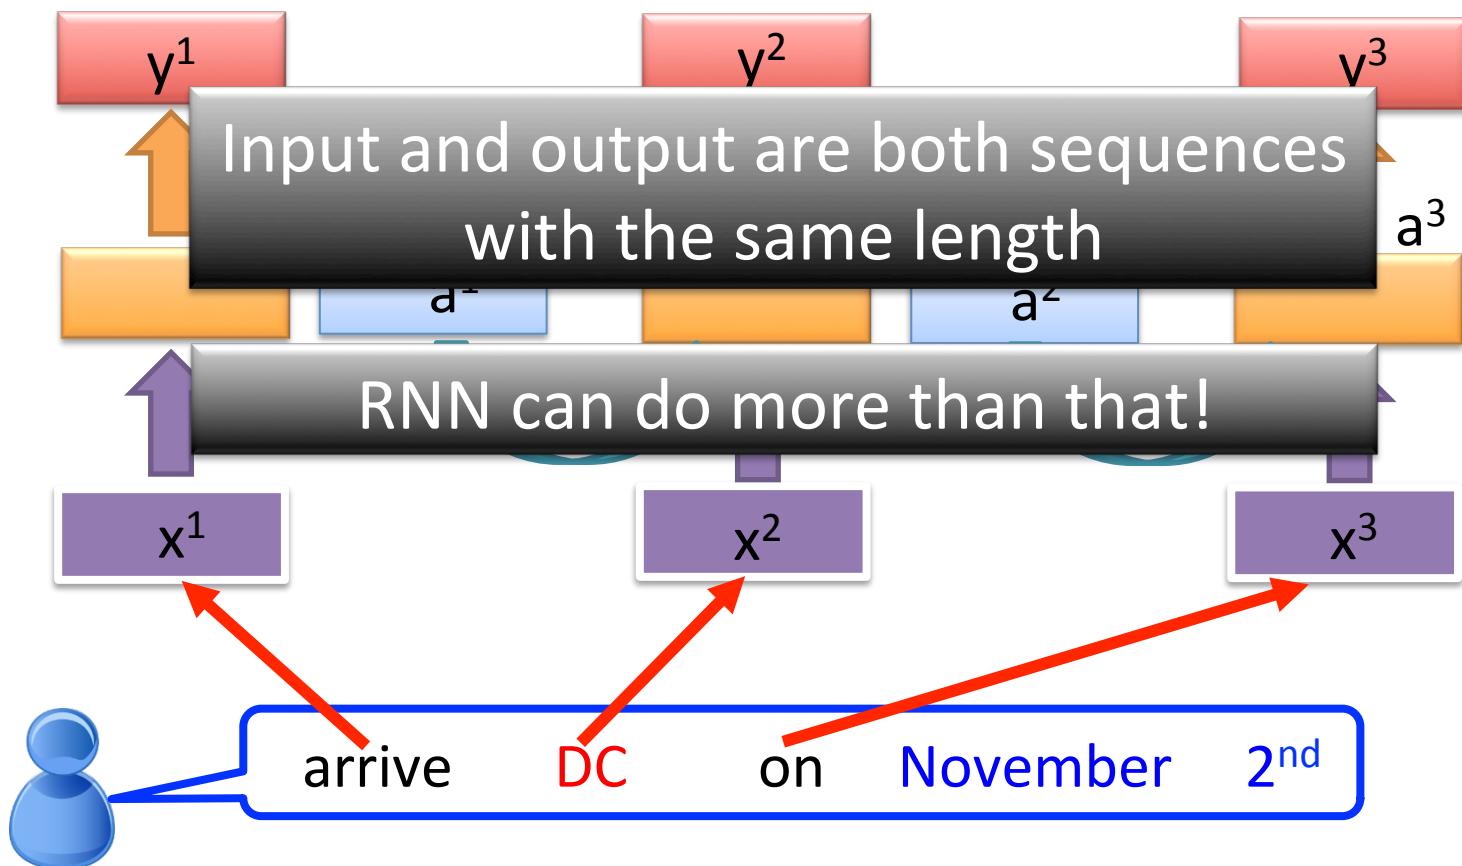
- Outperform or be comparable with LSTM in 4 different tasks

More Applications

Probability of
“arrive” in each slot

Probability of “DC”
in each slot

Probability of
“on” in each slot



Many to one

- Input is a vector sequence, but output is only one vector

Sentiment Analysis

This movie is great.....

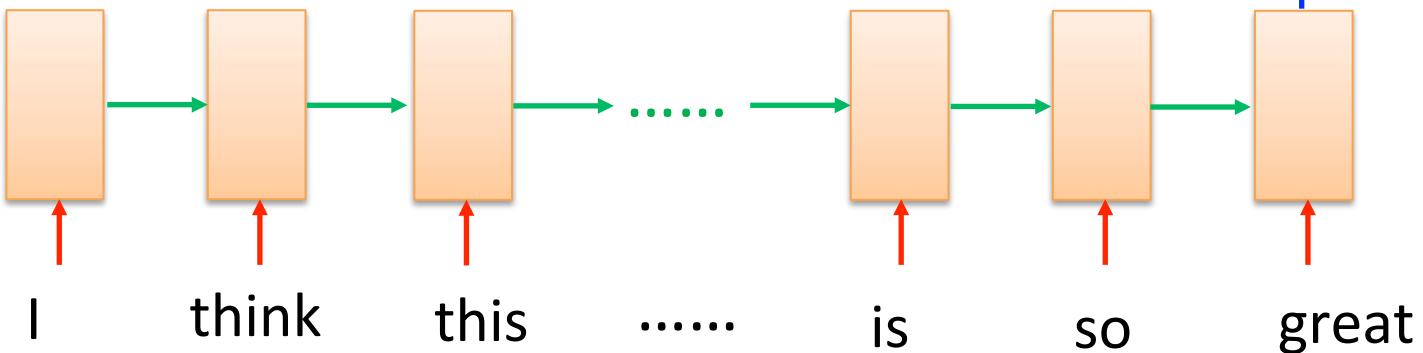
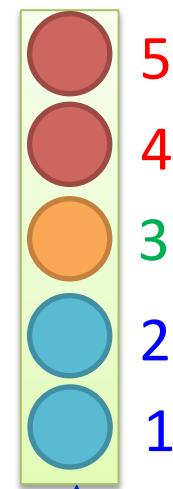
Positive

This movie sucks

Negative

This movie is awesome.....

Positive

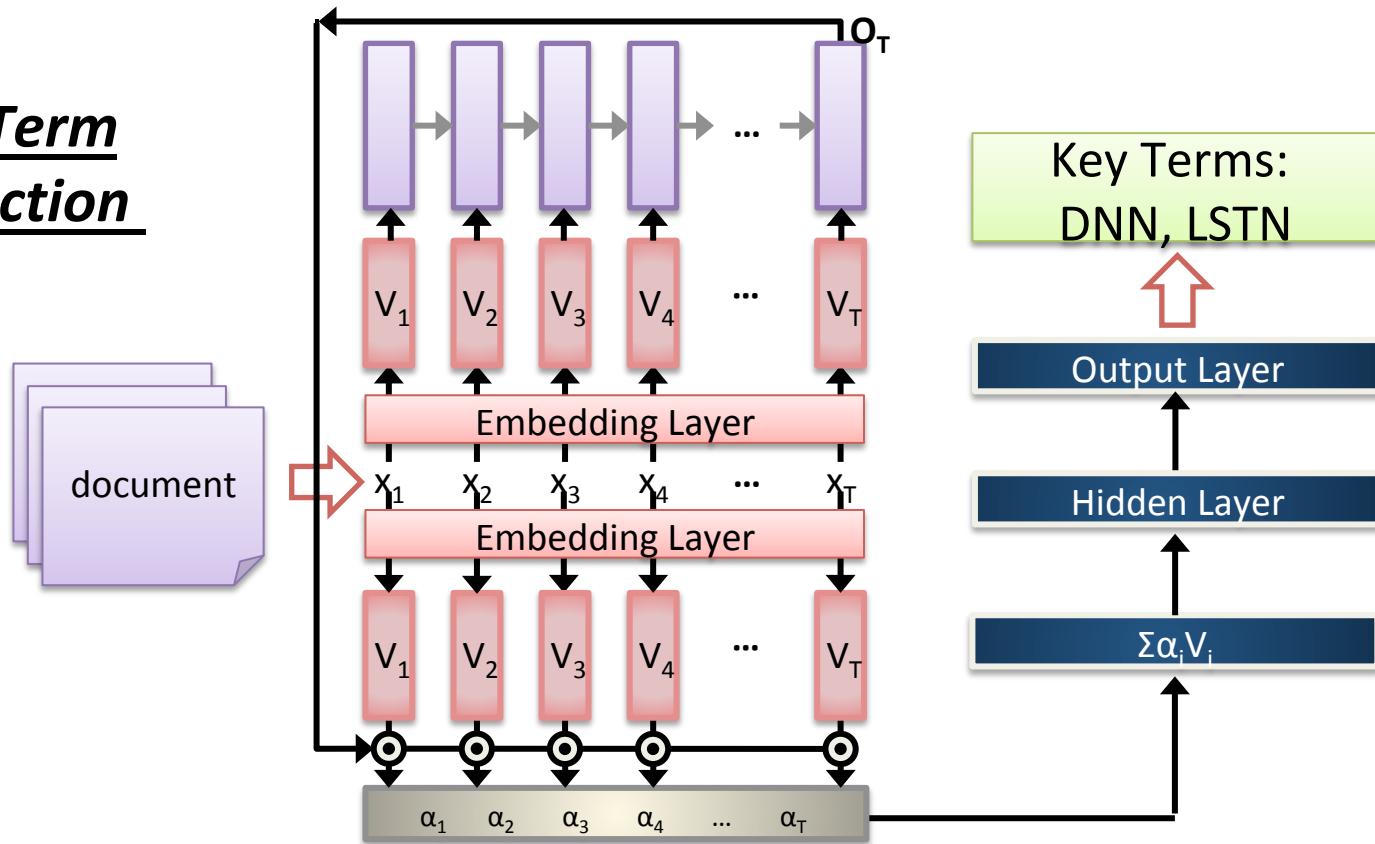


Many to one

[Shen & Lee, Interspeech 16]

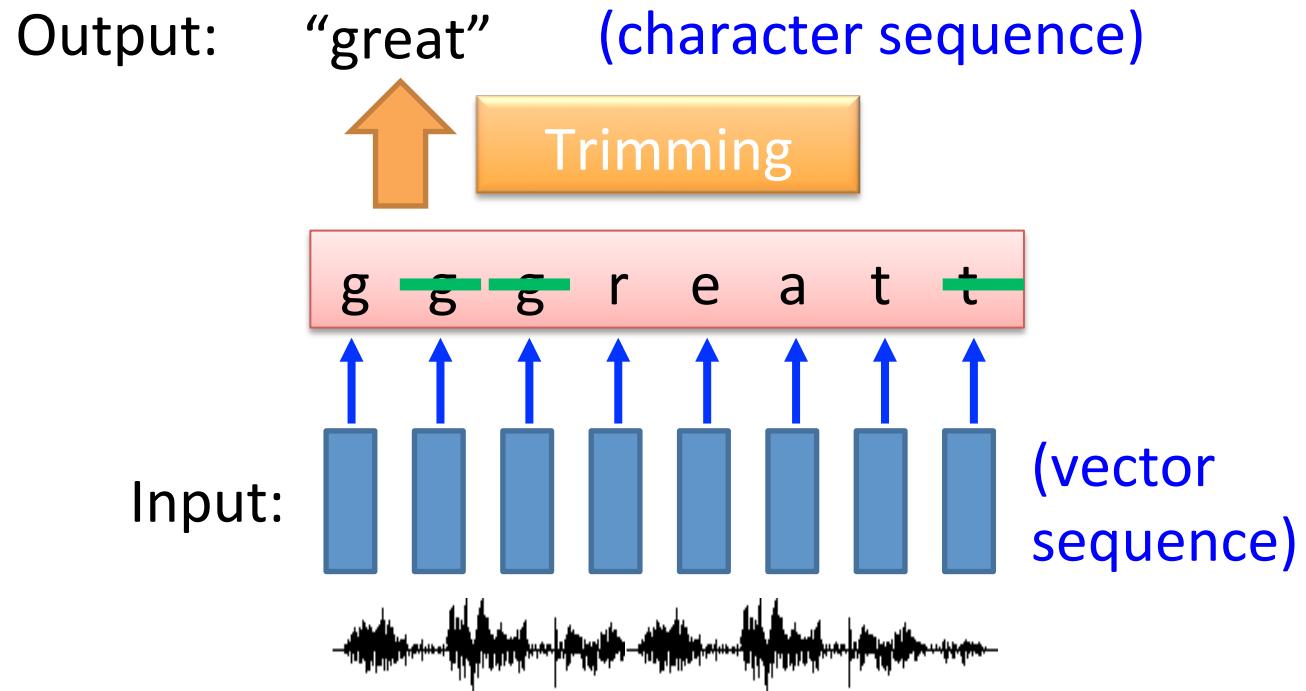
- Input is a vector sequence, but output is only one vector

Key Term Extraction



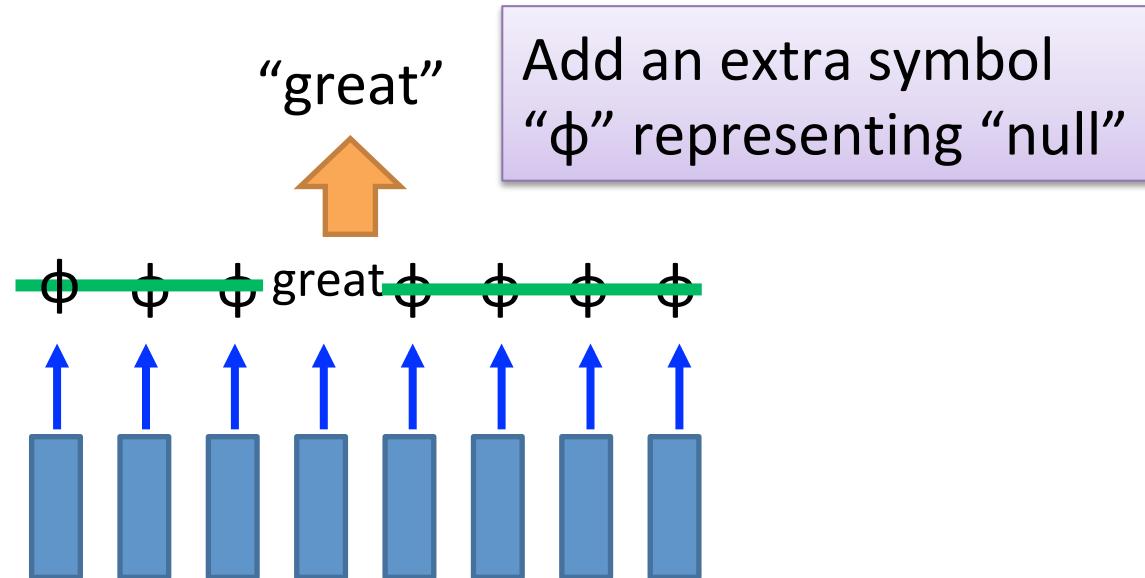
Many to Many (Output is shorter)

- Both input and output are both sequences, **but the output is shorter.**
 - E.g. **Speech Recognition**



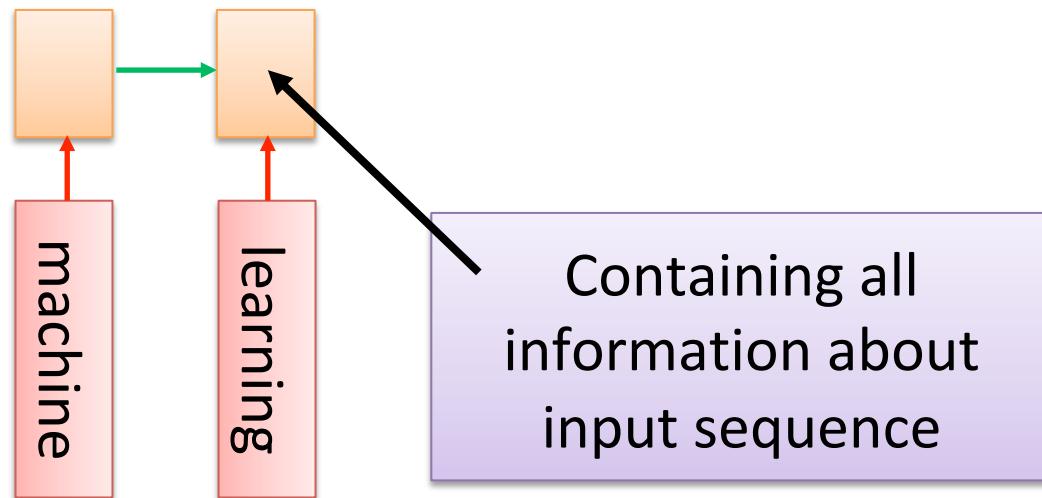
Many to Many (Output is shorter)

- Both input and output are both sequences, **but the output is shorter.**
- Connectionist Temporal Classification (CTC) [Alex Graves, ICML'06] [Alex Graves, ICML'14][Haşim Sak, Interspeech'15][Jie Li, Interspeech'15][Andrew Senior, ASRU'15]



Many to Many (No Limitation)

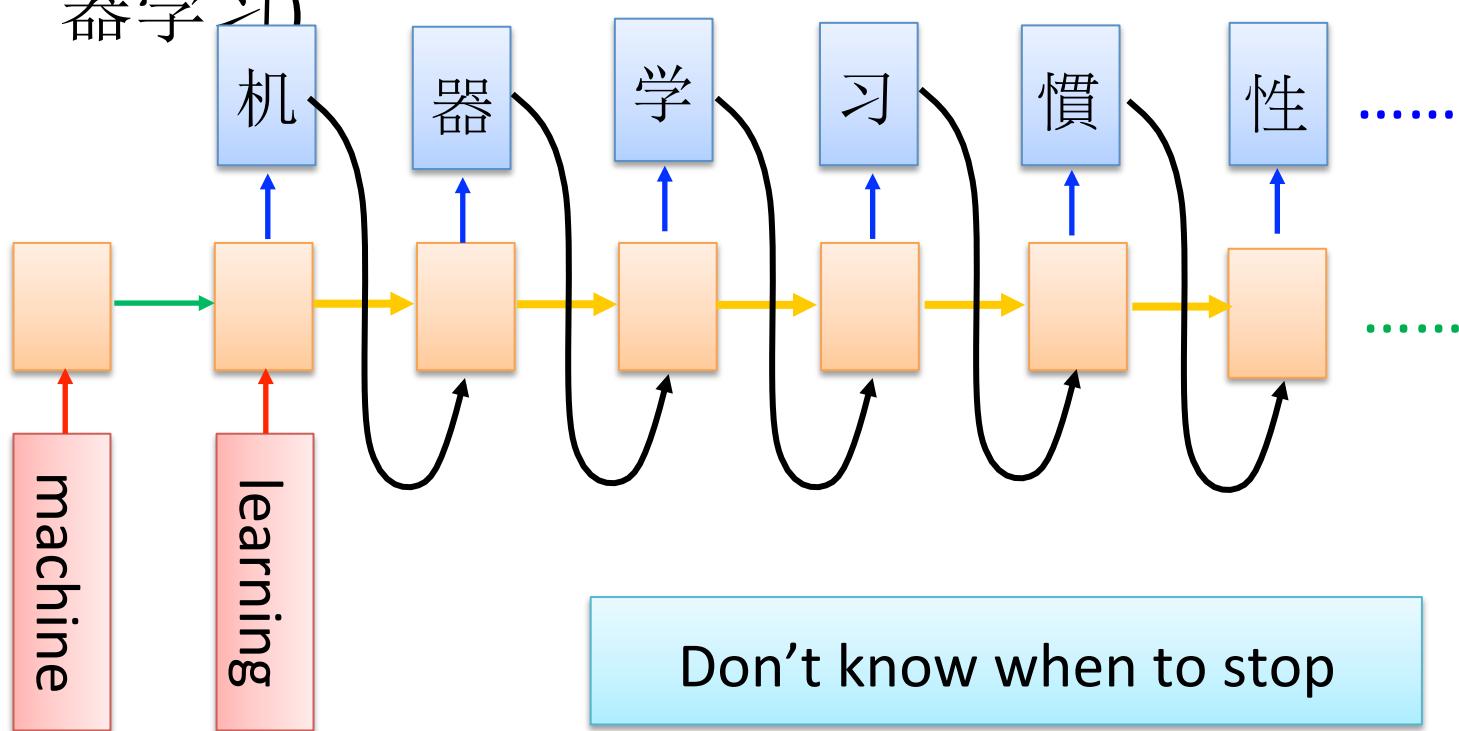
- Both input and output are both sequences *with different lengths*. → *Sequence to sequence learning*
 - E.g. *Machine Translation* (machine learning → 机器学习)



Many to Many (No Limitation)

- Both input and output are both sequences with different lengths. → Sequence to sequence learning

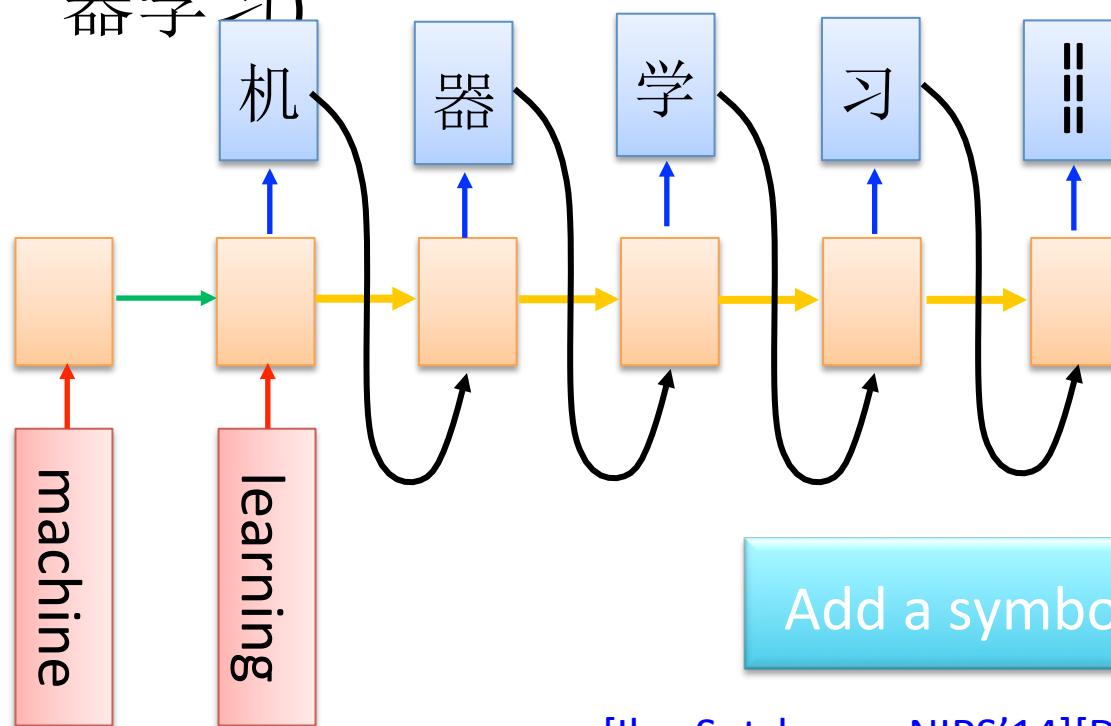
– E.g. Machine Translation (machine learning → 机器学习)



Many to Many (No Limitation)

- Both input and output are both sequences with different lengths. → Sequence to sequence learning

– E.g. Machine Translation (machine learning → 机 器 学 习)

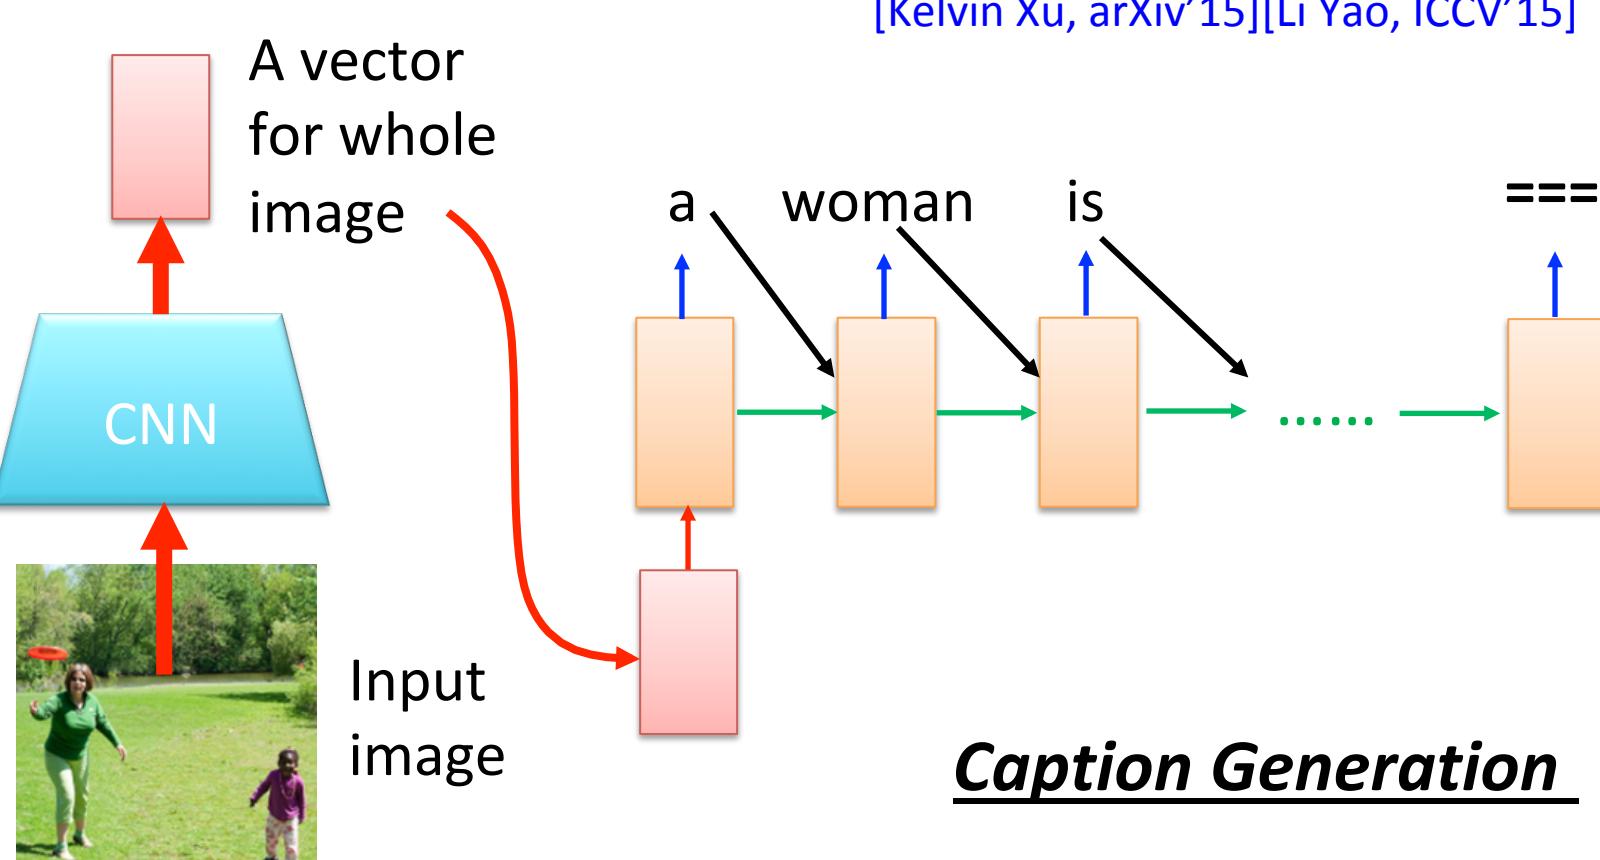


Add a symbol “==” (stop)

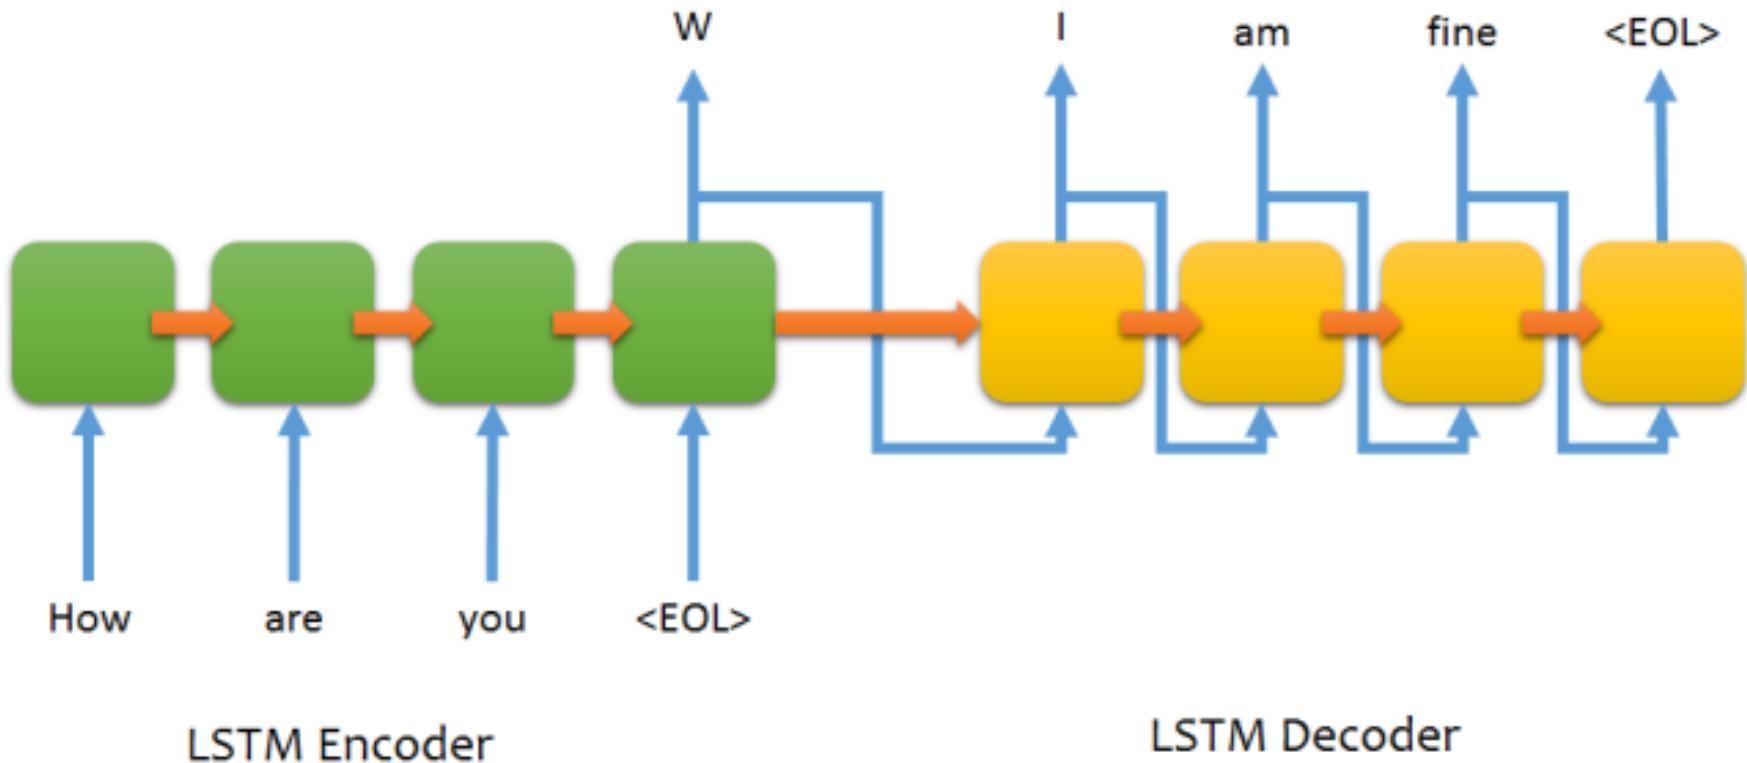
[Ilya Sutskever, NIPS'14][Dzmitry Bahdanau, arXiv'15]

Image Caption Generation

- Input an image, but output a sequence of words



Chat-bot



movie (~40,000 sentences), presidential debate...