

BIG DATA and AI for business

Lecture 4 (02/12, 02/14): Hadoop

Decisions, Operations & Information Technologies
Robert H. Smith School of Business
Spring, 2018



What we'll cover...

- Overview
- Architecture
 - Hadoop distributed file system (HDFS)
- Hands-on
 - Installation and configuration
 - Running Hadoop jobs

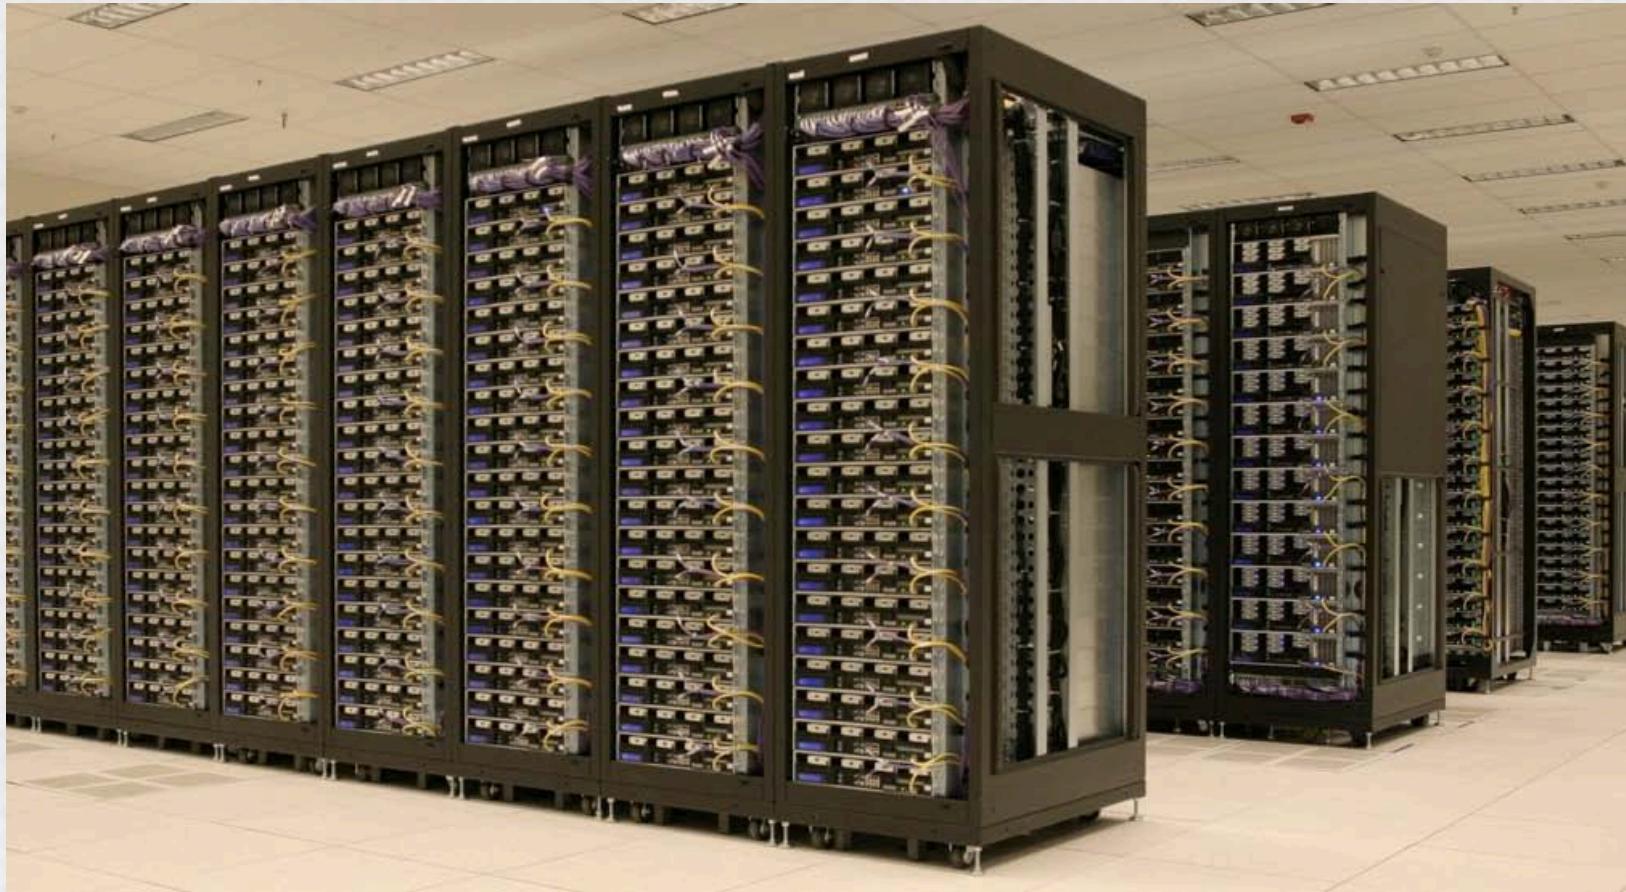
Distributed processing is non-trivial

- How to assign tasks to different workers in an efficient way?
- What happens if tasks fail?
- How do workers exchange results?
- How to synchronize distributed tasks allocated to different workers?

Big data storage is challenging

- Data **volumes** are massive
- Reliability of storing PBs of data is challenging
- All kinds of **failures**: Disk/Hardware/Network Failures
- Probability of failures simply increase with the number of machines ...

One popular solution: Hadoop



Hadoop Cluster at Yahoo!

Hadoop offers

- Redundant, Fault-tolerant data storage
- Parallel computation framework
- Job coordination



Hadoop offers

- Redundant, Fault-tolerant data storage
- Parallel computation framework
- Job coordination



Programmers

No longer need to
worry about



Q: Where file is
located?

Q: How to handle
failures & data lost?

Q: How to divide
computation?

Q: How to program
for scaling?

A little history on Hadoop

- Hadoop is an open-source implementation based on Google File System (GFS) and MapReduce from Google
- Hadoop was created by Doug Cutting and Mike Cafarella in 2005
- Hadoop was donated to Apache in 2006



Who uses Hadoop?

Social



User Tracking & Engagement



Homeland Security



eCommerce

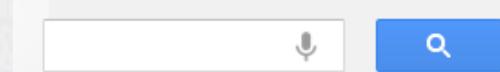


Financial Services



Real Time Search

Google



Who uses Hadoop?

2007

YAHOO!



last.fm

2008

Google™ ablegrape



ImageShack
online media hosting

Cascading

IBM®

facebook

ENORMO Every property. Everywhere.



krugle

rackspace.
HOSTING

Lookery Control freaks welcome

The New York Times joost

Zvents Discover Things To Do

FORMATION SCIENCES INSTITUTE

News Corporation

Cornell University Computing and Information Science

Visible MEASURES

LOTAME Locate, Target & Message with Social Media

NetSeer

parc®
Palo Alto Research Center

SECURITY ENHANCED DOMAIN NAME SYSTEM

veoh™

2009

AOL

cloudera

deepdyve

cooliris

eyearlike

TEXTMAP
THE ENTITY SEARCH ENGINE

PSG College of Technology

iterend

tailsweep

hulu™

RapLeaf

USCIMS

Ning quaxtcast

amazon web services

pressflip

detikSearch

WorldLingo

Systems@ETH Zürich

VK SOLUTIONS
Global Solutions Provider

TARAGANA
Innovation • Quality • Simplicity

HOSTING HABITAT

HOLA
SERVICES

Terrier

adknowledge

stampede

beta

2010

SAMSUNG

rubicon

BERKELEY LAB
LAWRENCE BERKELEY NATIONAL LABORATORY

VISIBLE
TECHNOLOGIES

APOLLO GROUP

ADSDAQ™

rackspace.
HOSTING

RapLeaf

wordnik

meBIGEN
Technological Revolution

comSCORE

trulia®
real estate search

Accela
COMMUNICATIONS

Forward3D

Linked in

Microsoft®

Infochimps

Pharm2Pharm

ADMELD

gumgum®

BrainPad

Pronux

The DataGraph Bbg

NETFLIX

mobileanalytics.tv

markt24.de

twitter™

media6degrees

BEEBLER

SLC Security
When Experience Matters...

eBay

Why HDFS?

- **Problem 1:** Data is too big to store on one machine.
- **HDFS:** Store the data on multiple machines!

Why HDFS?

- **Problem 2:** Very high end machines are too expensive
- **HDFS:** Run on commodity hardware!

Why HDFS?

- **Problem 3:** Commodity hardware can fail
- **HDFS:** Software is intelligent enough to handle hardware failure!

Why HDFS?

- **Problem 4:** What happens to the data if the machine storing the data fails?
- **HDFS:** Replicate the data!

Why HDFS?

- **Problem 5:** How can distributed machines organize the data in a coordinated way?
- **HDFS: Master-Slave Architecture!**

Another reason why Hadoop

- **Scan** 100TB datasets on 1000 node cluster
 - Remote storage @ 10MB/s = 165 mins
 - Local storage @ 50-200MB/s = 33-8 mins
- Moving computation is more efficient than moving data
- Need fault tolerant store with reasonable availability guarantees
 - Handle hardware faults transparently

Hadoop goals

- **Scalable**: Petabytes (10^{15} Bytes) of data on thousands of nodes
- **Economical**: Commodity components only
- **Reliable**: fault tolerance

Hadoop big picture

HADOOP 1.0

MapReduce
(cluster resource management
& data processing)

HDFS
(redundant, reliable storage)

HADOOP 2.0

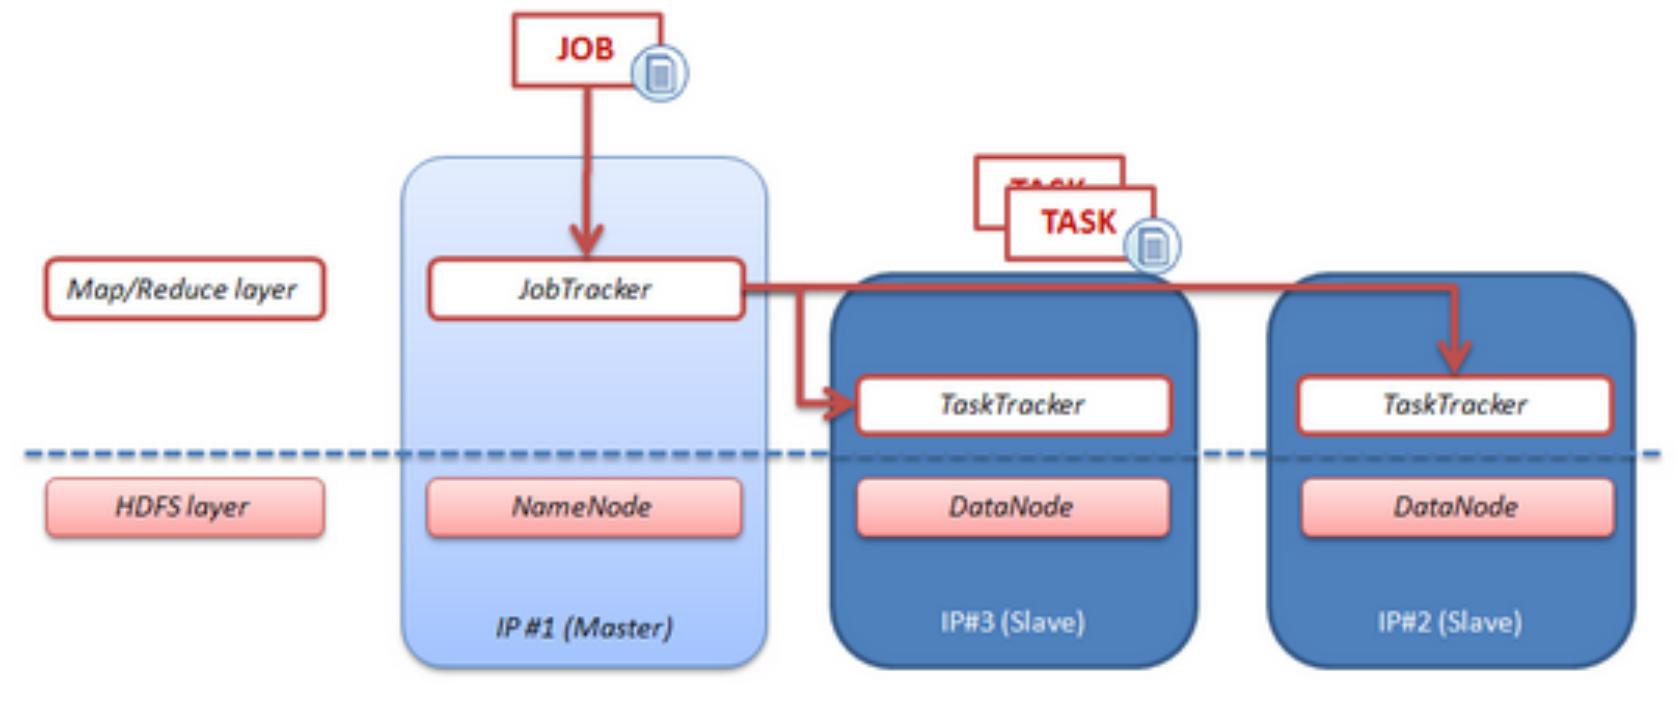
MapReduce
(data processing)

Others
(data processing)

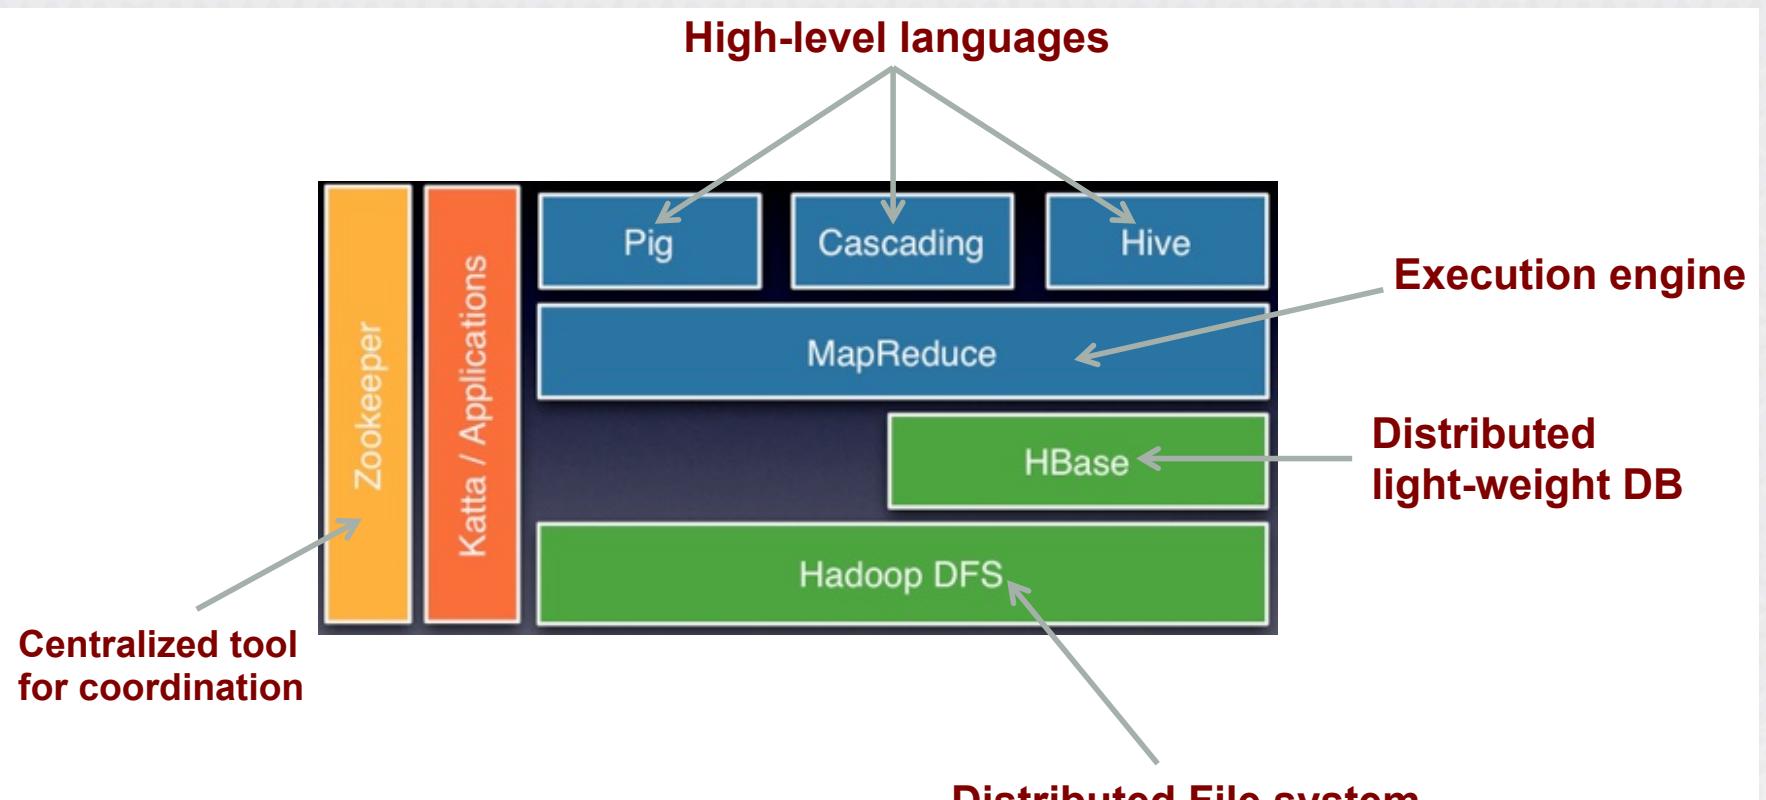
YARN
(cluster resource management)

HDFS
(redundant, reliable storage)

High-level architecture of Hadoop

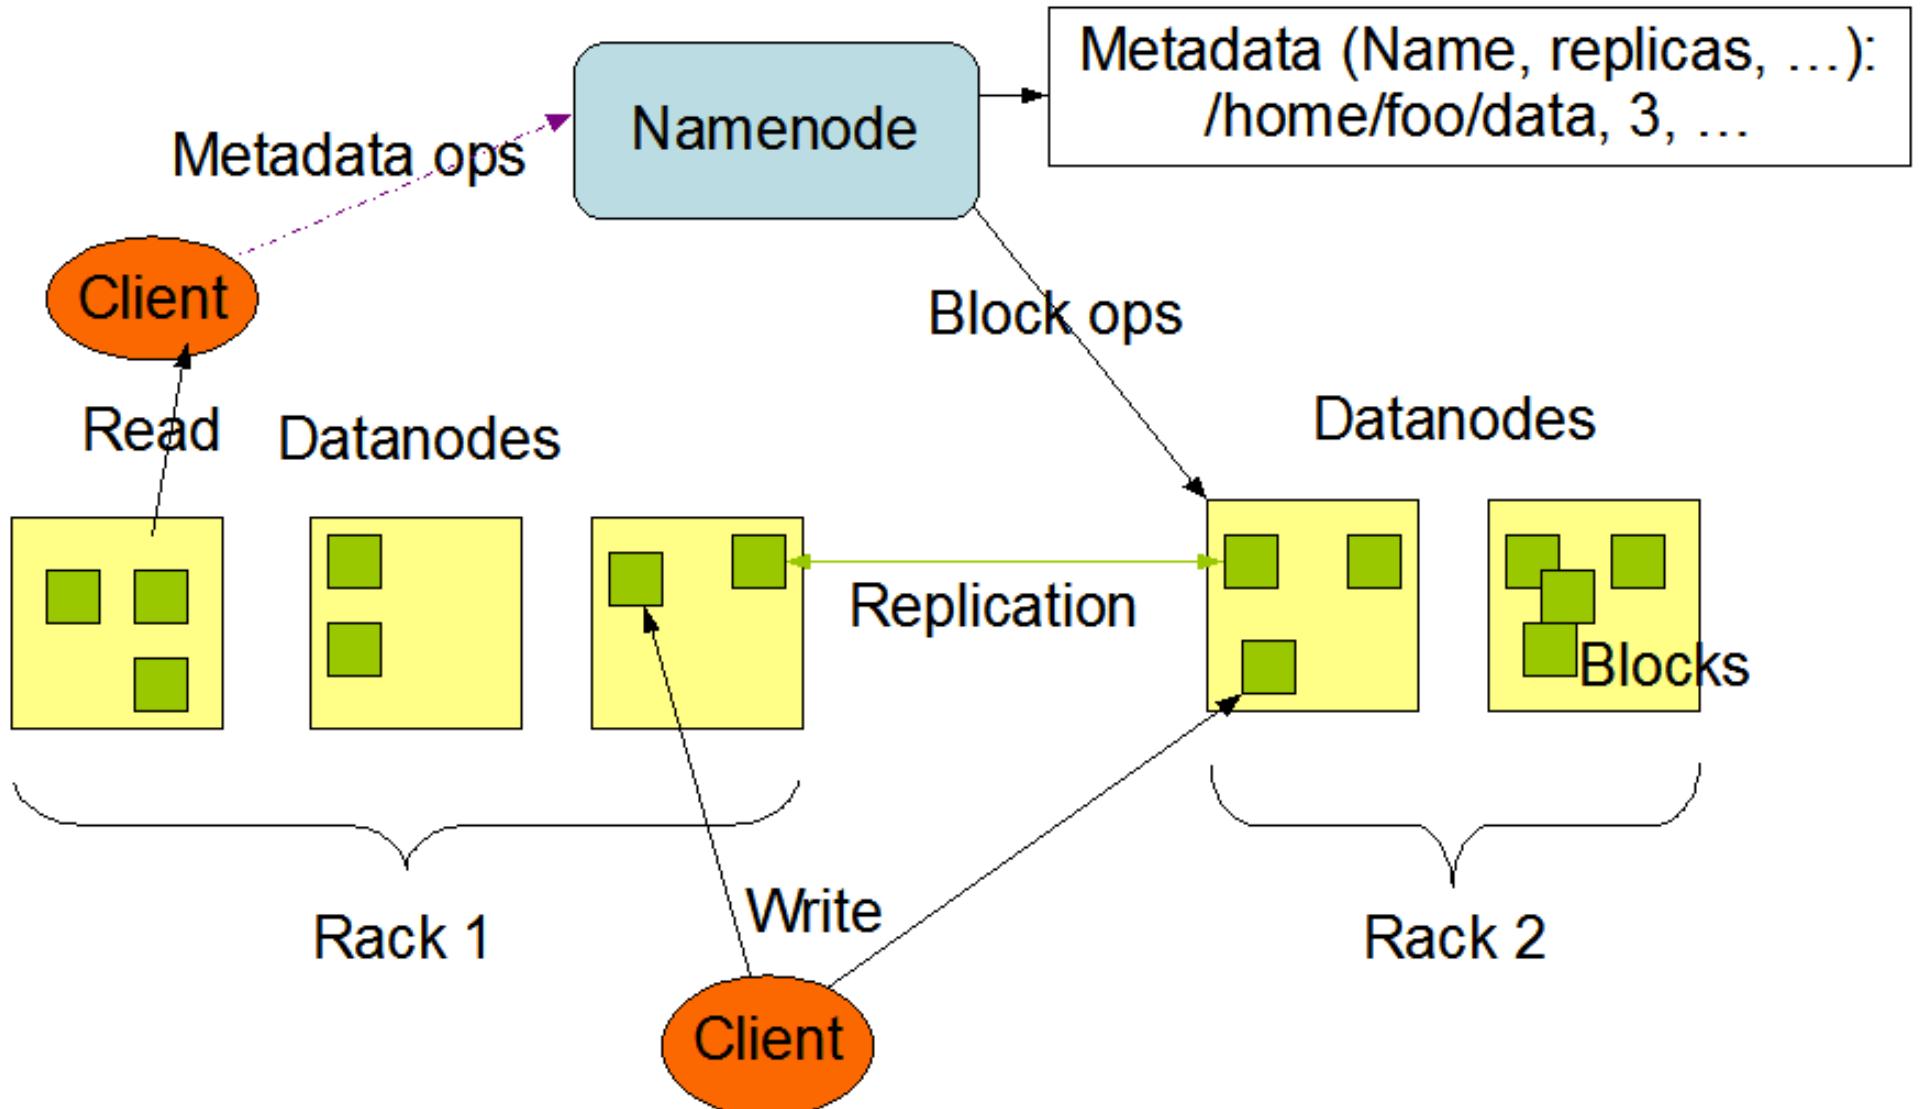


Hadoop big picture



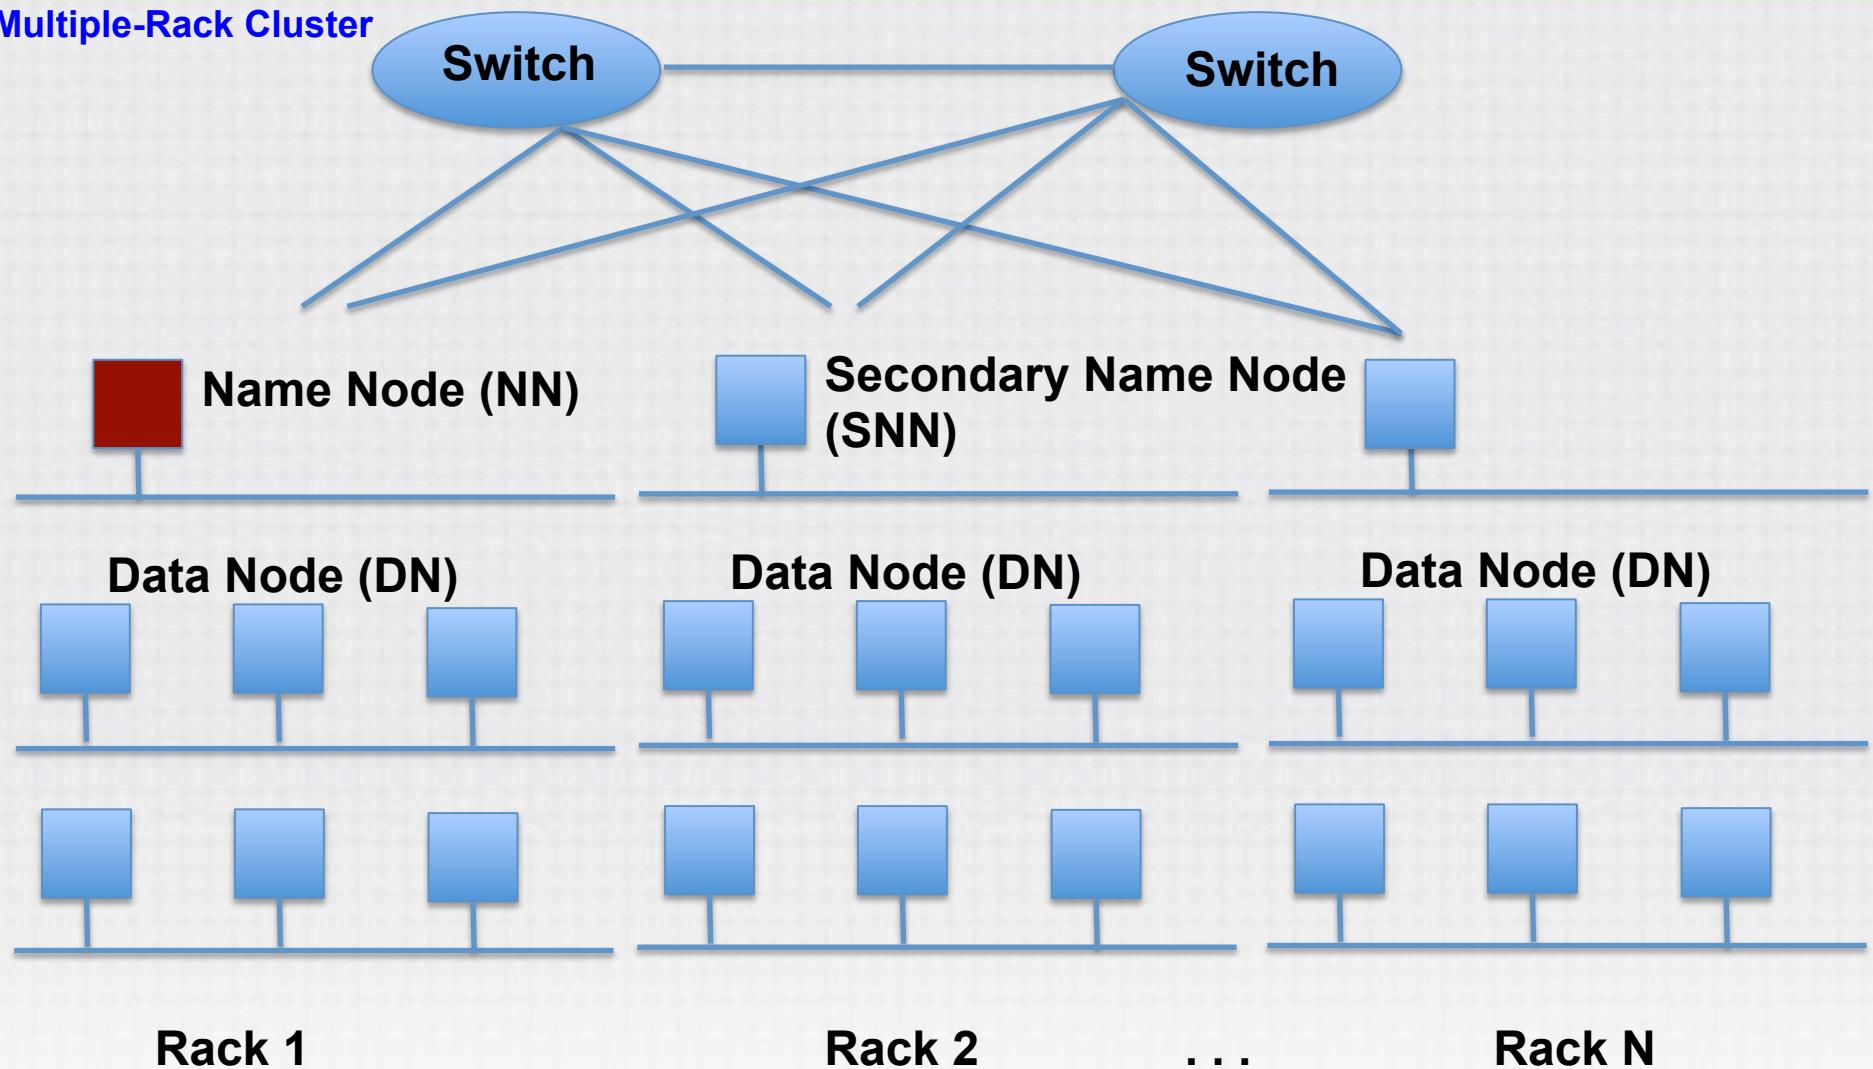
HDFS + MapReduce are enough to have things working

HDFS Architecture



HDFS architecture

Multiple-Rack Cluster



- **Master-Slave** architecture
- Single NameNode
 - Sometimes a backup: secondary NameNode
- Many (Thousands) DataNodes
- Files are split into fixed sized blocks and stored on data nodes (**default: 64MB**)
- Data blocks are replicated for fault tolerance and fast access (**default: 3**)

HDFS - Master (NameNode)

- Manages file system (FS) namespace
- File metadata
- Mapping file to list of blocks
- Authorization & Authentication
- Mapping of datanode to list of blocks
- Monitor datanode health
- Replicate missing blocks
- Keeps ALL namespace in memory

HDFS - Slave (DataNode)

- Handle block storage on multiple volumes & block integrity
- Clients access the blocks directly from data nodes
- Periodically send heartbeats and block reports to NameNode
- Blocks are stored as underlying OS's files

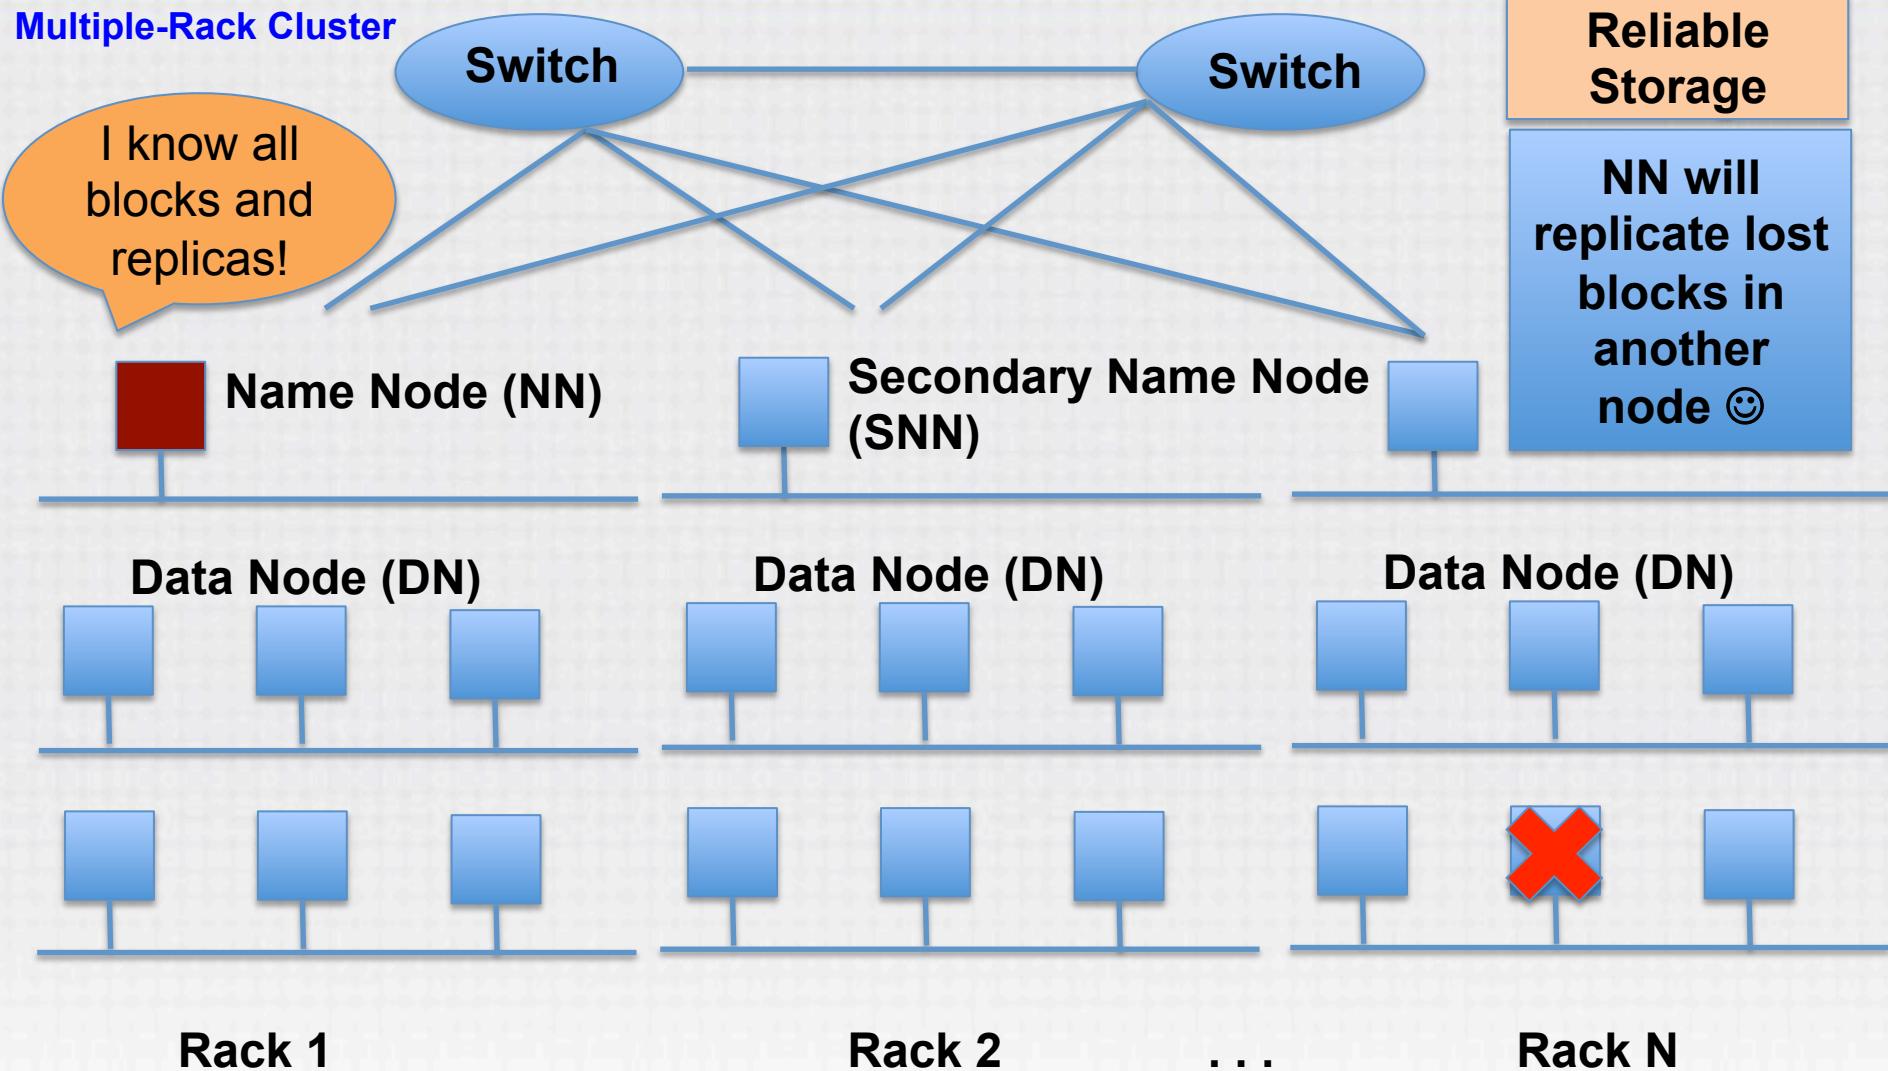
HDFS Name Node

Name Node**Snapshot of FS****Edit log: record
changes to FS**

Filename	Replication factor	Block ID
File 1	3	[1, 2, 3]
File 2	2	[4, 5, 6]
File 3	1	[7, 8]

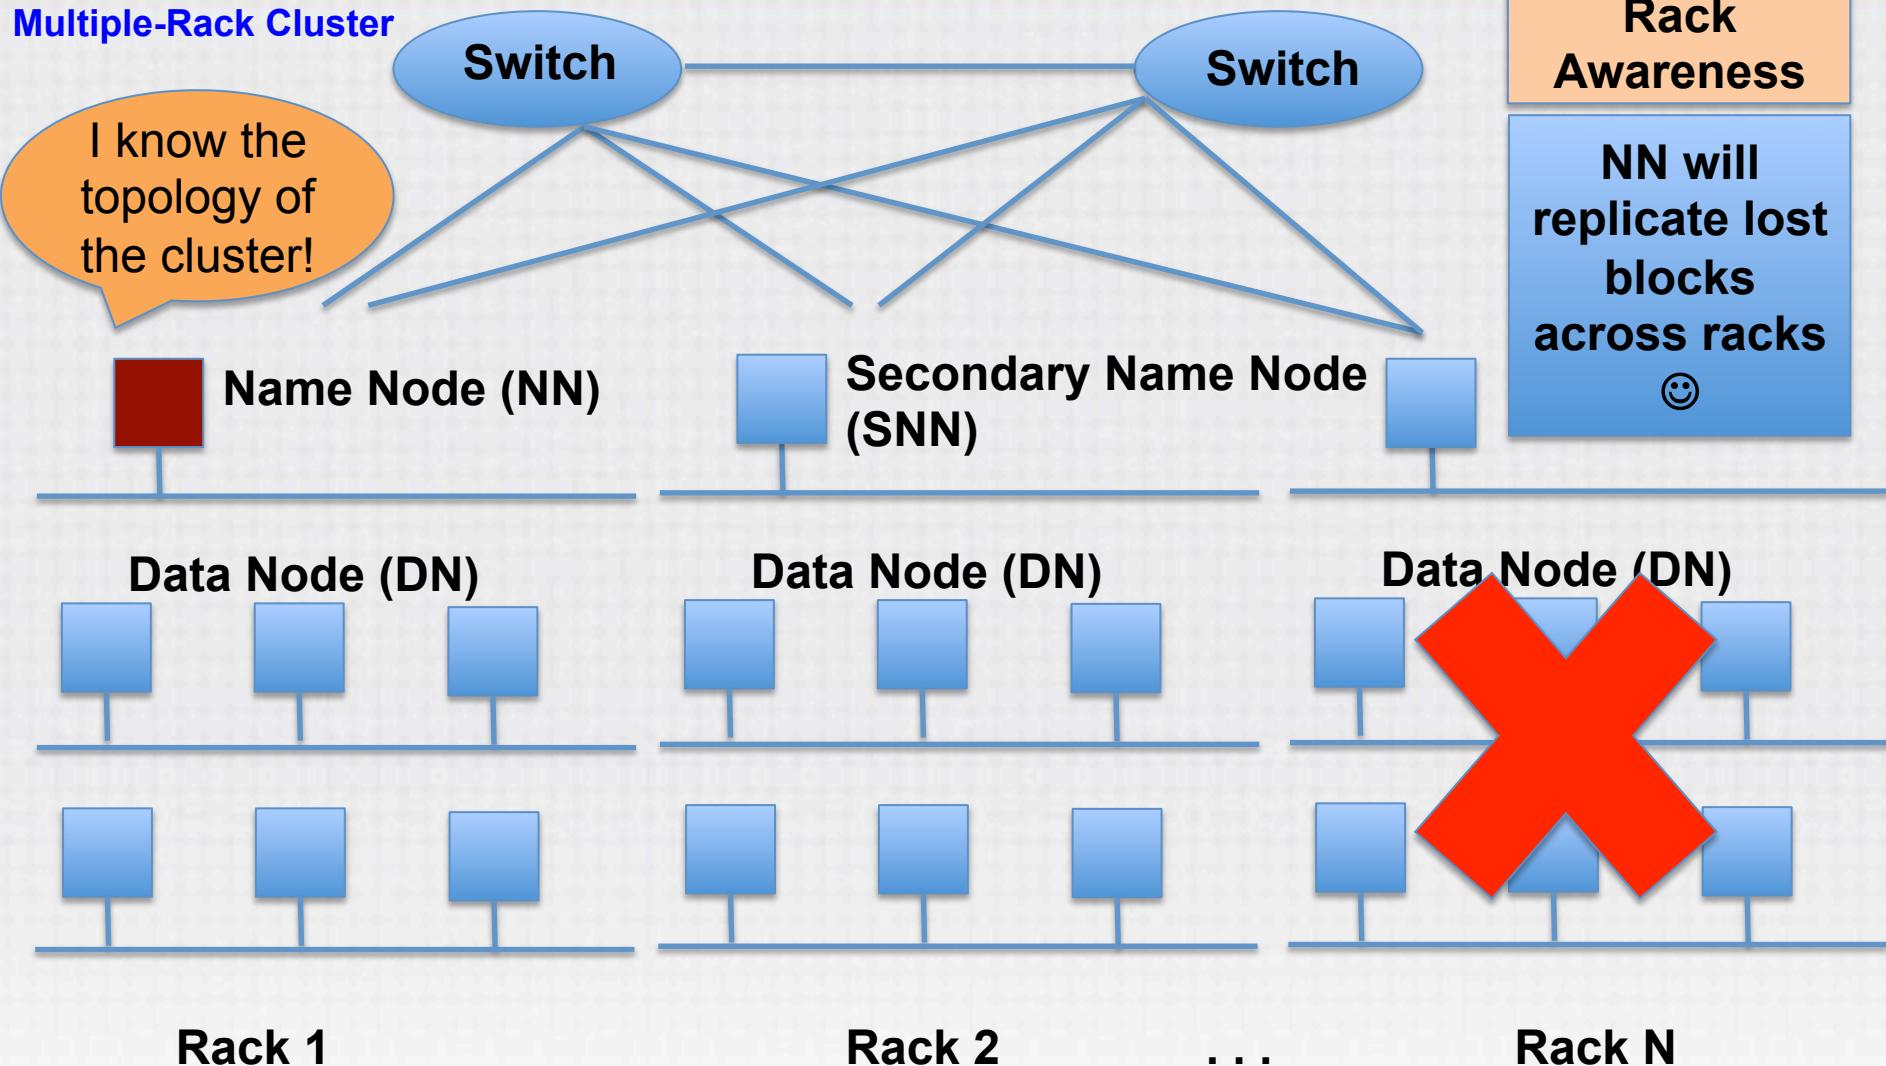
Data Nodes1, 2, 5, 7,
4, 31, 5, 3,
2, 8, 61, 4, 3,
2, 6

HDFS architecture



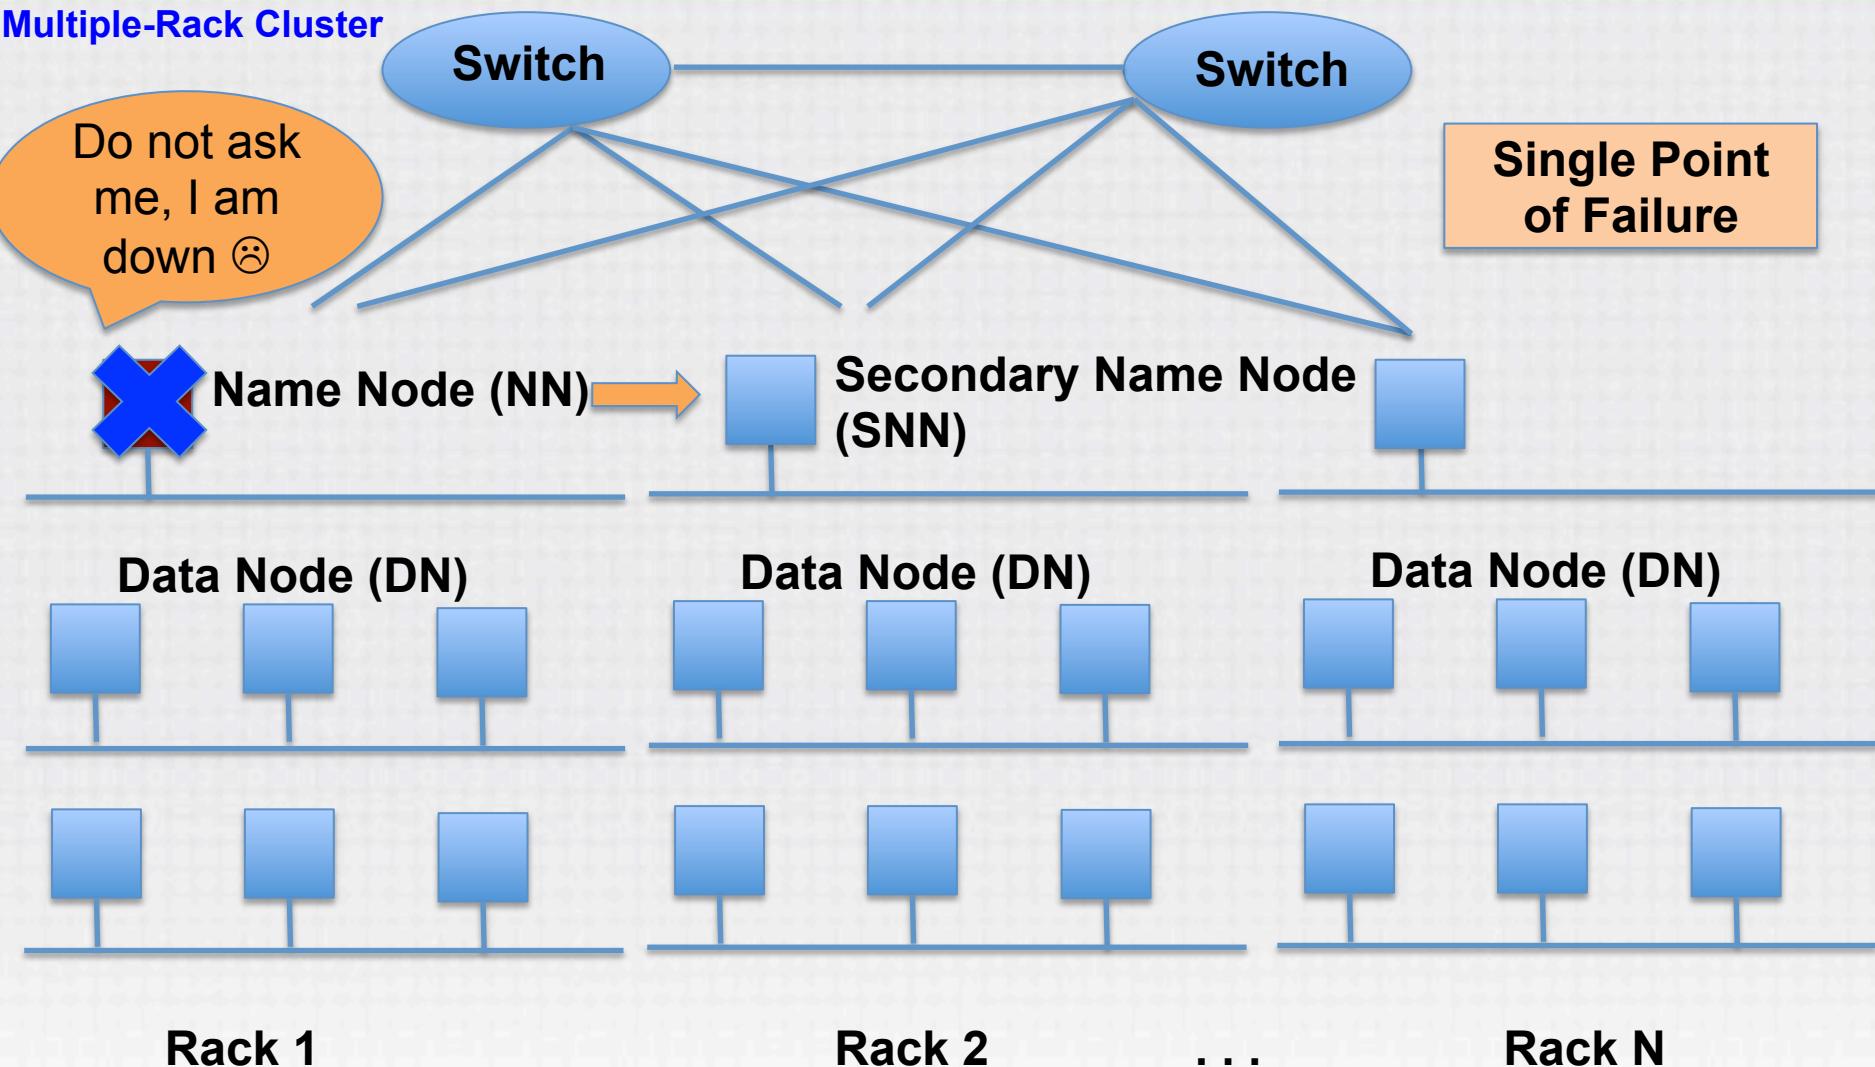
HDFS architecture

Multiple-Rack Cluster



HDFS architecture

Multiple-Rack Cluster

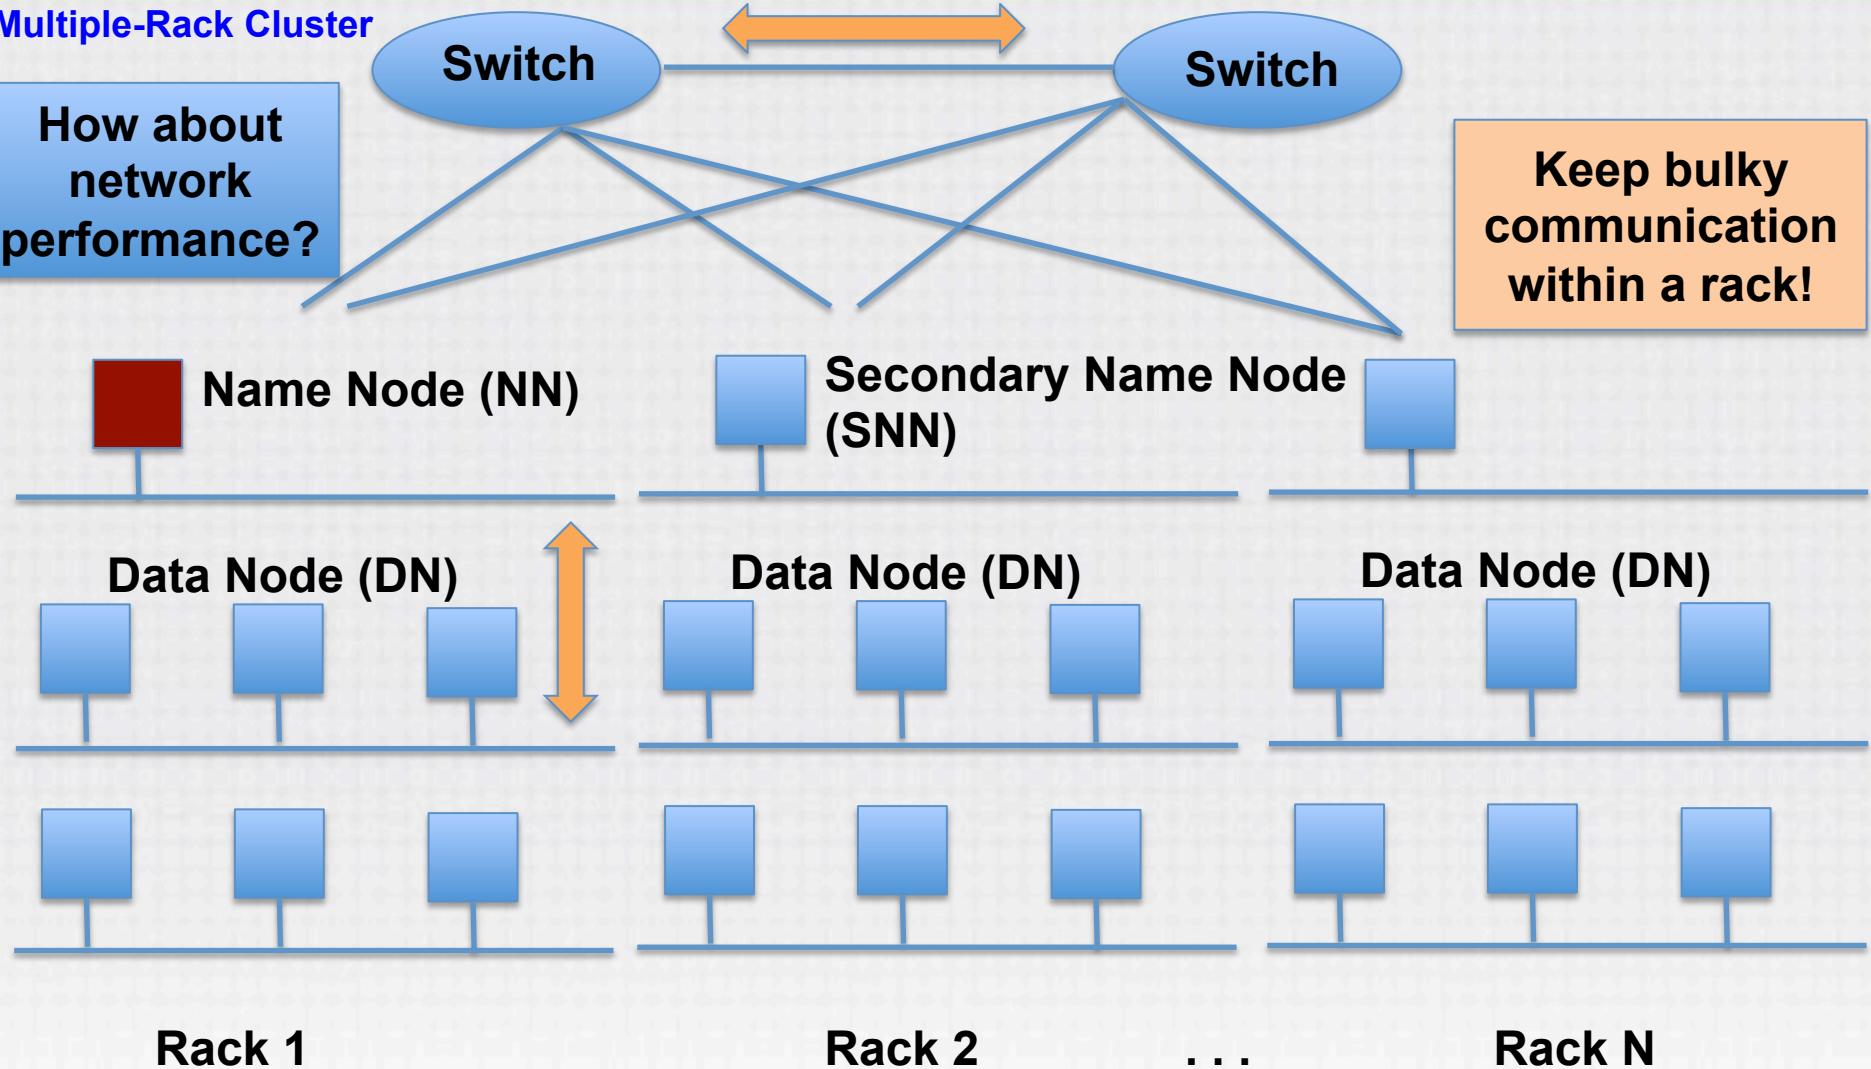


HDFS architecture

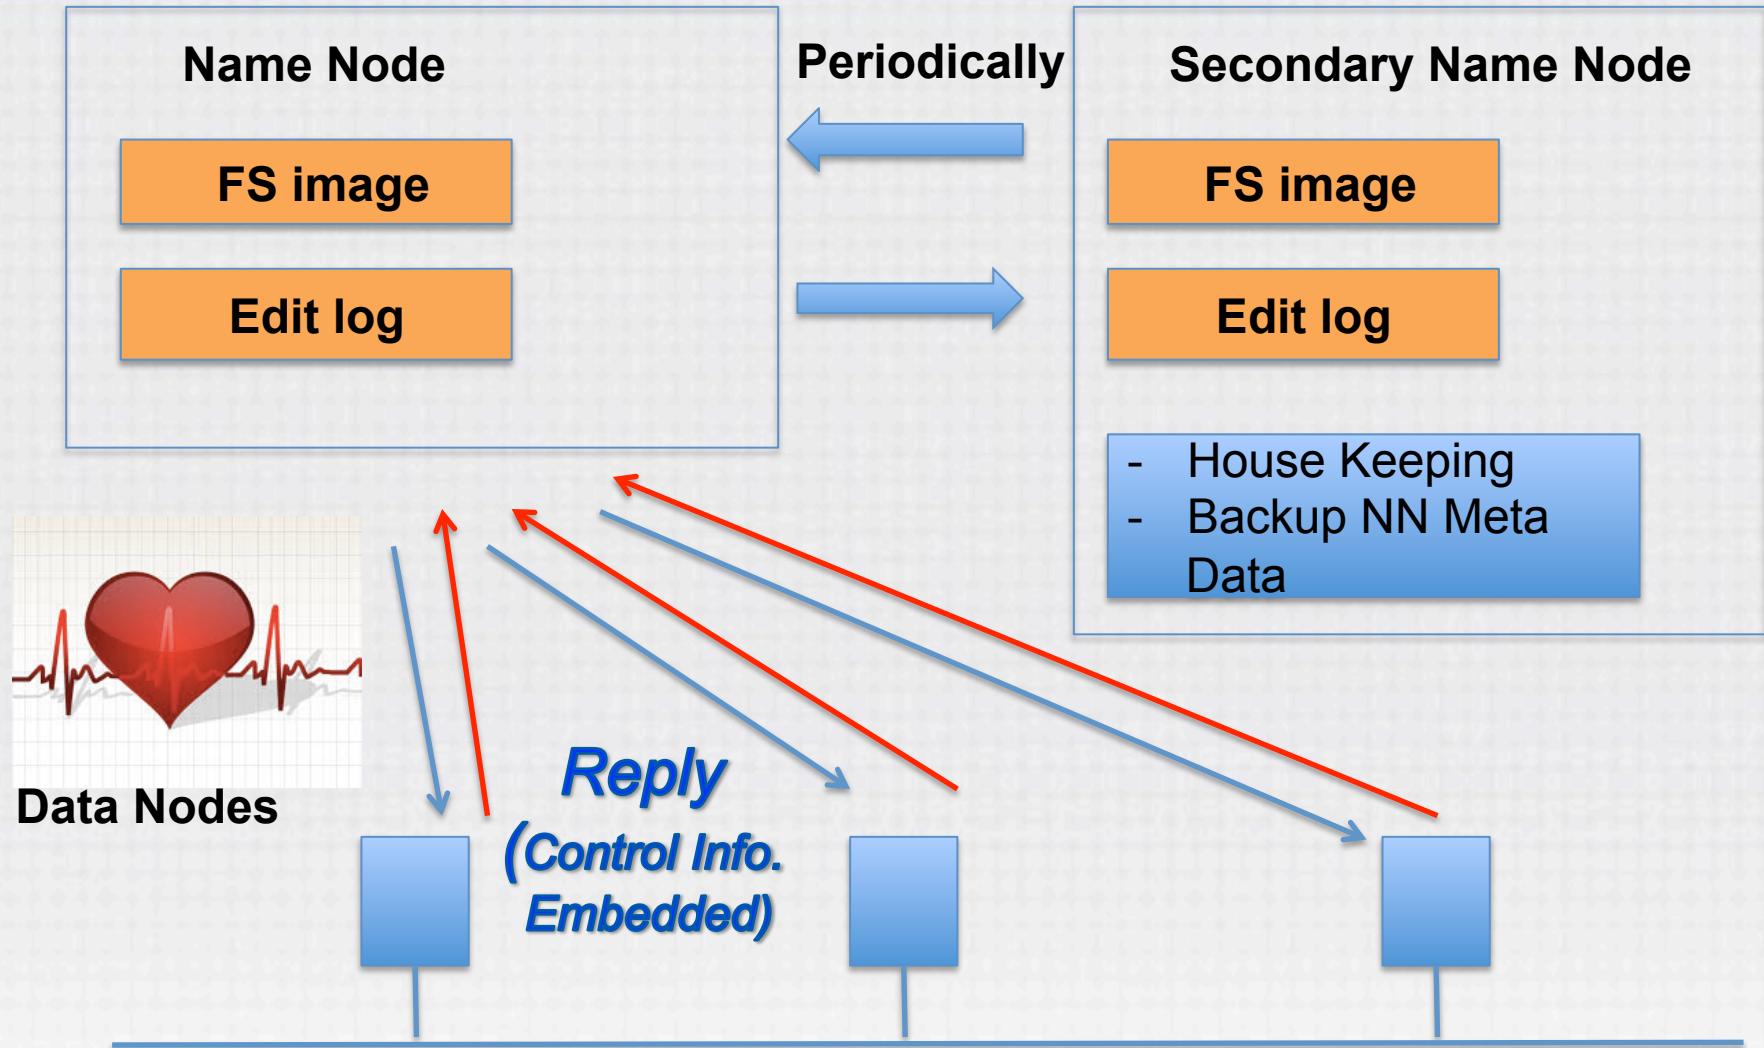
Multiple-Rack Cluster

How about
network
performance?

Keep bulky
communication
within a rack!



HDFS Inside: Name Node



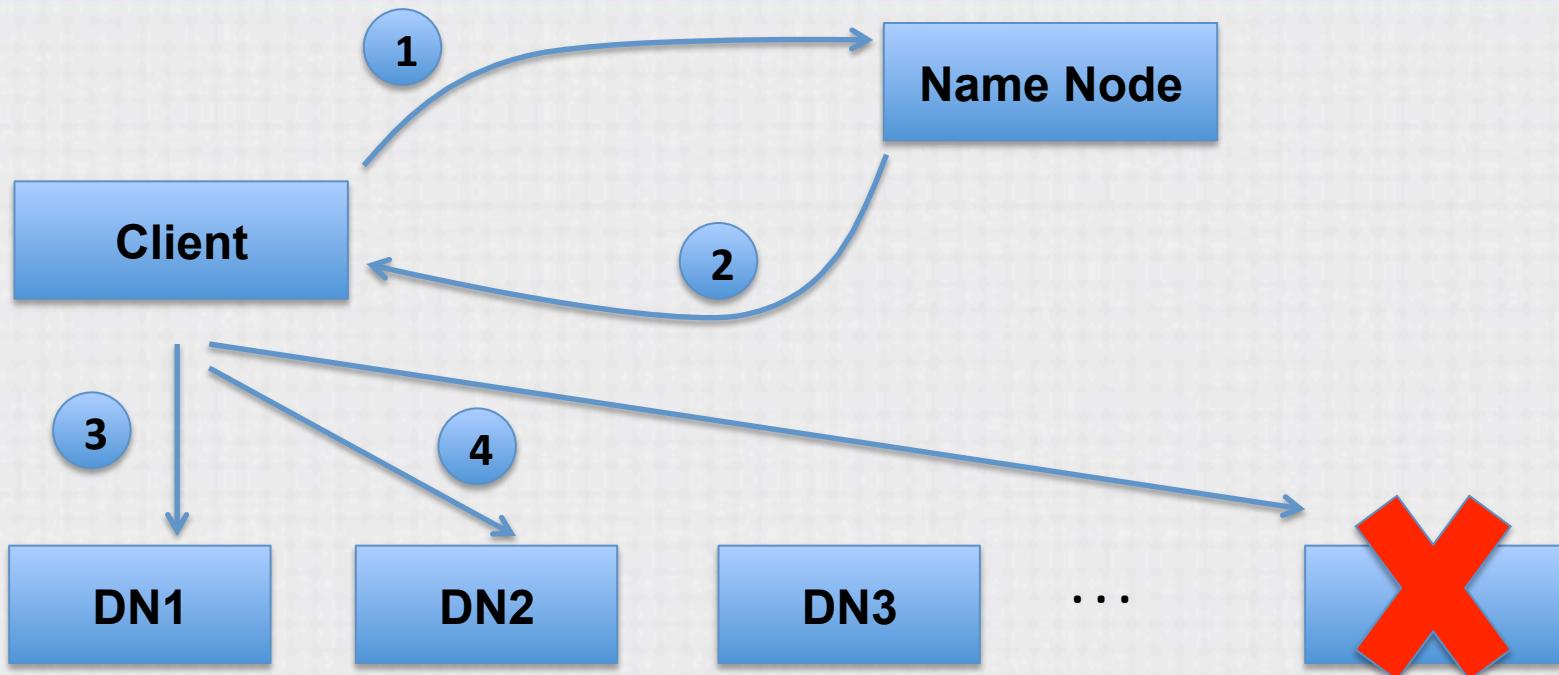
HDFS Inside: Blocks

- Q: Why do we need the abstraction “Blocks” in addition to “Files”?
- Reasons:
 - File can be larger than a single disk
 - Block is of fixed size, easy to manage and manipulate
 - Easy to replicate and do more fine grained load balancing

HDFS Inside: Blocks

- HDFS Block size is by default **64 MB**, why it is much larger than regular file system block?
- Reasons:
 - Minimize overhead: disk seek time is almost constant

HDFS Inside: Read



1. Client connects to NN to read data
2. NN tells client where to find the data blocks
3. Client reads blocks directly from data nodes (without going through NN)
4. In case of node failures, client connects to another node that serves the missing block

HDFS Inside: Read

- Q: Why does HDFS choose such a design for read? Why not ask client to read blocks through NN?
- Reasons:
 - Prevent NN from being the bottleneck of the cluster
 - Allow HDFS to scale to large number of concurrent clients
 - Spread the data traffic across the cluster

HDFS Inside: Read

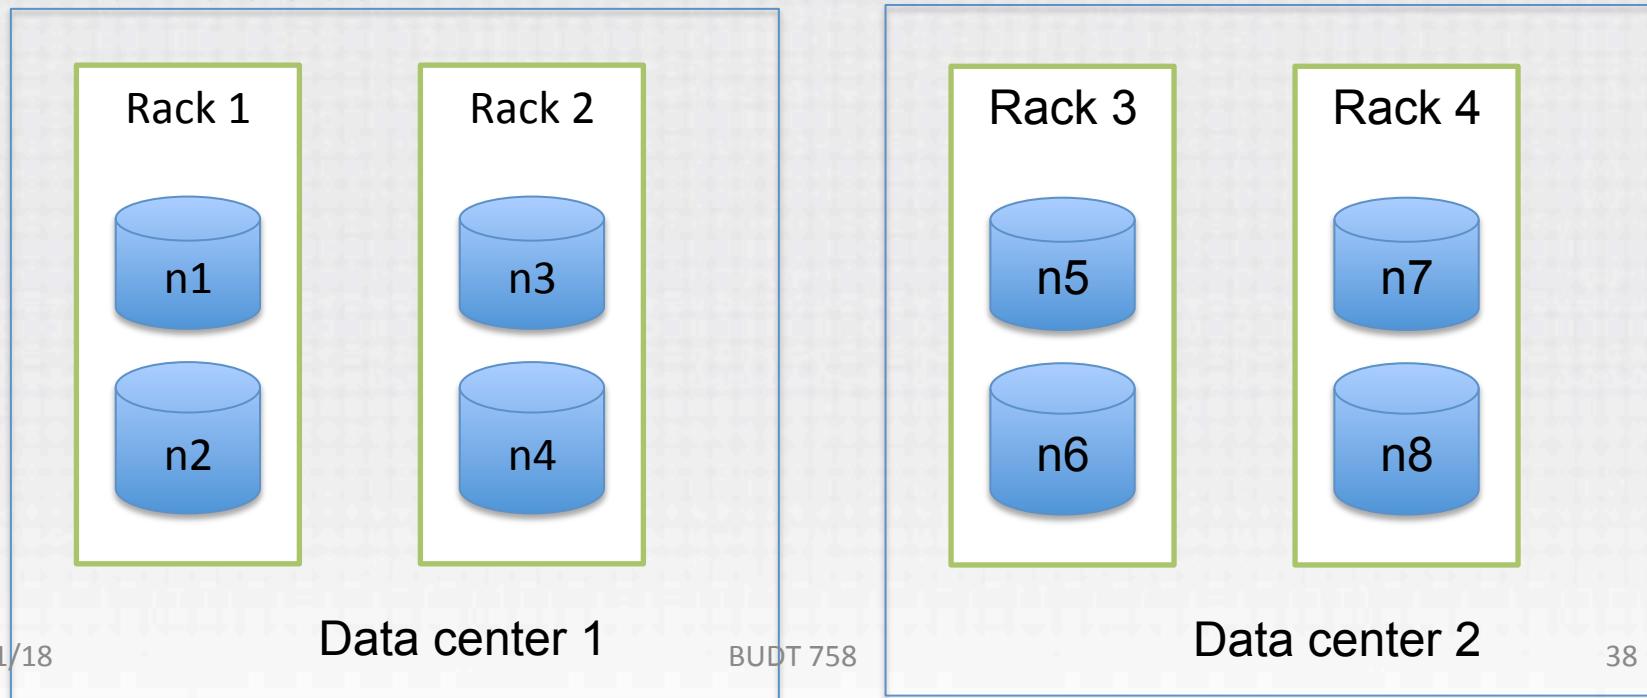
- Q: Given multiple replicas of the same block, how does NN decide which replica the client should read?
- **HDFS Solution:**
 - Rack awareness based on network topology

HDFS network topology

- The critical resource in HDFS is **bandwidth**, distance is defined based on that
 - Measuring bandwidths between any pair of nodes is too complex and **does not scale**
 - **Basic Idea:**
 - Processes on the same node
 - Different nodes on the same rack
 - Nodes on different racks in the same data center (cluster)
 - Nodes in different data centers
- 
- Bandwidth becomes less**

HDFS network topology

- HDFS takes a simple approach:
 - See the network as a tree
 - **Distance between two nodes is the sum of their distances to their closest common ancestor**



HDFS network topology

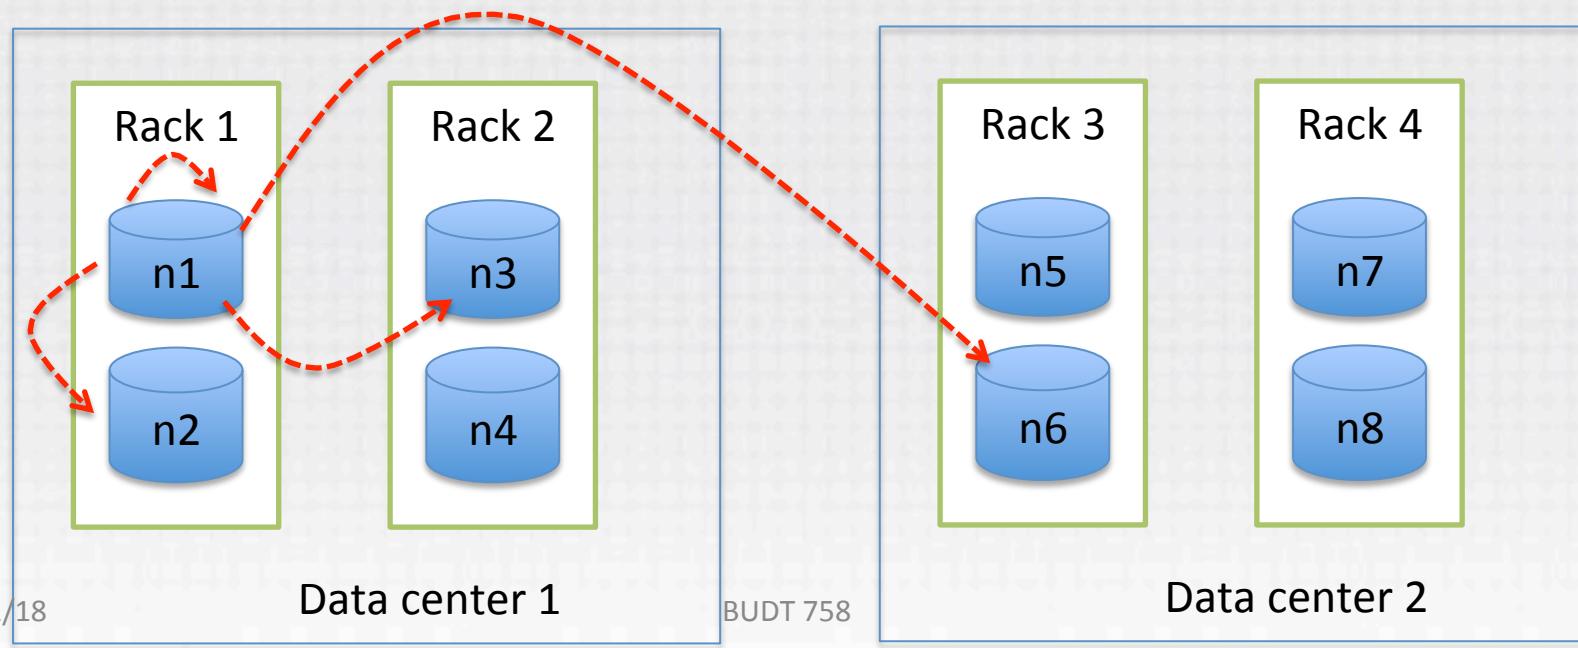
- What are the distance of the following pairs:

$\text{Dist}(d1/r1/n1, d1/r1/n1) = 0$

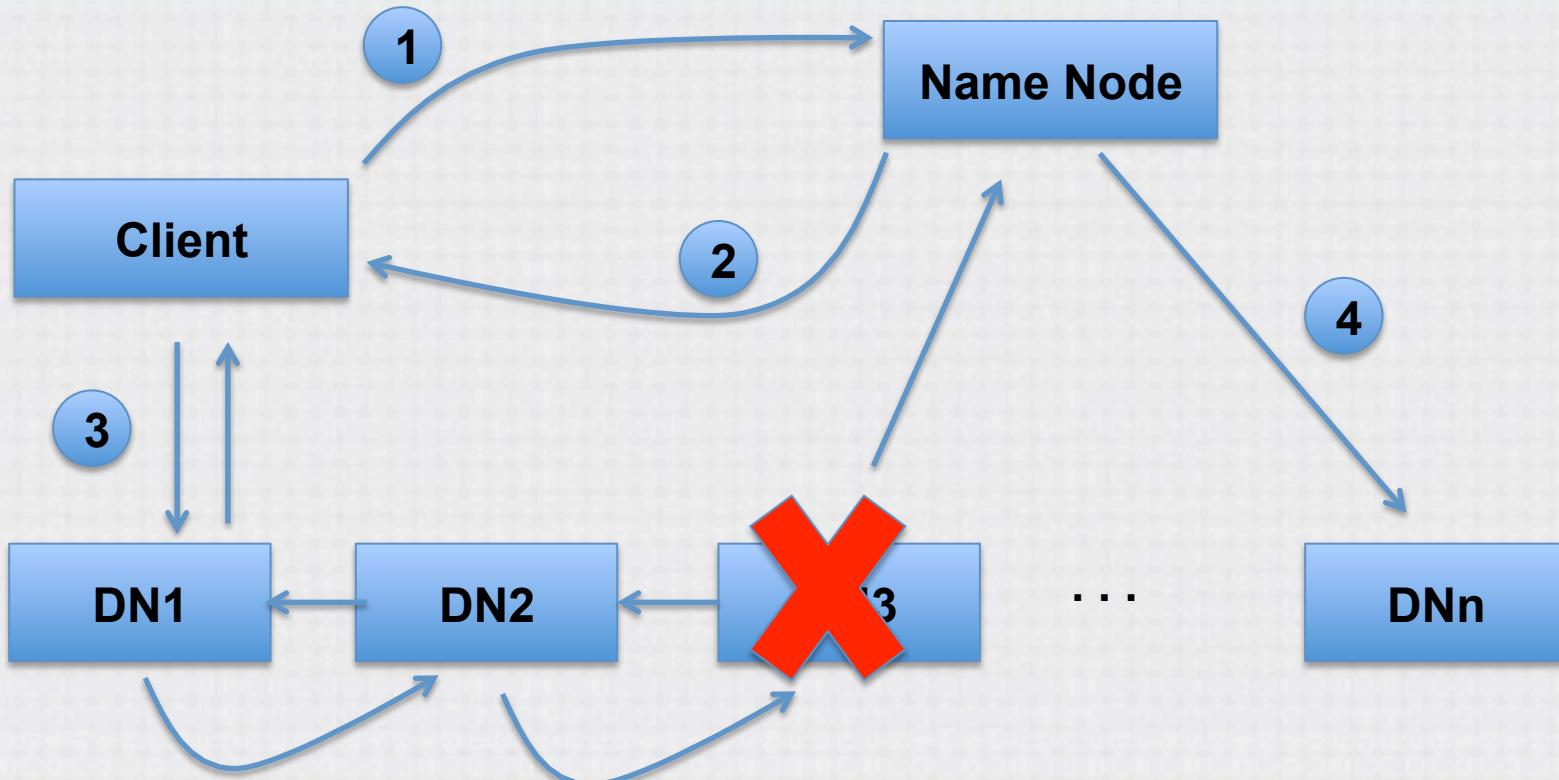
$\text{Dist}(d1/r1/n1, d1/r1/n2) = 2$

$\text{Dist}(d1/r1/n1, d1/r2/n3) = 4$

$\text{Dist}(d1/r1/n1, d2/r3/n6) = 6$



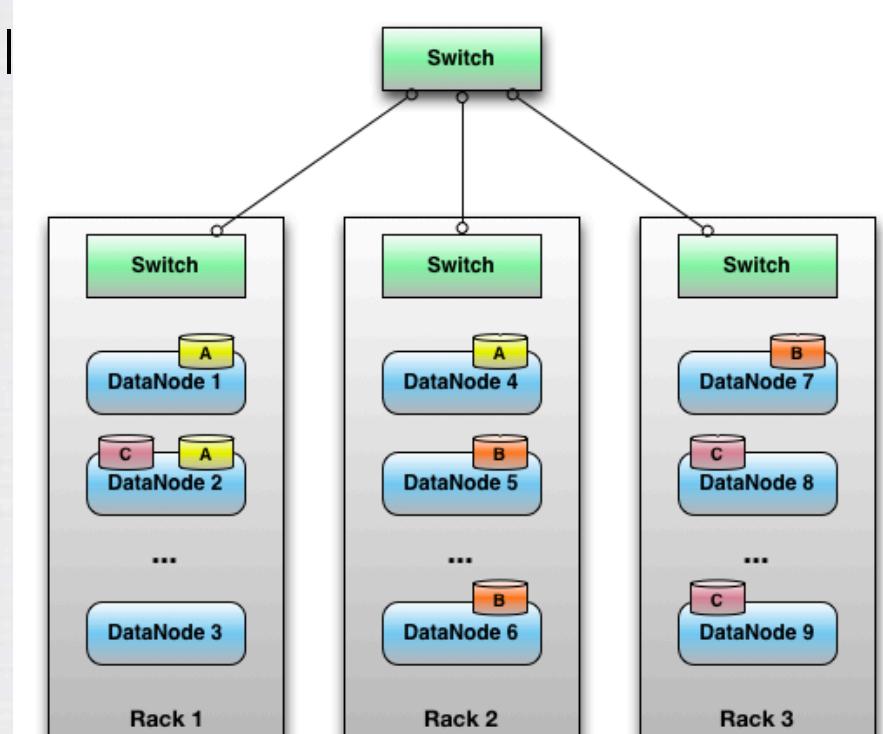
HDFS Inside: Write



1. Client connects to NN to write data
2. NN tells client write these data nodes
3. Client writes blocks directly to data nodes with desired replication factor
4. In case of node failures, NN will figure it out and replicate the missing blocks

Data replication

- **Frist copy** is written to the local node (write affinity).
- **Second copy** is written to a DataNode within a remote rack.
- **Third copy** is written to a DataNode in the same remote rack.
- **Additional** replicas are randomly placed.



Objectives: load balancing, fast access, fault tolerance.

HDFS Inside: Write

- Replication Strategy vs Tradeoffs

	Reliability	Write Bandwidth	Read Bandwidth
Put all replicas on one node			
Put all replicas on different racks			

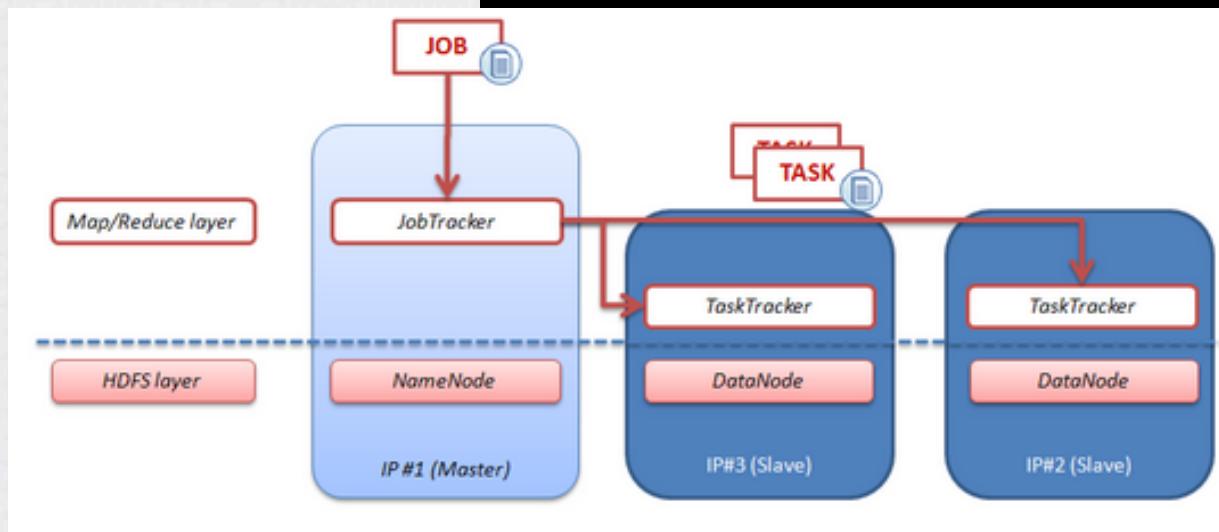
HDFS Inside: Write

- Replication Strategy vs Tradeoffs

	Reliability	Write Bandwidth	Read Bandwidth
Put all replicas on one node			
Put all replicas on different racks			
HDFS: 1-> same node as client 2-> a node on different rack 3-> a different node on the same rack as 2			

MapReduce: Hadoop execution layer

- **JobTracker** knows everything about submitted jobs
- Divides jobs into tasks and decides where to run each task
- Continuously communicating with TaskTracker



- **TaskTracker** execute task (multiple tasks per node)
- Monitors the execution of each task
- Continuously sending feedback to JobTracker

HDFS filesystem commands

1. List the contents of a directory
 - \$hadoop fs -ls
2. Create a directory in HDFS at given path(s)
 - \$hadoop fs -mkdir <directory name>
3. Upload and download a file in HDFS
 - Upload: \$hadoop fs -put <local file> <remote path>
 - Download: \$hadoop fs -get <file in HDFS> <local path>
4. See contents of a file
 - \$hadoop fs -cat <filename>
5. Delete a file/directory in HDFS
 - \$hadoop fs -rm/rmr <file or directory>

HDFS filesystem commands

6. Move file from source to destination
 - \$hadoop fs -mv <src> <dst>
7. Report the amount of space used and availability
 - \$hadoop fs -df hdfs:/
8. How much space a directory occupies
 - \$hadoop fs -du -s -h <dir name>
9. Change permission of files
 - \$sudo hadoop fs -chmod 600 <file>
10. Change owner and group of files
 - \$sudo hadoop fs -chown root:root <file>

HDFS admin commands

- DFSAdmin command
 - -report: reports basic statistics of HDFS
 - -safemode: though usually not required, an administrator can manually enter or leave safemode
 - enter, leave, get, wait
 - -refreshNodes: updates the set of hosts allowed to connect to namenode

Usage: hadoop dfsadmin [-report] [-safemode enter | leave | get | wait] [-refreshNodes] [-finalizeUpgrade] [-upgradeProgress status | details | force] [-metasave filename] [-setQuota <quota> <dirname>...<dirname>] [-clrQuota <dirname>...<dirname>] [-help [cmd]]

Installation and configuration

- Use the stable version, e.g., 1.0.3
- Local (standalone) mode
- **Fully-distributed mode**
- **Pseudo-distributed mode**

Fully distributed mode

- Assume we have three machines with following configurations

OS	Machine name	IP address
Ubuntu 13.04	Master.Hadoop	192.168.1.141
Ubuntu 9.11	Slave1.Hadoop	192.168.1.142
Fedora 17	Slave2.Hadoop	192.168.1.137

- Create a hadoop user account on for each machine

Modify machine names

- For different OS

- Ubuntu OS

```
hadoop@Master:~$ vi /etc/hostname
hadoop@Master:~$ hostname
Master.Hadoop
hadoop@Master:~$
```

- Fedora OS

```
[hadoop@Slave2 conf]$ vi /etc/sysconfig/network
[hadoop@Slave2 conf]$ hostname
Slave2.Hadoop
[hadoop@Slave2 conf]$
```

```
NETWORKING=yes
HOSTNAME=Slave2.Hadoop
NTPSERVERARGS=iburst
~
```

- Mac OS

- \$ sudo scutil --set HostName Master.Hadoop

Configure DNS

- Modify /etc/hosts under the master machine

```
hadoop@Master:~  
$ ping 192.168.1.142  
PING 192.168.1.142 (192.168.1.142) 56(84) bytes of data.  
64 bytes from 192.168.1.142: icmp_req=1 ttl=64 time=7.51 ms  
64 bytes from 192.168.1.142: icmp_req=2 ttl=64 time=0.974 ms  
^C  
--- 192.168.1.142 ping statistics ---  
2 packets transmitted, 2 received, 0% packet loss, time 1001ms  
rtt min/avg/max/mdev = 0.974/4.244/7.515/3.271 ms  
hadoop@Master:~  
$ ping Slave1.Hadoop  
ping: unknown host Slave1.Hadoop  
hadoop@Master:~
```

```
#127.0.0.1      localhost  
127.0.0.1      Master.Hadoop  
  
192.168.1.141  Master.Hadoop  
192.168.1.142  Slave1.Hadoop  
192.168.1.137  Slave2.Hadoop
```

Required software

1. JDK

Download:

<http://www.oracle.com/technetwork/java/javase/index.html>

e.g., jdk-7u25-linux-i586.tar.gz

2. Hadoop

Download:

<http://hadoop.apache.org/common/releases.html>

e.g., hadoop-1.0.3.tar.gz

3. SSH

- \$ sudo apt-get install openssh-server (for ubuntu OS)
- \$ yum install openssh-server (for Fedora OS)

Passwordless login for SSH (1)

- \$ ssh-keygen -t rsa -P '' (no spaces between single quotes; 'P' is uppercase)

```
root@Master:/usr# ssh-keygen -t rsa -P ''
Generating public/private rsa key pair.
Enter file in which to save the key (/root/.ssh/id_rsa):
Your identification has been saved in /root/.ssh/id_rsa.
Your public key has been saved in /root/.ssh/id_rsa.pub.
The key fingerprint is:
e5:bf:4a:ca:88:8b:34:a9:d4:0a:bc:ff:b1:51:74:dc root@Master.Hadoop
The key's randomart image is:
+---[ RSA 2048]----+
|                               |
|                               |
|       . o.E                 |
|       ..o                   |
|       .S .                  |
|. . . . . . .                |
|..+.. o . . .                |
|ooooo .+o o . .              |
| .oo.++. o ...               |
+-----+
```

- \$ cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
- \$ chmod 600 ~/.ssh/authorized_keys

Passwordless login for SSH (2)

- \$ sudo vi /etc/ssh/sshd_config
- Make sure the following lines are uncommented

```
# Authentication:  
LoginGraceTime 120  
PermitRootLogin yes  
StrictModes yes  
  
RSAAuthentication yes  
PubkeyAuthentication yes  
AuthorizedKeysFile      %h/.ssh/authorized_keys
```

- Restart ssh service
 - \$ sudo service ssh restart

Passwordless login for SSH (3)

- Do the `ssh-add` if ssh is still not working

```
seed@Master:~/ssh
$ ssh localhost
Agent admitted failure to sign using the key.
seed@localhost's password:

seed@Master:~/ssh
$ ssh-add ~/ssh/id_rsa
Identity added: /home/seed/.ssh/id_rsa (/home/seed/.ssh/id_rsa)
seed@Master:~/ssh
$ ssh localhost
Welcome to Ubuntu 11.04 (GNU/Linux 2.6.38-8-generic i686)

 * Documentation:  https://help.ubuntu.com/

New release 'oneiric' available.
Run 'do-release-upgrade' to upgrade to it.

Last login: Tue Aug 27 10:31:16 2013 from localhost
```

Passwordless login for SSH (4)

- Copy the generated public key to all slave machines in order passwordless access from the master to slaves

```
hadoop@Master:~/ssh
$ ssh-copy-id hadoop@Slave1.Hadoop
hadoop@slave1.hadoop's password: Input slave1.hadoop's passwd
Now try logging into the machine, with "ssh 'hadoop@Slave1.Hadoop'", and check to
n:

~/ssh/authorized_keys

to make sure we haven't added extra keys that you weren't expecting.

hadoop@Master:~/ssh
```

Passwordless login for SSH (5)

- Test if it is successful to access slave machines without password

```
hadoop@Master:~/ssh
$ ssh hadoop@Slave1.Hadoop
Linux Slave1.Hadoop 2.6.28-11-generic #42-Ubuntu SMP Fri Apr 17 01:57:59 UTC 200
9 i686

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

To access official Ubuntu documentation, please visit:
http://help.ubuntu.com/

105 packages can be updated.
55 updates are security updates.

Last login: Tue Aug 27 21:47:57 2013 from slave3.hadoop
hadoop@Slave1:~$
```

Passwordless login for SSH (6)

- To make passwordless access from slave machines to the master machine, we need to add public keys of slave machines into 'authorized_keys' on the master machine under the folder of .ssh using the following command
- \$ cat ~/.ssh/id_rsa.pub | ssh hadoop@Master.Hadoop 'cat >> ~/.ssh/authorized_keys'
- It is equal to two commands:
 - \$ scp ~/.ssh/id_rsa.pub [hadoop@Master.Hadoop:~/](#)
 - \$ cat ~/id_rsa.pub >> ~/.ssh/authorized_keys

```
[root@Slave3 root]# cat ~/.ssh/id_rsa.pub | ssh seed@Master.Hadoop 'cat >> ~/.ssh/authorized_keys'
The authenticity of host 'master.hadoop (192.168.1.141)' can't be established.
RSA key fingerprint is fa:ae:87:8c:d8:5f:03:a9:f8:13:0a:99:86:a9:bd:f1.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added 'master.hadoop,192.168.1.141' (RSA) to the list of known hosts.
seed@master.hadoop's password:
[root@Slave3 root]# ssh seed@Master.Hadoop
Welcome to Ubuntu 11.04 (GNU/Linux 2.6.38-8-generic i686)

 * Documentation:  https://help.ubuntu.com/

 New release 'oneiric' available.
 Run 'do-release-upgrade' to upgrade to it.

Last login: Tue Aug 27 21:26:03 2013 from slave2.hadoop
seed@Master:~
```

Set Java environment variables

- Edit /etc/profile (Ubuntu) or ~/.bashrc_profile (Mac)

```
umask 022

# set java environment
export JAVA_HOME=/usr/java/jdk1.7.0_25/
export JRE_HOME=/usr/java/jdk1.7.0_25/jre
export CLASSPATH=.:${CLASSPATH}:$JAVA_HOME/lib:$JRE_HOME/lib
export PATH=$PATH:$JAVA_HOME/bin:$JRE_HOME/bin
```

- Make it work through
 - \$ source /etc/profile or ./etc/profile

```
root@Master:/usr/java# vi /etc/profile
root@Master:/usr/java# . /etc/profile
root@Master:/usr/java#
```

Install Hadoop

```
cd /usr
tar -xzvf hadoop-1.1.2.tar.gz
mv hadoop-1.1.2 hadoop
chown -R hadoop:hadoop hadoop
rm -rf hadoop-1.1.2.tar.gz
```

- Modify /etc/profile to add some environment variables

```
umask 022

# set java environment
export JAVA_HOME=/usr/java/jdk1.7.0_25/
export JRE_HOME=/usr/java/jdk1.7.0_25/jre
export CLASSPATH=.:$CLASSPATH:$JAVA_HOME/lib:$JRE_HOME/lib
export PATH=$PATH:$JAVA_HOME/bin:$JRE_HOME/bin

# set Hadoop environment
export HADOOP_HOME=/usr/hadoop
export PATH=$PATH:$HADOOP_HOME/bin
:
```

Configure Hadoop (1)

- Edit /usr/hadoop/conf/hadoop-env.sh

```
# Set Hadoop-specific environment variables here.

# The only required environment variable is JAVA_HOME.  A
# optional.  When running a distributed configuration it
# set JAVA_HOME in this file, so that it is correctly def
# remote nodes.

# The java implementation to use.  Required.
export JAVA_HOME=/usr/java/jdk1.7.0_25
```

Configure Hadoop (2)

- Edit /usr/hadoop/conf/core-site.xml

```
<?xml version="1.0"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>

<!-- Put site-specific property overrides in this file. --&gt;

&lt;configuration&gt;
  &lt;property&gt;
    &lt;name&gt;hadoop.tmp.dir&lt;/name&gt;
    &lt;value&gt;/usr/hadoop/tmp&lt;/value&gt;
    &lt;description&gt;A base for other temporary directories.&lt;/description&gt;
  &lt;/property&gt;
  &lt!-- file system properties --&gt;
  &lt;property&gt;
    &lt;name&gt;fs.default.name&lt;/name&gt;
    &lt;value&gt;hdfs://192.168.1.141:9000&lt;/value&gt;
  &lt;/property&gt;
&lt;/configuration&gt;</pre>
```

Configure Hadoop (3)

- Edit /usr/hadoop/conf/hdfs-site.xml

```
<?xml version="1.0"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>

<!-- Put site-specific property overrides in this file. --&gt;

&lt;configuration&gt;
  &lt;property&gt;
    &lt;name&gt;dfs.replication&lt;/name&gt;
    &lt;value&gt;1&lt;/value&gt;
  &lt;/property&gt;
&lt;/configuration&gt;
~</pre>
```

Configure Hadoop (4)

- Edit /usr/hadoop/conf/mapred-site.xml

```
<?xml version="1.0"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>

<!-- Put site-specific property overrides in this file. --&gt;

&lt;configuration&gt;
  &lt;property&gt;
    &lt;name&gt;mapred.job.tracker&lt;/name&gt;
    &lt;value&gt;http://192.168.1.141:9001&lt;/value&gt;
  &lt;/property&gt;
&lt;/configuration&gt;
~</pre>
```

Configure Hadoop (5)

- Edit 'masters' file

```
root@Master:/usr/hadoop/conf# vi masters
root@Master:/usr/hadoop/conf# cat masters
192.168.1.141
root@Master:/usr/hadoop/conf#
```

- Configure 'slaves' file (only needed on the master machine)

```
root@Master:/usr/hadoop/conf# vi slaves
root@Master:/usr/hadoop/conf# cat slaves
192.168.1.142
192.168.1.137
```

Start Hadoop (1)

```
hadoop@Master:/usr/hadoop
$ hadoop namenode -format
Warning: $HADOOP_HOME is deprecated.

13/08/28 03:31:47 INFO namenode.NameNode: STARTUP_MSG:
/*****
STARTUP_MSG: Starting NameNode
STARTUP_MSG:   host = Master.Hadoop/127.0.0.1
STARTUP_MSG:   args = [-format]
STARTUP_MSG:   version = 1.1.2
STARTUP_MSG:   build = https://svn.apache.org/repos/asf/hadoop/common/branches/b
ranch-1.1 -r 1440782; compiled by 'hortonfo' on Thu Jan 31 02:03:24 UTC 2013
*****
13/08/28 03:31:47 INFO util.GSet: VM type      = 32-bit
13/08/28 03:31:47 INFO util.GSet: 2% max memory = 19.33375 MB
13/08/28 03:31:47 INFO util.GSet: capacity      = 2^22 = 4194304 entries
13/08/28 03:31:47 INFO util.GSet: recommended=4194304, actual=4194304
13/08/28 03:31:49 INFO namenode.FSNamesystem: fsOwner=hadoop
13/08/28 03:31:50 INFO namenode.FSNamesystem: supergroup=supergroup
13/08/28 03:31:50 INFO namenode.FSNamesystem: isPermissionEnabled=true
13/08/28 03:31:50 INFO namenode.FSNamesystem: dfs.block.invalidate.limit=100
13/08/28 03:31:50 INFO namenode.FSNamesystem: isAccessTokenEnabled=false accessK
eyUpdateInterval=0 min(s), accessTokenLifetime=0 min(s)
13/08/28 03:31:50 INFO namenode.NameNode: Caching file names occurring more than
10 times
13/08/28 03:31:51 INFO common.Storage: Image file of size 112 saved in 0 seconds
.
13/08/28 03:31:51 INFO namenode.FSEditLog: closing edit log: position=4, editlog
=/usr/hadoop/tmp/dfs/name/current/edits
13/08/28 03:31:51 INFO namenode.FSEditLog: close success: truncate to 4, editlog
=/usr/hadoop/tmp/dfs/name/current/edits
13/08/28 03:31:51 INFO common.Storage: Storage directory /usr/hadoop/tmp/dfs/nam
e has been successfully formatted.
13/08/28 03:31:51 INFO namenode.NameNode: SHUTDOWN_MSG:
/*****
SHUTDOWN_MSG: Shutting down NameNode at Master.Hadoop/127.0.0.1
*****
```

Start Hadoop (2)

```
hadoop@Master:/usr/hadoop
$ start-all.sh
Warning: $HADOOP_HOME is deprecated.

starting namenode, logging to /usr/hadoop/libexec/../logs/hadoop-hadoop-namenode
-Master.Hadoop.out
192.168.1.137: starting datanode, logging to /usr/hadoop/libexec/../logs/hadoop-
hadoop-datanode-Slave2.Hadoop.out
192.168.1.142: starting datanode, logging to /usr/hadoop/libexec/../logs/hadoop-
hadoop-datanode-Slave1.Hadoop.out
192.168.1.141: starting secondarynamenode, logging to /usr/hadoop/libexec/../log
s/hadoop-hadoop-secondarynamenode-Master.Hadoop.out
starting jobtracker, logging to /usr/hadoop/libexec/../logs/hadoop-hadoop-jobtra
cker-Master.Hadoop.out
192.168.1.142: starting tasktracker, logging to /usr/hadoop/libexec/../logs/hado
op-hadoop-tasktracker-Slave1.Hadoop.out
192.168.1.137: starting tasktracker, logging to /usr/hadoop/libexec/../logs/hado
op-hadoop-tasktracker-Slave2.Hadoop.out
hadoop@Master:/usr/hadoop
```

Check Hadoop status

```
hadoop@Master:/usr/hadoop/tmp
$ jps
6231 Jps
5753 SecondaryNameNode
5834 JobTracker
5549 NameNode
hadoop@Master:/usr/hadoop/tmp
$
```

```
[hadoop@Slave2 tmp]$ jps
3454 DataNode
3636 Jps
3532 TaskTracker
[hadoop@Slave2 tmp]$
```

```
hadoop@Master:/usr/hadoop
$ hadoop dfsadmin -report
Warning: $HADOOP_HOME is deprecated.

Configured Capacity: 24690470912 (22.99 GB)
Present Capacity: 15462825999 (14.4 GB)
DFS Remaining: 15462768640 (14.4 GB)
DFS Used: 57359 (56.01 KB)
DFS Used%: 0%
Under replicated blocks: 0
Blocks with corrupt replicas: 0
Missing blocks: 0

-----
Datanodes available: 2 (2 total, 0 dead)

Name: 192.168.1.137:50010
Decommission Status : Normal
Configured Capacity: 16651051008 (15.51 GB)
DFS Used: 28687 (28.01 KB)
Non DFS Used: 5241806833 (4.88 GB)
DFS Remaining: 11409215488(10.63 GB)
DFS Used%: 0%
DFS Remaining%: 68.52%
Last contact: Wed Aug 28 09:20:28 EDT 2013

Name: 192.168.1.142:50010
Decommission Status : Normal
Configured Capacity: 8039419904 (7.49 GB)
DFS Used: 28672 (28 KB)
Non DFS Used: 3985838080 (3.71 GB)
DFS Remaining: 4053553152(3.78 GB)
DFS Used%: 0%
DFS Remaining%: 50.42%
Last contact: Wed Aug 28 09:20:29 EDT 2013
```

Web access Hadoop

<http://192.168.1.141:50030>

The screenshot shows the "Master Hadoop Map/Reduce Administration" page. At the top, it displays system status: State: RUNNING, Started: Wed Aug 28 09:18:03 EDT 2013, Version: 1.2.2, r14460782, Compiled: Thu Jan 31 02:03:24 UTC 2013 by hortoninfo, Identifier: 201308280917, and SafeMode: OFF.

Cluster Summary (Heap Size is 8.7 MB/966.69 MB)

Running Map Tasks	Running Reduce Tasks	Total Submissions	Nodes	Occupied Map Slots	Occupied Reduce Slots	Reserved Map Slots	Reserved Reduce Slots	Map Task Capacity	Reduce Task Capacity	Avg. Tasks/Node	Blacklisted Nodes
0	0	0	2	0	0	0	0	4	4	4.00	0

Scheduling Information

Queue Name	State	Scheduling Information
default	running	N/A

Filter (Jobid, Priority, User, Name): Example: 'user smith 3200' will filter by 'smith' only in the user field and '3200' in all fields

Running Jobs

Retired Jobs

Web access Hadoop

http://192.168.1.141:50070

The screenshot shows a web browser window with the URL <http://192.168.1.141:50070/dfshealth.jsp> in the address bar. The page title is "NameNode 'Master.Hadoop:9000'".

Started: Wed Aug 28 09:17:50 EDT 2013
Version: 1.1.2, r1440782
Compiled: Thu Jan 31 02:03:24 UTC 2013 by hortonfo
Upgrades: There are no upgrades in progress.

[Browse the filesystem](#)
[Namenode Logs](#)

Cluster Summary

7 files and directories, 1 blocks = 8 total. Heap Size is 49.38 MB / 966.69 MB (5%)

Configured Capacity	22.99 GB
DFS Used	68 KB
Non DFS Used	8.59 GB
DFS Remaining	14.4 GB
DFS Used%	0 %
DFS Remaining%	62.63 %
Live Nodes	2
Dead Nodes	0
Decommissioning Nodes	0
Number of Under-Replicated Blocks	0

NameNode Storage:

Storage Directory	Type	State
/usr/hadoop/tmp/dfs/name	IMAGE_AND_EDITS	Active

This is Apache Hadoop release 1.1.2

Installation and configuration

- Local (standalone) mode
- Fully-distributed mode
- **Pseudo-distributed mode**

Please also refer to PDF file
Hadoop_installation_and_configuration.pdf