

# **BIG DATA and AI for business**

---

Hadoop

**Decisions, Operations & Information Technologies**  
**Robert H. Smith School of Business**  
**Fall, 2020**



# **Distributed processing is non-trivial**

- How to assign tasks to different workers in an efficient way?
- What happens if tasks fail?
- How do workers exchange results?
- How to synchronize distributed tasks allocated to different workers?

# **Big data storage is challenging**

- Data **volumes** are massive
- **Reliability** of storing PBs of data is challenging
- All kinds of **failures**: Disk/Hardware/Network Failures
- Probability of failures simply increase with the number of machines ...

# One popular solution: Hadoop



Hadoop Cluster at Yahoo!

# Hadoop offers

- Redundant, Fault-tolerant data storage
- Parallel computation framework
- Job coordination



# Hadoop offers

- Redundant, Fault-tolerant data storage
- Parallel computation framework
- Job coordination



Programmers

No longer need to  
worry about



Q: Where file is located?

Q: How to handle failures & data lost?

Q: How to divide computation?

Q: How to program for scaling?

# A little history on Hadoop

- Hadoop is an open-source implementation based on Google File System (GFS) and MapReduce from Google
- Hadoop was created by Doug Cutting and Mike Cafarella in 2005
- Hadoop was donated to Apache in 2006



# Who uses Hadoop?

Social



The image shows the Twitter for Business landing page. It features a green header with the Twitter logo and the text "Twitter for Business". Below the header, there's a section titled "Small business advertisers, learn how to get \$100 in free Twitter advertising from American Express." It includes a "Learn the Basics" section with links to "What is Twitter?", "Twitter Glossary", "Best Practices", and "Twitter on the Go". There's also an "Optimize" section with a "Case Resource API" link. The main visual is a collage of icons related to business, including a chalkboard with "ABCs", a bar chart, a pie chart, and the Facebook logo.

User Tracking & Engagement



Homeland Security



eCommerce



Financial Services



Real Time Search

Google



# Who uses Hadoop?

2007

**YAHOO!**



**last.fm**

2008

**Google** ablegrape



**Cascading**

**IBM**

**facebook**

**enORMO** Every property. Everywhere.



**krugle**

**rackspace** HOSTING

**Lookery** Control freaks welcome

**The New York Times** joost

**Zvents**

FORMATION SCIENCES INSTITUTE

**News Corporation**

Cornell University  
Computing and Information Science

Visible MEASURES

**LOTAME**  
Locate, Target & Message with Social Media

**NetSeer**

**parc**  
Palo Alto Research Center

SECURITY ENHANCED  
DOMAIN NAME SYSTEM

**veeh**™

2009

**AOL** cloudera

**deepdyve** BETA

**cooliris**

**eyearlike**

**TEXTMAP**  
THE ENTITY SEARCH ENGINE

PSG College of Technology

**iterend**

**tailsweep**

**hulu**™

**RapLeaf**

**USC CMS**

**Ning**

**quantcast**

**amazon** web services™

**pressflip**

**detikSearch**



**Systems@ETH** Zürich

**VK SOLUTIONS**  
Global Solutions Provider

**TARAGANA**  
Innovation • Quality • Simplicity



**HOSTING HABITAT**

**HOLA** SERVERS

**Terrier**

**adknowledge**



**stampede** beta

2010

**SAMSUNG**

**rubicon** PROJECT

**BERKELEY LAB**  
LAWRENCE BERKELEY NATIONAL LABORATORY

**VISIBLE** TECHNOLOGIES

**APOLLO GROUP**

**ADSDAQ**

**rackspace** HOSTING

**RapLeaf**

**wordnik**

**mBIGEN** network revolution

**comSCORE**

**trulia** real estate search

**Accela** COMMUNICATIONS

**Forward3D**

**LinkedIn** ®

**Microsoft**

**Infochimps**

**PharmZphork**  
Find the world's data

**ADMELD**

**gumgum**®

**Pronux**

**The DataGraph Bbg**

**NETFLIX**

**mobileanalytics.tv**

**markt24.de**

**twitter**™

**media6degrees**

**BEEBLER**

**SLC Security**  
When Experience Matters...

**eBay**®

# Why HDFS?

- **Problem 1:** Data is too big to store on one machine.
- **HDFS:** Store the data on multiple machines!

# Why HDFS?

- **Problem 2:** Very high end machines are too expensive
- **HDFS:** Run on commodity hardware!

# Why HDFS?

- **Problem 3:** Commodity hardware can fail
- **HDFS:** Software is intelligent enough to handle hardware failure!

# Why HDFS?

- **Problem 4:** What happens to the data if the machine storing the data fails?
- **HDFS:** Replicate the data!

# Why HDFS?

- **Problem 5:** How can distributed machines organize the data in a coordinated way?
- **HDFS: Master-Slave Architecture!**

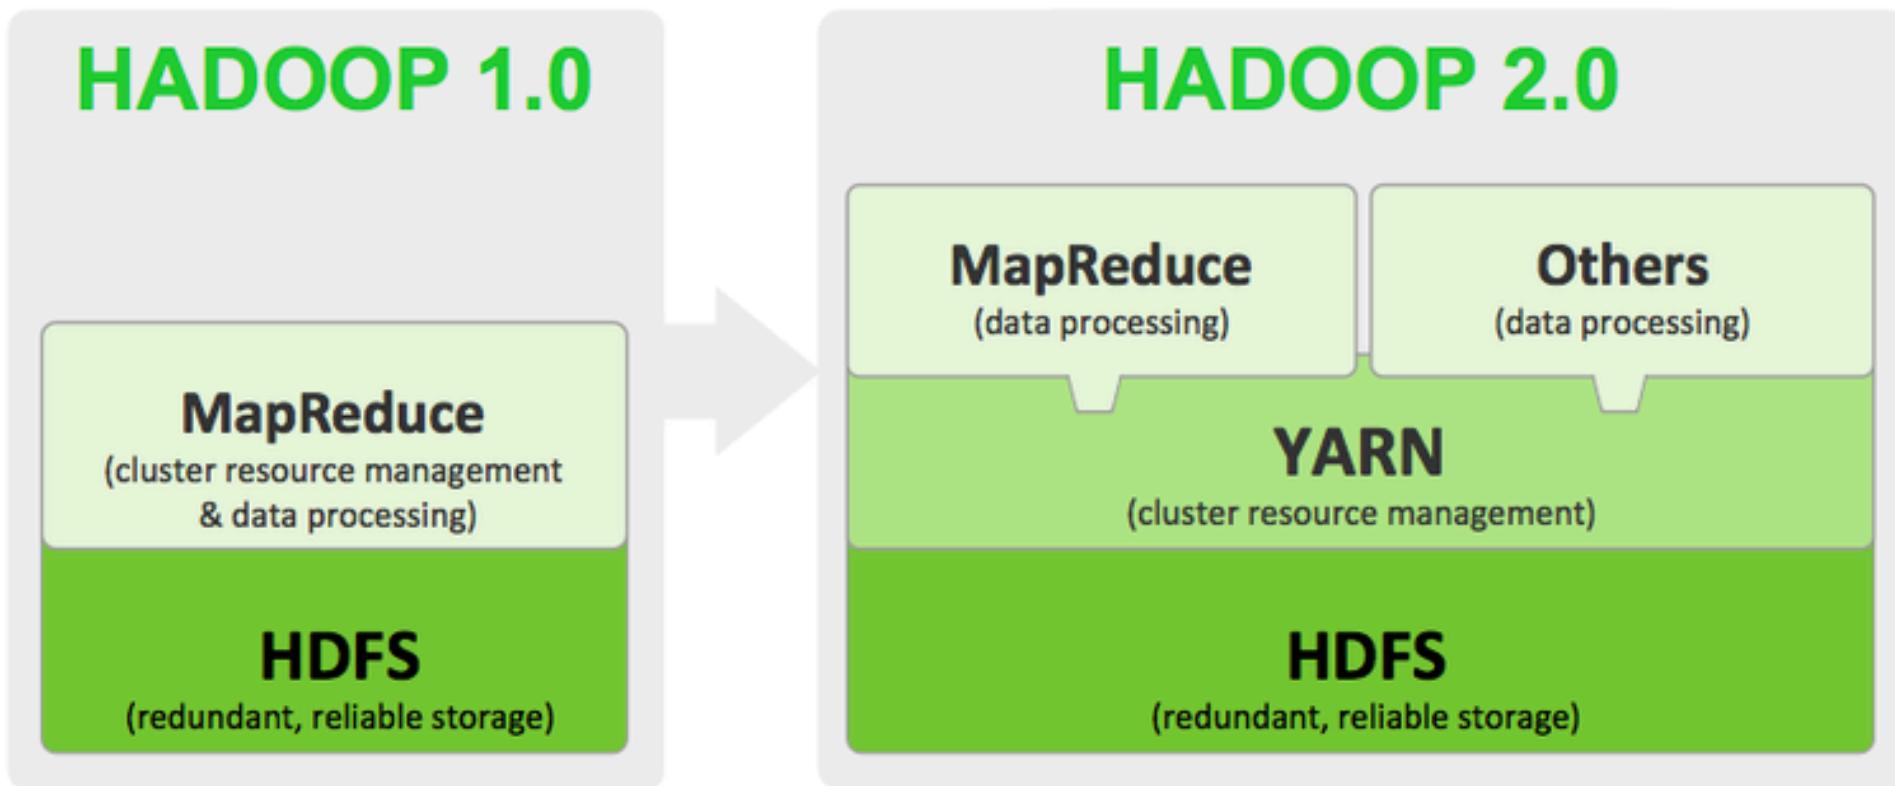
# Another reason why Hadoop

- **Scan** 100TB datasets on a 1000-node cluster
  - Remote storage @ 10MB/s = 165 mins
  - Local storage @ 50-200MB/s = 33-8 mins
- Moving computation is more efficient than moving data
- Need fault tolerant store with reasonable availability guarantees
  - Handle hardware faults transparently

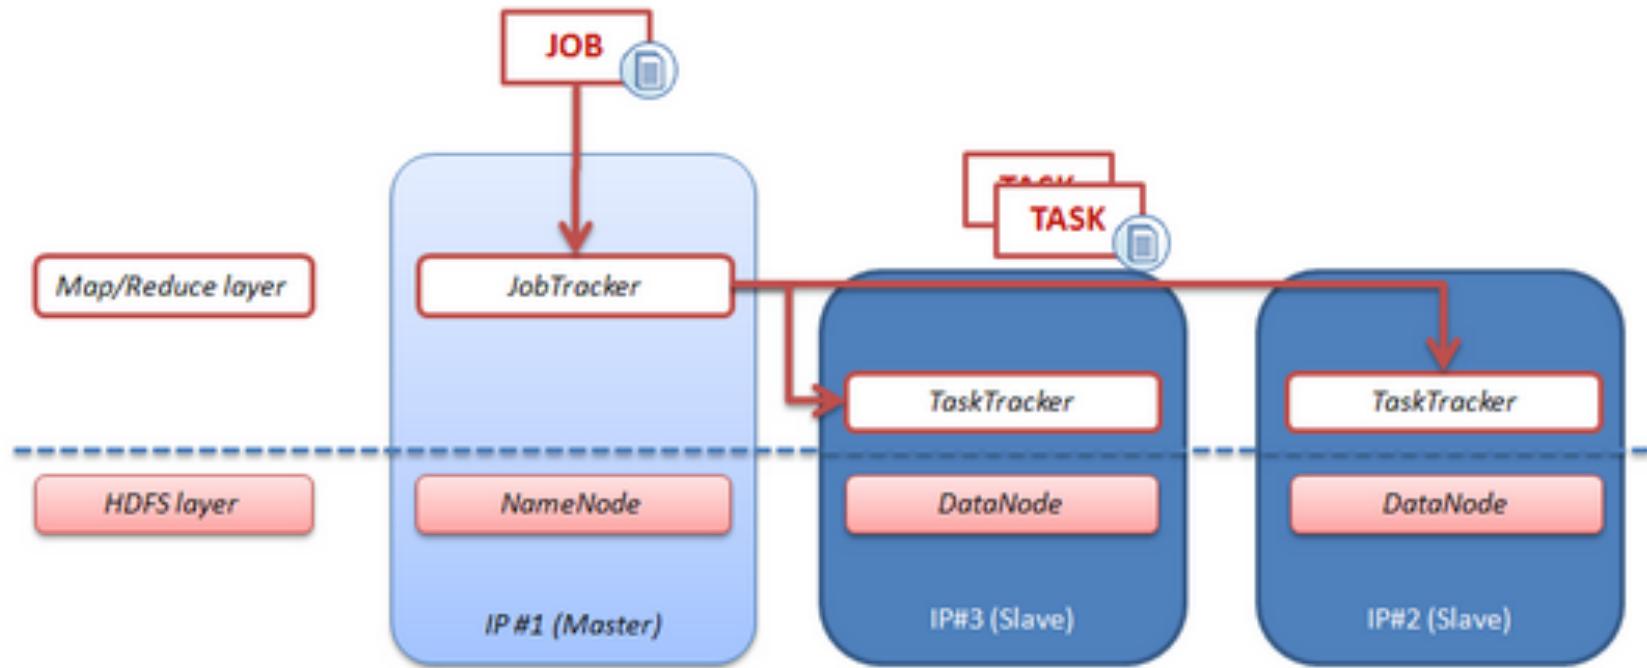
# Hadoop goals

- **Scalable**: Petabytes ( $10^{15}$  Bytes) of data on thousands of nodes
- **Economical**: Commodity components only
- **Reliable**: fault tolerance

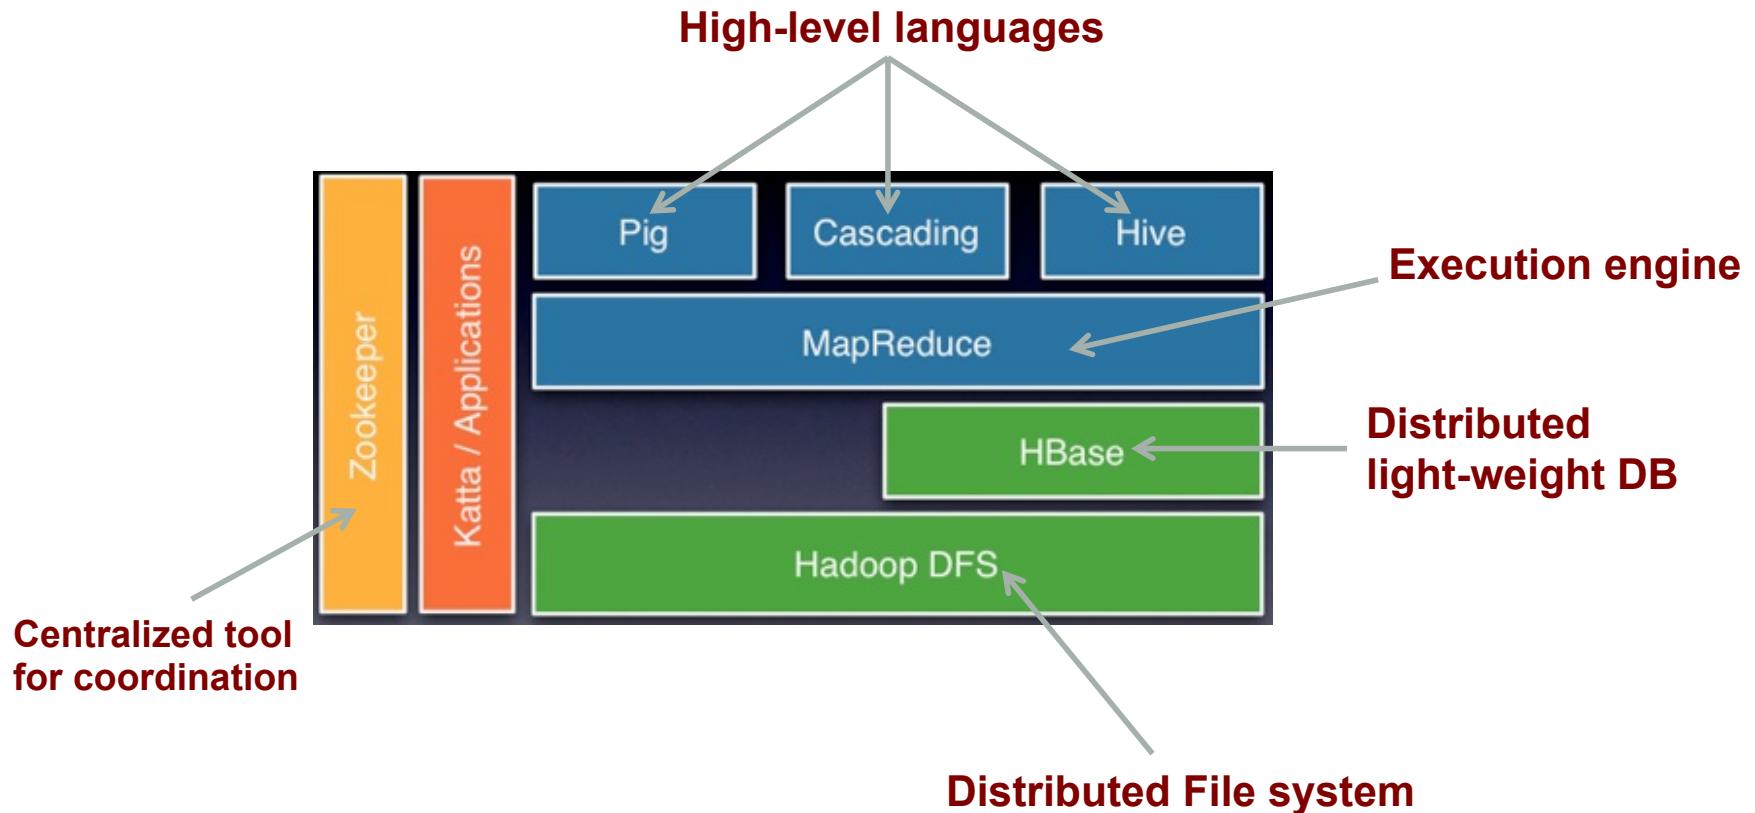
# Hadoop big picture



# High-level architecture of Hadoop

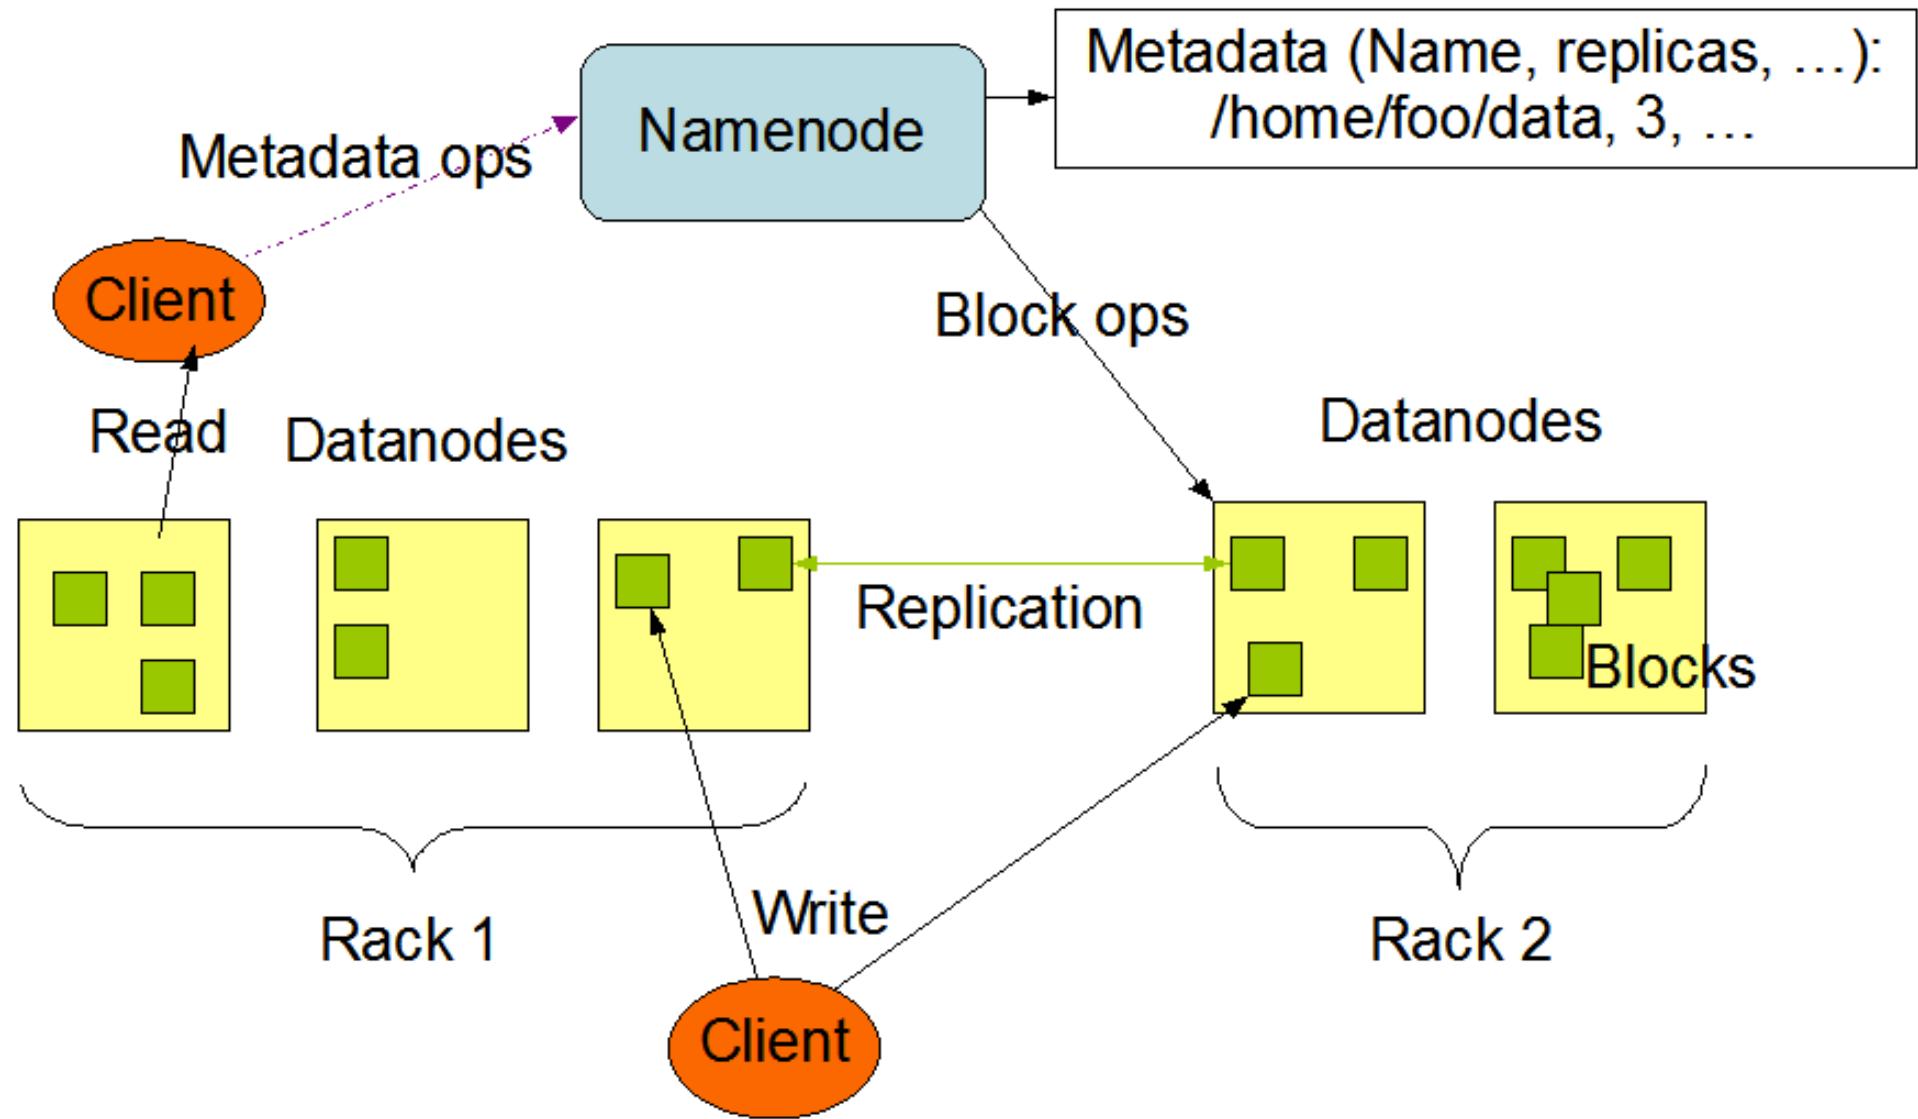


# Hadoop big picture

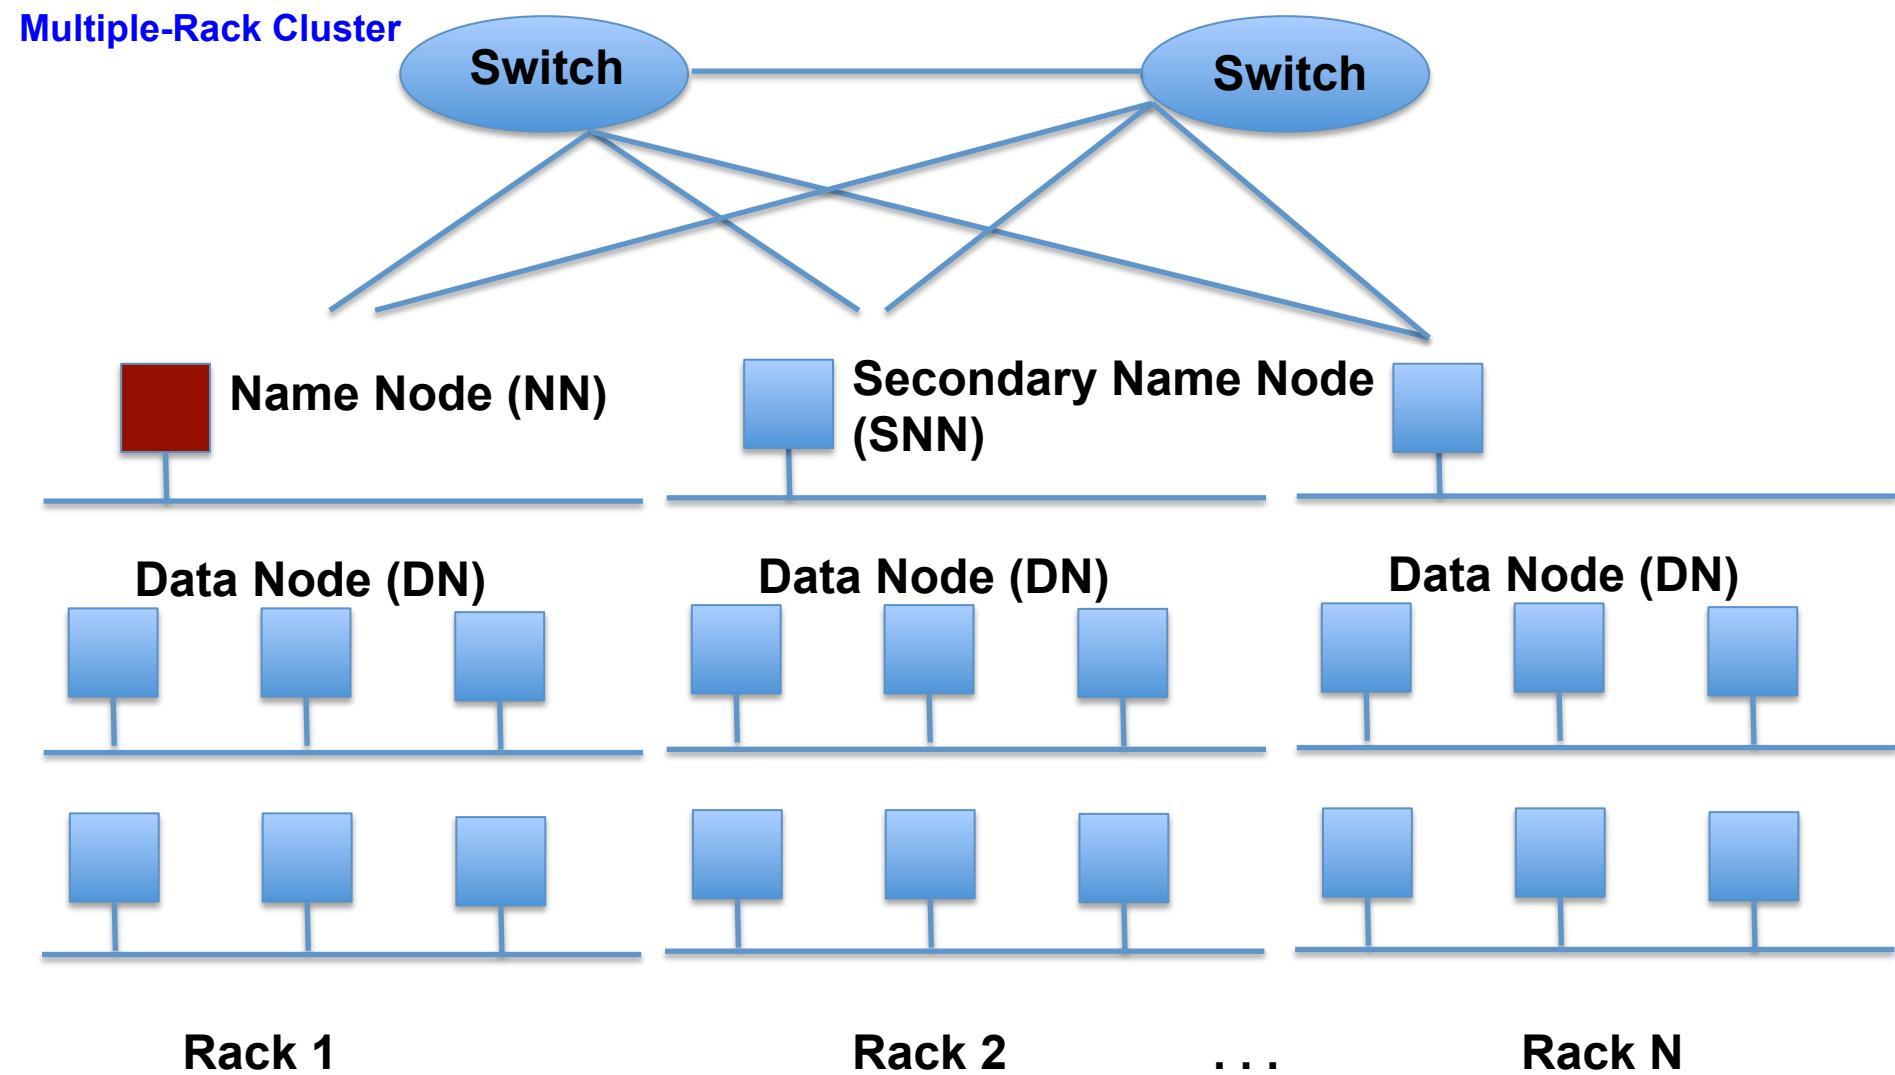


HDFS + MapReduce are enough to have things working

# HDFS Architecture



# HDFS architecture



# HDFS

- **Master-Slave** architecture
- Single NameNode
  - Sometimes a backup: secondary NameNode
- Many (Thousands) DataNodes
- Files are split into fixed sized blocks and stored on data nodes
- Data blocks are replicated for fault tolerance and fast access (**default: 3**)

# HDFS - Master (NameNode)

- Manages file system (FS) namespace
- File metadata
- Mapping file to list of blocks
- Authorization & Authentication
- Mapping of datanode to list of blocks
- Monitor datanode health
- Replicate missing blocks
- Keeps ALL namespace in memory

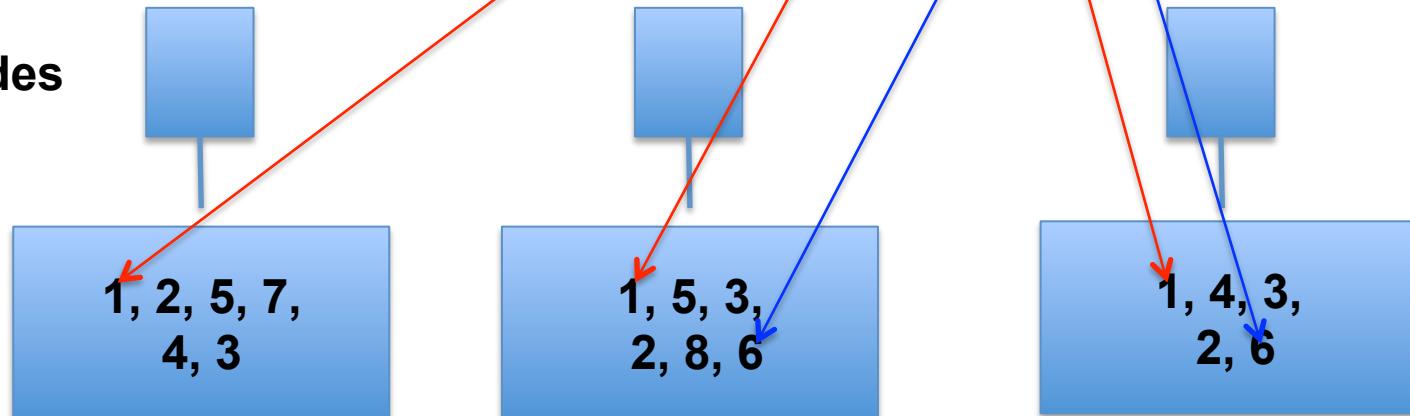
# HDFS - Slave (DataNode)

- Handle block storage on multiple volumes & block integrity
- Clients access the blocks directly from data nodes
- Periodically send heartbeats and block reports to NameNode
- Blocks are stored as underlying OS's files

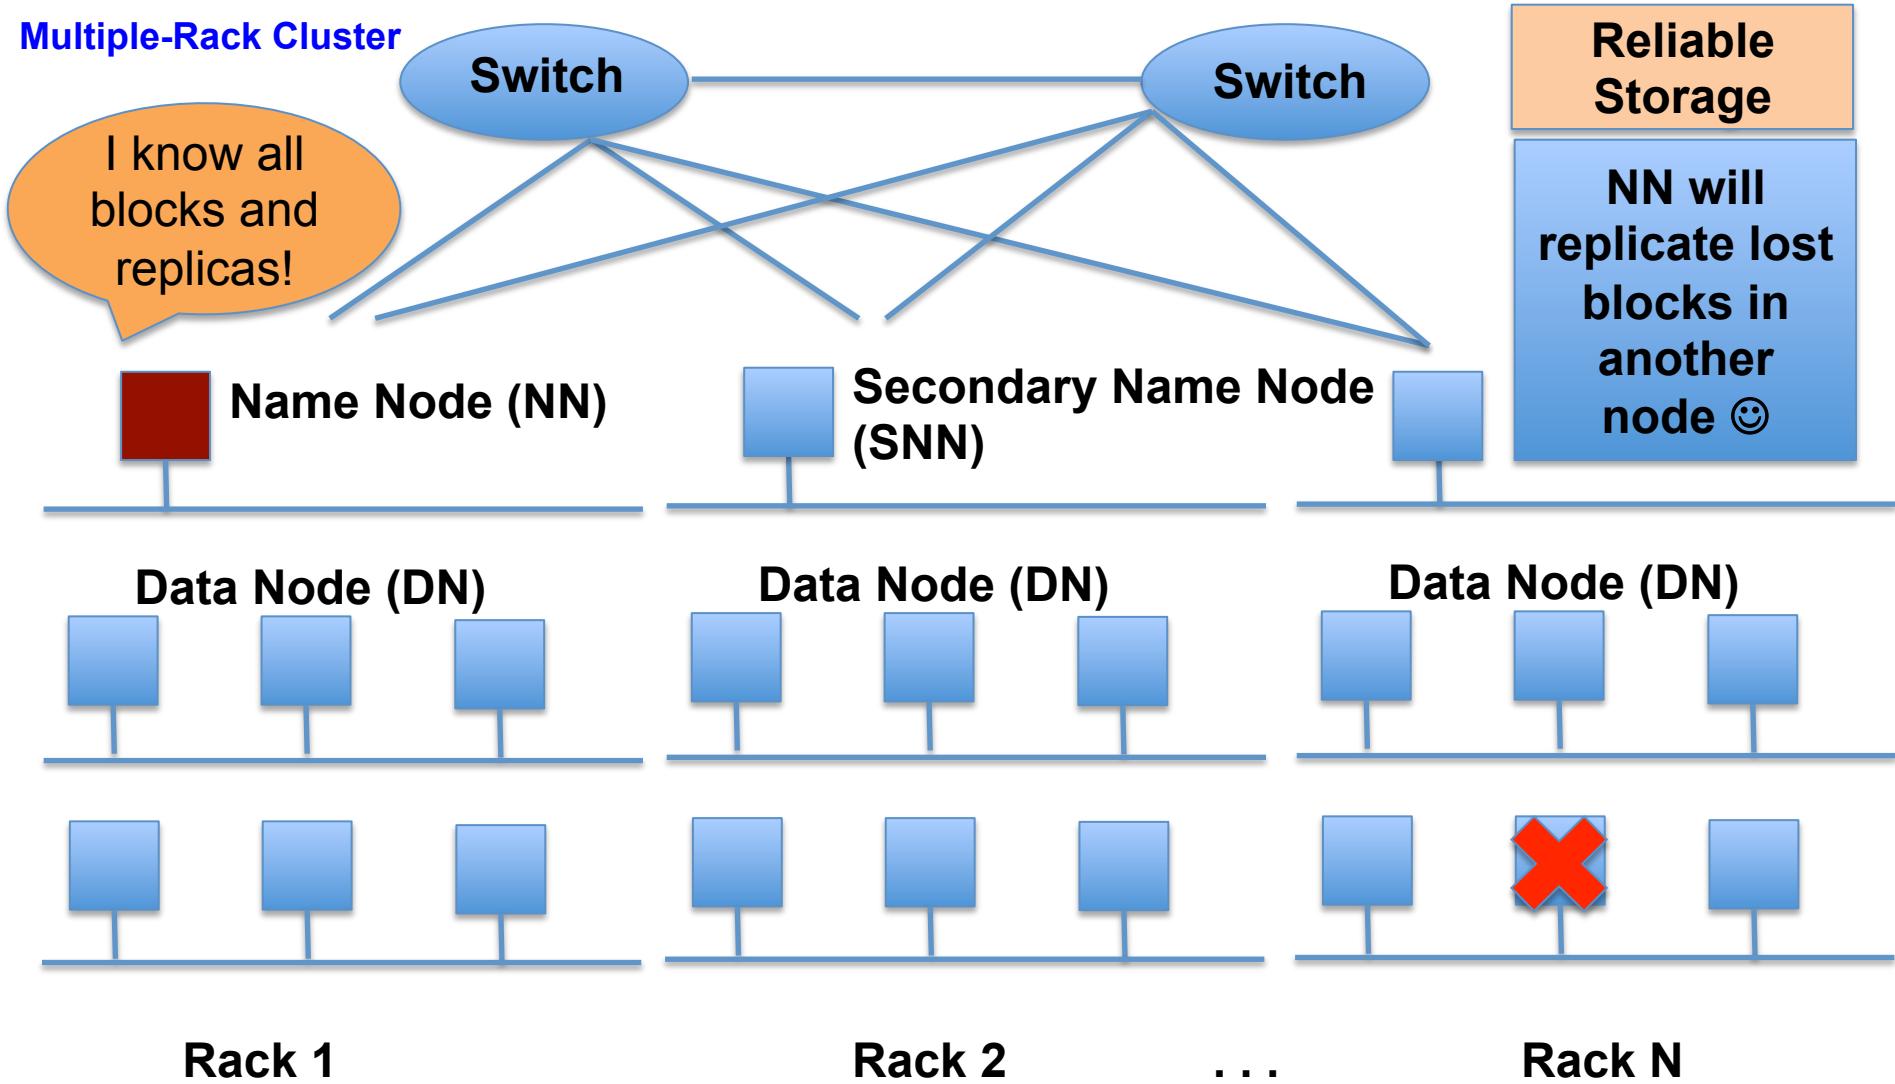
# HDFS Name Node

Name Node	Snapshot of FS	Edit log: record changes to FS
Filename	Replication factor	Block ID
File 1	3	[1, 2, 3]
File 2	2	[4, 5, 6]
File 3	1	[7, 8]

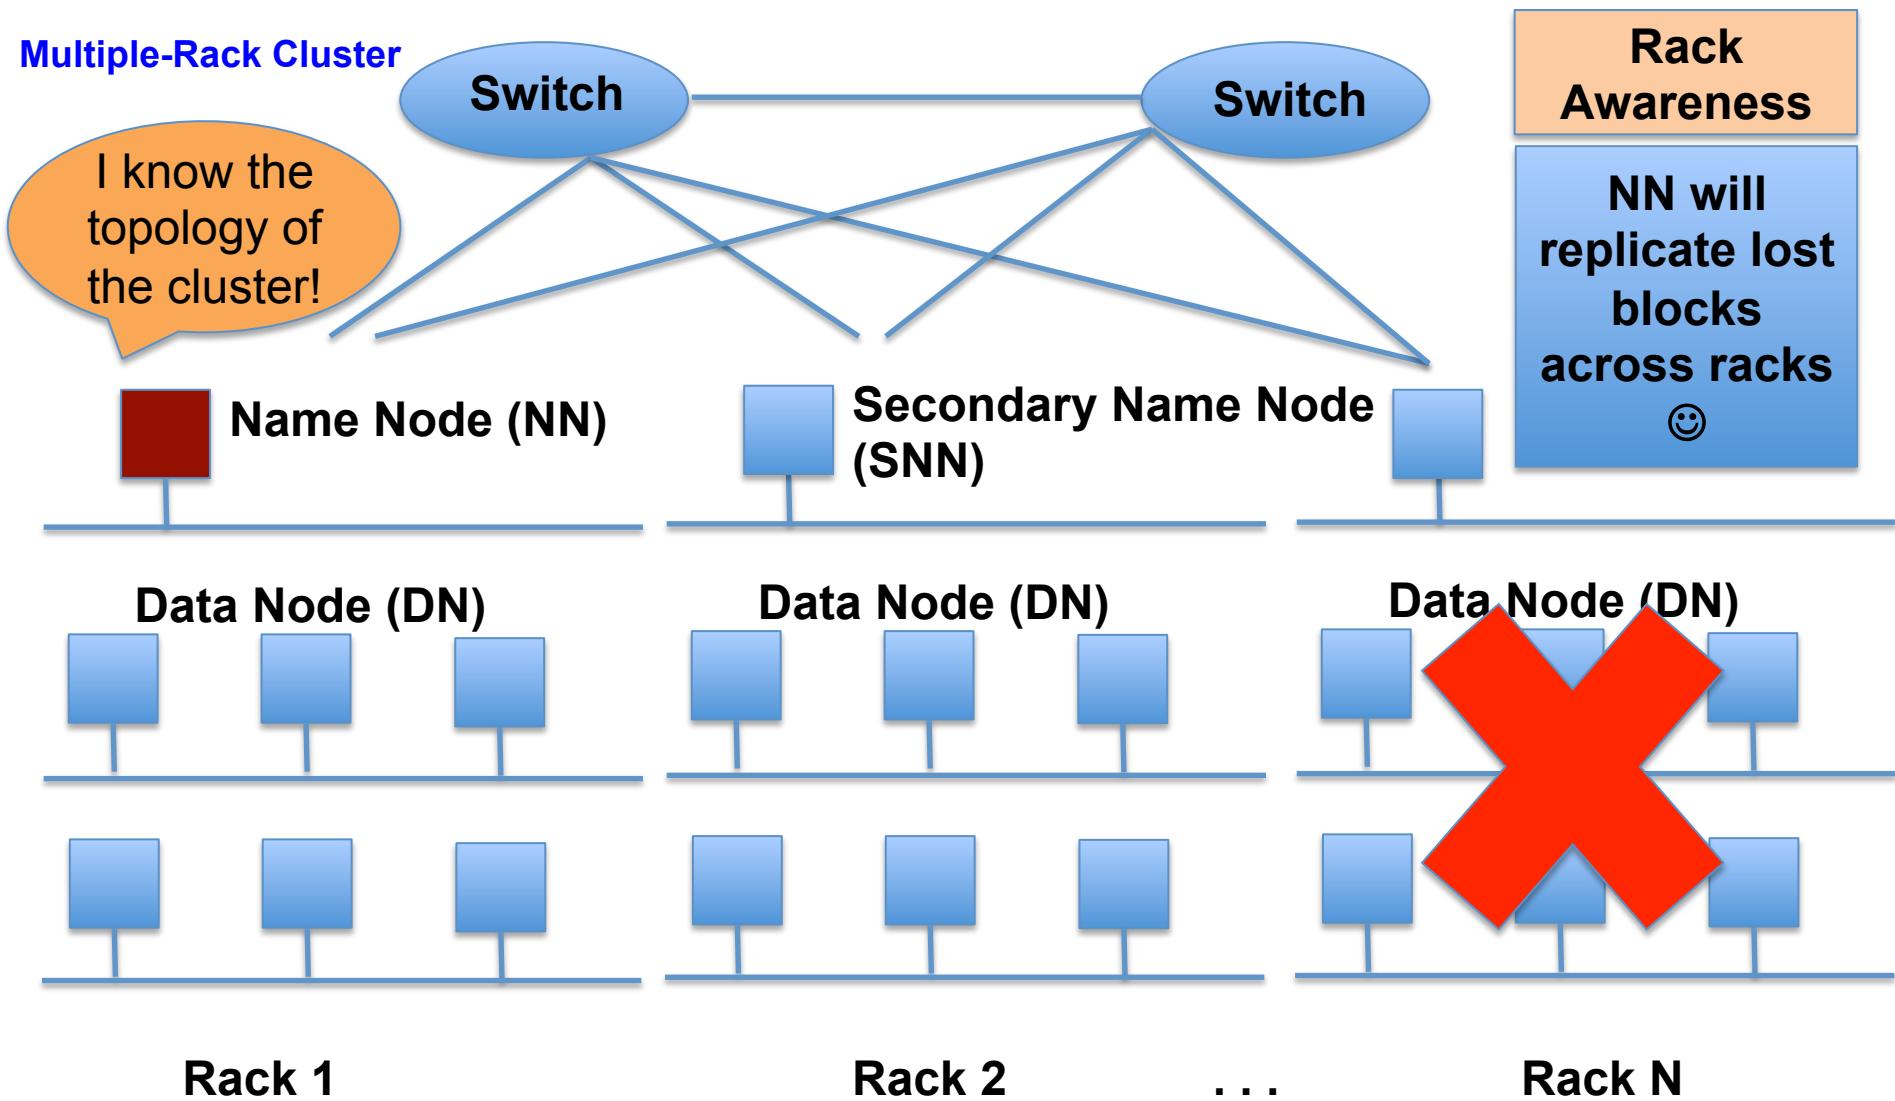
Data Nodes



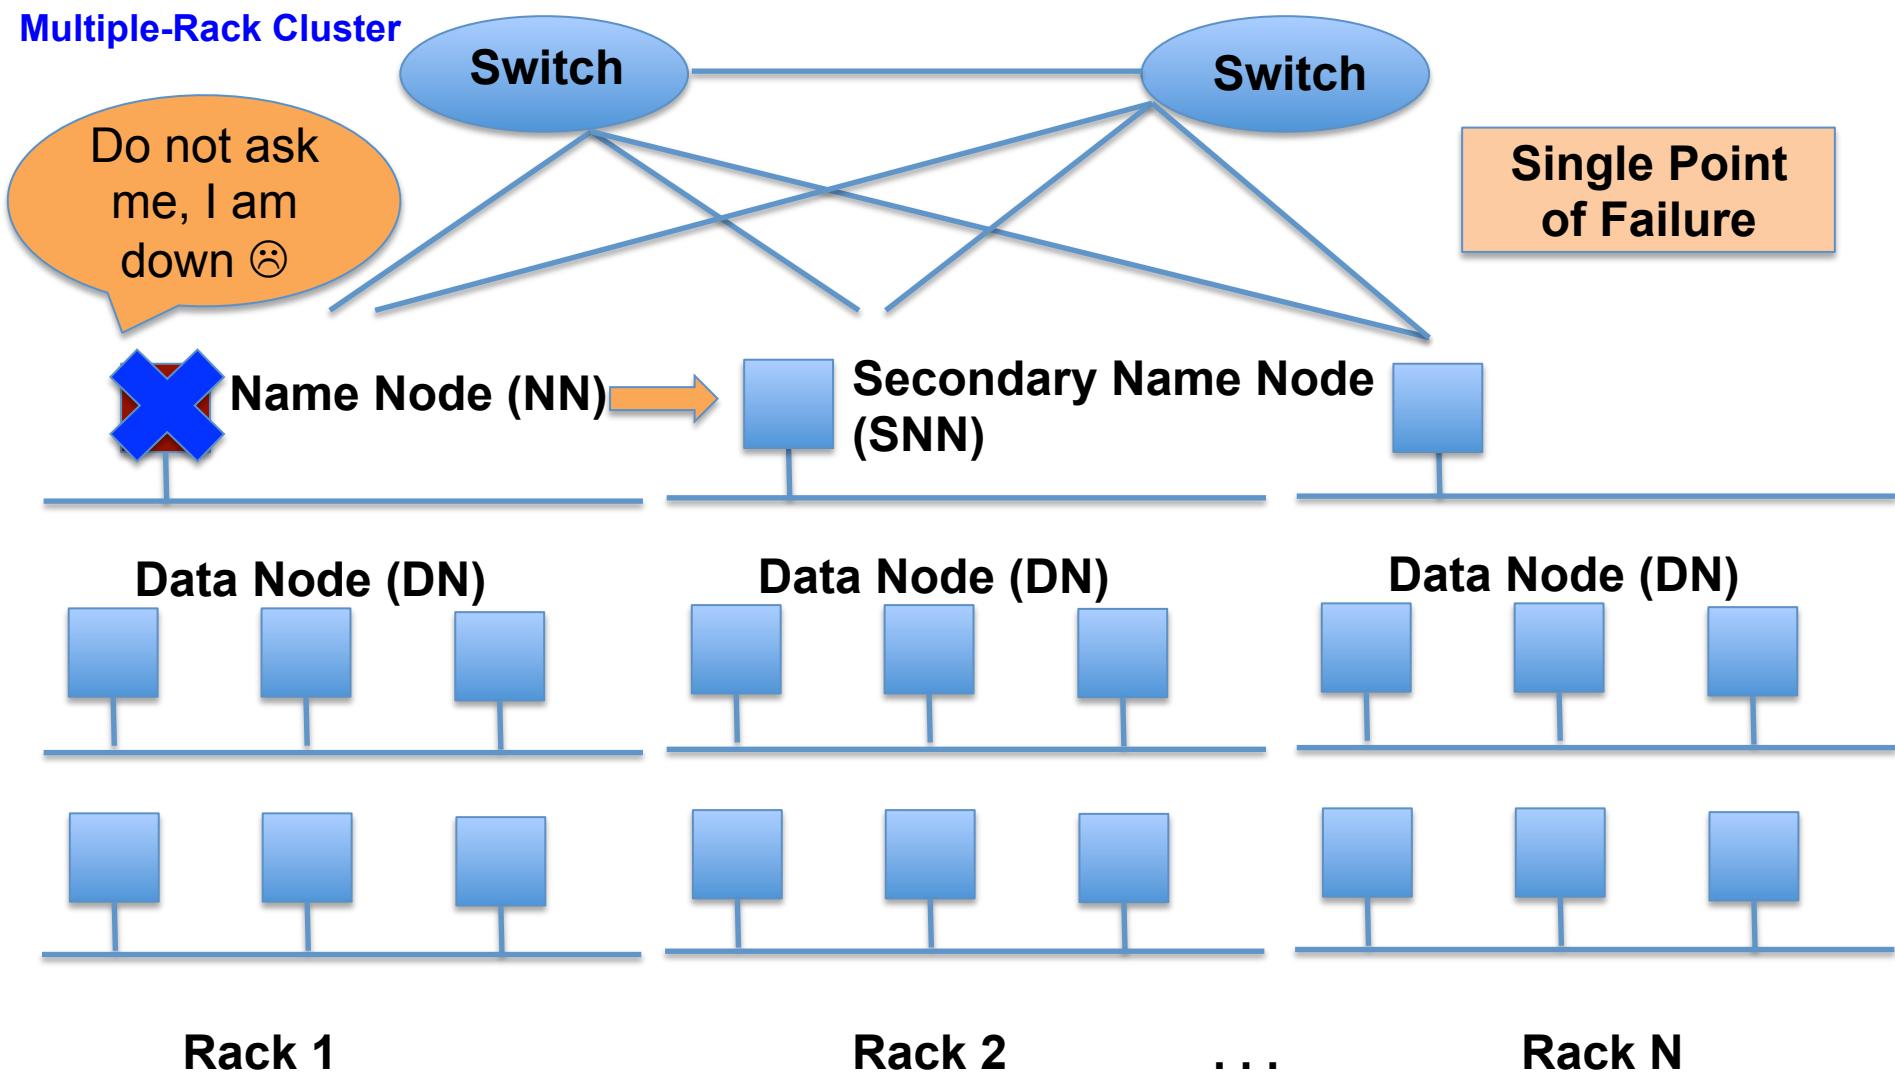
# HDFS architecture



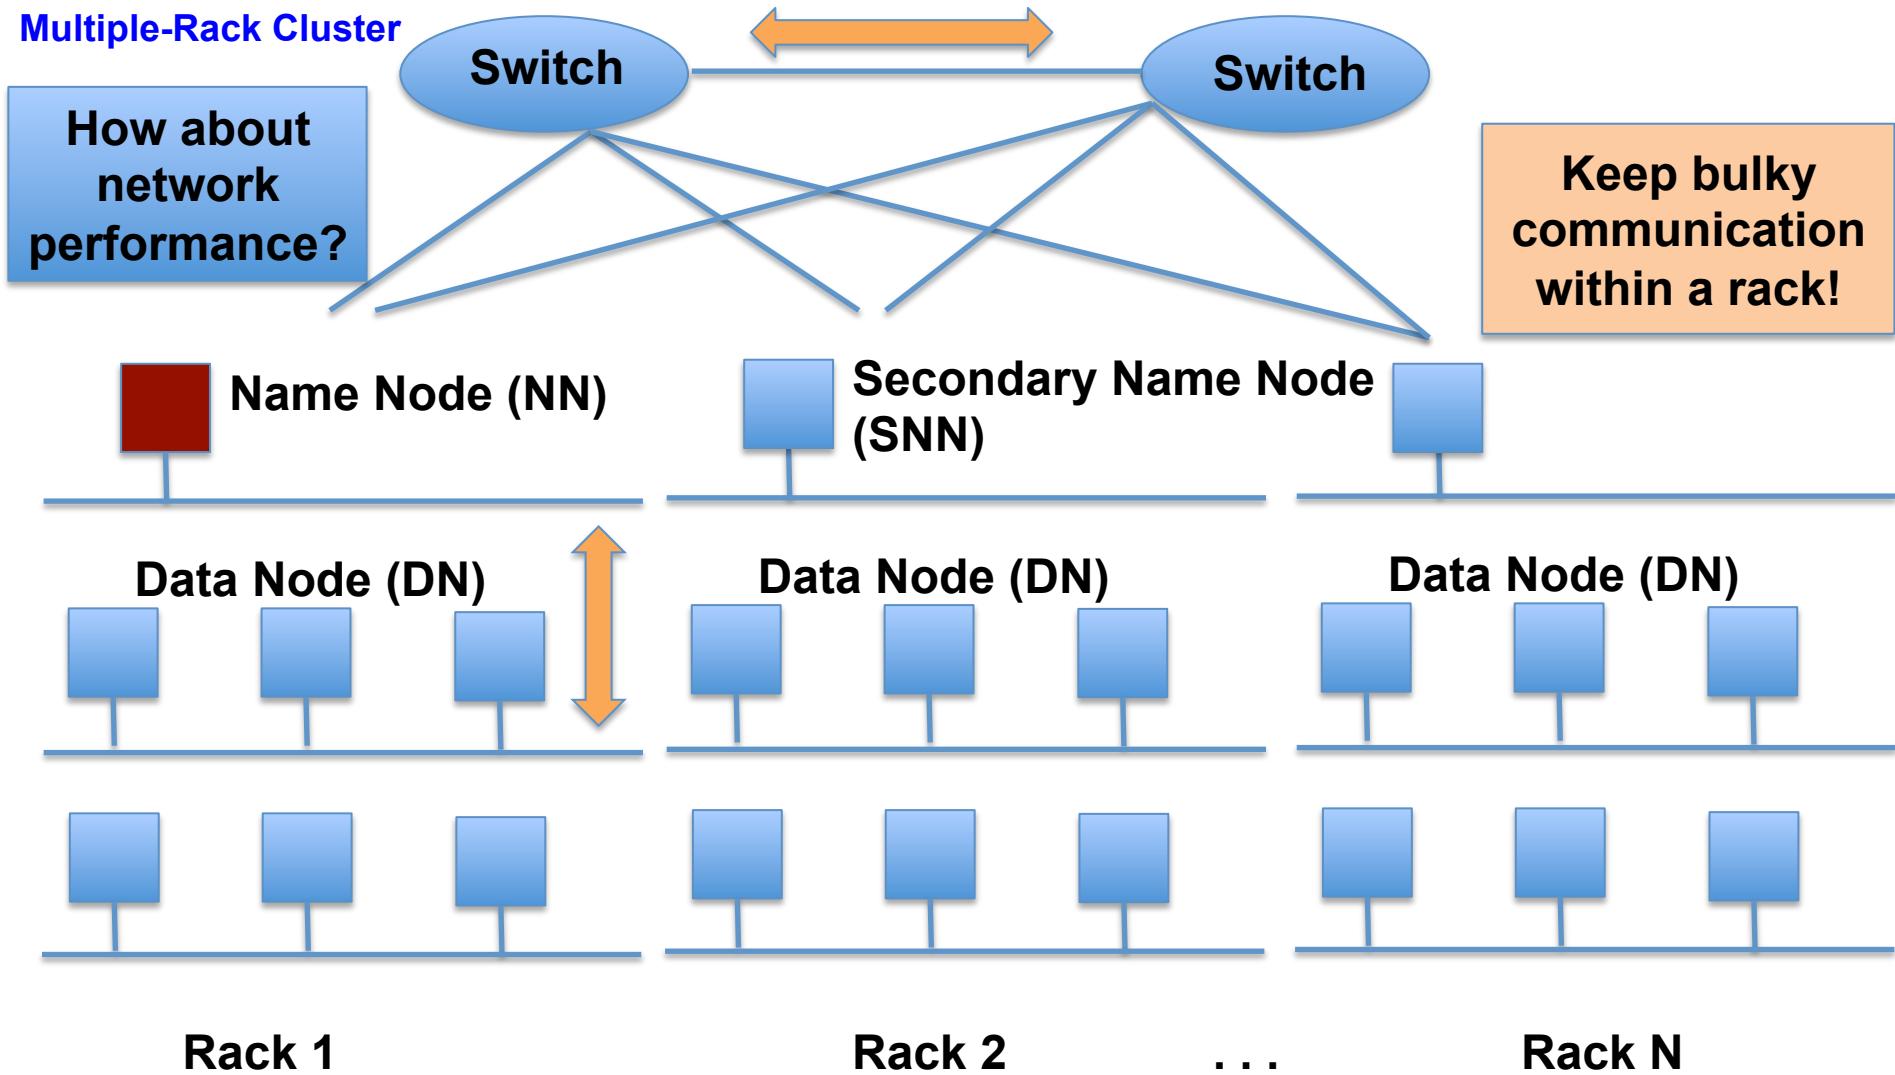
# HDFS architecture



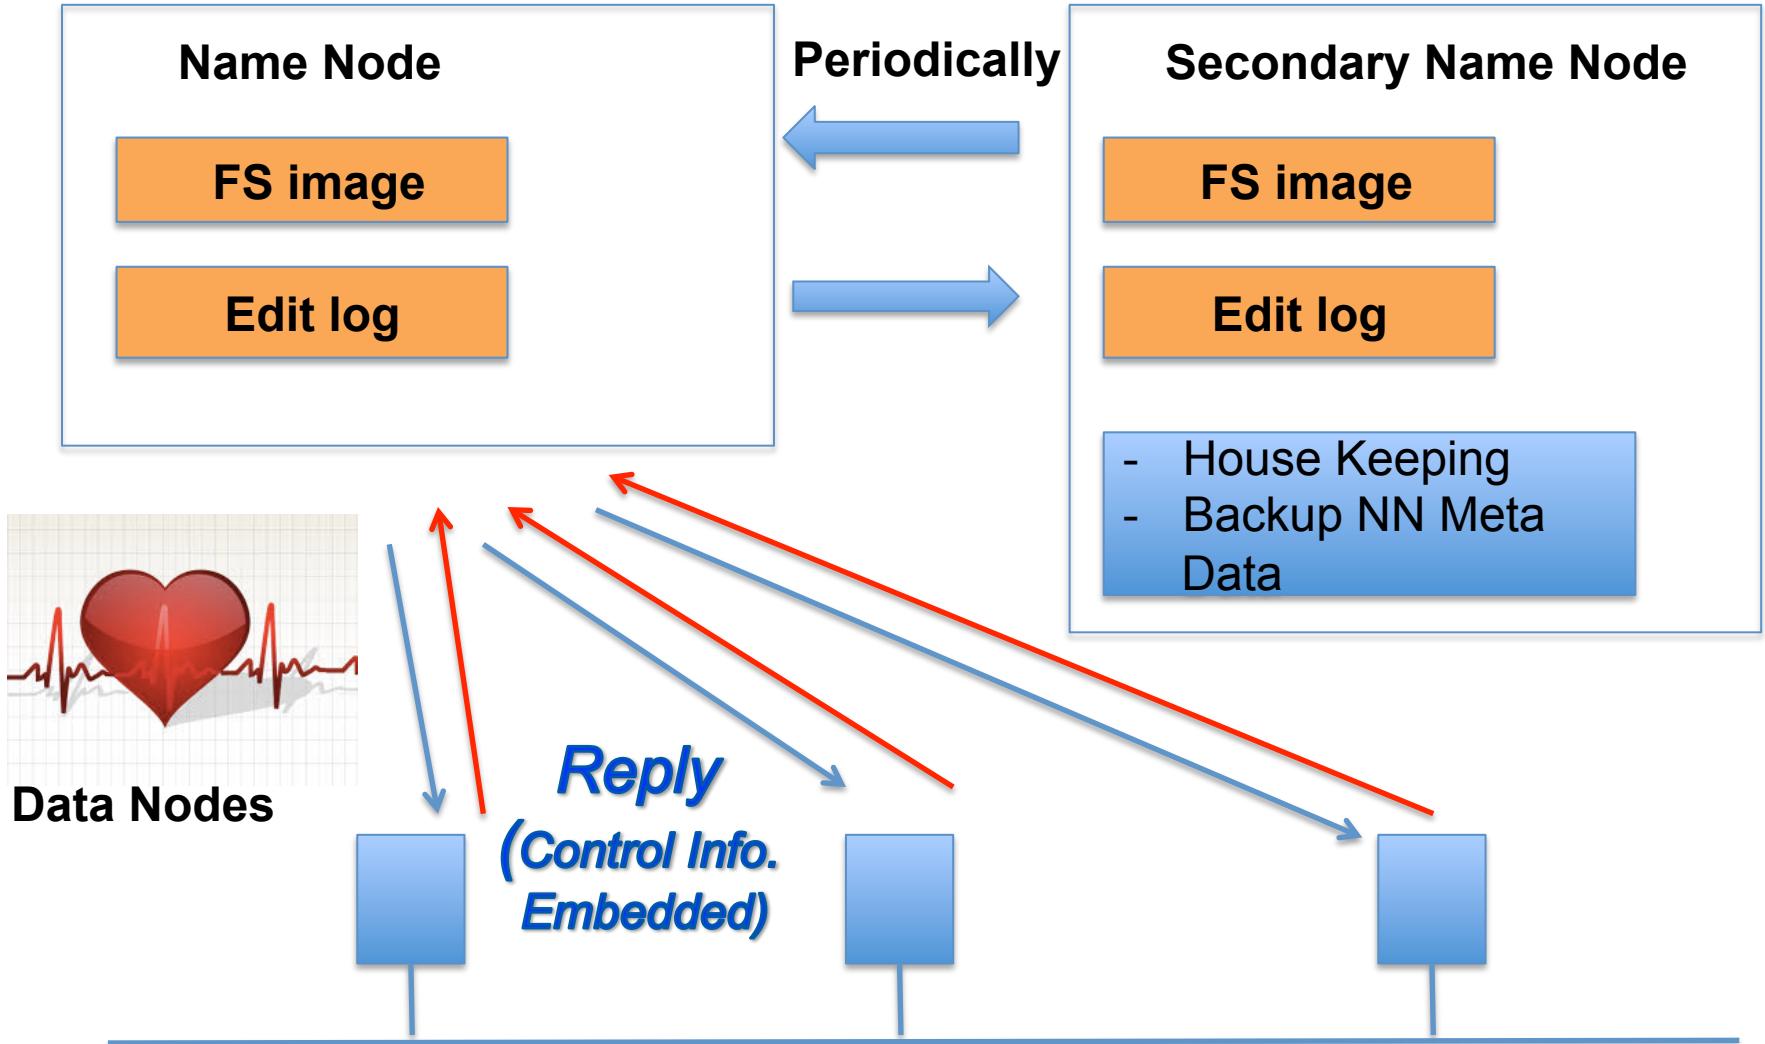
# HDFS architecture



# HDFS architecture



# HDFS Inside: Name Node



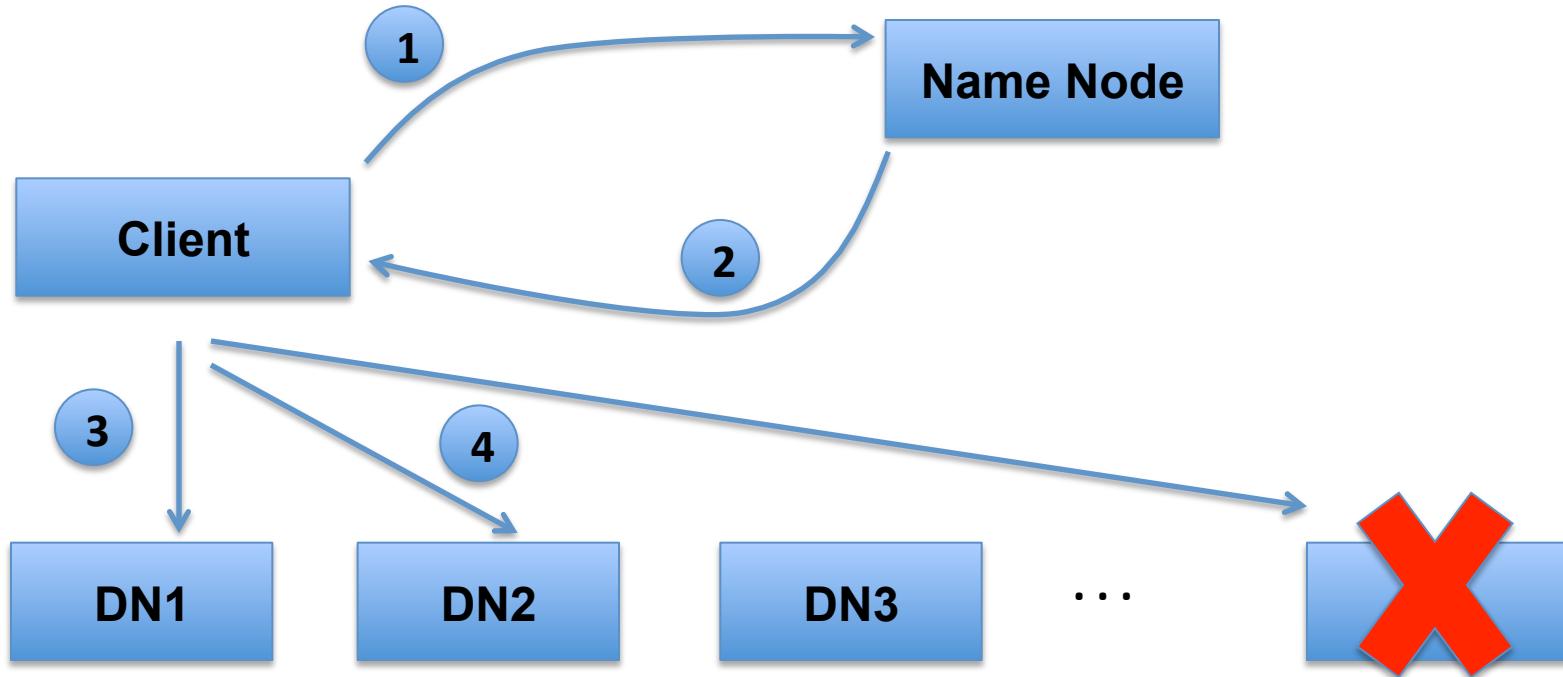
# HDFS Inside: Blocks

- Q: Why do we need the abstraction “Blocks” in addition to “Files”?
- Reasons:
  - File can be larger than a single disk
  - Block is of fixed size, easy to manage and manipulate
  - Easy to replicate and do more fine grained load balancing

# HDFS Inside: Blocks

- HDFS Block size is by default **64 MB**, why it is much larger than regular file system block?
- Reasons:
  - Minimize overhead: disk seek time is almost constant

# HDFS Inside: Read



1. Client connects to NN to read data
2. NN tells client where to find the data blocks
3. Client reads blocks directly from data nodes (without going through NN)
4. In case of node failures, client connects to another node that serves the missing block

# HDFS Inside: Read

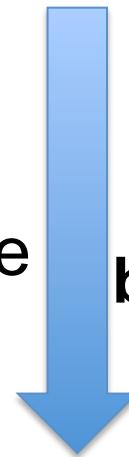
- Q: Why does HDFS choose such a design for read? Why not ask client to read blocks through NN?
- Reasons:
  - Prevent NN from being the bottleneck of the cluster
  - Allow HDFS to scale to large number of concurrent clients
  - Spread the data traffic across the cluster

# HDFS Inside: Read

- Q: Given multiple replicas of the same block, how does NN decide which replica the client should read?
- HDFS Solution:
  - Rack awareness based on network topology

# HDFS network topology

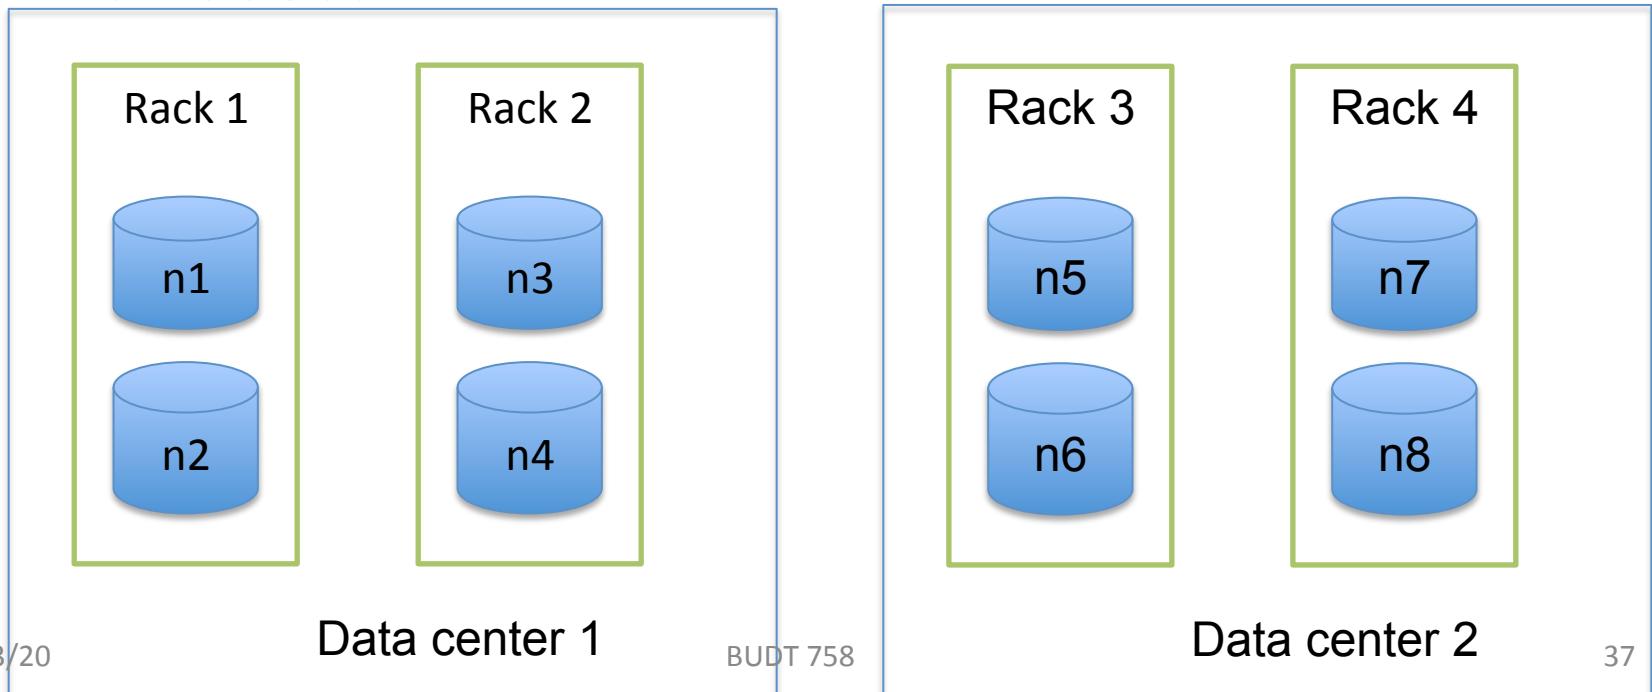
- The critical resource in HDFS is **bandwidth**, distance is defined based on that
- Measuring bandwidths between any pair of nodes is too complex and **does not scale**
- **Basic Idea:**
  - Processes on the same node
  - Different nodes on the same rack
  - Nodes on different racks in the same data center (cluster)
  - Nodes in different data centers



**Bandwidth becomes less**

# HDFS network topology

- HDFS takes a simple approach:
  - See the network as a tree
  - **Distance between two nodes is the sum of their distances to their closest common ancestor**



# HDFS network topology

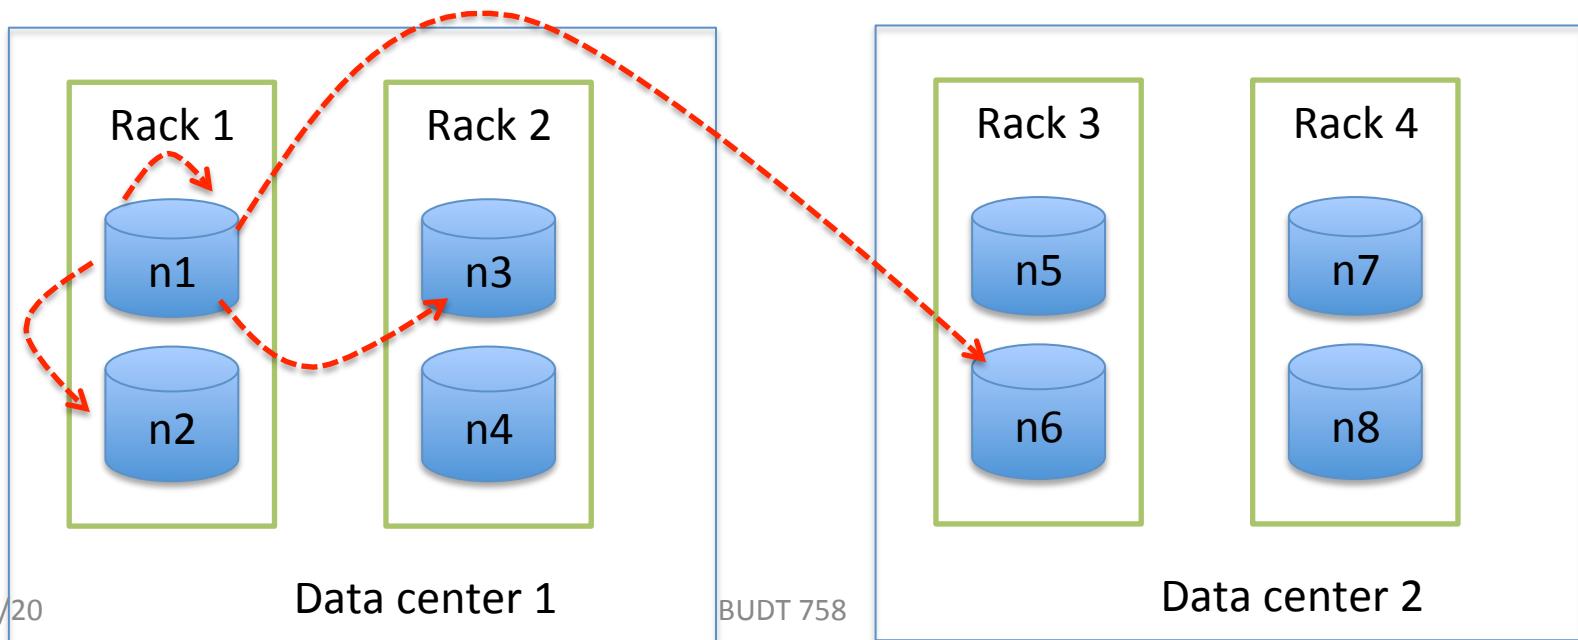
- What are the distance of the following pairs:

$\text{Dist}(d1/r1/n1, d1/r1/n1) = 0$

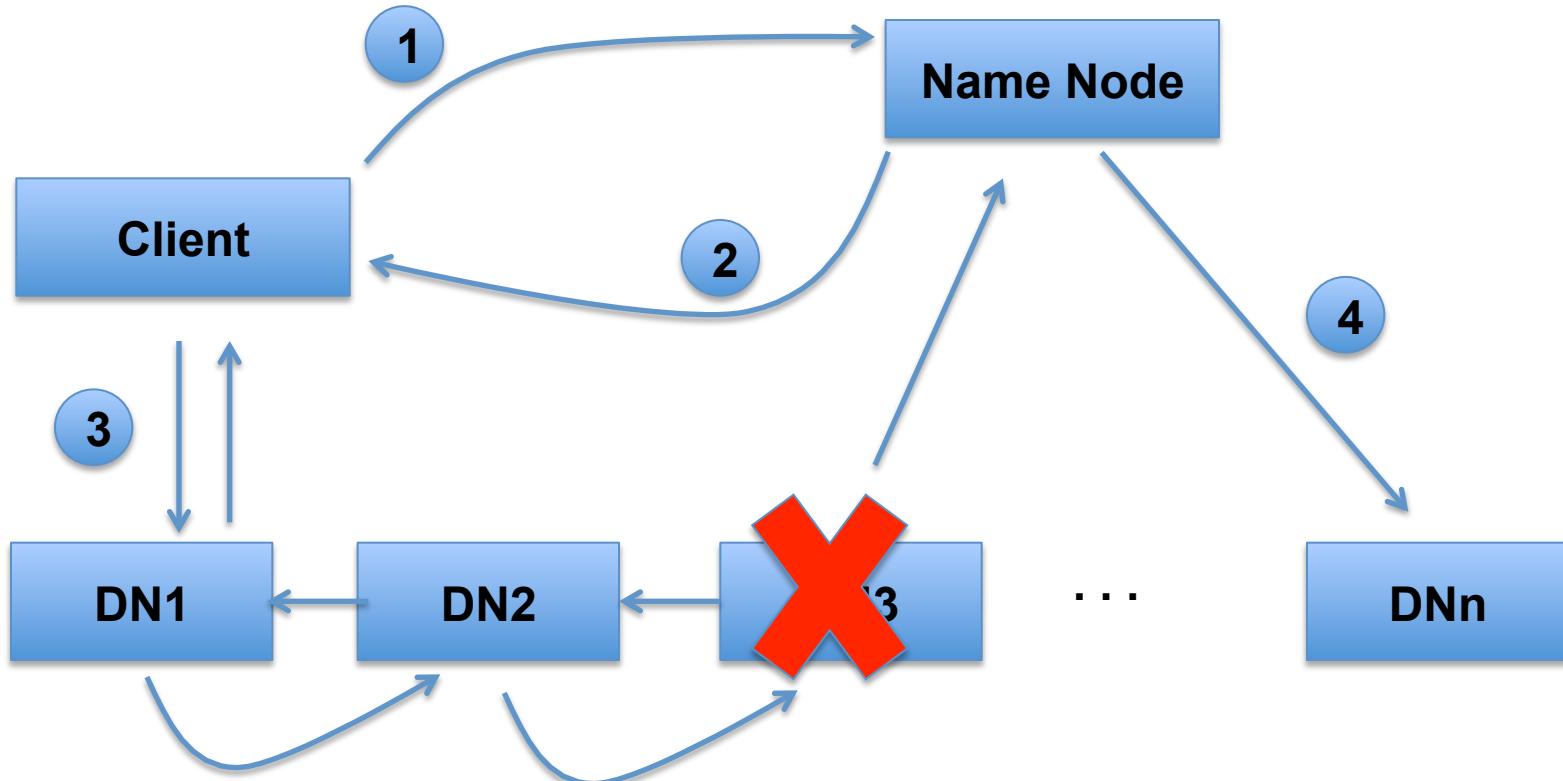
$\text{Dist}(d1/r1/n1, d1/r1/n2) = 2$

$\text{Dist}(d1/r1/n1, d1/r2/n3) = 4$

$\text{Dist}(d1/r1/n1, d2/r3/n6) = 6$



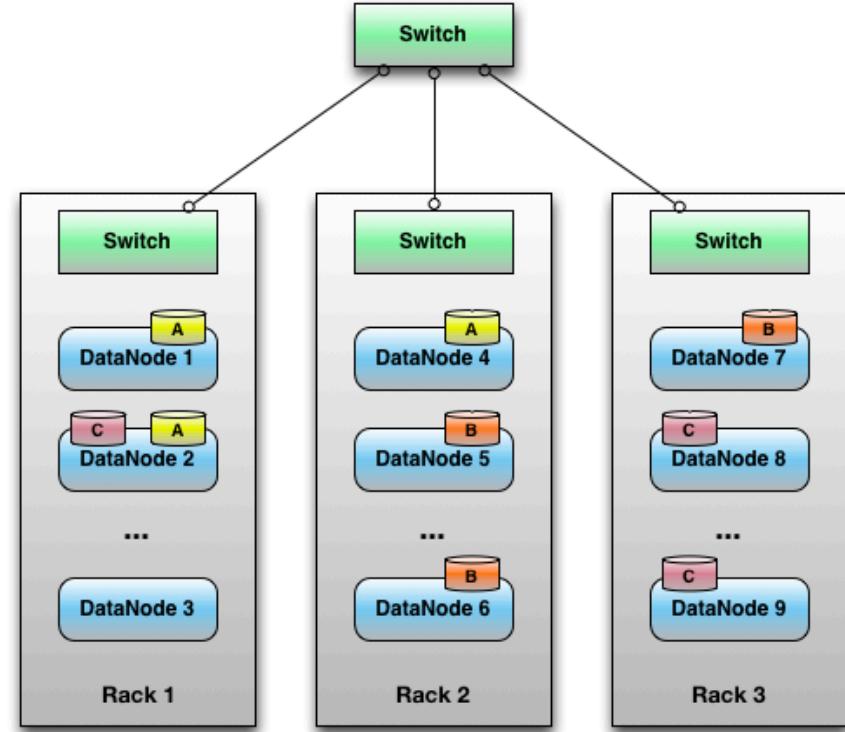
# HDFS Inside: Write



1. Client connects to NN to write data
2. NN tells client write these data nodes
3. Client writes blocks directly to data nodes with desired replication factor
4. In case of node failures, NN will figure it out and replicate the missing blocks

# Data replication

- **Frist copy** is written to the local node (write affinity).
- **Second copy** is written to a DataNode within a remote rack.
- **Third copy** is written to a DataNode in the same remote rack.
- **Additional** replicas are randomly placed.



**Objectives:** **load balancing, fast access, fault tolerance.**

# HDFS Inside: Write

- Replication Strategy vs Tradeoffs

	Reliability	Write Bandwidth	Read Bandwidth
Put all replicas on one node			
Put all replicas on different racks			

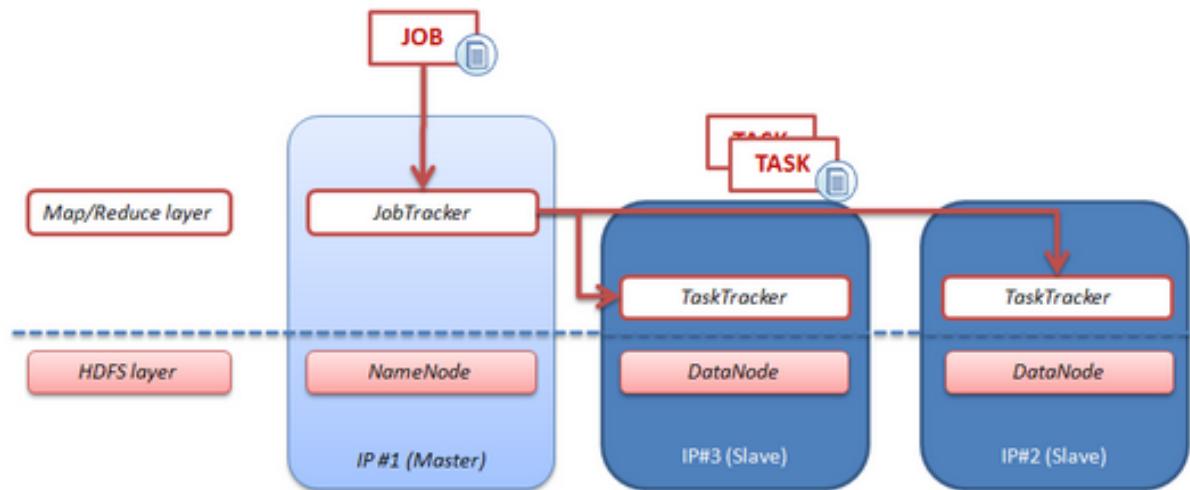
# HDFS Inside: Write

- Replication Strategy vs Tradeoffs

	Reliability	Write Bandwidth	Read Bandwidth
Put all replicas on one node			
Put all replicas on different racks			
HDFS: 1-> same node as client 2-> a node on different rack 3-> a different node on the same rack as 2			

# MapReduce: Hadoop execution layer

- **JobTracker** knows everything about submitted jobs
- Divides jobs into tasks and decides where to run each task
- Continuously communicating with TaskTracker



- **TaskTracker** execute task (multiple tasks per node)
- Monitors the execution of each task
- Continuously sending feedback to JobTracker